

Finding your way in R

Data Wrangling, Session 2

Kieran Healy

Code Horizons

July 22, 2024

Writing documents
Using Quarto to
produce and
reproduce your
work

Where we want to end up

Covid Cases

Kieran Healy

4/18/2021

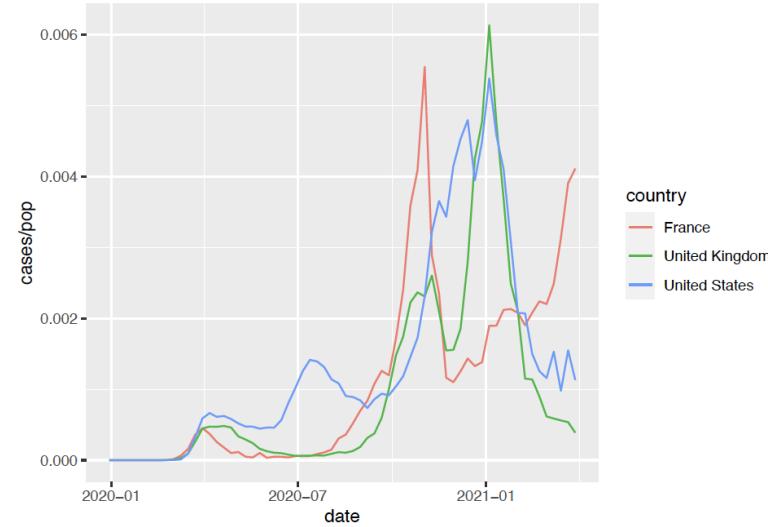
LOREM IPSUM

LOREM IPSUM dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Table 1: Total cases in three countries.

country	cases
United States	30706129
France	4822470
United Kingdom	4359388

LOREM IPSUM dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.



PDF out

Where we want to end up

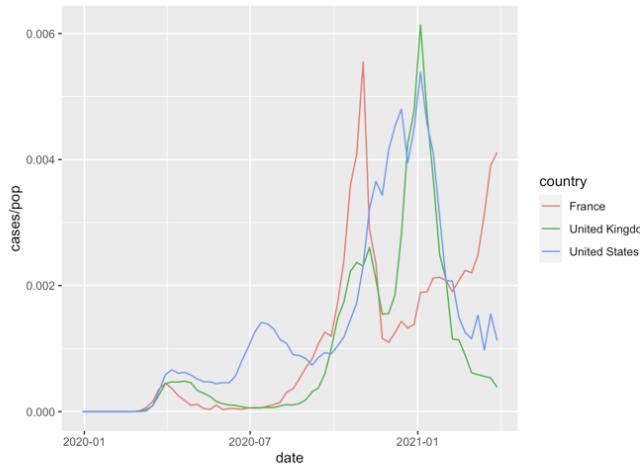
Lorem Ipsum

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Total cases in three countries.

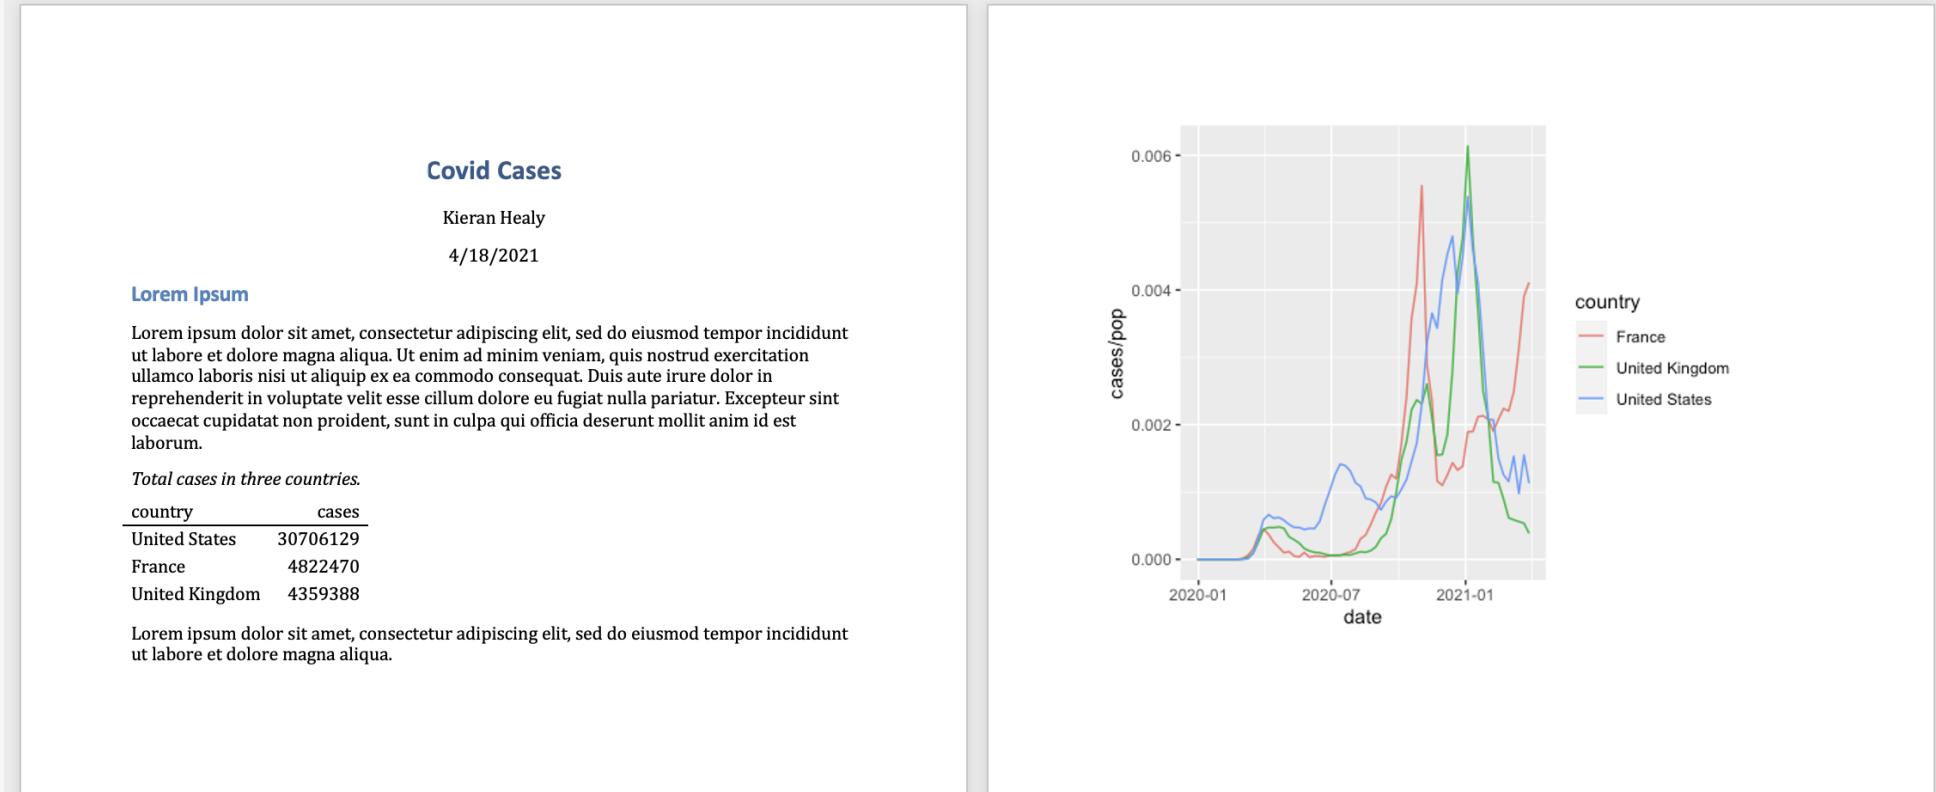
country	cases
United States	30706129
France	4822470
United Kingdom	4359388

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua.



HTML out

Where we want to end up



Word out

How to get there?

```
# COVID      covidcases.R  
  
# Get data from ECDC  
# FIXME Write a fn to  
# do this  
data_raw <- read_csv[url]  
  
# Clean it  
# Notes on the cleaning  
# process.  
  
covid <- data_raw %>%  
  mutate[...] %>%  
  select[...]  
  
# Make some plots  
covid %>%  
  ggplot [...] +  
  geom_line[...]
```

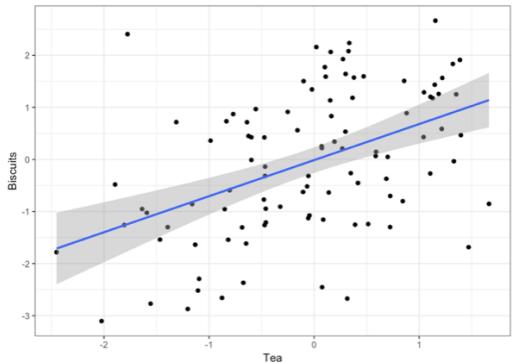
Write an R script with some notes inside. Create some figures and tables, paste them into our document.

This will work, but we can do better.

We can make this ...

1. Lorem Ipsum

 Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

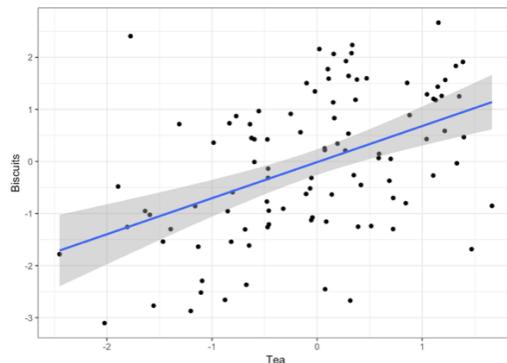


 Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

... by writing this

1. Lorem Ipsum

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enimad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.



Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Lorem Ipsum

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor* incididunt ut labore et dolore magna aliqua. Ut enimad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

```
library(ggplot2)
tea <- rnorm(100)
biscuits <- tea + rnorm(100, 0, 1.3)
data <- data.frame(tea, biscuits)
p <- ggplot(data, aes(x = tea, y = biscuits)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(x = "Tea", y = "Biscuits") + theme_bw()
print(p)
```

Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

The code gets replaced by its output

Lorem Ipsum

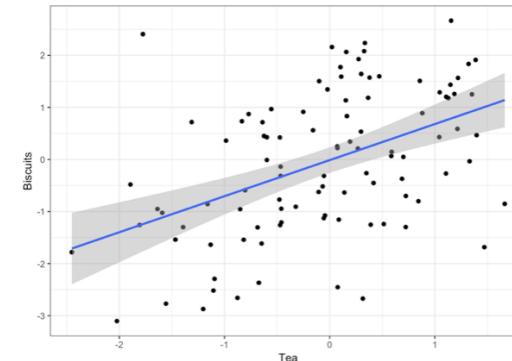
 Lorem ipsum dolor sit amet, consectetur adipisicing elit,
 sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enimad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat.

```
library(ggplot2)
tea <- rnorm(100)
biscuits <- tea + rnorm(100, 0, 1.3)
data <- data.frame(tea, biscuits)
p <- ggplot(data, aes(x = tea, y = biscuits)) +
  geom_point() +
  geom_smooth(method = "lm") +
  labs(x = "Tea", y = "Biscuits") + theme_bw()
print(p)
```

Duis aute irure dolor in reprehenderit in voluptate velit esse
cillum dolore eu fugiat nulla pariatur. Excepteur sint
occaecat cupidatat non proident, sunt in culpa qui officia
deserunt mollit anim id est laborum.

1. Lorem Ipsum

 Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do
 eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut
 enimad minim veniam, quis nostrud exercitation ullamco laboris
 nisi ut aliquip ex ea commodo consequat.



Duis aute irure dolor in reprehenderit in voluptate velit esse
cillum dolore eu fugiat nulla pariatur. Excepteur sint
occaecat cupidatat non proident, sunt in culpa qui officia
deserunt mollit anim id est laborum.

```
---
```

```
title: "Covid Cases"
author: "Kieran Healy"
date: "4/18/2021"
output: html_document
---
```

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(echo = FALSE)
```

## Introduction

We'll be looking at some COVID case data.

```{r libraries, message = FALSE}
library(tidyverse)
library(here)
library(janitor)
library(socviz)
```

```{r load-the-data, message = FALSE}
covid_cases <- read_csv("data/national_cases.csv")
```

## Lorem Ipsum



Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.



```{r case-table}
covid_cases %>%
 filter(country %in% c("United States", "United Kingdom", "France")) %>%
 group_by(country) %>%
 summarize(cases = sum(cases)) %>%
 arrange(desc(cases)) %>%
 kable(caption = "Total cases in three countries.")
```

```

```
---
```

```
title: "Covid Cases"
author: "Kieran Healy"
date: "4/18/2021"
output: html_document
```

```
--
```

```
```{r setup, include=FALSE}
knitr::opts_chunk$set(echo = FALSE)
```

```
Introduction
```

```
We'll be looking at some COVID case data.
```

```
```{r libraries, message = FALSE}
```

```
library(tidyverse)
library(here)
library(janitor)
library(socviz)
```

```
```{r load-the-data, message = FALSE}
covid_cases <- read_csv("data/national_cases.csv")
```
```

```
## Lorem Ipsum
```

```
Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.
```

```
```{r case-table}
covid_cases %>%
 filter(country %in% c("United States", "United Kingdom", "France")) %>%
 group_by(country) %>%
 summarize(cases = sum(cases)) %>%
 arrange(desc(cases)) %>%
 kable(caption = "Total cases in three countries.")
```

```

Header section with metadata

Code chunk

In RStudio, code chunks can be "played" one at a time

Code chunks can have their own labels and options

Text with Markdown formatting

Chunks are replaced by their output when the document is knitted

Covid Cases

Kieran Healy

4/18/2021

Introduction

We'll be looking at some COVID case data.

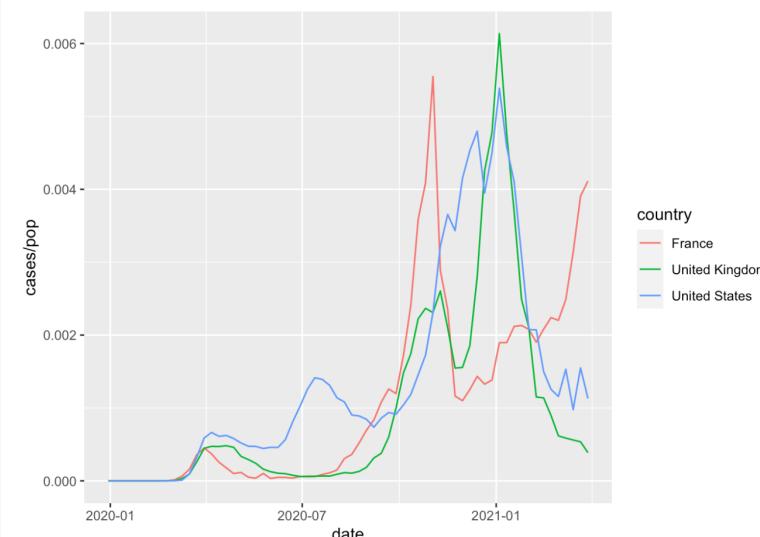
Lorem Ipsum

Lorem ipsum dolor sit amet, consectetur adipiscing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

Total cases in three countries.

| country | cases |
|----------------|----------|
| United States | 30706129 |
| France | 4822470 |
| United Kingdom | 4359388 |

Placeholder text for the chart area.



This approach has its limitations, but it's *very* useful.

Report notes.Rmd

We can see this *relationship* in a scatterplot.

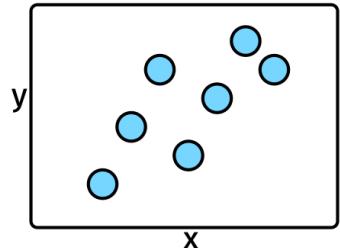
```
```{r my-code}
p <- ggplot(data, mapping)
p + geom_point()
````
```

As you can see, this plot looks pretty nice.

knit in R

Report notes.html

We can see this *relationship* in a scatterplot.



As you can see, this plot looks pretty nice.

When learning these workflows, stick with the defaults at the beginning. Later, you can customize the look of the output in all kinds of ways.

The right frame of mind

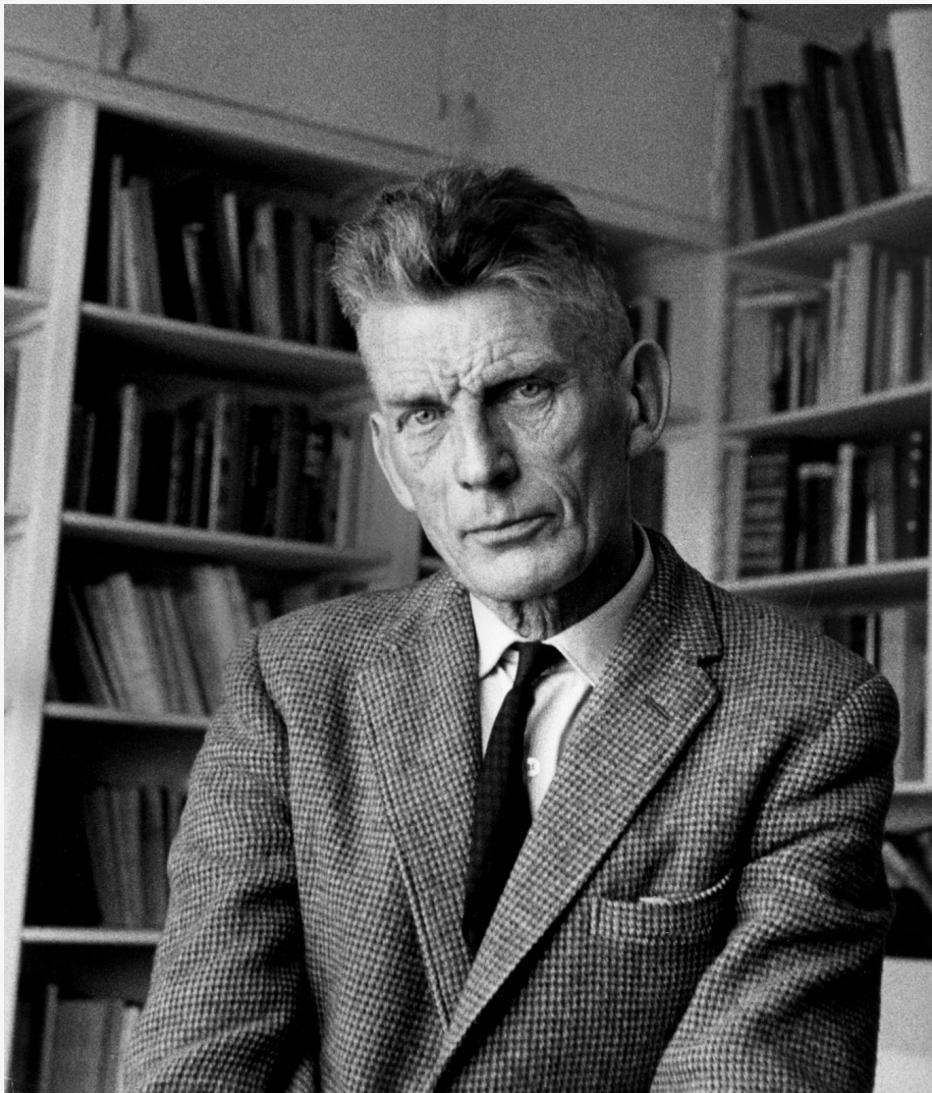
This is like learning how to drive a car, or how to cook in a kitchen ... or learning to speak a language.

After some orientation to what's where, you will learn best by *doing*.

Software is a pain, but you won't crash the car or burn your house down.

TYPE OUT YOUR CODE BY HAND

By far the best way to get a feel for how R works, and to get comfortable with the cycle of try-fail-cry-wail-retry that is a permanent part of writing any code.



Samuel Beckett

Ever tried.
Ever failed.
No matter.
Try again.
Fail again.
Fail better.

Samuel Beckett,
early data analyst

Getting Oriented

Loading the tidyverse libraries

```
library(tidyverse)
-- Attaching core tidyverse packages ━━━━━━━━━━━━━━ tidyverse 2.0.0 ━━━━━
✓ dplyr     1.1.3      ✓ readr     2.1.4
✓forcats   1.0.0      ✓ stringr   1.5.0
✓ ggplot2   3.4.4      ✓ tibble    3.2.1
✓ lubridate 1.9.3      ✓ tidyverse  1.3.0
✓ purrr    1.0.2
-- Conflicts ━━━━━━━━━━━━━━ tidyverse_conflicts() ━━━
✖ dplyr::filter() masks stats::filter()
✖ dplyr::lag()   masks stats::lag()
ℹ Use the conflicted package to force all conflicts to become errors
```

The tidyverse has several components.

We'll return to this message about Conflicts later.

Tidyverse components

```
library(tidyverse)
```

Call the package and ...

```
Loading tidyverse: ggplot2
```

◀ Draw graphs

```
Loading tidyverse: tibble
```

◀ Nicer data tables

```
Loading tidyverse: tidyr
```

◀ Tidy your data

```
Loading tidyverse: readr
```

◀ Get data into R

```
Loading tidyverse: purrr
```

◀ Fancy Iteration

```
Loading tidyverse: dplyr
```

◀ Action verbs for tables

What R looks like

Code you can type and run:

```
## Inside code chunks, lines beginning with a # character are comments
## Comments are ignored by R

my_numbers ← c(1, 1, 2, 4, 1, 3, 1, 5) # Anything after a # character is ignored as well
```

Output:

```
my_numbers
[1] 1 1 2 4 1 3 1 5
```

This is equivalent to running the code above, typing `my_numbers` at the console, and hitting enter.

What R looks like

By convention, code output in documents is prefixed by `##`

Also by convention, outputting vectors, etc, gets a counter keeping track of the number of elements. For example,

```
letters
```

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s"  
[20] "t" "u" "v" "w" "x" "y" "z"
```

SOME THINGS TO KNOW ABOUT R

0. It's a calculator

Arithmetic

```
(31 * 12) / 2^4
```

```
[1] 23.25
```

```
sqrt(25)
```

```
[1] 5
```

```
log(100)
```

```
[1] 4.60517
```

```
log10(100)
```

```
[1] 2
```

0. It's a calculator

Arithmetic

```
(31 * 12) / 2^4
```

```
[1] 23.25
```

```
sqrt(25)
```

```
[1] 5
```

```
log(100)
```

```
[1] 4.60517
```

```
log10(100)
```

```
[1] 2
```

Logic

```
4 < 10
```

```
[1] TRUE
```

```
4 > 2 & 1 > 0.5 # The "&" means "and"
```

```
[1] TRUE
```

```
4 < 2 | 1 > 0.5 # The "/" means "or"
```

```
[1] TRUE
```

```
4 < 2 | 1 < 0.5
```

```
[1] FALSE
```

Boolean and Logical operators

Logical equality and inequality (yielding a `TRUE` or `FALSE` result) is done with `=` and `!=`. Other logical operators include `<`, `>`, `≤`, `≥`, and `!` for negation.

```
## A logical test  
2 = 2 # Write `=` twice
```

```
[1] TRUE
```

```
## This will cause an error, because R will think you are trying to assign a value  
2 = 2
```

```
## Error in 2 = 2 : invalid (do_set) left-hand side to assignment
```

```
3 != 7 # Write `!` and then `=` to make `!=`
```

```
[1] TRUE
```

Watch out!

Here's a gotcha. You might think you could write `3 < 5 & 7` and have it be interpreted as "Three is less than five and also less than seven [True or False?]" :

```
3 < 5 & 7
```

```
[1] TRUE
```

It seems to work!

Watch out!

But now try `3 < 5 & 1`, where your intention is “Three is less than five and also less than one [True or False?]”

```
3 < 5 & 1
```

```
[1] TRUE
```

What's happening is that `3 < 5` is evaluated first, and resolves to `TRUE`, leaving us with the expression `TRUE & 1`.

R interprets this as `TRUE & as.logical(1)`.

In Boolean algebra, `1` resolves to `TRUE`. Any other number is `FALSE`. So,

Watch out!

```
TRUE & as.logical(1)
```

```
[1] TRUE
```

```
3 < 5 & 3 < 1
```

```
[1] FALSE
```

You have to make your comparisons explicit.

Logic and floating point arithmetic

Let's evaluate $0.6 + 0.2 = 0.8$

Logic and floating point arithmetic

Let's evaluate $0.6 + 0.2 = 0.8$

```
0.6 + 0.2 = 0.8
```

```
[1] TRUE
```

Logic and floating point arithmetic

Let's evaluate $0.6 + 0.2 = 0.8$

```
0.6 + 0.2 = 0.8
```

```
[1] TRUE
```

Now let's try $0.6 + 0.3 = 0.9$

Logic and floating point arithmetic

Let's evaluate $0.6 + 0.2 = 0.8$

```
0.6 + 0.2 = 0.8
```

```
[1] TRUE
```

Now let's try $0.6 + 0.3 = 0.9$

```
0.6 + 0.3 = 0.9
```

```
[1] FALSE
```

Er. That's not right.

Welcome to floating point math!

In Base 10, you can't precisely express fractions like $\frac{1}{3}$ and $\frac{1}{9}$. They come out as repeating decimals: 0.3333... or 0.1111... You *can* cleanly represent fractions that use a prime factor of the base, which in the case of Base 10 are 2 and 5.

Welcome to floating point math!

In Base 10, you can't precisely express fractions like $\frac{1}{3}$ and $\frac{1}{9}$. They come out as repeating decimals: 0.3333... or 0.1111... You *can* cleanly represent fractions that use a prime factor of the base, which in the case of Base 10 are 2 and 5.

Computers represent numbers as binary (i.e. Base 2) floating-points. In Base 2, the only prime factor is 2. So $\frac{1}{5}$ or $\frac{1}{10}$ in binary would be repeating.

Logic and floating point arithmetic

When you do binary math on repeating numbers and convert back to decimals you get tiny leftovers, and this can mess up *logical* comparisons of equality. The `all.equal()` function exists for this purpose.

```
print(.1 + .2)
```

```
[1] 0.3
```

```
print(.1 + .2, digits=18)
```

```
[1] 0.30000000000000044
```

```
all.equal(.1 + .2, 0.3)
```

```
[1] TRUE
```

See e.g. <https://0.3000000000000004.com>

More later on why
this might bite
you, and how to
deal with it

For now, “Be very careful about doing logical comparisons on floating-point numbers” is not a bad rule.

1. Everything in R has a name

```
my_numbers # We created this a few minutes ago
```

```
[1] 1 1 2 4 1 3 1 5
```

```
letters # This one is built-in
```

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s"  
[20] "t" "u" "v" "w" "x" "y" "z"
```

```
pi # Also built-in
```

```
[1] 3.141593
```

Some names are forbidden

Or it's a *really* bad idea to try to use them

```
TRUE  
FALSE  
Inf  
NaN  
NA  
NULL
```

```
for  
if  
while  
break  
function
```

2. Everything is an object

There are a few built-in objects:

```
letters
```

```
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o" "p" "q" "r" "s"  
[20] "t" "u" "v" "w" "x" "y" "z"
```

```
pi
```

```
[1] 3.141593
```

```
LETTERS
```

```
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S"  
[20] "T" "U" "V" "W" "X" "Y" "Z"
```

3. You can create objects

In fact, this is mostly what we will be doing.

Objects are created by **assigning** a thing to a name:

```
## name ... gets ... this stuff  
my_numbers ← c(1, 2, 3, 1, 3, 5, 25, 10)  
  
## name ... gets ... the output of the function `c()`  
your_numbers ← c(5, 31, 71, 1, 3, 21, 6, 52)
```

The **c()** function *combines* or *concatenates* things

The assignment operator

The assignment operator performs the action of creating objects

Use a keyboard shortcut to write it:

Press **option** and **-** on a Mac

Press **alt** and **-** on Windows

Assignment with =

You can use = as well as ← for assignment.

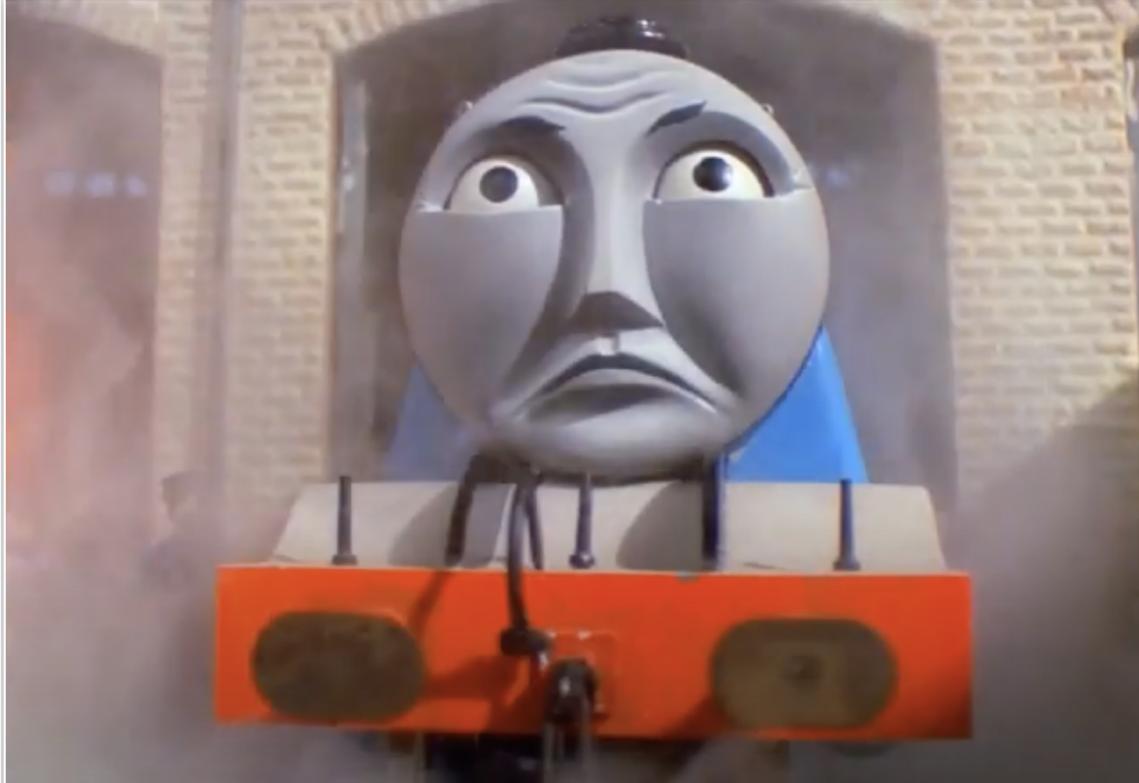
```
my_numbers = c(1, 2, 3, 1, 3, 5, 25)  
my_numbers  
[1] 1 2 3 1 3 5 25
```

On the other hand, = has a different meaning when used in functions.

I'm going to use ← for assignment throughout.

Be consistent either way.

Assignment with =



***It isn't wrong,
but we just
don't do it.***

4. Do things to objects with functions

```
## this object ... gets ... the output of this function  
my_numbers <- c(1, 2, 3, 1, 3, 5, 25, 10)
```

```
your_numbers <- c(5, 31, 71, 1, 3, 21, 6, 52)
```

```
my_numbers
```

```
[1] 1 2 3 1 3 5 25 10
```

4. Do things to objects with functions

Functions can be identified by the parentheses after their names.

```
my_numbers
```

```
[1] 1 2 3 1 3 5 25 10
```

```
## If you run this you'll get an error  
mean()
```

What functions usually do

They take **inputs** to **arguments**

They perform **actions**

They produce, or return, **outputs**

mean(x = my_numbers)

What functions usually do

They take **inputs** to **arguments**

They perform **actions**

They produce, or return, **outputs**

mean(x = my_numbers)

[1] **6.25**

What functions usually do

```
## Get the mean of what? Of x.  
## You need to tell the function what x is  
mean(x = my_numbers)
```

```
[1] 6.25
```

```
mean(x = your_numbers)
```

```
[1] 23.75
```

What functions usually do

If you don't *name* the arguments, R assumes you are providing them in the order the function expects.

```
mean(your_numbers)
```

```
[1] 23.75
```

What functions usually do

What arguments? Which order? Read the function's help page

```
help(mean)  
## quicker  
?mean
```

How to read an R help page?

What functions usually do

Arguments often tell the function what to do in specific circumstances

```
missing_numbers ← c(1:10, NA, 20, 32, 50, 104, 32, 147, 99, NA, 45)
```

```
mean(missing_numbers)
```

```
[1] NA
```

```
mean(missing_numbers, na.rm = TRUE)
```

```
[1] 32.44444
```

Or select from one of several options

```
## Look at ?mean to see what `trim` does  
mean(missing_numbers, na.rm = TRUE, trim = 0.1)
```

```
[1] 27.25
```

What functions usually do

There are all kinds of functions. They return different things.

```
summary(my_numbers)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|------|---------|--------|------|---------|-------|
| 1.00 | 1.75 | 3.00 | 6.25 | 6.25 | 25.00 |

What functions usually do

You can assign the output of a function to a name, which turns it into an object. (Otherwise it'll send its output to the console.)

```
my_summary ← summary(my_numbers)
```

```
my_summary
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|------|---------|--------|------|---------|-------|
| 1.00 | 1.75 | 3.00 | 6.25 | 6.25 | 25.00 |

What functions usually do

Objects hang around in your work environment until they are overwritten by you, or are deleted.

```
## rm() function removes objects
rm(my_summary)

my_summary

## Error: object 'my_summary' not found
```

Functions can be nested

```
c(1:20)
```

```
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
```

```
mean(c(1:20))
```

```
[1] 10.5
```

```
summary(mean(c(1:20)))
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|------|---------|--------|------|---------|------|
| 10.5 | 10.5 | 10.5 | 10.5 | 10.5 | 10.5 |

```
names(summary(mean(c(1:20))))
```

```
[1] "Min."      "1st Qu."   "Median"    "Mean"      "3rd Qu."   "Max."
```

```
length(names(summary(mean(c(1:20)))))
```

```
[1] 6
```

Nested functions are evaluated from the inside out.

Use the pipe operator: |>

Instead of deeply nesting functions in parentheses, we can use the *pipe operator*:

```
c(1:20) %> mean() %> summary() %> names() %> length()
```

```
[1] 6
```

Read this operator as “*and then*”

Use the pipe operator: |>

Better, vertical space is free in R:

```
c(1:20) |>  
  mean() |>  
  summary() |>  
  names() |>  
  length()
```

```
[1] 6
```

Pipelines make code more **readable**

Not great, Bob:

```
serve(stir(pour_in_pan(whisk(crack_eggs(get_from_fridge(eggs), into = "bowl"), len = 40), temp = "med-high")))
```

Notice how the first thing you read is the last operation performed.

Pipelines make code more **readable**

We can use vertical space and indents, but it's really not much better:

```
serve(
    stir(
        pour_in_pan(
            whisk(
                crack_eggs(
                    get_from_fridge(eggs),
                    into = "bowl"),
                len = 40),
            temp = "med-high")
    )
)
```

Pipelines make code more **readable**

Much nicer:

```
eggs ▷  
  get_from_fridge() ▷  
  crack_eggs(into = "bowl") ▷  
  whisk(len = 40) ▷  
  pour_in_pan(temp = "med-high") ▷  
  stir() ▷  
  serve()
```

We'll still use nested parentheses quite a bit, often in the context of a function working inside a pipeline. But it's good not to have too many levels of nesting.

The other pipe: %>%

The Base R pipe operator, `>` is a relatively recent addition to R.

Piping operations were originally introduced in a package called `magrittr`, where it took the form `%>%`

The other pipe: %>%

The Base R pipe operator, `>` is a relatively recent addition to R.

Piping operations were originally introduced in a package called `magrittr`, where it took the form `%>%`

It's been so successful, a version of it has been incorporated into Base R. It *mostly* but does not *quite* work the same way as `%>%` in every case.

The other pipe: %>%

The Base R pipe operator, `>` is a relatively recent addition to R.

Piping operations were originally introduced in a package called `magrittr`, where it took the form `%>%`

It's been so successful, a version of it has been incorporated into Base R. It *mostly* but does not *quite* work the same way as `%>%` in every case. We'll use the Base R pipe in this course, but you'll see the Magrittr pipe a lot out in the world.

With the Base R pipe, you can only pass an object to the *first* argument in a function. This is fine for most tidyverse pipelines, where the first argument is usually (implicitly) the data. But it does mean that most Base R functions will continue not to be easily piped, as most of them do not follow the convention of passing the current data as the first argument

Functions are bundled into packages

Packages are loaded into your working environment using the `library()` function:

```
## A package containing a dataset rather than functions
library(gapminder)
```

```
gapminder
```

```
# A tibble: 1,704 × 6
  country continent year lifeExp      pop gdpPercap
  <fct>     <fct>   <int>    <dbl>    <int>      <dbl>
1 Afghanistan Asia     1952     28.8  8425333      779.
2 Afghanistan Asia     1957     30.3  9240934      821.
3 Afghanistan Asia     1962     32.0  10267083     853.
4 Afghanistan Asia     1967     34.0  11537966     836.
5 Afghanistan Asia     1972     36.1  13079460     740.
6 Afghanistan Asia     1977     38.4  14880372     786.
7 Afghanistan Asia     1982     39.9  12881816     978.
8 Afghanistan Asia     1987     40.8  13867957     852.
9 Afghanistan Asia     1992     41.7  16317921     649.
10 Afghanistan Asia    1997     41.8  22227415     635.
# i 1,694 more rows
```

Functions are bundled into packages

You need only *install* a package once (and occasionally update it):

```
## Do at least once for each package. Once done, not needed each time.  
install.packages("palmerpenguins", repos = "http://cran.rstudio.com")  
  
## Needed sometimes, especially after an R major version upgrade.  
update.packages(repos = "http://cran.rstudio.com")
```

Functions are bundled into packages

But you must *load* the package in each R session before you can access its contents:

```
## To load a package, usually at the start of your RMarkdown document or script file
library(palmerpenguins)
penguins

# A tibble: 344 × 8
  species   island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
  <fct>     <fct>        <dbl>          <dbl>            <int>        <int>
1 Adelie    Torgersen      39.1           18.7            181        3750
2 Adelie    Torgersen      39.5           17.4            186        3800
3 Adelie    Torgersen      40.3            18              195        3250
4 Adelie    Torgersen       NA             NA              NA         NA
5 Adelie    Torgersen      36.7           19.3            193        3450
6 Adelie    Torgersen      39.3           20.6            190        3650
7 Adelie    Torgersen      38.9           17.8            181        3625
8 Adelie    Torgersen      39.2           19.6            195        4675
9 Adelie    Torgersen      34.1           18.1            193        3475
10 Adelie   Torgersen       42              20.2            190        4250
# i 334 more rows
# i 2 more variables: sex <fct>, year <int>
```

Grabbing a single function with ::

“Reach in” to an unloaded package and grab a function directly, using
`<package> :: <function>`

Grabbing a single function with ::

```
## A little glimpse of what we'll do soon
penguins %>
  select(species, body_mass_g, sex) %>
  gtsummary::tbl_summary(by = species)
```

| Characteristic | Adelie, N = 152 | Chinstrap, N = 68 | Gentoo, N = 124 |
|---------------------------|-----------------------|-----------------------|-----------------------|
| body_mass_g, Median (IQR) | 3,700 (3,350 – 4,000) | 3,700 (3,488 – 3,950) | 5,000 (4,700 – 5,500) |
| Unknown | 1 | 0 | 1 |
| sex, n (%) | | | |
| female | 73 (50) | 34 (50) | 58 (49) |
| male | 73 (50) | 34 (50) | 61 (51) |
| Unknown | 6 | 0 | 5 |

Remember those conflicts?

```
library(tidyverse)

## ━━ Attaching packages ━━━━━━━━━━━━━━━━ tidyverse 1.3.0 ━━

## ✓ ggplot2 3.3.3      ✓ purrr  0.3.4
## ✓ tibble  3.1.0      ✓ dplyr   1.0.5
## ✓ tidyr   1.1.3      ✓ stringr 1.4.0
## ✓ readr   1.4.0      ✓ forcats 0.5.1

## ━━ Conflicts ━━━━━━━━━━━━━━━━ tidyverse_conflicts() ━━

## x dplyr::filter()  masks stats::filter()
## x purrr::is_null() masks testthat::is_null()
## x dplyr::lag()     masks stats::lag()
## x dplyr::matches() masks tidyrr::matches(), testthat::matches()
```

Notice how some functions in different packages have the same names.

Related concepts of *namespaces* and *environments*.

The scope of names

```
x ← c(1:10)
y ← c(90:100)
```

```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
y
```

```
[1] 90 91 92 93 94 95 96 97 98 99 100
```

```
mean()
```

```
## Error in mean.default() : argument "x" is missing, with no default
```

The scope of names

```
mean(x) # argument names are internal to functions
```

```
[1] 5.5
```

```
mean(x = x)
```

```
[1] 5.5
```

```
mean(x = y)
```

```
[1] 95
```

```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
y
```

```
[1] 90 91 92 93 94 95 96 97 98 99 100
```

5. Objects come in **types** and **classes**

I'm going to speak somewhat loosely here for now, and gloss over some distinctions between object classes and data structures, as well as kinds of objects and their attributes.

5. Objects come in **types** and **classes**

I'm going to speak somewhat loosely here for now, and gloss over some distinctions between object classes and data structures, as well as kinds of objects and their attributes.

The object inspector in RStudio is your friend.

You can ask an object what it is at the console, too:

```
class(my_numbers)
```

```
[1] "numeric"
```

```
typeof(my_numbers)
```

```
[1] "double"
```

5. Objects come in **types** and **classes**

Objects can have more than one (nested) class:

```
summary(my_numbers)
```

| Min. | 1st Qu. | Median | Mean | 3rd Qu. | Max. |
|------|---------|--------|------|---------|-------|
| 1.00 | 1.75 | 3.00 | 6.25 | 6.25 | 25.00 |

```
my_smry ← summary(my_numbers) # remember, outputs can be assigned to a name, creating an object
```

```
class(summary(my_numbers)) # functions can be nested, and are evaluated from the inside out
```

```
[1] "summaryDefault" "table"
```

```
class(my_smry) # equivalent to the previous line
```

```
[1] "summaryDefault" "table"
```

5. Objects come in **types** and **classes**

```
typeof(my_smry)
```

```
[1] "double"
```

```
attributes(my_smry)
```

```
$names  
[1] "Min."     "1st Qu." "Median"   "Mean"      "3rd Qu." "Max."
```

```
$class  
[1] "summaryDefault" "table"
```

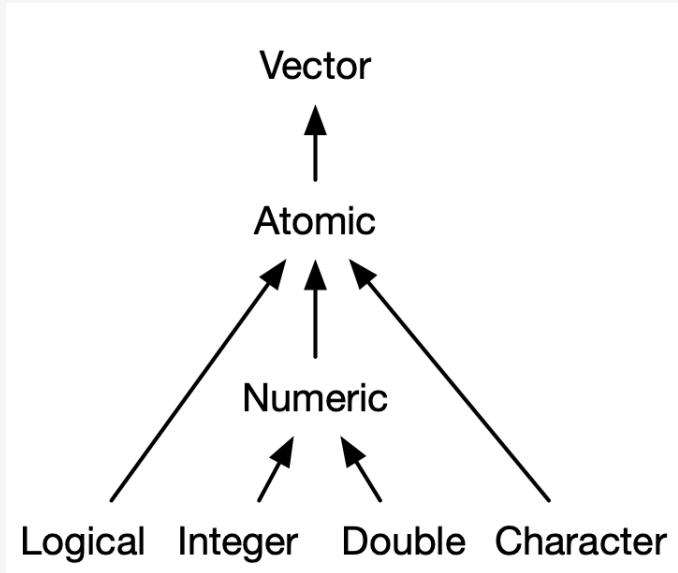
```
## In this case, the functions extract the corresponding attribute  
class(my_smry)
```

```
[1] "summaryDefault" "table"
```

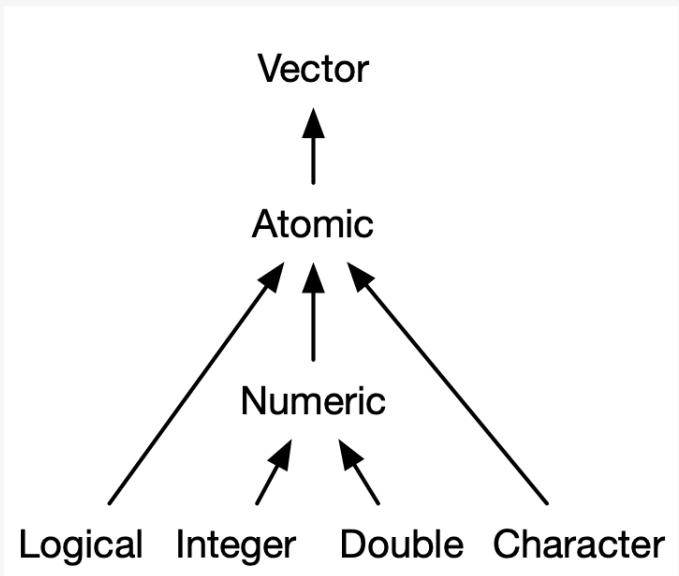
```
names(my_smry)
```

```
[1] "Min."     "1st Qu." "Median"   "Mean"      "3rd Qu." "Max."
```

Kinds of vector



Kinds of vector



```
my_int <- c(1, 3, 5, 6, 10)  
is.integer(my_int)
```

```
[1] FALSE
```

```
is.double(my_int)
```

```
[1] TRUE
```

```
my_int <- as.integer(my_int)  
is.integer(my_int)
```

```
[1] TRUE
```

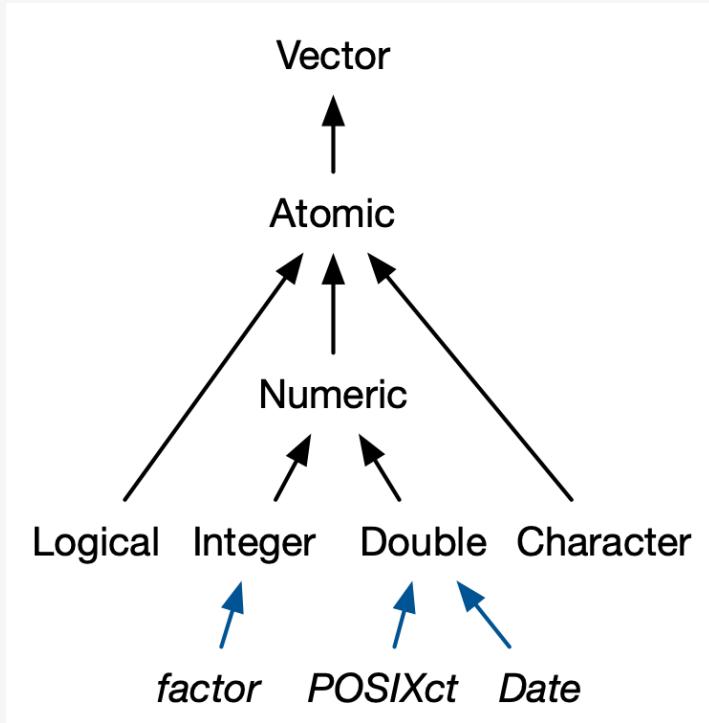
```
my_chr <- c("Mary", "had", "a", "little", "lamb")  
is.character(my_chr)
```

```
[1] TRUE
```

```
my_lgl <- c(TRUE, FALSE, TRUE)  
is.logical(my_lgl)
```

```
[1] TRUE
```

Kinds of vector



```
## Factors are for storing unordered or ordered categorical variables
```

```
x <- factor(c("Yes", "No", "No", "Maybe", "Yes", "Yes", "Yes", "No"))
x
```

```
[1] Yes   No    No    Maybe  Yes   Yes   Yes   No
Levels: Maybe No Yes
```

```
summary(x) # Alphabetical order by default
```

```
Maybe    No    Yes
1       3       4
```

```
typeof(x)      # Underneath, a factor is a type of integer ...
```

```
[1] "integer"
```

```
attributes(x) # ... with labels for its numbers, or "levels"
```

```
$levels
[1] "Maybe" "No"     "Yes"
```

```
$class
[1] "factor"
```

```
levels(x)
```

```
[1] "Maybe" "No"     "Yes"
```

```
is.ordered(x)
```

Vector types can't be heterogenous

Objects can be manually or automatically coerced from one class to another. Take care.

```
class(my_numbers)
```

```
[1] "numeric"
```

```
my_new_vector ← c(my_numbers, "Apple")
```

```
my_new_vector # vectors are homogeneous/atomic
```

```
[1] "1"      "2"      "3"      "1"      "3"      "5"      "25"     "10"     "Apple"
```

```
class(my_new_vector)
```

```
[1] "character"
```

Vector types can't be heterogenous

Objects can be manually or automatically coerced from one class to another. Take care.

```
my_dbl ← c(2.1, 4.77, 30.111, 3.14519)  
is.double(my_dbl)
```

```
[1] TRUE
```

```
my_dbl ← as.integer(my_dbl)  
  
my_dbl
```

```
[1] 2 4 30 3
```

A table of data is a kind of **list**

```
gapminder # tibbles and data frames can contain vectors of different types
```

```
# A tibble: 1,704 × 6
  country   continent   year lifeExp     pop gdpPercap
  <fct>     <fct>     <int>   <dbl>   <int>     <dbl>
1 Afghanistan Asia      1952    28.8  8425333     779.
2 Afghanistan Asia      1957    30.3  9240934     821.
3 Afghanistan Asia      1962    32.0  10267083    853.
4 Afghanistan Asia      1967    34.0  11537966    836.
5 Afghanistan Asia      1972    36.1  13079460    740.
6 Afghanistan Asia      1977    38.4  14880372    786.
7 Afghanistan Asia      1982    39.9  12881816    978.
8 Afghanistan Asia      1987    40.8  13867957    852.
9 Afghanistan Asia      1992    41.7  16317921    649.
10 Afghanistan Asia     1997    41.8  22227415    635.
# i 1,694 more rows
```

```
class(gapminder)
```

```
[1] "tbl_df"     "tbl"        "data.frame"
```

```
typeof(gapminder) # hmm
```

```
[1] "list"
```

A table of data is a kind of **list**

Lists *can* be heterogenous. Underneath, most complex R objects are some kind of list with different components.

A *data frame* is a list of vectors of the same length, where the vectors can be of different types (e.g. numeric, character, logical, etc)

A *tibble* is an enhanced data frame

Some classes are versions of others

Base R's trusty `data.frame`

```
library(socviz)
titanic

  fate    sex    n percent
1 perished male 1364    62.0
2 perished female 126     5.7
3 survived male  367    16.7
4 survived female 344    15.6
```

```
class(titanic)

[1] "data.frame"
```

```
## The ` `$` idiom picks out a named column here;
## more generally, the named element of a list
titanic$percent
```

```
[1] 62.0 5.7 16.7 15.6
```

Some classes are versions of others

Base R's trusty `data.frame`

```
library(socviz)
titanic

  fate   sex     n percent
1 perished male 1364    62.0
2 perished female 126     5.7
3 survived male  367    16.7
4 survived female 344    15.6
```

```
class(titanic)
```

```
[1] "data.frame"
```

```
## The ` `$` idiom picks out a named column here;
## more generally, the named element of a list
titanic$percent
```

```
[1] 62.0  5.7 16.7 15.6
```

The Tidyverse's enhanced `tibble`

```
## tibbles are build on data frames
titanic_tb ← as_tibble(titanic)
titanic_tb
```

```
# A tibble: 4 × 4
  fate     sex     n percent
  <fct>   <fct>   <dbl>   <dbl>
1 perished male    1364    62
2 perished female   126     5.7
3 survived male    367    16.7
4 survived female   344    15.6
```

```
class(titanic_tb)
```

```
[1] "tbl_df"      "tbl"        "data.frame"
```

A data frame and a tibble are both fundamentally a list of vectors of the same length, where the vectors can be of different types (e.g. numeric, character, logical, etc)

All of this will be clearer in use

gss_sm

```
# A tibble: 2,867 × 32
  year   id ballot      age child� sibs degree race   sex   region income16
  <dbl> <dbl> <labelled> <dbl> <dbl> <labe> <fct> <fct> <fct> <fct> <fct>
1 2016     1 1           47     3 2    Bach... White Male  New E... $170000...
2 2016     2 2           61     0 3    High ... White Male  New E... $50000 ...
3 2016     3 3           72     2 3    Bach... White Male  New E... $75000 ...
4 2016     4 1           43     4 3    High ... White Fema... New E... $170000...
5 2016     5 3           55     2 2    Gradu... White Fema... New E... $170000...
6 2016     6 2           53     2 2    Junio... White Fema... New E... $60000 ...
7 2016     7 1           50     2 2    High ... White Male  New E... $170000...
8 2016     8 3           23     3 6    High ... Other Fema... Middl... $30000 ...
9 2016     9 1           45     3 5    High ... Black Male  Middl... $60000 ...
10 2016    10 3          71     4 1   Junio... White Male  Middl... $60000 ...
# i 2,857 more rows
# i 21 more variables: relig <fct>, marital <fct>, padeg <fct>, madeg <fct>,
# partyid <fct>, polviews <fct>, happy <fct>, partners <fct>, grass <fct>,
# zodiac <fct>, pres12 <labelled>, wtssall <dbl>, income_rc <fct>,
# agegrp <fct>, ageq <fct>, siblings <fct>, kids <fct>, religion <fct>,
# bigregion <fct>, partners_rc <fct>, obama <dbl>
```

Tidyverse tools are generally *type safe*, meaning their functions return the same type of thing every time, or fail if they cannot do this. So it's good to know about the various data types.

6. Arithmetic on vectors

In R, all numbers are vectors of different sorts. Even single numbers (“scalars”) are conceptually vectors of length 1.

Arithmetic on vectors (and arrays generally) follows a series of *recycling rules* that favor ease of expression of vectorized, “elementwise” operations.

See if you can predict what the following operations do:

6. Arithmetic on vectors

```
my_numbers
```

```
[1] 1 2 3 1 3 5 25 10
```

```
result1 ← my_numbers + 1
```

6. Arithmetic on vectors

```
my_numbers
```

```
[1] 1 2 3 1 3 5 25 10
```

```
result1 ← my_numbers + 1
```

```
result1
```

```
[1] 2 3 4 2 4 6 26 11
```

6. Arithmetic on vectors

```
result2 ← my_numbers + my_numbers
```

6. Arithmetic on vectors

```
result2 ← my_numbers + my_numbers
```

```
result2
```

```
[1]  2  4  6  2  6 10 50 20
```

6. Arithmetic on vectors

```
two_nums ← c(5, 10)  
result3 ← my_numbers + two_nums
```

6. Arithmetic on vectors

```
two_nums ← c(5, 10)  
result3 ← my_numbers + two_nums  
result3  
[1]  6 12  8 11  8 15 30 20
```

6. Arithmetic on vectors

```
three_nums ← c(1, 5, 10)  
  
result4 ← my_numbers + three_nums
```

Warning in my_numbers + three_nums: longer object length is not a multiple of
shorter object length

6. Arithmetic on vectors

```
three_nums ← c(1, 5, 10)  
  
result4 ← my_numbers + three_nums
```

Warning in my_numbers + three_nums: longer object length is not a multiple of shorter object length

```
result4  
  
[1]  2  7 13  2  8 15 26 15
```

Note that you get a *warning* here. It'll still do it, though! Don't ignore warnings until you understand what they mean.

7. R will be frustrating

The IDE tries its best to help you. Learn to attend to what it is trying to say.

```
Warning message:  
In my_numbers + two_nums :  
  longer object length is not a multiple of shorter object length
```

```
38   p <- ggplot(data = gapminder  
39   p <- ggplot(data = gapminder  
40     mapping = aes(x = gdpPercap,  
41                   y = lifeExp))  
42 |
```

```
39 p <- ggplot(data = gapminder,  
40                 mapping = aes(x = gdpPercap,  
41                               y = lifeExp)))  
42 `unexpected token ')'`  
43 ^`|
```