

Manipulating Tables with **dplyr** (contd)

Data Wrangling: Session 3

Kieran Healy

Statistical Horizons, December 2022

Window functions and moving averages

Load our libraries

```
library(here)      # manage file paths  
library(socviz)    # data and some useful functions  
library(tidyverse) # your friend and mine
```

dp1yr's **window** functions

Ranking and cumulation within groups.

```
## Data on COVID-19
```

```
library(covdata)
```

```
covnat_weekly
```

```
## # A tibble: 4,020 × 11
```

```
##   date      year_week cname  iso3    pop cases deaths cu_ca...1 cu_de...2 r14_c...3
##   <date>    <chr>    <chr> <chr>  <dbl> <dbl>  <dbl>  <dbl>    <dbl>    <dbl>
## 1 2019-12-30 2020-01  Austr... AUT   8.93e6    0    0      0      0    NA
## 2 2020-01-06 2020-02  Austr... AUT   8.93e6    0    0      0      0    0
## 3 2020-01-13 2020-03  Austr... AUT   8.93e6    0    0      0      0    0
## 4 2020-01-20 2020-04  Austr... AUT   8.93e6    0    0      0      0    0
## 5 2020-01-27 2020-05  Austr... AUT   8.93e6    0    0      0      0    0
## 6 2020-02-03 2020-06  Austr... AUT   8.93e6    0    0      0      0    0
## 7 2020-02-10 2020-07  Austr... AUT   8.93e6    0    0      0      0    0
## 8 2020-02-17 2020-08  Austr... AUT   8.93e6    0    0      0      0    0
## 9 2020-02-24 2020-09  Austr... AUT   8.93e6   12    0     12      0  0.134
## 10 2020-03-02 2020-10  Austr... AUT   8.93e6  114    0    126      0  1.41
```

```
## # ... with 4,010 more rows, 1 more variable: r14_deaths <dbl>, and abbreviated
```

```
## #   variable names 1cu_cases, 2cu_deaths, 3r14_cases
```

dp1yr's **window** functions

cumsum() gives cumulative sums

```
covnat_weekly |>
  filter(iso3 == "FRA") |>
  select(date, cname, iso3, cases) |>
  mutate(cumulative = cumsum(cases))
```

```
## # A tibble: 134 × 5
##   date      cname iso3  cases cumulative
##   <date>    <chr> <chr> <dbl>      <dbl>
## 1 2019-12-30 France FRA      0          0
## 2 2020-01-06 France FRA      0          0
## 3 2020-01-13 France FRA      0          0
## 4 2020-01-20 France FRA      3          3
## 5 2020-01-27 France FRA      3          6
## 6 2020-02-03 France FRA      6         12
## 7 2020-02-10 France FRA      0         12
## 8 2020-02-17 France FRA      4         16
## 9 2020-02-24 France FRA     133        149
## 10 2020-03-02 France FRA     981       1130
## # ... with 124 more rows
```

dp1yr's **window** functions

cume_dist() gives the proportion of values less than or equal to the current value.

```
covnat_weekly |>
  select(date, cname, iso3, deaths) |>
  filter(iso3 == "FRA") |>
  filter(cume_dist(desc(deaths)) < 0.1) # i.e. Top 10%
```

```
## # A tibble: 13 × 4
##   date      cname iso3  deaths
##   <date>    <chr> <chr>  <dbl>
## 1 2020-10-26 France FRA    3517
## 2 2020-11-02 France FRA    5281
## 3 2020-11-09 France FRA    6018
## 4 2020-11-16 France FRA    6208
## 5 2020-11-23 France FRA    5215
## 6 2020-11-30 France FRA    4450
## 7 2020-12-07 France FRA    4257
## 8 2020-12-14 France FRA    3786
## 9 2020-12-21 France FRA    3560
## 10 2021-01-04 France FRA    3851
## 11 2021-01-11 France FRA    3833
## 12 2021-01-18 France FRA    3754
## 13 2021-01-25 France FRA    3535
```

The dp1yr vignette on Window functions is good.

An application

```
covus |>
  filter(measure == "death") |>
  group_by(state) |>
  arrange(state, desc(date)) |>
  filter(state %in% "NY")
```

```
## # A tibble: 371 × 7
## # Groups:   state [1]
##   date      state fips data_quality_grade measure count measure_label
##   <date>     <chr> <chr> <lgl>           <chr>   <dbl> <chr>
## 1 2021-03-07 NY     36    NA             death  39029 Deaths
## 2 2021-03-06 NY     36    NA             death  38970 Deaths
## 3 2021-03-05 NY     36    NA             death  38891 Deaths
## 4 2021-03-04 NY     36    NA             death  38796 Deaths
## 5 2021-03-03 NY     36    NA             death  38735 Deaths
## 6 2021-03-02 NY     36    NA             death  38660 Deaths
## 7 2021-03-01 NY     36    NA             death  38577 Deaths
## 8 2021-02-28 NY     36    NA             death  38497 Deaths
## 9 2021-02-27 NY     36    NA             death  38407 Deaths
## 10 2021-02-26 NY     36    NA             death  38321 Deaths
## # ... with 361 more rows
```

Here the count measure is *cumulative* deaths. What if we want to recover the daily count for all the states in the data?

An application

`dplyr` has `lead()` and `lag()` functions. These allow you to access the previous and next values in a vector. You can calculate offsets this way.

```
my_vec <- c(1:20)
my_vec
```

```
## [1]  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
```

```
lag(my_vec) # first element has no lag
```

```
## [1] NA  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19
```

```
my_vec - lag(my_vec)
```

```
## [1] NA  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1  1
```


An application

We can write the expression directly:

```
covus |>
  select(-data_quality_grade) |>
  filter(measure == "death") |>
  group_by(state) |>
  arrange(date) |>
  mutate(deaths_daily = count - lag(count, order_by = date)) |>
  arrange(state, desc(date)) |>
  filter(state %in% "NY")

## # A tibble: 371 × 7
## # Groups:   state [1]
##   date      state fips  measure count measure_label deaths_daily
##   <date>    <chr> <chr> <chr>    <dbl> <chr>          <dbl>
## 1 2021-03-07 NY     36    death  39029 Deaths         59
## 2 2021-03-06 NY     36    death  38970 Deaths         79
## 3 2021-03-05 NY     36    death  38891 Deaths         95
## 4 2021-03-04 NY     36    death  38796 Deaths         61
## 5 2021-03-03 NY     36    death  38735 Deaths         75
## 6 2021-03-02 NY     36    death  38660 Deaths         83
## 7 2021-03-01 NY     36    death  38577 Deaths         80
## 8 2021-02-28 NY     36    death  38497 Deaths         90
## 9 2021-02-27 NY     36    death  38407 Deaths         86
## 10 2021-02-26 NY     36    death  38321 Deaths         94
## # ... with 361 more rows
```

Writing our own functions

But we could also write a function to do this.

We write functions using the special `function()` function.*

```
my_fun <- function(x) {  
  x + 1  
}
```

```
my_fun # we've created the function; it's just an object
```

```
## function(x) {  
##   x + 1  
## }
```

```
my_fun(x = 1) # But we can supply it with an input!
```

```
## [1] 2
```

```
my_fun(10)
```

```
## [1] 11
```

*Nerds love this sort of stuff.

Writing our own **functions**

We write our function. It's just the expression we originally wrote, wrapped up.

```
get_daily_count <- function(count, date){  
  count - lag(count, order_by = date)  
}
```

This function has no generality, error-handling, or anything else. It's a once-off.

Writing our own **functions**

Now we can use it like any other:

```
covus |>
  filter(measure == "death") |>
  select(-data_quality_grade) |>
  group_by(state) |>
  arrange(date) |>
  mutate(deaths_daily = get_daily_count(count, date)) |>
  arrange(state, desc(date)) |>
  filter(state %in% "NY")

## # A tibble: 371 × 7
## # Groups:   state [1]
##   date      state fips  measure count measure_label deaths_daily
##   <date>    <chr> <chr> <chr>    <dbl> <chr>          <dbl>
## 1 2021-03-07 NY     36    death  39029 Deaths         59
## 2 2021-03-06 NY     36    death  38970 Deaths         79
## 3 2021-03-05 NY     36    death  38891 Deaths         95
## 4 2021-03-04 NY     36    death  38796 Deaths         61
## 5 2021-03-03 NY     36    death  38735 Deaths         75
## 6 2021-03-02 NY     36    death  38660 Deaths         83
## 7 2021-03-01 NY     36    death  38577 Deaths         80
## 8 2021-02-28 NY     36    death  38497 Deaths         90
## 9 2021-02-27 NY     36    death  38407 Deaths         86
## 10 2021-02-26 NY     36    death  38321 Deaths         94
## # ... with 361 more rows
```

Not super-useful quite yet, but if our task had more steps ...

Tidy moving averages with **slider**

dplyr's window functions don't include moving averages.

There are several options, notably **RcppRoll**

We'll use the **slider** package.

```
# install.packages("slider")  
library(slider)
```

Tidy moving averages with **slider**

```
covus |>
  filter(measure == "death") |>
  select(-data_quality_grade) |>
  group_by(state) |>
  arrange(date) |>
  mutate(
    deaths_daily = get_daily_count(count, date),
    deaths7 = slide_mean(deaths_daily,
                        before = 7,
                        na_rm = TRUE)) |>
  arrange(state, desc(date)) |>
  filter(state %in% "NY")
```

```
## # A tibble: 371 × 8
## # Groups:   state [1]
##   date      state fips  measure count measure_label deaths_daily deaths7
##   <date>    <chr> <chr> <chr>    <dbl> <chr>          <dbl>    <dbl>
## 1 2021-03-07 NY     36    death  39029 Deaths        59      77.8
## 2 2021-03-06 NY     36    death  38970 Deaths        79      81.1
## 3 2021-03-05 NY     36    death  38891 Deaths        95       83
## 4 2021-03-04 NY     36    death  38796 Deaths        61     82.6
## 5 2021-03-03 NY     36    death  38735 Deaths        75       88
## 6 2021-03-02 NY     36    death  38660 Deaths        83     89.9
## 7 2021-03-01 NY     36    death  38577 Deaths        80     90.8
## 8 2021-02-28 NY     36    death  38497 Deaths        90     90.1
## 9 2021-02-27 NY     36    death  38407 Deaths        86     91.5
## 10 2021-02-26 NY     36    death  38321 Deaths        94     95.6
## # ... with 361 more rows
```

Tidy moving averages with **slider**

```
deaths7 = slide_mean(deaths_daily,  
  before = 7,  
  na_rm = TRUE)) |>
```

Notice the Tidyverse-style `na_rm` argument rather than the usual base `na.rm`

The package provides a lot of different functions, from general-purpose **`slide_max()`**, **`slide_min()`** to more specialized sliding functions. In particular note e.g. **`slide_index_mean()`** that addresses some subtleties in averaging over dates with gaps.

Tidy up after yourself with `relocate()`

```
gss_sm
```

```
## # A tibble: 2,867 × 32
##   year    id ballot  age childs sibs degree race  sex  region incom...1 relig
##   <dbl> <dbl> <labe> <dbl>  <dbl> <lab> <fct>  <fct> <fct> <fct>  <fct>  <fct>
## 1  2016     1 1      47     3 2    Bache... White Male  New E... $17000... None
## 2  2016     2 2      61     0 3    High ... White Male  New E... $50000... None
## 3  2016     3 3      72     2 3    Bache... White Male  New E... $75000... Cath...
## 4  2016     4 1      43     4 3    High ... White Fema... New E... $17000... Cath...
## 5  2016     5 3      55     2 2    Gradu... White Fema... New E... $17000... None
## 6  2016     6 2      53     2 2    Junio... White Fema... New E... $60000... None
## 7  2016     7 1      50     2 2    High ... White Male  New E... $17000... None
## 8  2016     8 3      23     3 6    High ... Other Fema... Middl... $30000... Cath...
## 9  2016     9 1      45     3 5    High ... Black Male  Middl... $60000... Prot...
## 10 2016    10 3      71     4 1    Junio... White Male  Middl... $60000... None
## # ... with 2,857 more rows, 20 more variables: marital <fct>, padeg <fct>,
## #   madeg <fct>, partyid <fct>, polviews <fct>, happy <fct>, partners <fct>,
## #   grass <fct>, zodiac <fct>, pres12 <labelled>, wtssall <dbl>,
## #   income_rc <fct>, agegrp <fct>, ageq <fct>, siblings <fct>, kids <fct>,
## #   religion <fct>, bigregion <fct>, partners_rc <fct>, obama <dbl>, and
## #   abbreviated variable name 1income16
```


Shuffle columns around

gss_sm

```
## # A tibble: 2,867 × 32
##   year    id ballot  age childs sibs degree race  sex  region incom...1 relig
##   <dbl> <dbl> <labe> <dbl>  <dbl> <lab> <fct>  <fct> <fct> <fct>  <fct>  <fct>
## 1  2016     1  1      47     3  2  Bache... White Male  New E... $17000... None
## 2  2016     2  2      61     0  3  High ... White Male  New E... $50000... None
## 3  2016     3  3      72     2  3  Bache... White Male  New E... $75000... Cath...
## 4  2016     4  1      43     4  3  High ... White Fema... New E... $17000... Cath...
## 5  2016     5  3      55     2  2  Gradu... White Fema... New E... $17000... None
## 6  2016     6  2      53     2  2  Junio... White Fema... New E... $60000... None
## 7  2016     7  1      50     2  2  High ... White Male  New E... $17000... None
## 8  2016     8  3      23     3  6  High ... Other Fema... Middl... $30000... Cath...
## 9  2016     9  1      45     3  5  High ... Black Male  Middl... $60000... Prot...
## 10 2016    10  3      71     4  1  Junio... White Male  Middl... $60000... None
## # ... with 2,857 more rows, 20 more variables: marital <fct>, padeg <fct>,
## #   madeg <fct>, partyid <fct>, polviews <fct>, happy <fct>, partners <fct>,
## #   grass <fct>, zodiac <fct>, pres12 <labelled>, wtssall <dbl>,
## #   income_rc <fct>, agegrp <fct>, ageq <fct>, siblings <fct>, kids <fct>,
## #   religion <fct>, bigregion <fct>, partners_rc <fct>, obama <dbl>, and
## #   abbreviated variable name 1income16
```

Shuffle columns around

```
gss_sm |>
```

```
  select(region, bigregion, year,  
         id:region,  
         starts_with("p"),  
         contains("income"))
```

```
## # A tibble: 2,867 × 19  
##   region bigre...1 year   id ballot   age childs sibs  degree race  sex  padeg  
##   <fct>  <fct>    <dbl> <dbl> <labe> <dbl> <dbl> <lab> <fct>  <fct> <fct> <fct>  
## 1 New E... Northe... 2016     1 1      47     3 2    Bache... White Male Grad...  
## 2 New E... Northe... 2016     2 2      61     0 3    High ... White Male Lt H...  
## 3 New E... Northe... 2016     3 3      72     2 3    Bache... White Male High...  
## 4 New E... Northe... 2016     4 1      43     4 3    High ... White Fema... <NA>  
## 5 New E... Northe... 2016     5 3      55     2 2    Gradu... White Fema... Bach...  
## 6 New E... Northe... 2016     6 2      53     2 2    Junio... White Fema... <NA>  
## 7 New E... Northe... 2016     7 1      50     2 2    High ... White Male High...  
## 8 Middl... Northe... 2016     8 3      23     3 6    High ... Other Fema... Lt H...  
## 9 Middl... Northe... 2016     9 1      45     3 5    High ... Black Male Lt H...  
## 10 Middl... Northe... 2016    10 3      71     4 1    Junio... White Male High...  
## # ... with 2,857 more rows, 7 more variables: partyid <fct>, polviews <fct>,  
## #   partners <fct>, pres12 <labelled>, partners_rc <fct>, income16 <fct>,  
## #   income_rc <fct>, and abbreviated variable name 1bigregion
```

Shuffle columns around

```
gss_sm |>
  select(region, bigregion, year,
         id:region,
         starts_with("p"),
         contains("income")) |>
  rename(children = child,
         siblings = sibs)
```

```
## # A tibble: 2,867 × 19
##   region    bigre...1 year    id ballot   age child...2 sibli...3 degree race  sex
##   <fct>      <fct>   <dbl> <dbl> <labe> <dbl>   <dbl> <label> <fct> <fct> <fct>
## 1 New Engl... Northe... 2016     1 1      47      3 2      Bache... White Male
## 2 New Engl... Northe... 2016     2 2      61      0 3      High ... White Male
## 3 New Engl... Northe... 2016     3 3      72      2 3      Bache... White Male
## 4 New Engl... Northe... 2016     4 1      43      4 3      High ... White Fema...
## 5 New Engl... Northe... 2016     5 3      55      2 2      Gradu... White Fema...
## 6 New Engl... Northe... 2016     6 2      53      2 2      Junio... White Fema...
## 7 New Engl... Northe... 2016     7 1      50      2 2      High ... White Male
## 8 Middle A... Northe... 2016     8 3      23      3 6      High ... Other Fema...
## 9 Middle A... Northe... 2016     9 1      45      3 5      High ... Black Male
## 10 Middle A... Northe... 2016    10 3      71      4 1      Junio... White Male
## # ... with 2,857 more rows, 8 more variables: padeg <fct>, partyid <fct>,
## #   polviews <fct>, partners <fct>, pres12 <labelled>, partners_rc <fct>,
## #   income16 <fct>, income_rc <fct>, and abbreviated variable names 1bigregion,
## #   2children, 3siblings
```

Shuffle columns around

```
gss_sm |>
  select(region, bigregion, year,
         id:region,
         starts_with("p"),
         contains("income")) |>
  rename(children = child,
         siblings = sibs) |>
  relocate(id)
```

```
## # A tibble: 2,867 × 19
##       id region    bigre...1 year ballot age child...2 sibli...3 degree race sex
##   <dbl> <fct>    <fct>    <dbl> <lab> <dbl> <dbl> <label> <fct> <fct> <fct>
## 1     1 New Engl... Northe... 2016 1      47      3 2 Bache... White Male
## 2     2 New Engl... Northe... 2016 2      61      0 3 High ... White Male
## 3     3 New Engl... Northe... 2016 3      72      2 3 Bache... White Male
## 4     4 New Engl... Northe... 2016 1      43      4 3 High ... White Fema...
## 5     5 New Engl... Northe... 2016 3      55      2 2 Gradu... White Fema...
## 6     6 New Engl... Northe... 2016 2      53      2 2 Junio... White Fema...
## 7     7 New Engl... Northe... 2016 1      50      2 2 High ... White Male
## 8     8 Middle A... Northe... 2016 3      23      3 6 High ... Other Fema...
## 9     9 Middle A... Northe... 2016 1      45      3 5 High ... Black Male
## 10    10 Middle A... Northe... 2016 3      71      4 1 Junio... White Male
## # ... with 2,857 more rows, 8 more variables: padeg <fct>, partyid <fct>,
## # polviews <fct>, partners <fct>, pres12 <labelled>, partners_rc <fct>,
## # income16 <fct>, income_rc <fct>, and abbreviated variable names 1bigregion,
## # 2children, 3siblings
```

Shuffle columns around

```
gss_sm |>
  select(region, bigregion, year,
         id:region,
         starts_with("p"),
         contains("income")) |>
  rename(children = child,
         siblings = sibs) |>
  relocate(id) |>
  select(-ballot)
```

```
## # A tibble: 2,867 × 18
##       id region    bigre...1 year  age child...2 sibli...3 degree race sex padeg
##   <dbl> <fct>    <fct>    <dbl> <dbl> <dbl> <label> <fct> <fct> <fct> <fct>
## 1     1 New Engla... Northe... 2016    47     3 2 Bache... White Male Grad...
## 2     2 New Engla... Northe... 2016    61     0 3 High ... White Male Lt H...
## 3     3 New Engla... Northe... 2016    72     2 3 Bache... White Male High...
## 4     4 New Engla... Northe... 2016    43     4 3 High ... White Fema... <NA>
## 5     5 New Engla... Northe... 2016    55     2 2 Gradu... White Fema... Bach...
## 6     6 New Engla... Northe... 2016    53     2 2 Junio... White Fema... <NA>
## 7     7 New Engla... Northe... 2016    50     2 2 High ... White Male High...
## 8     8 Middle At... Northe... 2016    23     3 6 High ... Other Fema... Lt H...
## 9     9 Middle At... Northe... 2016    45     3 5 High ... Black Male Lt H...
## 10    10 Middle At... Northe... 2016    71     4 1 Junio... White Male High...
## # ... with 2,857 more rows, 7 more variables: partyid <fct>, polviews <fct>,
## #   partners <fct>, pres12 <labelled>, partners_rc <fct>, income16 <fct>,
## #   income_rc <fct>, and abbreviated variable names 1bigregion, 2children,
## #   3siblings
```

Shuffle columns around

```
gss_sm |>
  select(region, bigregion, year,
         id:region,
         starts_with("p"),
         contains("income")) |>
  rename(children = child,
         siblings = sibs) |>
  relocate(id) |>
  select(-ballot) |>
  relocate(where(is.numeric),
           .before = where(is.factor))
```

```
## # A tibble: 2,867 × 18
##       id year age children siblings pres12 region bigre...1 degree race sex
##   <dbl> <dbl> <dbl>   <dbl> <labelle> <labe> <fct> <fct> <fct> <fct> <fct>
## 1     1   2016  47     3 2         3 New E... Northe... Bache... White Male
## 2     2   2016  61     0 3         1 New E... Northe... High ... White Male
## 3     3   2016  72     2 3         2 New E... Northe... Bache... White Male
## 4     4   2016  43     4 3         2 New E... Northe... High ... White Fema...
## 5     5   2016  55     2 2         1 New E... Northe... Gradu... White Fema...
## 6     6   2016  53     2 2         1 New E... Northe... Junio... White Fema...
## 7     7   2016  50     2 2        NA New E... Northe... High ... White Male
## 8     8   2016  23     3 6        NA Middl... Northe... High ... Other Fema...
## 9     9   2016  45     3 5        NA Middl... Northe... High ... Black Male
## 10    10   2016  71     4 1         2 Middl... Northe... Junio... White Male
## # ... with 2,857 more rows, 7 more variables: padeg <fct>, partyid <fct>,
## #   polviews <fct>, partners <fct>, partners_rc <fct>, income16 <fct>,
## #   income_rc <fct>, and abbreviated variable name 1bigregion
```

Shuffle columns around

```
gss_sm |>
  select(region, bigregion, year,
         id:region,
         starts_with("p"),
         contains("income")) |>
  rename(children = child,
         siblings = sibs) |>
  relocate(id) |>
  select(-ballot) |>
  relocate(where(is.numeric),
           .before = where(is.factor)) |>
  relocate(contains("region"),
           .after = year)
```

```
## # A tibble: 2,867 × 18
##       id year region    bigre...1 age child...2 sibli...3 pres12 degree race sex
##   <dbl> <dbl> <fct>    <fct>    <dbl>    <dbl> <label> <label> <fct> <fct> <fct>
## 1     1   2016 New Engl... Northe...    47         3 2         3 Bache... White Male
## 2     2   2016 New Engl... Northe...    61         0 3         1 High ... White Male
## 3     3   2016 New Engl... Northe...    72         2 3         2 Bache... White Male
## 4     4   2016 New Engl... Northe...    43         4 3         2 High ... White Fema...
## 5     5   2016 New Engl... Northe...    55         2 2         1 Gradu... White Fema...
## 6     6   2016 New Engl... Northe...    53         2 2         1 Junio... White Fema...
## 7     7   2016 New Engl... Northe...    50         2 2        NA High ... White Male
## 8     8   2016 Middle A... Northe...    23         3 6        NA High ... Other Fema...
## 9     9   2016 Middle A... Northe...    45         3 5        NA High ... Black Male
## 10    10   2016 Middle A... Northe...    71         4 1         2 Junio... White Male
## # ... with 2,857 more rows, 7 more variables: padeg <fct>, partyid <fct>,
## # polviews <fct>, partners <fct>, partners_rc <fct>, income16 <fct>,
## # income_rc <fct>, and abbreviated variable names 1bigregion, 2children,
## # 3siblings
```

Example: UK Election Data

```
library(ukelection2019)
```

```
ukvote2019
```

```
## # A tibble: 3,320 × 13
```

```
##   cid      const...1 elect...2 party...3 candi...4 votes vote_...5 vote_...6 total...7 vrank
##   <chr>      <chr>      <int> <chr>      <chr>      <int>  <dbl>  <dbl>  <int> <int>
## 1 W07000049 Aberav...  50747 Labour  Stephe... 17008    53.8   -14.3   31598     1
## 2 W07000049 Aberav...  50747 Conser... Charlo...  6518    20.6     2.9   31598     2
## 3 W07000049 Aberav...  50747 The Br... Glenda...  3108     9.8     9.8   31598     3
## 4 W07000049 Aberav...  50747 Plaid ... Nigel ...  2711     8.6     0.3   31598     4
## 5 W07000049 Aberav...  50747 Libera... Sheila...  1072     3.4     1.6   31598     5
## 6 W07000049 Aberav...  50747 Indepe... Captai...   731     2.3     2.3   31598     6
## 7 W07000049 Aberav...  50747 Green   Giorgi...   450     1.4     1.4   31598     7
## 8 W07000058 Aberco...  44699 Conser... Robin ... 14687    46.1     1.5   31865     1
## 9 W07000058 Aberco...  44699 Labour  Emily ... 12653    39.7    -2.9   31865     2
## 10 W07000058 Aberco...  44699 Plaid ... Lisa G...  2704     8.5    -1.4   31865     3
## # ... with 3,310 more rows, 3 more variables: turnout <dbl>, fname <chr>,
## #   lname <chr>, and abbreviated variable names 1constituency, 2electorate,
## #   3party_name, 4candidate, 5vote_share_percent, 6vote_share_change,
## #   7total_votes_cast
```


Example: UK Election Data

Use `sample_n()` to sample `n` rows of your tibble.

```
library(ukelection2019)
```

```
ukvote2019 |>  
  sample_n(10)
```

```
## # A tibble: 10 × 13  
##   cid      const...1 elect...2 party...3 candi...4 votes vote_...5 vote_...6 total...7 vrank  
##   <chr>      <chr>      <int> <chr>      <chr>      <int>  <dbl>  <dbl>    <int> <int>  
## 1 E14000699 Fareham    78337 Green    Nick L...  2412    4.2    1.9    57250    4  
## 2 E14000702 Filton...  74016 Conser... Jack L... 26293   48.9   -1.1    53752    1  
## 3 E14000952 Norfol...  78455 Conser... Liz Tr... 35507    69     6.2    51466    1  
## 4 E14001053 Workin...  61370 Conser... Mark J... 20488   49.3    7.5    41599    1  
## 5 E14000617 Cambri...  79951 Labour  Daniel... 25776    48    -3.9    53729    1  
## 6 E14000864 Norwic...  77845 Labour  Clive ... 27766   53.7   -7.2    51673    1  
## 7 E14000552 Bedford  71579 Green   Adrian...  960     2     0     47301    4  
## 8 E14001029 West B...  62111 The Br... Christ...  1475    4.1    4.1    35975    3  
## 9 E14000792 Lincoln  74778 Labour  Karen ... 20753    41   -6.9    50629    2  
## 10 E14000632 Cheste...  70994 Conser... Leigh ... 16720    37    2.2    45186    2  
## # ... with 3 more variables: turnout <dbl>, fname <chr>, lname <chr>, and  
## # abbreviated variable names 1constituency, 2electorate, 3party_name,  
## # 4candidate, 5vote_share_percent, 6vote_share_change, 7total_votes_cast
```

Example: UK Election Data

A vector of unique constituency names

```
ukvote2019 |>  
  distinct(constituency)
```

```
## # A tibble: 650 × 1  
##   constituency  
##   <chr>  
## 1 Aberavon  
## 2 Aberconwy  
## 3 Aberdeen North  
## 4 Aberdeen South  
## 5 Aberdeenshire West & Kincardine  
## 6 Airdrie & Shotts  
## 7 Aldershot  
## 8 Aldridge-Brownhills  
## 9 Altrincham & Sale West  
## 10 Alyn & Deeside  
## # ... with 640 more rows
```

Example: UK Election Data

Tally them up

```
ukvote2019 |>  
  distinct(constituency) |>  
  tally()
```

```
## # A tibble: 1 × 1  
##       n  
##   <int>  
## 1   650
```

```
# Base R / non-pipeline version
```

```
length(unique(ukvote2019$constituency))
```

```
## [1] 650
```

Example: UK Election Data

Which parties fielded the most candidates?

```
ukvote2019 |>  
  count(party_name) |>  
  arrange(desc(n))
```

```
## # A tibble: 69 × 2  
##   party_name      n  
##   <chr>        <int>  
## 1 Conservative    636  
## 2 Labour          631  
## 3 Liberal Democrat 611  
## 4 Green           497  
## 5 The Brexit Party  275  
## 6 Independent      224  
## 7 Scottish National Party 59  
## 8 UKIP             44  
## 9 Plaid Cymru      36  
## 10 Christian Peoples Alliance 29  
## # ... with 59 more rows
```

Example: UK Election Data

Top 5

```
ukvote2019 |>  
  count(party_name) |>  
  slice_max(order_by = n, n = 5)
```

```
## # A tibble: 5 × 2  
##   party_name      n  
##   <chr>      <int>  
## 1 Conservative    636  
## 2 Labour          631  
## 3 Liberal Democrat 611  
## 4 Green           497  
## 5 The Brexit Party 275
```

Example: UK Election Data

Top 5

```
ukvote2019 |>
  count(party_name) |>
  slice_max(order_by = n, n = 5)
```

```
## # A tibble: 5 × 2
##   party_name      n
##   <chr>      <int>
## 1 Conservative    636
## 2 Labour          631
## 3 Liberal Democrat 611
## 4 Green           497
## 5 The Brexit Party 275
```

Bottom 5

```
ukvote2019 |>
  count(party_name) |>
  slice_min(order_by = n, n = 5)
```

```
## # A tibble: 25 × 2
##   party_name      n
##   <chr>      <int>
## 1 Ashfield Independents    1
## 2 Best for Luton           1
## 3 Birkenhead Social Justice Party    1
## 4 British National Party    1
## 5 Burnley & Padiham Independent Party 1
## 6 Church of the Militant Elvis Party  1
## 7 Citizens Movement Party UK    1
## 8 CumbriaFirst              1
## 9 Heavy Woollen District Independents 1
## 10 Independent Network        1
## # ... with 15 more rows
```

Example: UK Election Data

How many constituencies are there?

```
ukvote2019 |>  
  count(constituency)
```

```
## # A tibble: 650 × 2  
##   constituency      n  
##   <chr>          <int>  
## 1 Aberavon        7  
## 2 Aberconwy        4  
## 3 Aberdeen North   6  
## 4 Aberdeen South   4  
## 5 Aberdeenshire West & Kincardine 4  
## 6 Airdrie & Shotts  5  
## 7 Aldershot        4  
## 8 Aldridge-Brownhills 5  
## 9 Altrincham & Sale West 6  
## 10 Alyn & Deeside    5  
## # ... with 640 more rows
```

Example: UK Election Data

How many constituencies are there?

```
ukvote2019 |>  
  count(constituency)
```

```
## # A tibble: 650 × 2  
##   constituency      n  
##   <chr>          <int>  
## 1 Aberavon        7  
## 2 Aberconwy        4  
## 3 Aberdeen North   6  
## 4 Aberdeen South   4  
## 5 Aberdeenshire West & Kincardine 4  
## 6 Airdrie & Shotts  5  
## 7 Aldershot        4  
## 8 Aldridge-Brownhills 5  
## 9 Altrincham & Sale West 6  
## 10 Alyn & Deeside   5  
## # ... with 640 more rows
```

```
ukvote2019 |>  
  distinct(constituency) |>  
  count()
```

```
## # A tibble: 1 × 1  
##       n  
##   <int>  
## 1   650
```

```
# Base R style ...  
length(unique(ukvote2019$constituency))  
  
## [1] 650
```


Counting Twice Over

```
ukvote2019 |>  
  count(constituency) |>  
  count(n)
```

```
## # A tibble: 8 × 2  
##       n     nn  
##   <int> <int>  
## 1     3     21  
## 2     4    194  
## 3     5    226  
## 4     6    139  
## 5     7     49  
## 6     8     18  
## 7     9      2  
## 8    12      1
```

Counting Twice Over

ukvote2019

```
## # A tibble: 3,320 × 13
##   cid      const...1 elect...2 party...3 candi...4 votes vote_...5 vote_...6 total...7 vrank
##   <chr>      <chr>      <int> <chr>      <chr>      <int>      <dbl>      <dbl>      <int> <int>
## 1 W07000049 Aberav...  50747 Labour    Stephe... 17008      53.8      -14.3     31598     1
## 2 W07000049 Aberav...  50747 Conser... Charlo...  6518      20.6       2.9     31598     2
## 3 W07000049 Aberav...  50747 The Br... Glenda...  3108       9.8       9.8     31598     3
## 4 W07000049 Aberav...  50747 Plaid ... Nigel ...  2711       8.6       0.3     31598     4
## 5 W07000049 Aberav...  50747 Libera... Sheila...  1072       3.4       1.6     31598     5
## 6 W07000049 Aberav...  50747 Indepe... Captai...   731       2.3       2.3     31598     6
## 7 W07000049 Aberav...  50747 Green    Giorgi...   450       1.4       1.4     31598     7
## 8 W07000058 Aberco...  44699 Conser... Robin ... 14687      46.1       1.5     31865     1
## 9 W07000058 Aberco...  44699 Labour    Emily ... 12653      39.7      -2.9     31865     2
## 10 W07000058 Aberco...  44699 Plaid ... Lisa G...  2704       8.5      -1.4     31865     3
## # ... with 3,310 more rows, 3 more variables: turnout <dbl>, fname <chr>,
## #   lname <chr>, and abbreviated variable names 1constituency, 2electorate,
## #   3party_name, 4candidate, 5vote_share_percent, 6vote_share_change,
## #   7total_votes_cast
```

Counting Twice Over

```
ukvote2019 |>
```

```
count(constituency, name = "n_cands")
```

```
## # A tibble: 650 × 2
##   constituency n_cands
##   <chr>        <int>
## 1 Aberavon      7
## 2 Aberconwy     4
## 3 Aberdeen North 6
## 4 Aberdeen South 4
## 5 Aberdeenshire West & Kincardine 4
## 6 Airdrie & Shotts 5
## 7 Aldershot     4
## 8 Aldridge-Brownhills 5
## 9 Altrincham & Sale West 6
## 10 Alyn & Deeside 5
## # ... with 640 more rows
```

Counting Twice Over

```
ukvote2019 |>  
  count(constituency, name = "n_cands") |>  
  count(n_cands, name = "n_const")
```

```
## # A tibble: 8 × 2  
##   n_cands n_const  
##   <int>   <int>  
## 1      3      21  
## 2      4     194  
## 3      5     226  
## 4      6     139  
## 5      7      49  
## 6      8      18  
## 7      9       2  
## 8     12       1
```

Recapping Yesterday and Looking Ahead

Data Wrangling

Kieran Healy

Statistical Horizons, December 2022

Recapping and Looking Ahead

Coding as gardening

Working in RStudio with RMarkdown documents

Recapping and Looking Ahead

Core **dp1yr** verbs

Subset your table: `filter()` rows, `select()` columns

Logically `group_by()` one or more columns

Add columns with `mutate()`

Summarize (by group, or the whole table) with `summarize()`

Recapping and Looking Ahead

Expand your **dp1yr** actions

Count up rows with `n()`, `tally()` or `count()`

Calculate quantities with `sum()`, `mean()`, `min()`, etc

Subset rows with logical expressions or `slice` functions

Conditionally select columns by name directly, with `%in%` or `%nin%`, or with tidy selectors like `starts_with()`, `ends_with()`, `contains()`

Conditionally select columns by *type* with `where()` and some criterion, e.g. `where(is.numeric)`

Conditionally select and then *act* on columns with `across(where(<condition>), <action>)`

Recapping and Looking Ahead

Expand your **dp1yr** actions

Tidy up columns with `relocate()` and `rename()`

Tidy up rows with `arrange()`

Recapping and Looking Ahead

Today's sessions

Two dplyr gotchas

1. Comparisons filtering on proportions

Let's say you are working with proportions ...

```
df
```

```
## # A tibble: 4 × 3
##   id      prop1 prop2
##   <chr> <dbl> <dbl>
## 1 A      0.1   0.2
## 2 B      0.1   0.21
## 3 C      0.11  0.2
## 4 D      0.1   0.1
```

1. Comparisons filtering on proportions

And you want to focus on cases where `prop1` *plus* `prop2` is greater than 0.3:

1. Comparisons filtering on proportions

And you want to focus on cases where prop1 *plus* prop2 is greater than 0.3:

```
df |>  
  filter(prop1 + prop2 > 0.3)
```

```
## # A tibble: 3 × 3  
##   id    prop1 prop2  
##   <chr> <dbl> <dbl>  
## 1 A      0.1   0.2  
## 2 B      0.1  0.21  
## 3 C      0.11  0.2
```

The row with id **A** shouldn't have been included there.

1. Comparisons filtering on proportions

And you want to focus on cases where `prop1` *plus* `prop2` is greater than 0.3:

```
df |>
  filter(prop1 + prop2 > 0.3)
```

```
## # A tibble: 3 × 3
##   id      prop1 prop2
##   <chr> <dbl> <dbl>
## 1 A      0.1   0.2
## 2 B      0.1   0.21
## 3 C      0.11  0.2
```

The row with `id` **A** shouldn't have been included there.

This is not `dlpyr`'s fault. It's our floating point friend again.

1. Comparisons filtering on proportions

```
df |>  
  filter(prop1 + prop2 == 0.3)
```

```
## # A tibble: 0 × 3  
## # ... with 3 variables: id <chr>, prop1 <dbl>, prop2 <dbl>
```

The row with id **A** *should* have been included here!

1. Comparisons filtering on proportions

This won't give the right behavior either:

```
df |>
  mutate(prop3 = prop1 + prop2) |>
  filter(prop3 == 0.3)

## # A tibble: 0 × 4
## # ... with 4 variables: id <chr>, prop1 <dbl>, prop2 <dbl>, prop3 <dbl>
```

1. Comparisons filtering on proportions

So, beware.

```
df |>
  filter(prop1*100 + prop2*100 == 0.3*100)
```

```
## # A tibble: 1 × 3
##   id    prop1 prop2
##   <chr> <dbl> <dbl>
## 1 A      0.1   0.2
```

Better:

```
df |>
  filter(near(prop1 + prop2, 0.3))
```

```
## # A tibble: 1 × 3
##   id    prop1 prop2
##   <chr> <dbl> <dbl>
## 1 A      0.1   0.2
```

2. Zero Counts in dplyr

```
df <- read_csv(here("data", "first_terms.csv"))
```

```
df
```

```
## # A tibble: 280 × 4
##   pid start_year party    sex
##   <dbl> <date>    <chr>  <chr>
## 1  3160 2013-01-03 Republican M
## 2  3161 2013-01-03 Democrat  F
## 3  3162 2013-01-03 Democrat  M
## 4  3163 2013-01-03 Republican M
## 5  3164 2013-01-03 Democrat  M
## 6  3165 2013-01-03 Republican M
## 7  3166 2013-01-03 Republican M
## 8  3167 2013-01-03 Democrat  F
## 9  3168 2013-01-03 Republican M
## 10 3169 2013-01-03 Democrat  M
## # ... with 270 more rows
```

2. Zero Counts in dplyr

```
df |>
  group_by(start_year, party, sex) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N))
```

A tibble: 14 × 5

Groups: start_year, party [8]

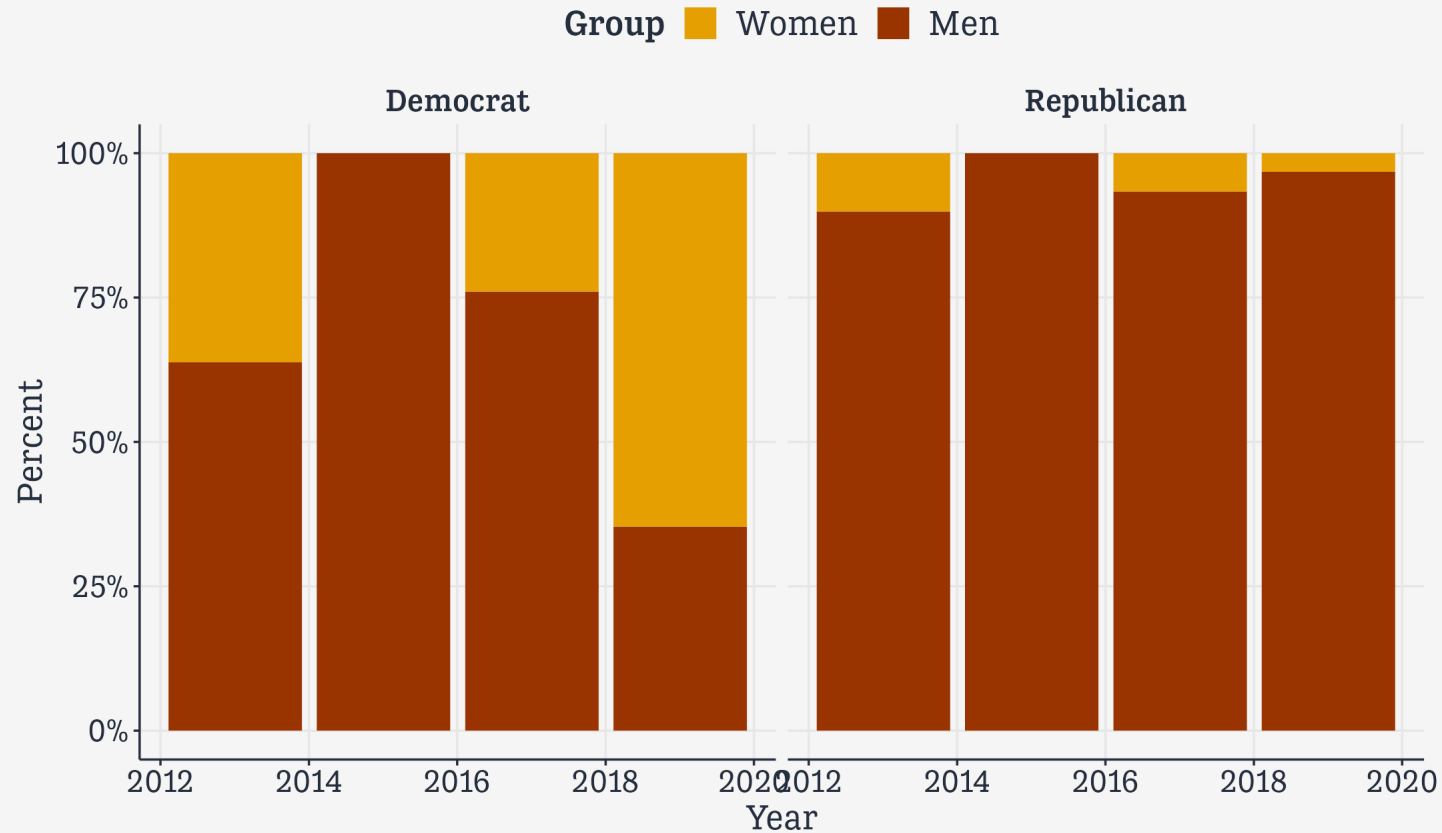
##	start_year	party	sex	N	freq
##	<date>	<chr>	<chr>	<int>	<dbl>
## 1	2013-01-03	Democrat	F	21	0.362
## 2	2013-01-03	Democrat	M	37	0.638
## 3	2013-01-03	Republican	F	8	0.101
## 4	2013-01-03	Republican	M	71	0.899
## 5	2015-01-03	Democrat	M	1	1
## 6	2015-01-03	Republican	M	5	1
## 7	2017-01-03	Democrat	F	6	0.24
## 8	2017-01-03	Democrat	M	19	0.76
## 9	2017-01-03	Republican	F	2	0.0667
## 10	2017-01-03	Republican	M	28	0.933
## 11	2019-01-03	Democrat	F	33	0.647
## 12	2019-01-03	Democrat	M	18	0.353
## 13	2019-01-03	Republican	F	1	0.0323
## 14	2019-01-03	Republican	M	30	0.968

2. Zero Counts in dplyr

```
p_col <- df |>
  group_by(start_year, party, sex) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N)) |>
  ggplot(aes(x = start_year,
             y = freq,
             fill = sex)) +
  geom_col() +
  scale_y_continuous(labels = scales::percent) +
  scale_fill_manual(values = sex_colors, labels = c("Women", "Men")) +
  labs(x = "Year", y = "Percent", fill = "Group") +
  facet_wrap(~ party)
```

2. Zero Counts in dplyr

p_col

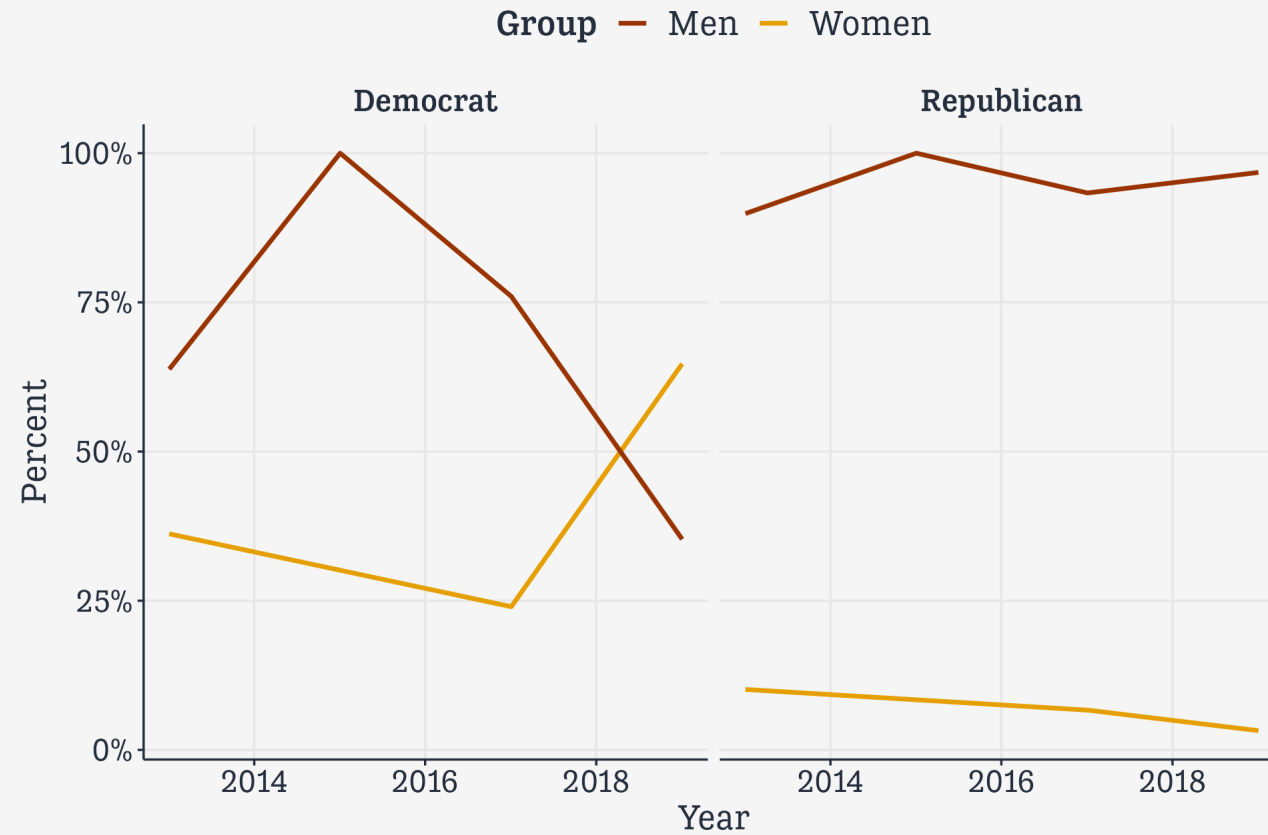


Zero Counts in dplyr

```
p_line <- df |>
  group_by(start_year, party, sex) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N)) |>
  ggplot(aes(x = start_year,
             y = freq,
             color = sex)) +
  geom_line(size = 1.1) +
  scale_y_continuous(labels = scales::percent) +
  scale_color_manual(values = sex_colors, labels = c("Women", "Men")) +
  guides(color = guide_legend(reverse = TRUE)) +
  labs(x = "Year", y = "Percent", color = "Group") +
  facet_wrap(~ party)
```

2. Zero Counts in dplyr

p_line



Option 1: **factors** and **. drop**

Factors are for categorical variables and are stored differently from characters.

This can matter when modeling, and also now.

```
df_f <- df |>  
  mutate(party_f = factor(party))
```

```
df_f
```

```
## # A tibble: 280 × 5  
##   pid start_year party    sex party_f  
##   <dbl> <date>    <chr>  <chr> <fct>  
## 1  3160 2013-01-03 Republican M    Republican  
## 2  3161 2013-01-03 Democrat  F    Democrat  
## 3  3162 2013-01-03 Democrat  M    Democrat  
## 4  3163 2013-01-03 Republican M    Republican  
## 5  3164 2013-01-03 Democrat  M    Democrat  
## 6  3165 2013-01-03 Republican M    Republican  
## 7  3166 2013-01-03 Republican M    Republican  
## 8  3167 2013-01-03 Democrat  F    Democrat  
## 9  3168 2013-01-03 Republican M    Republican  
## 10 3169 2013-01-03 Democrat  M    Democrat  
## # ... with 270 more rows
```

Option 1: **factors** and **. drop**

```
df_f |>
  group_by(party_f) |>
  tally()
```

```
## # A tibble: 2 × 2
##   party_f      n
##   <fct>    <int>
## 1 Democrat    135
## 2 Republican  145
```

Factors are integer values with named labels, or *levels*:

```
typeof(df_f$party_f)
```

```
## [1] "integer"
```

```
levels(df_f$party_f)
```

```
## [1] "Democrat" "Republican"
```

Option 1: **factors** and **. drop**

By default, unused levels won't display:

```
df_f <- df |>
  mutate(party_f = factor(party,
                          levels = c("Democrat",
                                      "Republican",
                                      "Libertarian")))

df_f |>
  group_by(party_f) |>
  tally()
```

```
## # A tibble: 2 × 2
##   party_f      n
##   <fct>    <int>
## 1 Democrat    135
## 2 Republican  145
```

```
levels(df_f$party_f)
```

```
## [1] "Democrat" "Republican" "Libertarian"
```

Option 1: **factors** and **. drop**

By default, unused levels won't display:

```
df |>
  mutate(across(where(is.character), as_factor)) |>
  group_by(start_year, party, sex) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N))
```

```
## # A tibble: 14 × 5
## # Groups:   start_year, party [8]
##   start_year party    sex      N  freq
##   <date>      <fct>    <fct> <int> <dbl>
## 1 2013-01-03 Republican M      71 0.899
## 2 2013-01-03 Republican F       8 0.101
## 3 2013-01-03 Democrat  M      37 0.638
## 4 2013-01-03 Democrat  F      21 0.362
## 5 2015-01-03 Republican M       5 1
## 6 2015-01-03 Democrat  M       1 1
## 7 2017-01-03 Republican M      28 0.933
## 8 2017-01-03 Republican F       2 0.0667
## 9 2017-01-03 Democrat  M      19 0.76
## 10 2017-01-03 Democrat  F       6 0.24
## 11 2019-01-03 Republican M      30 0.968
## 12 2019-01-03 Republican F       1 0.0323
## 13 2019-01-03 Democrat  M      18 0.353
## 14 2019-01-03 Democrat  F      33 0.647
```

Option 1: **factors** and **.drop**

You can make dplyr keep empty factor levels though:

```
df |>
  mutate(across(where(is.character), as_factor)) |>
  group_by(start_year, party, sex, .drop = FALSE) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N))
```

```
## # A tibble: 16 × 5
## # Groups:   start_year, party [8]
##   start_year party    sex      N  freq
##   <date>      <fct>    <fct> <int> <dbl>
## 1 2013-01-03 Republican M      71 0.899
## 2 2013-01-03 Republican F       8 0.101
## 3 2013-01-03 Democrat  M      37 0.638
## 4 2013-01-03 Democrat  F      21 0.362
## 5 2015-01-03 Republican M       5 1
## 6 2015-01-03 Republican F       0 0
## 7 2015-01-03 Democrat  M       1 1
## 8 2015-01-03 Democrat  F       0 0
## 9 2017-01-03 Republican M      28 0.933
## 10 2017-01-03 Republican F       2 0.0667
## 11 2017-01-03 Democrat  M      19 0.76
## 12 2017-01-03 Democrat  F       6 0.24
## 13 2019-01-03 Republican M      30 0.968
## 14 2019-01-03 Republican F       1 0.0323
## 15 2019-01-03 Democrat  M      18 0.353
## 16 2019-01-03 Democrat  F      33 0.647
```

Option 2: **ungroup()** and **complete()**

Maybe you don't want to deal with factors.

```
df_c <- df |>
  group_by(start_year, party, sex) |>
  summarize(N = n()) |>
  mutate(freq = N / sum(N)) |>
  ungroup() |>
  complete(start_year, party, sex,
           fill = list(N = 0, freq = 0))
```

Option 2: **ungroup()** and **complete()**

df_c

```
## # A tibble: 16 × 5
##   start_year party    sex      N   freq
##   <date>      <chr>  <chr> <int> <dbl>
## 1 2013-01-03 Democrat  F      21 0.362
## 2 2013-01-03 Democrat  M      37 0.638
## 3 2013-01-03 Republican F       8 0.101
## 4 2013-01-03 Republican M      71 0.899
## 5 2015-01-03 Democrat  F       0 0
## 6 2015-01-03 Democrat  M       1 1
## 7 2015-01-03 Republican F       0 0
## 8 2015-01-03 Republican M       5 1
## 9 2017-01-03 Democrat  F       6 0.24
## 10 2017-01-03 Democrat  M      19 0.76
## 11 2017-01-03 Republican F       2 0.0667
## 12 2017-01-03 Republican M      28 0.933
## 13 2019-01-03 Democrat  F      33 0.647
## 14 2019-01-03 Democrat  M      18 0.353
## 15 2019-01-03 Republican F       1 0.0323
## 16 2019-01-03 Republican M      30 0.968
```

Option 2: **ungroup()** and **complete()**

```
p_out <- df_c |>
  ggplot(aes(x = start_year,
             y = freq,
             color = sex)) +
  geom_line(size = 1.1) +
  scale_y_continuous(labels = scales::percent) +
  scale_color_manual(values = sex_colors, labels = c("Women", "Men")) +
  guides(color = guide_legend(reverse = TRUE)) +
  labs(x = "Year", y = "Percent", color = "Group") +
  facet_wrap(~ party)
```


Option 2: **ungroup()** and **complete()**

p_out

