Reading in data with readr and haven

Session 6

Kieran Healy

Statistical Horizons, September 2021

Load the packages, as always

```
library(here)      # manage file paths
```

```
## here() starts at /Users/kjhealy/Documents/courses/data_wrangling
```

```
library(socviz)    # data and some useful functions
```

```
##
## Attaching package: 'socviz'

## The following object is masked from 'package:kjhutils':
##
##     %nin%
```

```
library(tidyverse) # your friend and mine
```

```
## ── Attaching packages ─────────────────────────────────────── tidyverse 1.3.1 ──

## ✓ ggplot2 3.3.5     ✓ purrr   0.3.4
## ✓ tibble  3.1.4     ✓ dplyr   1.0.7
## ✓ tidyr   1.1.3     ✓ stringr 1.4.0
## ✓ readr   2.0.1     ✓ forcats 0.5.1

## ── Conflicts ────────────────────────────────────────── tidyverse_conflicts() ──
## x readr::edition_get()   masks testthat::edition_get()
## x dplyr::filter()        masks stats::filter()
## x purrr::is_null()       masks testthat::is_null()
## x dplyr::lag()           masks stats::lag()
## x readr::local_edition() masks testthat::local_edition()
## x dplyr::matches()       masks tidyr::matches(), testthat::matches()
```

```
library(haven)     # for Stata, SAS, and SPSS files
```

We've put a lot of pieces in place at this point

Including several things we haven't fully exploited yet

Data we want to get into R

Data we want to get into R

Nice, clean CSV files.

Data we want to get into R

Nice, clean CSV files.

More troublesome CSVs.

Data we want to get into R

Nice, clean CSV files.

More troublesome CSVs.

Other plain-text formats.

Data we want to get into R

Nice, clean CSV files.

More troublesome CSVs.

Other plain-text formats.

Foreign formats, like Stata.

Data we want to get into R

Nice, clean CSV files.

More troublesome CSVs.

Other plain-text formats.

Foreign formats, like Stata.

Quite messy things like tables on web pages.

Data we want to get into R

Nice, clean CSV files.

More troublesome CSVs.

Other plain-text formats.

Foreign formats, like Stata.

Quite messy things like tables on web pages.

... and more besides.

Reading in CSV files

CSV is not really a proper format at all!

Reading in CSV files

CSV is not really a proper format at all!

Base R has `read.csv()`

Reading in CSV files

CSV is not really a proper format at all!

Base R has `read.csv()`

As is often the case, the tidyverse has a corresponding "underscored" version, `read_csv()`. It is *much* pickier and more talkative than the Base R version.

Where's the data? Using here()

If we're loading a file, it's coming from *somewhere*.

If it's on our local disk somewhere, we will need to interact with the file system. We should try to do this in a way that avoids *absolute* file paths.

```
# This is not portable
df <- read_csv("/Users/kjhealy/Documents/data/misc/project/data/mydata.csv")
```

Where's the data? Using here()

If we're loading a file, it's coming from *somewhere*.

If it's on our local disk somewhere, we will need to interact with the file system. We should try to do this in a way that avoids *absolute* file paths.

```
# This is not portable
df <- read_csv("/Users/kjhealy/Documents/data/misc/project/data/mydata.csv")
```

We should also do it in a way that is *platform independent*.

This makes it easier to share your work, move it around, etc. Projects should be self-contained.

Where's the data? Using here()

The `here` package, and `here()` function builds paths relative to the top level of your R project.

```r
here() # this path will be different for you
```

```
## [1] "/Users/kjhealy/Documents/courses/data_wrangling"
```

Where's the data? Using here()

This seminar's files all live in an RStudio project. It looks like this:

```
## /Users/kjhealy/Documents/courses/data_wrangling
## ├── LICENSE
## ├── README.Rmd
## ├── README.md
## ├── code
## ├── course_notes.Rmd
## ├── course_notes.html
## ├── data
## ├── data_wrangling.Rproj
## ├── docs
## ├── dot_r.html
## ├── dot_r.md
## ├── office
## ├── pdf_slides
## ├── r_code
## ├── scratch.Rmd
## ├── scratch.docx
## ├── scratch.html
## ├── scratch.log
## ├── scratch.pdf
## ├── scratch.tex
## ├── scratch_kh.Rmd
## ├── slides
```

Where's the data? Using here()

So:

```
## Load the file relative to the path from the top of the project, without separators, etc
organs <- read_csv(file = here("data", "organdonation.csv"))
```

Where's the data? Using here()

So:

```
## Load the file relative to the path from the top of the project, without separators, etc
organans <- read_csv(file = here("data", "organdonation.csv"))
```

```
organs
```

```
## # A tibble: 238 × 21
##    country    year donors   pop pop.dens   gdp gdp.lag health health.lag pubhealth
##    <chr>     <dbl>  <dbl> <dbl>    <dbl> <dbl>   <dbl>  <dbl>      <dbl>     <dbl>
##  1 Australia    NA  NA    17065    0.220 16774   16591   1300       1224       4.8
##  2 Australia  1991  12.1  17284    0.223 17171   16774   1379       1300       5.4
##  3 Australia  1992  12.4  17495    0.226 17914   17171   1455       1379       5.4
##  4 Australia  1993  12.5  17667    0.228 18883   17914   1540       1455       5.4
##  5 Australia  1994  10.2  17855    0.231 19849   18883   1626       1540       5.4
##  6 Australia  1995  10.2  18072    0.233 21079   19849   1737       1626       5.5
##  7 Australia  1996  10.6  18311    0.237 21923   21079   1846       1737       5.6
##  8 Australia  1997  10.3  18518    0.239 22961   21923   1948       1846       5.7
##  9 Australia  1998  10.5  18711    0.242 24148   22961   2077       1948       5.9
## 10 Australia  1999   8.67 18926    0.244 25445   24148   2231       2077       6.1
## # … with 228 more rows, and 11 more variables: roads <dbl>, cerebvas <dbl>,
## #   assault <dbl>, external <dbl>, txp.pop <dbl>, world <chr>, opt <chr>,
## #   consent.law <chr>, consent.practice <chr>, consistent <chr>, ccode <chr>
```

And there it is.

Where's the data? Using here()

Get in the habit of putting this at the top of your files:

```
here::i_am("analysis.Rmd") # or whatever your Rmd or R file is called
```

See the `here` project page for more details.

`read_csv()` comes in different varieties

`read_csv()` Field separator is ,

```
organs <- read_csv(file = here("data", "organdonation.csv"))
```

`read_csv2()` Field separator is ;

```
# Example only
my_data <- read_csv2(file = here("data", "my_euro_file.csv))
```

Both are special cases of `read_delim()`

Other species are also catered to

`read_tsv()` Tab separated.

`read_fwf()` Fixed-width files.

`read_log()` Log files (i.e. computer log files).

`read_lines()` Just read in lines, without trying to parse them.

Also often useful ...

`read_table()`

Data that's separated by one (or more) columns of space.

You can read files remotely, too

You can give all of these functions local files, or they can point to URLs.

Compressed files will be automatically uncompressed.

(Be careful what you download from remote locations!)

```
organ_remote <- read_csv("http://kjhealy.co/organdonation.csv")

organ_remote
```

```
## # A tibble: 238 × 21
##    country    year donors   pop pop.dens   gdp gdp.lag health health.lag pubhealth
##    <chr>     <dbl>  <dbl> <dbl>    <dbl> <dbl>   <dbl>  <dbl>      <dbl>     <dbl>
##  1 Australia    NA  NA    17065    0.220 16774   16591   1300       1224       4.8
##  2 Australia  1991  12.1  17284    0.223 17171   16774   1379       1300       5.4
##  3 Australia  1992  12.4  17495    0.226 17914   17171   1455       1379       5.4
##  4 Australia  1993  12.5  17667    0.228 18883   17914   1540       1455       5.4
##  5 Australia  1994  10.2  17855    0.231 19849   18883   1626       1540       5.4
##  6 Australia  1995  10.2  18072    0.233 21079   19849   1737       1626       5.5
##  7 Australia  1996  10.6  18311    0.237 21923   21079   1846       1737       5.6
##  8 Australia  1997  10.3  18518    0.239 22961   21923   1948       1846       5.7
##  9 Australia  1998  10.5  18711    0.242 24148   22961   2077       1948       5.9
## 10 Australia  1999   8.67 18926    0.244 25445   24148   2231       2077       6.1
## # … with 228 more rows, and 11 more variables: roads <dbl>, cerebvas <dbl>,
## #   assault <dbl>, external <dbl>, txp.pop <dbl>, world <chr>, opt <chr>,
## #   consent.law <chr>, consent.practice <chr>, consistent <chr>, ccode <chr>
```

An example: read_table()

```
England and Wales, Total Population, Death rates (period 1x1),  Last modified: 02 Apr
2018;  Methods Protocol: v6 (2017)

  Year          Age             Female          Male            Total
  1841            0             0.136067        0.169189        0.152777
  1841            1             0.059577        0.063208        0.061386
  1841            2             0.036406        0.036976        0.036689
  1841            3             0.024913        0.026055        0.025480
  1841            4             0.018457        0.019089        0.018772
  1841            5             0.013967        0.014279        0.014123
  1841            6             0.010870        0.011210        0.011040
  1841            7             0.008591        0.008985        0.008788
  1841            8             0.006860        0.007246        0.007053
  1841            9             0.005772        0.006050        0.005911
  1841           10             0.005303        0.005382        0.005343
  1841           11             0.005114        0.005002        0.005057
  1841           12             0.005145        0.004856        0.004999
  1841           13             0.005455        0.004955        0.005202
```

```
  1841          105             0.576967        1.727848        0.700373
  1841          106             0.677711        6.000000        0.795287
  1841          107             0.900000               .        0.900000
  1841          108             1.388430               .        1.388430
  1841          109                    .               .               .
  1841          110+                   .               .               .
  1842            0             0.148491        0.184007        0.166481
  1842            1             0.063038        0.066596        0.064818
  1842            2             0.035203        0.035854        0.035527
```

An example: read_table()

```
England and Wales, Total Population, Death rates (period 1x1),  Last modified: 02 Apr
2018;  Methods Protocol: v6 (2017)

  Year          Age             Female          Male            Total
  1841            0             0.136067        0.169189        0.152777
  1841            1             0.059577        0.063208        0.061386
  1841            2             0.036406        0.036976        0.036689
  1841            3             0.024913        0.026055        0.025480
  1841            4             0.018457        0.019089        0.018772
  1841            5             0.013967        0.014279        0.014123
  1841            6             0.010870        0.011210        0.011040
  1841            7             0.008591        0.008985        0.008788
  1841            8             0.006860        0.007246        0.007053
  1841            9             0.005772        0.006050        0.005911
  1841           10             0.005303        0.005382        0.005343
  1841           11             0.005114        0.005002        0.005057
  1841           12             0.005145        0.004856        0.004999
  1841           13             0.005455        0.004955        0.005202
```

```
  1841          105             0.576967        1.727848        0.700373
  1841          106             0.677711        6.000000        0.795287
  1841          107             0.900000               .        0.900000
  1841          108             1.388430               .        1.388430
  1841          109                    .               .               .
  1841          110+                   .               .               .
  1842            0             0.148491        0.184007        0.166481
  1842            1             0.063038        0.066596        0.064818
  1842            2             0.035203        0.035854        0.035527
```

```r
engmort <- read_table(here("data", "mortality.txt"),
                      skip = 2, na = ".")

engmort
```

```
## # A tibble: 222 × 5
##     Year Age    Female    Male   Total
##    <dbl> <chr>   <dbl>   <dbl>   <dbl>
##  1  1841 0     0.136   0.169   0.153
##  2  1841 1     0.0596  0.0632  0.0614
##  3  1841 2     0.0364  0.0370  0.0367
##  4  1841 3     0.0249  0.0261  0.0255
##  5  1841 4     0.0185  0.0191  0.0188
##  6  1841 5     0.0140  0.0143  0.0141
##  7  1841 6     0.0109  0.0112  0.0110
##  8  1841 7     0.00859 0.00898 0.00879
##  9  1841 8     0.00686 0.00725 0.00705
## 10  1841 9     0.00577 0.00605 0.00591
## # … with 212 more rows
```

Pay attention to the column specification

```
engmort <- read_table(here("data", "mortality.txt"),
                      skip = 2, na = ".")
```

```
##
## ── Column specification ────────────────────────────────────────────────────────
## cols(
##   Year = col_double(),
##   Age = col_character(),
##   Female = col_double(),
##   Male = col_double(),
##   Total = col_double()
## )
```

The column specification tells you what the read function did. That is, how it interpreted each of the columns. It will also report if things don't go as expected.

Pay attention to the column specification

```r
engmort <- read_table(here("data", "mortality.txt"),
                      skip = 2, na = ".")
```

```
##
## ── Column specification ────────────────────────────────────────────────────────
## cols(
##   Year = col_double(),
##   Age = col_character(),
##   Female = col_double(),
##   Male = col_double(),
##   Total = col_double()
## )
```

The column specification tells you what the read function did. That is, how it interpreted each of the columns. It will also report if things don't go as expected.

Why is `age` imported in `character` format?

Pay attention to the **column specification**

Absent you giving them a column specification, the `read_` functions try to *guess* what the type of each column is. They do this by looking at the first thousand rows of each column.

They may guess incorrectly!

Normalizing names and recoding

```r
read_table(here("data", "mortality.txt"),
           skip = 2, na = ".")
```

```
## # A tibble: 222 × 5
##     Year Age    Female    Male   Total
##    <dbl> <chr>   <dbl>   <dbl>   <dbl>
##  1  1841 0     0.136   0.169   0.153
##  2  1841 1     0.0596  0.0632  0.0614
##  3  1841 2     0.0364  0.0370  0.0367
##  4  1841 3     0.0249  0.0261  0.0255
##  5  1841 4     0.0185  0.0191  0.0188
##  6  1841 5     0.0140  0.0143  0.0141
##  7  1841 6     0.0109  0.0112  0.0110
##  8  1841 7     0.00859 0.00898 0.00879
##  9  1841 8     0.00686 0.00725 0.00705
## 10  1841 9     0.00577 0.00605 0.00591
## # … with 212 more rows
```

Normalizing names and recoding

```r
read_table(here("data", "mortality.txt"),
           skip = 2, na = ".") %>%
  janitor::clean_names()
```

```
## # A tibble: 222 × 5
##     year age    female    male   total
##    <dbl> <chr>   <dbl>   <dbl>   <dbl>
##  1  1841 0     0.136   0.169   0.153  
##  2  1841 1     0.0596  0.0632  0.0614 
##  3  1841 2     0.0364  0.0370  0.0367 
##  4  1841 3     0.0249  0.0261  0.0255 
##  5  1841 4     0.0185  0.0191  0.0188 
##  6  1841 5     0.0140  0.0143  0.0141 
##  7  1841 6     0.0109  0.0112  0.0110 
##  8  1841 7     0.00859 0.00898 0.00879
##  9  1841 8     0.00686 0.00725 0.00705
## 10  1841 9     0.00577 0.00605 0.00591
## # … with 212 more rows
```

Normalizing names and recoding

```r
read_table(here("data", "mortality.txt"),
           skip = 2, na = ".") %>%
  janitor::clean_names() %>%
  mutate(age = as.integer(recode(age, "110+" = "110")))
```

```
## # A tibble: 222 × 5
##     year   age  female    male   total
##    <dbl> <int>   <dbl>   <dbl>   <dbl>
##  1  1841     0 0.136   0.169   0.153
##  2  1841     1 0.0596  0.0632  0.0614
##  3  1841     2 0.0364  0.0370  0.0367
##  4  1841     3 0.0249  0.0261  0.0255
##  5  1841     4 0.0185  0.0191  0.0188
##  6  1841     5 0.0140  0.0143  0.0141
##  7  1841     6 0.0109  0.0112  0.0110
##  8  1841     7 0.00859 0.00898 0.00879
##  9  1841     8 0.00686 0.00725 0.00705
## 10  1841     9 0.00577 0.00605 0.00591
## # … with 212 more rows
```

Normalizing names and recoding

```r
read_table(here("data", "mortality.txt"),
           skip = 2, na = ".") %>%
  janitor::clean_names() %>%
  mutate(age = as.integer(recode(age, "110+" = "110")))
```

```
## # A tibble: 222 × 5
##     year   age  female    male   total
##    <dbl> <int>   <dbl>   <dbl>   <dbl>
##  1  1841     0 0.136   0.169   0.153
##  2  1841     1 0.0596  0.0632  0.0614
##  3  1841     2 0.0364  0.0370  0.0367
##  4  1841     3 0.0249  0.0261  0.0255
##  5  1841     4 0.0185  0.0191  0.0188
##  6  1841     5 0.0140  0.0143  0.0141
##  7  1841     6 0.0109  0.0112  0.0110
##  8  1841     7 0.00859 0.00898 0.00879
##  9  1841     8 0.00686 0.00725 0.00705
## 10  1841     9 0.00577 0.00605 0.00591
## # … with 212 more rows
```

The `janitor` package is very handy!

The main cost of normalizing names comes with, e.g., data where there is a codebook you need to consult. But in general it's worth it.

More on column specifications

CDC/NCHS data: Provisional COVID-19 Death Counts by Sex, Age, and State

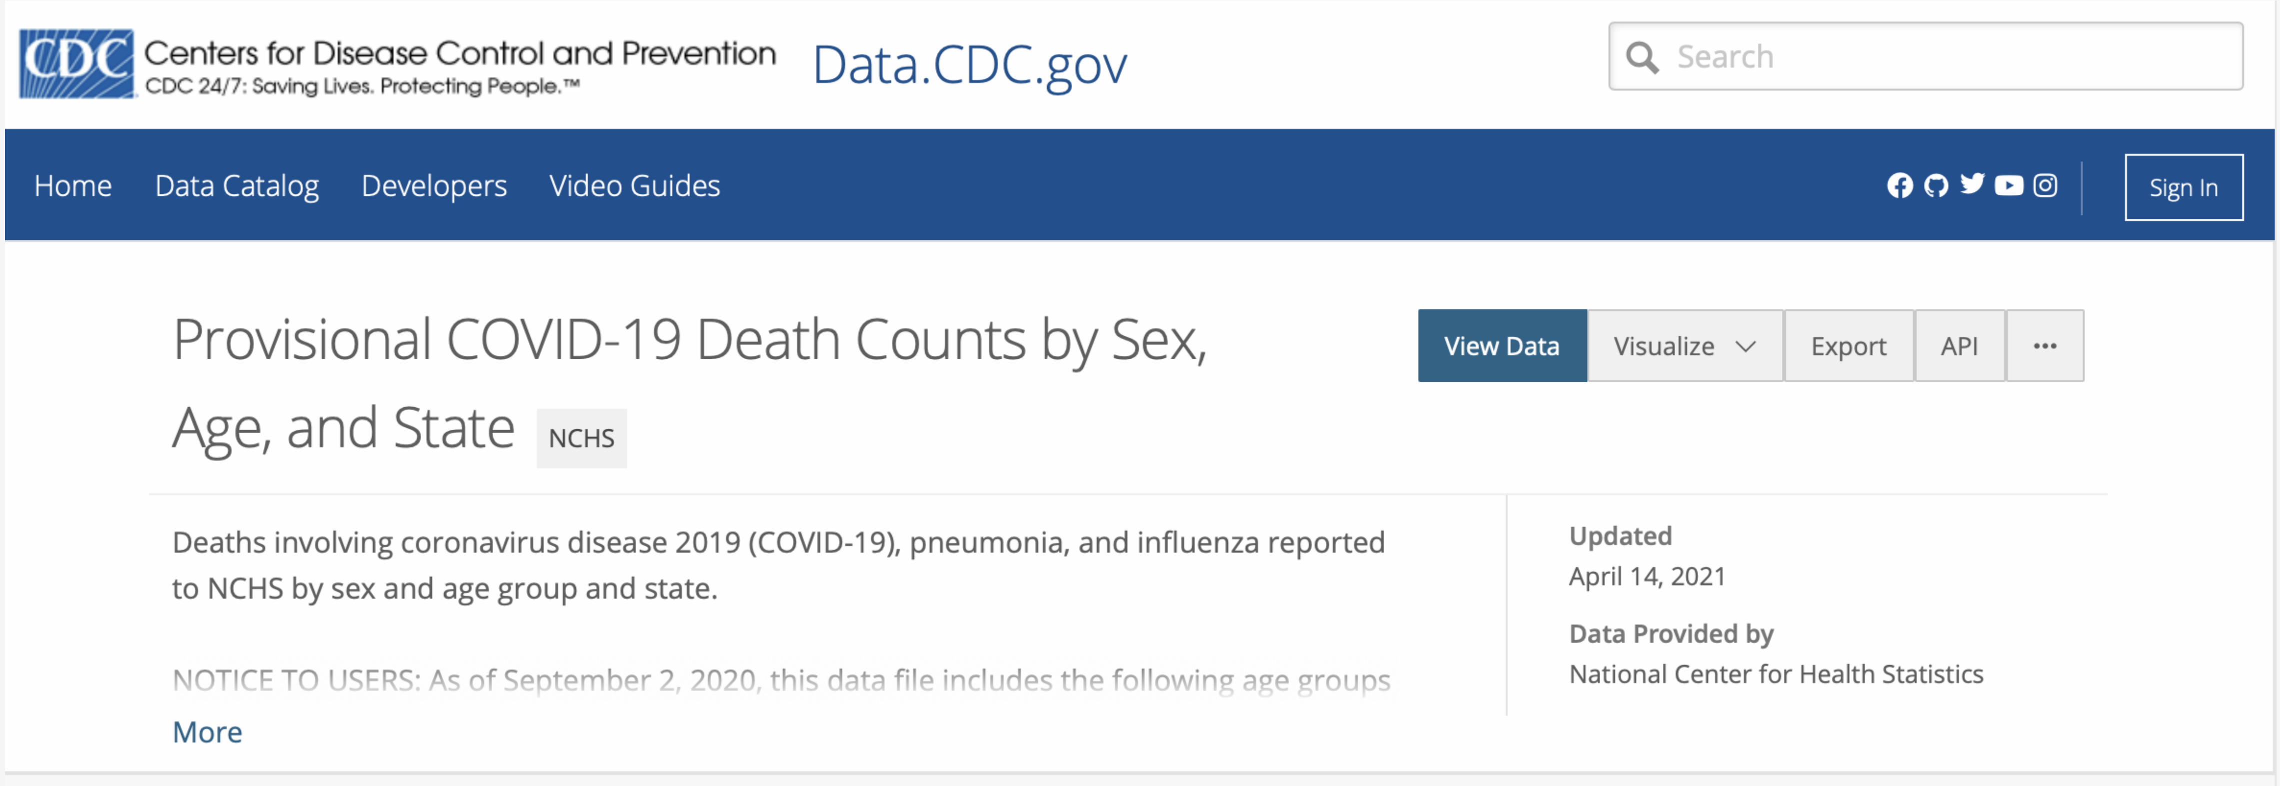

More on column specifications

What's in this Dataset?

Rows	Columns
52.3K	16

Columns in this Dataset

Column Name	Description	Type
Data As Of	Date of analysis	Date & Time
Start Date	First date of data period	Date & Time
End Date	Last date of data period	Date & Time
Group	Indicator of whether data measured by Month, by Year, or ...	Plain Text
Year	Year in which death occurred	Number
Month	Month in which death occurred	Number

Let's try to load it

```
nchs <- with_edition(1, read_csv(here("data", "SAS_on_2021-04-13.csv")))
```

```
## Warning: 88128 parsing failures.
##  row  col           expected actual                                                                        file
## 2755 Year 1/0/T/F/TRUE/FALSE   2020 '/Users/kjhealy/Documents/courses/data_wrangling/data/SAS_on_2021-04-13.csv'
## 2756 Year 1/0/T/F/TRUE/FALSE   2020 '/Users/kjhealy/Documents/courses/data_wrangling/data/SAS_on_2021-04-13.csv'
## 2757 Year 1/0/T/F/TRUE/FALSE   2020 '/Users/kjhealy/Documents/courses/data_wrangling/data/SAS_on_2021-04-13.csv'
## 2758 Year 1/0/T/F/TRUE/FALSE   2020 '/Users/kjhealy/Documents/courses/data_wrangling/data/SAS_on_2021-04-13.csv'
## 2759 Year 1/0/T/F/TRUE/FALSE   2020 '/Users/kjhealy/Documents/courses/data_wrangling/data/SAS_on_2021-04-13.csv'
## .... .... .................. ...... ...........................................................................
## See problems(...) for more details.
```

Let's try to load it

```
problems(nchs)
```

```
## # A tibble: 88,128 × 5
##      row col   expected           actual file
##    <int> <chr> <chr>              <chr>  <chr>
##  1  2755 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  2  2756 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  3  2757 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  4  2758 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  5  2759 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  6  2760 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  7  2761 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  8  2762 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  9  2763 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## 10  2764 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## # … with 88,118 more rows
```

Let's try to load it

```
problems(nchs)
```

```
## # A tibble: 88,128 × 5
##      row col   expected           actual file
##    <int> <chr> <chr>              <chr>  <chr>
##  1  2755 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  2  2756 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  3  2757 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  4  2758 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  5  2759 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  6  2760 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  7  2761 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  8  2762 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  9  2763 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## 10  2764 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## # … with 88,118 more rows
```

Problems are stored as an attribute of the `nchs` object, so we can revisit them.

Let's try to load it

```
problems(nchs)
```

```
## # A tibble: 88,128 × 5
##      row col   expected           actual file
##    <int> <chr> <chr>              <chr>  <chr>
##  1  2755 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  2  2756 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  3  2757 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  4  2758 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  5  2759 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  6  2760 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  7  2761 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  8  2762 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
##  9  2763 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## 10  2764 Year  1/0/T/F/TRUE/FALSE 2020   '/Users/kjhealy/Documents/courses/data…
## # … with 88,118 more rows
```

Problems are stored as an attribute of the `nchs` object, so we can revisit them.

Parsing failures tend to cascade. Our data only has 56k rows but we got 88k failures.

Take a look with head()

```
head(nchs)
```

```
## # A tibble: 6 × 16
##   `Data As Of` `Start Date` `End Date` Group    Year  Month State Sex   `Age Group`
##   <chr>        <chr>        <chr>      <chr>    <lgl> <lgl> <chr> <chr> <chr>
## 1 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … All Ages
## 2 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … Under 1 ye…
## 3 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … 0-17 years
## 4 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … 1-4 years
## 5 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … 5-14 years
## 6 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Unit… All … 15-24 years
## # … with 7 more variables: COVID-19 Deaths <dbl>, Total Deaths <dbl>,
## #   Pneumonia Deaths <dbl>, Pneumonia and COVID-19 Deaths <dbl>,
## #   Influenza Deaths <dbl>, Pneumonia, Influenza, or COVID-19 Deaths <dbl>,
## #   Footnote <chr>
```

Take a look with tail()

```
tail(nchs)
```

```
## # A tibble: 6 × 16
##   `Data As Of` `Start Date` `End Date` Group    Year  Month State Sex   `Age Group`
##   <chr>        <chr>        <chr>      <chr>    <lgl> <lgl> <chr> <chr> <chr>
## 1 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 45-54 years
## 2 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 50-64 years
## 3 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 55-64 years
## 4 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 65-74 years
## 5 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 75-84 years
## 6 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Puer… Fema… 85 years a…
## # … with 7 more variables: COVID-19 Deaths <dbl>, Total Deaths <dbl>,
## #   Pneumonia Deaths <dbl>, Pneumonia and COVID-19 Deaths <dbl>,
## #   Influenza Deaths <dbl>, Pneumonia, Influenza, or COVID-19 Deaths <dbl>,
## #   Footnote <chr>
```

Take a look with slice_sample()

```
nchs %>%
  slice_sample(n = 10)
```

```
## # A tibble: 10 × 16
##    `Data As Of` `Start Date` `End Date` Group    Year  Month State         Sex  
##    <chr>        <chr>        <chr>      <chr>    <lgl> <lgl> <chr>         <chr>
##  1 04/07/2021   02/01/2020   02/29/2020 By Month NA    NA    Alaska        All …
##  2 04/07/2021   01/01/2020   01/31/2020 By Month NA    TRUE  Georgia       Fema…
##  3 04/07/2021   11/01/2020   11/30/2020 By Month NA    NA    California    All …
##  4 04/07/2021   04/01/2021   04/03/2021 By Month NA    NA    Rhode Island  Male 
##  5 04/07/2021   04/01/2020   04/30/2020 By Month NA    NA    Illinois      All …
##  6 04/07/2021   01/01/2021   04/03/2021 By Year  NA    NA    Nebraska      Male 
##  7 04/07/2021   02/01/2020   02/29/2020 By Month NA    NA    United States Fema…
##  8 04/07/2021   10/01/2020   10/31/2020 By Month NA    NA    Alabama       All …
##  9 04/07/2021   01/01/2021   01/31/2021 By Month NA    TRUE  New Hampshire Fema…
## 10 04/07/2021   01/01/2021   01/31/2021 By Month NA    TRUE  Kansas        Fema…
## # … with 8 more variables: Age Group <chr>, COVID-19 Deaths <dbl>,
## #   Total Deaths <dbl>, Pneumonia Deaths <dbl>,
## #   Pneumonia and COVID-19 Deaths <dbl>, Influenza Deaths <dbl>,
## #   Pneumonia, Influenza, or COVID-19 Deaths <dbl>, Footnote <chr>
```

Aside: one that happened earlier ...

```
nchs %>%
  slice_sample(n = 10)
```

```
## # A tibble: 10 x 16
##    `Data As Of` `Start Date` `End Date` Group   Year  Month State        Sex
##    <chr>        <chr>        <chr>      <chr>   <lgl> <lgl> <chr>        <chr>
##  1 04/07/2021   01/01/2020   04/03/2021 By Tot… NA    NA    Minnesota    Male
##  2 04/07/2021   02/01/2020   02/29/2020 By Mon… NA    NA    Georgia      Male
##  3 04/07/2021   02/01/2021   02/28/2021 By Mon… NA    NA    Maine        Male
##  4 04/07/2021   11/01/2020   11/30/2020 By Mon… NA    NA    New Jersey   Female
##  5 04/07/2021   01/01/2020   12/31/2020 By Year NA    NA    Rhode Island All Se…
##  6 04/07/2021   01/01/2020   01/31/2020 By Mon… NA    TRUE  New York     All Se…
##  7 04/07/2021   05/01/2020   05/31/2020 By Mon… NA    NA    District of… Male
##  8 04/07/2021   04/01/2021   04/03/2021 By Mon… NA    NA    North Carol… Female
##  9 04/07/2021   03/01/2021   03/31/2021 By Mon… NA    NA    Kentucky     Male
## 10 04/07/2021   04/01/2021   04/03/2021 By Mon… NA    NA    New Mexico   Female
## # … with 8 more variables: Age Group <chr>, COVID-19 Deaths <dbl>,
## #   Total Deaths <dbl>, Pneumonia Deaths <dbl>,
## #   Pneumonia and COVID-19 Deaths <dbl>, Influenza Deaths <dbl>,
## #   Pneumonia, Influenza, or COVID-19 Deaths <dbl>, Footnote <chr>
```

Take a look with slice()

Let's look at the rows `read_csv()` complained about.

```
nchs %>%
  slice(2750:2760)
```

```
## # A tibble: 11 × 16
##    `Data As Of` `Start Date` `End Date` Group    Year  Month State         Sex
##    <chr>        <chr>        <chr>      <chr>    <lgl> <lgl> <chr>         <chr>
##  1 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Puerto Rico   Fema…
##  2 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Puerto Rico   Fema…
##  3 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Puerto Rico   Fema…
##  4 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Puerto Rico   Fema…
##  5 04/07/2021   01/01/2020   04/03/2021 By Total NA    NA    Puerto Rico   Fema…
##  6 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
##  7 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
##  8 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
##  9 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
## 10 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
## 11 04/07/2021   01/01/2020   12/31/2020 By Year  NA    NA    United States All …
## # … with 8 more variables: Age Group <chr>, COVID-19 Deaths <dbl>,
## #   Total Deaths <dbl>, Pneumonia Deaths <dbl>,
## #   Pneumonia and COVID-19 Deaths <dbl>, Influenza Deaths <dbl>,
## #   Pneumonia, Influenza, or COVID-19 Deaths <dbl>, Footnote <chr>
```

Take a look with slice()

```
nchs %>%
  slice(2750:2760) %>%
  select(Year, Month, State)
```

```
## # A tibble: 11 × 3
##    Year  Month State
##    <lgl> <lgl> <chr>
##  1 NA    NA    Puerto Rico
##  2 NA    NA    Puerto Rico
##  3 NA    NA    Puerto Rico
##  4 NA    NA    Puerto Rico
##  5 NA    NA    Puerto Rico
##  6 NA    NA    United States
##  7 NA    NA    United States
##  8 NA    NA    United States
##  9 NA    NA    United States
## 10 NA    NA    United States
## 11 NA    NA    United States
```

Hm, something to do with the transition to national numbers maybe?

Take a look with select() and filter()

```
nchs %>%
  select(Year, Month, State) %>%
  filter(State == "New York")
```

```
## # A tibble: 969 × 3
##    Year  Month State
##    <lgl> <lgl> <chr>
##  1 NA    NA    New York
##  2 NA    NA    New York
##  3 NA    NA    New York
##  4 NA    NA    New York
##  5 NA    NA    New York
##  6 NA    NA    New York
##  7 NA    NA    New York
##  8 NA    NA    New York
##  9 NA    NA    New York
## 10 NA    NA    New York
## # … with 959 more rows
```

Take a look with is.na()

```
nchs %>%
  select(Year, Month, State) %>%
  filter(!is.na(Year))
```

```
## # A tibble: 0 × 3
## # … with 3 variables: Year <lgl>, Month <lgl>, State <chr>
```

It really has been read in as a completely empty column.

That doesn't seem like it can be right.

Take a look with distinct()

```
nchs %>%
  select(Year) %>%
  distinct(Year)
```

```
## # A tibble: 1 × 1
##   Year
##   <lgl>
## 1 NA
```

Again, it's been read in as a completely empty column.

Take a look with read_lines()

Time to reach for a different kitchen knife.

```
read_lines(here("data", "SAS_on_2021-04-13.csv"), n_max = 10)
```

```
##  [1] "Data As Of,Start Date,End Date,Group,Year,Month,State,Sex,Age Group,COVID-19 Deaths,Total Deaths,Pneumonia Death
##  [2] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,All Ages,539723,4161167,466437,263147,9037,7
##  [3] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,Under 1 year,59,22626,246,10,21,316,"
##  [4] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,0-17 years,251,39620,667,46,179,1051,"
##  [5] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,1-4 years,31,4069,137,5,61,224,"
##  [6] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,5-14 years,89,6578,195,19,76,341,"
##  [7] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,15-24 years,804,42596,930,317,81,1493,"
##  [8] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,18-29 years,1996,75339,2184,884,150,3434,"
##  [9] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,25-34 years,3543,88196,3493,1617,237,5638,"
## [10] "04/07/2021,01/01/2020,04/03/2021,By Total,,,United States,All Sexes,30-39 years,5792,107348,5276,2658,318,8706,"
```

We can get the whole thing this way

```
raw_file <- read_lines(here("data", "SAS_on_2021-04-13.csv"))
```

This imports the data as a long, long character vector, with each element being a line.

```
# reminder: indexing 1D vectors
letters[5:6]
```

```
## [1] "e" "f"
```

Now we're just looking at lines in a file

```r
# This is not a tibble; we have to index it the basic way
raw_file[2753:2758]
```

```
## [1] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,65-74 years,203,2650,410,151,,466,One or more data
## [2] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,75-84 years,234,4274,656,154,16,751,"
## [3] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,85 years and over,222,6164,795,136,29,909,"
## [4] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,All Ages,380949,3372967,349667,178222,8779
## [5] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,Under 1 year,48,19356,224,9,21,284,"
## [6] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,0-17 years,189,33808,598,35,178,930,"
```

Now we're just looking at lines in a file

```r
# This is not a tibble; we have to index it the basic way
raw_file[2753:2758]
```

```
## [1] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,65-74 years,203,2650,410,151,,466,One or more data
## [2] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,75-84 years,234,4274,656,154,16,751,"
## [3] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,85 years and over,222,6164,795,136,29,909,"
## [4] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,All Ages,380949,3372967,349667,178222,8779
## [5] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,Under 1 year,48,19356,224,9,21,284,"
## [6] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,0-17 years,189,33808,598,35,178,930,"
```

There you are, you little bastard.

Now we're just looking at lines in a file

```r
# This is not a tibble; we have to index it the basic way
raw_file[2753:2758]
```

```
## [1] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,65-74 years,203,2650,410,151,,466,One or more data
## [2] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,75-84 years,234,4274,656,154,16,751,"
## [3] "04/07/2021,01/01/2020,04/03/2021,By Total,,,Puerto Rico,Female,85 years and over,222,6164,795,136,29,909,"
## [4] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,All Ages,380949,3372967,349667,178222,8779
## [5] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,Under 1 year,48,19356,224,9,21,284,"
## [6] "04/07/2021,01/01/2020,12/31/2020,By Year,2020,,United States,All Sexes,0-17 years,189,33808,598,35,178,930,"
```

There you are, you little bastard.

In this case, this is due to the kind of data this is, mixing multiple reporting levels and totals. That is, it's not a mistake in the *data*, but rather in the *parsing*.

OK, let's go back to the colspec!

```
nchs <- with_edition(1, read_csv(here("data", "SAS_on_2021-04-13.csv")))
```

```
##
## ── Column specification ────────────────────────────────────────────────────────
## cols(
##   `Data As Of` = col_character(),
##   `Start Date` = col_character(),
##   `End Date` = col_character(),
##   Group = col_character(),
##   Year = col_logical(),
##   Month = col_logical(),
##   State = col_character(),
##   Sex = col_character(),
##   `Age Group` = col_character(),
##   `COVID-19 Deaths` = col_double(),
##   `Total Deaths` = col_double(),
##   `Pneumonia Deaths` = col_double(),
##   `Pneumonia and COVID-19 Deaths` = col_double(),
##   `Influenza Deaths` = col_double(),
##   `Pneumonia, Influenza, or COVID-19 Deaths` = col_double(),
##   Footnote = col_character()
## )
```

We can just copy it from the console output! It's valid code.

We use it with col_types

```
nchs <- with_edition(1, read_csv(here("data", "SAS_on_2021-04-13.csv"),
           col_types = cols(
  `Data As Of` = col_character(),
  `Start Date` = col_character(),
  `End Date` = col_character(),
  Group = col_character(),
  Year = col_logical(),
  Month = col_logical(),
  State = col_character(),
  Sex = col_character(),
  `Age Group` = col_character(),
  `COVID-19 Deaths` = col_double(),
  `Total Deaths` = col_double(),
  `Pneumonia Deaths` = col_double(),
  `Pneumonia and COVID-19 Deaths` = col_double(),
  `Influenza Deaths` = col_double(),
  `Pneumonia, Influenza, or COVID-19 Deaths` = col_double(),
  Footnote = col_character()
)))
```

But we know we need to make some adjustments.

Fixes

```r
# Date format
us_style <-  "%m/%d/%Y"

nchs <- with_edition(1, read_csv(
  here("data", "SAS_on_2021-04-13.csv"),
  col_types = cols(
    `Data As Of` = col_date(format = us_style),
    `Start Date` = col_date(format = us_style),
    `End Date` = col_date(format = us_style),
    Group = col_character(),
    Year = col_character(),
    Month = col_character(),
    State = col_character(),
    Sex = col_character(),
    `Age Group` = col_character(),
    `COVID-19 Deaths` = col_integer(),
    `Total Deaths` = col_integer(),
    `Pneumonia Deaths` = col_integer(),
    `Pneumonia and COVID-19 Deaths` = col_integer(),
    `Influenza Deaths` = col_integer(),
    `Pneumonia, Influenza, or COVID-19 Deaths` = col_integer(),
    Footnote = col_character()
  )) %>%
  janitor::clean_names() %>%
  select(-footnote) %>%
  mutate(age_group = stringr::str_to_sentence(age_group)) %>%
  filter(!stringr::str_detect(state, "Total"))
)
```

Now let's look again

```
dim(nchs)
```

```
## [1] 52326    15
```

```
nchs %>%
  select(year, month, state) %>%
  filter(!is.na(year))
```

```
## # A tibble: 49,572 × 3
##    year  month state
##    <chr> <chr> <chr>
##  1 2020  <NA>  United States
##  2 2020  <NA>  United States
##  3 2020  <NA>  United States
##  4 2020  <NA>  United States
##  5 2020  <NA>  United States
##  6 2020  <NA>  United States
##  7 2020  <NA>  United States
##  8 2020  <NA>  United States
##  9 2020  <NA>  United States
## 10 2020  <NA>  United States
## # … with 49,562 more rows
```

Now let's look again

```
nchs %>%
  distinct(year)
```

```
## # A tibble: 3 × 1
##   year
##   <chr>
## 1 <NA>
## 2 2020
## 3 2021
```

Lessons learned

Lessons learned

I said at the start that it was no fun, but also weirdly satisfying.

Lessons learned

I said at the start that it was no fun, but also weirdly satisfying.

When `read_csv()` warns you of a parsing failure *don't ignore it*.

Lessons learned

I said at the start that it was no fun, but also weirdly satisfying.

When `read_csv()` warns you of a parsing failure *don't ignore it*.

`read_lines()` lets you get the file in a nearly unprocessed form.

Lessons learned

I said at the start that it was no fun, but also weirdly satisfying.

When `read_csv()` warns you of a parsing failure *don't ignore it*.

`read_lines()` lets you get the file in a nearly unprocessed form.

The `colspec` output is your friend.

If we wanted to ...

```r
library(stringr)  # it's back!
```

If we wanted to ...

```r
library(stringr)  # it's back!

nchs
```

```
## # A tibble: 52,326 × 15
##    data_as_of start_date end_date   group    year  month state   sex   ag
##    <date>     <date>     <date>     <chr>    <chr> <chr> <chr>   <chr> <c
##  1 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … Al
##  2 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … Un
##  3 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 0-
##  4 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 1-
##  5 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 5-
##  6 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 15
##  7 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 18
##  8 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 25
##  9 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 30
## 10 2021-04-07 2020-01-01 2021-04-03 By Total <NA>  <NA>  United… All … 35
## # … with 52,316 more rows, and 6 more variables: covid_19_deaths <int>,
## #   total_deaths <int>, pneumonia_deaths <int>,
## #   pneumonia_and_covid_19_deaths <int>, influenza_deaths <int>,
## #   pneumonia_influenza_or_covid_19_deaths <int>
```

If we wanted to ...

```r
library(stringr)  # it's back!

nchs %>%
  select(!(c(data_as_of:end_date, year, month)))
```

```
## # A tibble: 52,326 × 10
##    group    state         sex       age_group covid_19_deaths total_death
##    <chr>    <chr>         <chr>     <chr>               <int>        <int
##  1 By Total United States All Sexes All ages           539723      416116
##  2 By Total United States All Sexes Under 1 …              59        2262
##  3 By Total United States All Sexes 0-17 yea…             251        3962
##  4 By Total United States All Sexes 1-4 years              31         406
##  5 By Total United States All Sexes 5-14 yea…              89         657
##  6 By Total United States All Sexes 15-24 ye…             804        4259
##  7 By Total United States All Sexes 18-29 ye…            1996        7533
##  8 By Total United States All Sexes 25-34 ye…            3543        8819
##  9 By Total United States All Sexes 30-39 ye…            5792       10734
## 10 By Total United States All Sexes 35-44 ye…            9259       12684
## # … with 52,316 more rows, and 3 more variables:
## #   pneumonia_and_covid_19_deaths <int>, influenza_deaths <int>,
## #   pneumonia_influenza_or_covid_19_deaths <int>
```

If we wanted to ...

```r
library(stringr)  # it's back!

nchs %>%
  select(!(c(data_as_of:end_date, year, month))) %>%
  pivot_longer(covid_19_deaths:pneumonia_influenza_or_covid_19_deaths,
               names_to = "outcome",
               values_to = "n")
```

```
## # A tibble: 313,956 × 6
##    group    state         sex       age_group    outcome
##    <chr>    <chr>         <chr>     <chr>        <chr>
##  1 By Total United States All Sexes All ages     covid_19_deaths
##  2 By Total United States All Sexes All ages     total_deaths
##  3 By Total United States All Sexes All ages     pneumonia_deaths
##  4 By Total United States All Sexes All ages     pneumonia_and_covid_19_…
##  5 By Total United States All Sexes All ages     influenza_deaths
##  6 By Total United States All Sexes All ages     pneumonia_influenza_or_…
##  7 By Total United States All Sexes Under 1 year covid_19_deaths
##  8 By Total United States All Sexes Under 1 year total_deaths
##  9 By Total United States All Sexes Under 1 year pneumonia_deaths
## 10 By Total United States All Sexes Under 1 year pneumonia_and_covid_19_…
## # … with 313,946 more rows
```

If we wanted to ...

```r
library(stringr)  # it's back!

nchs %>%
  select(!(c(data_as_of:end_date, year, month))) %>%
  pivot_longer(covid_19_deaths:pneumonia_influenza_or_covid_19_deaths,
               names_to = "outcome",
               values_to = "n") %>%
  mutate(outcome = str_to_sentence(outcome),
         outcome = str_replace_all(outcome, "_", " "),
         outcome = str_replace(outcome, "(C|c)ovid 19", "COVID-19"))
```

```
## # A tibble: 313,956 × 6
##    group    state         sex       age_group    outcome
##    <chr>    <chr>         <chr>     <chr>        <chr>
##  1 By Total United States All Sexes All ages     COVID-19 deaths
##  2 By Total United States All Sexes All ages     Total deaths
##  3 By Total United States All Sexes All ages     Pneumonia deaths
##  4 By Total United States All Sexes All ages     Pneumonia and COVID-19 …
##  5 By Total United States All Sexes All ages     Influenza deaths
##  6 By Total United States All Sexes All ages     Pneumonia influenza or …
##  7 By Total United States All Sexes Under 1 year COVID-19 deaths
##  8 By Total United States All Sexes Under 1 year Total deaths
##  9 By Total United States All Sexes Under 1 year Pneumonia deaths
## 10 By Total United States All Sexes Under 1 year Pneumonia and COVID-19 …
## # … with 313,946 more rows
```

If we wanted to ...

```r
library(stringr)  # it's back!

nchs %>%
  select(!(c(data_as_of:end_date, year, month))) %>%
  pivot_longer(covid_19_deaths:pneumonia_influenza_or_covid_19_deaths,
               names_to = "outcome",
               values_to = "n") %>%
  mutate(outcome = str_to_sentence(outcome),
         outcome = str_replace_all(outcome, "_", " "),
         outcome = str_replace(outcome, "(C|c)ovid 19", "COVID-19"))
```

```
## # A tibble: 313,956 × 6
##    group    state         sex       age_group    outcome
##    <chr>    <chr>         <chr>     <chr>        <chr>
##  1 By Total United States All Sexes All ages     COVID-19 deaths
##  2 By Total United States All Sexes All ages     Total deaths
##  3 By Total United States All Sexes All ages     Pneumonia deaths
##  4 By Total United States All Sexes All ages     Pneumonia and COVID-19 …
##  5 By Total United States All Sexes All ages     Influenza deaths
##  6 By Total United States All Sexes All ages     Pneumonia influenza or …
##  7 By Total United States All Sexes Under 1 year COVID-19 deaths
##  8 By Total United States All Sexes Under 1 year Total deaths
##  9 By Total United States All Sexes Under 1 year Pneumonia deaths
## 10 By Total United States All Sexes Under 1 year Pneumonia and COVID-19 …
## # … with 313,946 more rows
```

Put this in `nchs_fmt`

... we could make a table or graph

```
nchs_fmt %>%
  select(state, age_group, outcome, n)
```

```
## # A tibble: 313,956 × 4
##    state         age_group   outcome                                    n
##    <chr>         <chr>       <chr>                                  <int>
##  1 United States All ages    COVID-19 deaths                       539723
##  2 United States All ages    Total deaths                         4161167
##  3 United States All ages    Pneumonia deaths                      466437
##  4 United States All ages    Pneumonia and COVID-19 deaths         263147
##  5 United States All ages    Influenza deaths                        9037
##  6 United States All ages    Pneumonia influenza or COVID-19 deaths  750804
##  7 United States Under 1 year COVID-19 deaths                          59
##  8 United States Under 1 year Total deaths                          22626
##  9 United States Under 1 year Pneumonia deaths                        246
## 10 United States Under 1 year Pneumonia and COVID-19 deaths            10
## # … with 313,946 more rows
```

Cleaned up (but not tidy)

```
nchs_fmt %>%
  distinct(group)
```

```
## # A tibble: 3 × 1
##   group
##   <chr>
## 1 By Total
## 2 By Year
## 3 By Month
```

Cleaned up (but not tidy)

```
nchs_fmt %>%
  distinct(group)
```

```
## # A tibble: 3 × 1
##   group
##   <chr>
## 1 By Total
## 2 By Year
## 3 By Month
```

```
nchs_fmt %>%
  distinct(age_group)
```

```
## # A tibble: 17 × 1
##    age_group
##    <chr>
##  1 All ages
##  2 Under 1 year
##  3 0-17 years
##  4 1-4 years
##  5 5-14 years
##  6 15-24 years
##  7 18-29 years
##  8 25-34 years
##  9 30-39 years
## 10 35-44 years
## 11 40-49 years
## 12 45-54 years
## 13 50-64 years
## 14 55-64 years
## 15 65-74 years
## 16 75-84 years
## 17 85 years and over
```

Make our plot

```r
p_out <- nchs_fmt %>%
  filter(group %in% "By Total",
         sex %in% "All Sexes",
         state %in% "United States",
         age_group %in% c("0-17 years",
                          "18-29 years",
                          "30-39 years",
                          "40-49 years",
                          "50-64 years",
                          "65-74 years",
                          "85 years and over"),
         outcome %in% "COVID-19 deaths") %>%
  mutate(age_group = str_replace(age_group, "years", "yrs"),
         age_group = str_replace(age_group, " and over", ""),
         age_group = str_replace(age_group, "85", "85+")) %>%
  ggplot(mapping = aes(x = n, y = age_group)) +
  geom_col() + scale_x_continuous(labels = scales::comma) +
  labs(x = "Deaths", y = NULL, title = "U.S. COVID-19 mortality totals by age group")
```

Result

```
print(p_out)
```

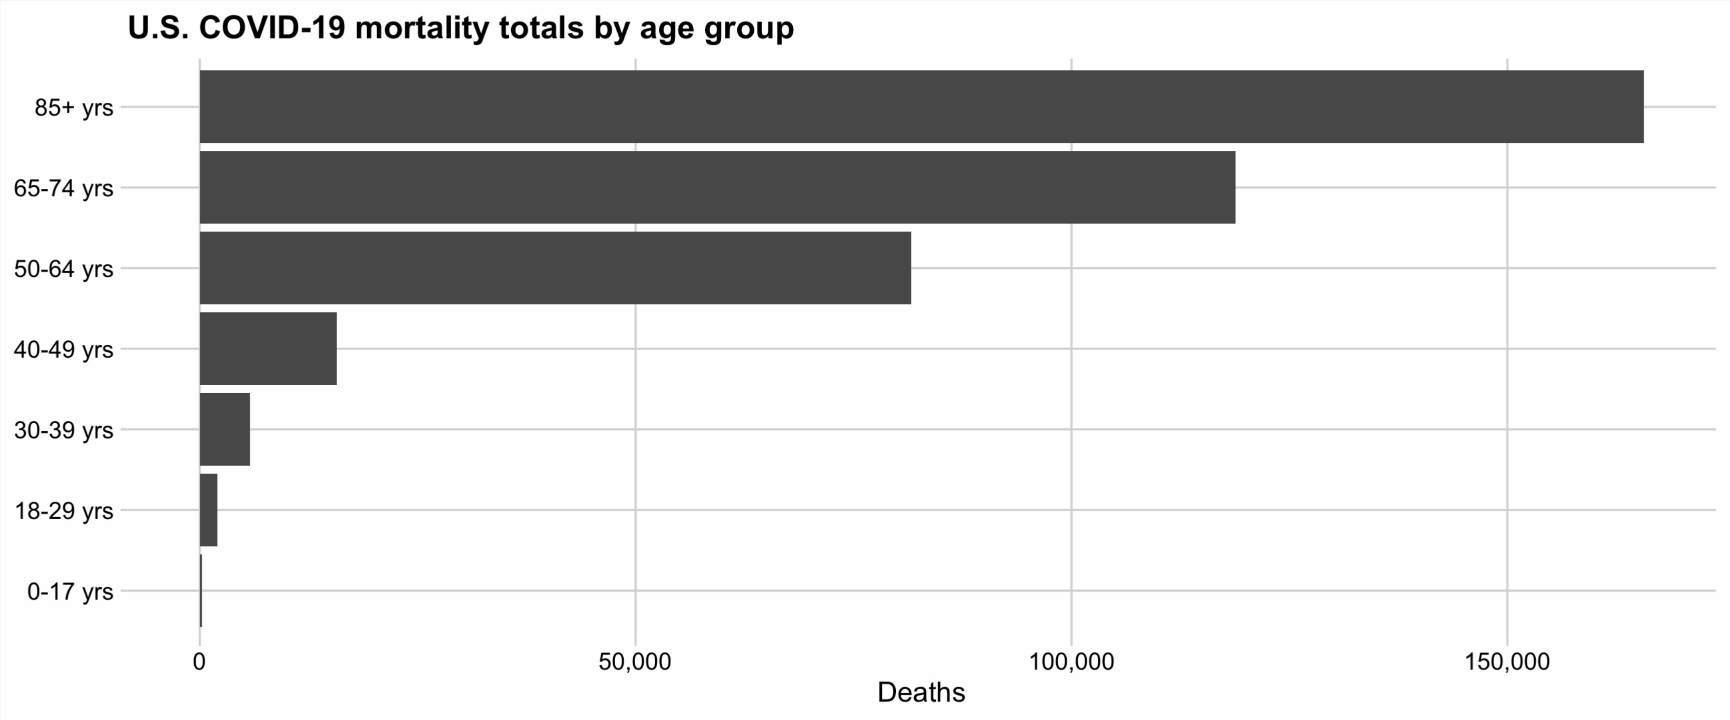

Every dataset is different

Dropping missing values: a quick demo

```r
df <- tribble(
  ~a, ~b, ~c,
  1, NA, 2,
  NA, NA, NA,
  2, 2, 2
)

df
```

```
## # A tibble: 3 × 3
##       a     b     c
##   <dbl> <dbl> <dbl>
## 1     1    NA     2
## 2    NA    NA    NA
## 3     2     2     2
```

Dropping **missing values**: a quick demo

```
# 1 Write it yourself
df %>%
  filter(complete.cases(.))
```

```
## # A tibble: 1 × 3
##       a     b     c
##   <dbl> <dbl> <dbl>
## 1     2     2     2
```

```
# 2 Convenience function
df %>%
  drop_na()
```

```
## # A tibble: 1 × 3
##       a     b     c
##   <dbl> <dbl> <dbl>
## 1     2     2     2
```

Both drop all rows with *any* missing cases.

Dropping missing values: a quick demo

What if we only want to drop all rows with *all* missing cases?

```
# 3
df %>%
  filter(!(is.na(.))) # Pronoun
```

```
## # A tibble: 2 × 3
##       a     b     c
##   <dbl> <dbl> <dbl>
## 1     1    NA     2
## 2     2     2     2
```

Read `.` as "the thing we're looking at" or "the thing we're computing on right now".

Cleaning a table

With that in mind … Some marketing data

SEGMENT	DESCRIPTION	R	F	M
Champions	Bought recently, buy often and spend the most	4– 5	4– 5	4– 5
Loyal Customers	Spend good money. Responsive to promotions	2– 5	3– 5	3– 5
Potential Loyalist	Recent customers, spent good amount, bought more than once	3– 5	1– 3	1– 3
New Customers	Bought more recently, but not often	4– 5	<= 1	<= 1
Promising	Recent shoppers, but haven't spent much	3– 4	<= 1	<= 1
Need Attention	Above average recency, frequency & monetary values	2– 3	2– 3	2– 3
About To Sleep	Below average recency, frequency & monetary values	2– 3	<= 2	<= 2
At Risk	Spent big money, purchased often but long time ago	<= 2	2– 5	2– 5
Can't Lose Them	Made big purchases and often, but long time ago	<= 1	4– 5	4– 5
Hibernating	Low spenders, low frequency, purchased long time ago	1– 2	1– 2	1– 2
Lost	Lowest recency, frequency & monetary scores	<= 2	<= 2	<= 2

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv"))
```

```
## # A tibble: 23 × 5
##    SEGMENT            DESCRIPTION                                R     F     M
##    <chr>              <chr>                                      <chr> <chr> <chr>
##  1 <NA>               <NA>                                       <NA>  <NA>  <NA>
##  2 Champions          Bought recently, buy often and spend th… 4- 5  4- 5  4- 5
##  3 <NA>               <NA>                                       <NA>  <NA>  <NA>
##  4 Loyal Customers    Spend good money. Responsive to promoti… 2- 5  3- 5  3- 5
##  5 <NA>               <NA>                                       <NA>  <NA>  <NA>
##  6 Potential Loyalist Recent customers, spent good amount, bo… 3- 5  1- 3  1- 3
##  7 <NA>               <NA>                                       <NA>  <NA>  <NA>
##  8 New Customers      Bought more recently, but not often        4- 5  <= 1  <= 1
##  9 <NA>               <NA>                                       <NA>  <NA>  <NA>
## 10 Promising          Recent shoppers, but haven't spent much    3- 4  <= 1  <= 1
## # … with 13 more rows
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names()
```

```
## # A tibble: 23 × 5
##    segment             description                               r     f     m    
##    <chr>               <chr>                                     <chr> <chr> <chr>
##  1 <NA>                <NA>                                      <NA>  <NA>  <NA> 
##  2 Champions           Bought recently, buy often and spend th… 4- 5  4- 5  4- 5 
##  3 <NA>                <NA>                                      <NA>  <NA>  <NA> 
##  4 Loyal Customers     Spend good money. Responsive to promoti… 2- 5  3- 5  3- 5 
##  5 <NA>                <NA>                                      <NA>  <NA>  <NA> 
##  6 Potential Loyalist Recent customers, spent good amount, bo… 3- 5  1- 3  1- 3 
##  7 <NA>                <NA>                                      <NA>  <NA>  <NA> 
##  8 New Customers       Bought more recently, but not often       4- 5  <= 1  <= 1 
##  9 <NA>                <NA>                                      <NA>  <NA>  <NA> 
## 10 Promising           Recent shoppers, but haven't spent much   3- 4  <= 1  <= 1 
## # … with 13 more rows
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.)))   #<<
```

```
## # A tibble: 11 × 5
##    segment            description                                r     f     m    
##    <chr>              <chr>                                      <chr> <chr> <chr>
##  1 Champions          Bought recently, buy often and spend th… 4- 5  4- 5  4- 5 
##  2 Loyal Customers    Spend good money. Responsive to promoti… 2- 5  3- 5  3- 5 
##  3 Potential Loyalist Recent customers, spent good amount, bo… 3- 5  1- 3  1- 3 
##  4 New Customers      Bought more recently, but not often      4- 5  <= 1  <= 1 
##  5 Promising          Recent shoppers, but haven't spent much  3- 4  <= 1  <= 1 
##  6 Need Attention     Above average recency, frequency & mone… 2- 3  2- 3  2- 3 
##  7 About To Sleep     Below average recency, frequency & mone… 2- 3  <= 2  <= 2 
##  8 At Risk            Spent big money, purchased often but lo… <= 2  2- 5  2- 5 
##  9 Can't Lose Them    Made big purchases and often, but long … <= 1  4- 5  4- 5 
## 10 Hibernating        Low spenders, low frequency, purchased … 1- 2  1- 2  1- 2 
## 11 Lost               Lowest recency, frequency & monetary sc… <= 2  <= 2  <= 2
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m)
```

```
## # A tibble: 33 × 4
##    segment            description                               name  value
##    <chr>              <chr>                                     <chr> <chr>
##  1 Champions          Bought recently, buy often and spend the most r     4- 5
##  2 Champions          Bought recently, buy often and spend the most f     4- 5
##  3 Champions          Bought recently, buy often and spend the most m     4- 5
##  4 Loyal Customers    Spend good money. Responsive to promotions    r     2- 5
##  5 Loyal Customers    Spend good money. Responsive to promotions    f     3- 5
##  6 Loyal Customers    Spend good money. Responsive to promotions    m     3- 5
##  7 Potential Loyalist Recent customers, spent good amount, bought m… r     3- 5
##  8 Potential Loyalist Recent customers, spent good amount, bought m… f     1- 3
##  9 Potential Loyalist Recent customers, spent good amount, bought m… m     1- 3
## 10 New Customers      Bought more recently, but not often           r     4- 5
## # … with 23 more rows
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m) %>%
  separate(col = value, into = c("lo", "hi"),
           remove = FALSE, convert = TRUE,
           fill = "left")
```

```
## # A tibble: 33 × 6
##    segment            description                          name  value    lo    hi
##    <chr>              <chr>                                <chr> <chr> <int> <int>
##  1 Champions          Bought recently, buy often and sp… r     4- 5      4     5
##  2 Champions          Bought recently, buy often and sp… f     4- 5      4     5
##  3 Champions          Bought recently, buy often and sp… m     4- 5      4     5
##  4 Loyal Customers    Spend good money. Responsive to p… r     2- 5      2     5
##  5 Loyal Customers    Spend good money. Responsive to p… f     3- 5      3     5
##  6 Loyal Customers    Spend good money. Responsive to p… m     3- 5      3     5
##  7 Potential Loyalist Recent customers, spent good amou… r     3- 5      3     5
##  8 Potential Loyalist Recent customers, spent good amou… f     1- 3      1     3
##  9 Potential Loyalist Recent customers, spent good amou… m     1- 3      1     3
## 10 New Customers      Bought more recently, but not oft… r     4- 5      4     5
## # … with 23 more rows
```

Cleaning a table

```
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m) %>%
  separate(col = value, into = c("lo", "hi"),
           remove = FALSE, convert = TRUE,
           fill = "left") %>%
  select(-value)
```

```
## # A tibble: 33 × 5
##    segment            description                            name     lo    hi
##    <chr>              <chr>                                  <chr> <int> <int>
##  1 Champions          Bought recently, buy often and spend th… r         4     5
##  2 Champions          Bought recently, buy often and spend th… f         4     5
##  3 Champions          Bought recently, buy often and spend th… m         4     5
##  4 Loyal Customers    Spend good money. Responsive to promoti… r         2     5
##  5 Loyal Customers    Spend good money. Responsive to promoti… f         3     5
##  6 Loyal Customers    Spend good money. Responsive to promoti… m         3     5
##  7 Potential Loyalist Recent customers, spent good amount, bo… r         3     5
##  8 Potential Loyalist Recent customers, spent good amount, bo… f         1     3
##  9 Potential Loyalist Recent customers, spent good amount, bo… m         1     3
## 10 New Customers      Bought more recently, but not often      r         4     5
## # … with 23 more rows
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m) %>%
  separate(col = value, into = c("lo", "hi"),
           remove = FALSE, convert = TRUE,
           fill = "left") %>%
  select(-value) %>%
  pivot_wider(names_from = name,
              values_from = lo:hi)
```

```
## # A tibble: 11 × 8
##    segment             description                lo_r  lo_f  lo_m  hi_r  hi_f  hi_m
##    <chr>               <chr>                     <int> <int> <int> <int> <int> <int>
##  1 Champions           Bought recently, buy …        4     4     4     5     5     5
##  2 Loyal Customers     Spend good money. Res…        2     3     3     5     5     5
##  3 Potential Loyalist  Recent customers, spe…        3     1     1     5     3     3
##  4 New Customers       Bought more recently,…        4    NA    NA     5     1     1
##  5 Promising           Recent shoppers, but …        3    NA    NA     4     1     1
##  6 Need Attention      Above average recency…        2     2     2     3     3     3
##  7 About To Sleep      Below average recency…        2    NA    NA     3     2     2
##  8 At Risk             Spent big money, purc…       NA     2     2     2     5     5
##  9 Can’t Lose Them     Made big purchases an…       NA     4     4     1     5     5
## 10 Hibernating         Low spenders, low fre…        1     1     1     2     2     2
## 11 Lost                Lowest recency, frequ…       NA    NA    NA     2     2     2
```

Cleaning a table

```r
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m) %>%
  separate(col = value, into = c("lo", "hi"),
           remove = FALSE, convert = TRUE,
           fill = "left") %>%
  select(-value) %>%
  pivot_wider(names_from = name,
              values_from = lo:hi) %>%
  mutate(across(where(is.integer), replace_na, 0))
```

```
## # A tibble: 11 × 8
##    segment            description                  lo_r  lo_f  lo_m  hi_r  hi_f  hi_m
##    <chr>              <chr>                       <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
##  1 Champions          Bought recently, buy …          4     4     4     5     5     5
##  2 Loyal Customers    Spend good money. Res…          2     3     3     5     5     5
##  3 Potential Loyalist Recent customers, spe…          3     1     1     5     3     3
##  4 New Customers      Bought more recently,…          4     0     0     5     1     1
##  5 Promising          Recent shoppers, but …          3     0     0     4     1     1
##  6 Need Attention     Above average recency…          2     2     2     3     3     3
##  7 About To Sleep     Below average recency…          2     0     0     3     2     2
##  8 At Risk            Spent big money, purc…          0     2     2     2     5     5
##  9 Can't Lose Them    Made big purchases an…          0     4     4     1     5     5
## 10 Hibernating        Low spenders, low fre…          1     1     1     2     2     2
## 11 Lost               Lowest recency, frequ…          0     0     0     2     2     2
```

Cleaning a table

```
read_csv(here("data", "rfm_table.csv")) %>%
  janitor::clean_names() %>%
  filter(!(is.na(.))) %>%
  pivot_longer(cols = r:m) %>%
  separate(col = value, into = c("lo", "hi"),
           remove = FALSE, convert = TRUE,
           fill = "left") %>%
  select(-value) %>%
  pivot_wider(names_from = name,
              values_from = lo:hi) %>%
  mutate(across(where(is.integer), replace_na, 0)) %>%
  select(segment,
         lo_r, hi_r,
         lo_f, hi_f,
         lo_m, hi_m,
         description)
```

```
## # A tibble: 11 × 8
##    segment             lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>              <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions              4     5     4     5     4     5 Bought recently, buy …
##  2 Loyal Customers        2     5     3     5     3     5 Spend good money. Res…
##  3 Potential Loyalist     3     5     1     3     1     3 Recent customers, spe…
##  4 New Customers          4     5     0     1     0     1 Bought more recently,…
##  5 Promising              3     4     0     1     0     1 Recent shoppers, but …
##  6 Need Attention         2     3     2     3     2     3 Above average recency…
##  7 About To Sleep         2     3     0     2     0     2 Below average recency…
##  8 At Risk                0     2     2     5     2     5 Spent big money, purc…
##  9 Can't Lose Them        0     1     4     5     4     5 Made big purchases an…
## 10 Hibernating            1     2     1     2     1     2 Low spenders, low fre…
## 11 Lost                   0     2     0     2     0     2 Lowest recency, frequ…
```

Maybe a candidate for rowwise?

```r
rfm_table
```

```
## # A tibble: 11 × 8
##    segment            lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>             <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions             4     5     4     5     4     5 Bought recently, buy …
##  2 Loyal Customers       2     5     3     5     3     5 Spend good money. Res…
##  3 Potential Loyalist    3     5     1     3     1     3 Recent customers, spe…
##  4 New Customers         4     5     0     1     0     1 Bought more recently,…
##  5 Promising             3     4     0     1     0     1 Recent shoppers, but …
##  6 Need Attention        2     3     2     3     2     3 Above average recency…
##  7 About To Sleep        2     3     0     2     0     2 Below average recency…
##  8 At Risk               0     2     2     5     2     5 Spent big money, purc…
##  9 Can't Lose Them       0     1     4     5     4     5 Made big purchases an…
## 10 Hibernating           1     2     1     2     1     2 Low spenders, low fre…
## 11 Lost                  0     2     0     2     0     2 Lowest recency, frequ…
```

Maybe a candidate for rowwise?

This does what we expect:

```
rfm_table %>%
  mutate(sum_lo = lo_r + lo_f + lo_m,
         sum_hi = hi_r + hi_f + hi_m) %>%
  select(segment, sum_lo, sum_hi, everything())
```

```
## # A tibble: 11 × 10
##    segment            sum_lo sum_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>               <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions              12     15     4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers         8     15     2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist      5     11     3     5     1     3     1     3 Recent cus…
##  4 New Customers           4      7     4     5     0     1     0     1 Bought mor…
##  5 Promising               3      6     3     4     0     1     0     1 Recent sho…
##  6 Need Attention          6      9     2     3     2     3     2     3 Above aver…
##  7 About To Sleep          2      7     2     3     0     2     0     2 Below aver…
##  8 At Risk                 4     12     0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them         8     11     0     1     4     5     4     5 Made big p…
## 10 Hibernating             3      6     1     2     1     2     1     2 Low spende…
## 11 Lost                    0      6     0     2     0     2     0     2 Lowest rec…
```

This adds each column, elementwise.

Maybe a candidate for rowwise?

But this does not:

```
rfm_table %>%
  mutate(sum_lo = sum(lo_r, lo_f, lo_m),
         sum_hi = sum(hi_r, hi_f, hi_m)) %>%
  select(segment, sum_lo, sum_hi, everything())
```

```
## # A tibble: 11 × 10
##    segment            sum_lo sum_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>               <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions              55    105     4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers        55    105     2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist     55    105     3     5     1     3     1     3 Recent cus…
##  4 New Customers          55    105     4     5     0     1     0     1 Bought mor…
##  5 Promising              55    105     3     4     0     1     0     1 Recent sho…
##  6 Need Attention         55    105     2     3     2     3     2     3 Above aver…
##  7 About To Sleep         55    105     2     3     0     2     0     2 Below aver…
##  8 At Risk                55    105     0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them        55    105     0     1     4     5     4     5 Made big p…
## 10 Hibernating            55    105     1     2     1     2     1     2 Low spende…
## 11 Lost                   55    105     0     2     0     2     0     2 Lowest rec…
```

Sum is taking all the columns, adding them up (into a single mumber), and putting that result in each row.

Maybe a candidate for rowwise?

Similarly, this will not give the answer we probably expect:

```
rfm_table %>%
  mutate(mean_lo = mean(c(lo_r, lo_f, lo_m)),
         mean_hi = mean(c(hi_r, hi_f, hi_m))) %>%
  select(segment, mean_lo, mean_hi, everything())
```

```
## # A tibble: 11 × 10
##    segment            mean_lo mean_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>                <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions             1.67    3.18     4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers       1.67    3.18     2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist    1.67    3.18     3     5     1     3     1     3 Recent cus…
##  4 New Customers         1.67    3.18     4     5     0     1     0     1 Bought mor…
##  5 Promising             1.67    3.18     3     4     0     1     0     1 Recent sho…
##  6 Need Attention        1.67    3.18     2     3     2     3     2     3 Above aver…
##  7 About To Sleep        1.67    3.18     2     3     0     2     0     2 Below aver…
##  8 At Risk               1.67    3.18     0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them       1.67    3.18     0     1     4     5     4     5 Made big p…
## 10 Hibernating           1.67    3.18     1     2     1     2     1     2 Low spende…
## 11 Lost                  1.67    3.18     0     2     0     2     0     2 Lowest rec…
```

Maybe a candidate for rowwise?

But this will:

```r
rfm_table %>%
  rowwise() %>%
  mutate(mean_lo = mean(c(lo_r, lo_f, lo_m)),
         mean_hi = mean(c(hi_r, hi_f, hi_m))) %>%
  select(segment, mean_lo, mean_hi, everything())
```

```
## # A tibble: 11 × 10
## # Rowwise:
##    segment            mean_lo mean_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>                <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions            4       5         4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers      2.67    5         2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist   1.67    3.67      3     5     1     3     1     3 Recent cus…
##  4 New Customers        1.33    2.33      4     5     0     1     0     1 Bought mor…
##  5 Promising            1       2         3     4     0     1     0     1 Recent sho…
##  6 Need Attention       2       3         2     3     2     3     2     3 Above aver…
##  7 About To Sleep       0.667   2.33      2     3     0     2     0     2 Below aver…
##  8 At Risk              1.33    4         0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them      2.67    3.67      0     1     4     5     4     5 Made big p…
## 10 Hibernating          1       2         1     2     1     2     1     2 Low spende…
## 11 Lost                 0       2         0     2     0     2     0     2 Lowest rec…
```

Rowwise operations aren't very efficient

In general, you'll want to see if some vectorized ("operating on columns, but elementwise") function exists, as it'll be faster.

And most of the time, R and the tidyverse "wants" you to work in vectorized, columnar terms ... hence your first move will often be to pivot the data into long format.

So, `rowwise()` is not likely to see a whole lot of further development.

You probably want `group_by()` instead

```r
rfm_table %>%
  group_by(segment) %>%
  mutate(mean_lo = mean(c(lo_r, lo_f, lo_m)),
         mean_hi = mean(c(hi_r, hi_f, hi_m))) %>%
  select(segment, mean_lo, mean_hi, everything())
```

```
## # A tibble: 11 × 10
## # Groups:   segment [11]
##    segment            mean_lo mean_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>                <dbl>   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions            4       5         4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers      2.67    5         2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist   1.67    3.67      3     5     1     3     1     3 Recent cus…
##  4 New Customers        1.33    2.33      4     5     0     1     0     1 Bought mor…
##  5 Promising            1       2         3     4     0     1     0     1 Recent sho…
##  6 Need Attention       2       3         2     3     2     3     2     3 Above aver…
##  7 About To Sleep       0.667   2.33      2     3     0     2     0     2 Below aver…
##  8 At Risk              1.33    4         0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them      2.67    3.67      0     1     4     5     4     5 Made big p…
## 10 Hibernating          1       2         1     2     1     2     1     2 Low spende…
## 11 Lost                 0       2         0     2     0     2     0     2 Lowest rec…
```

You probably want group_by() instead

```
rfm_table %>%
  group_by(segment) %>%
  mutate(sum_lo = sum(lo_r, lo_f, lo_m),
         sum_hi = sum(hi_r, hi_f, hi_m)) %>%
  select(segment, sum_lo, sum_hi, everything())
```

```
## # A tibble: 11 × 10
## # Groups:   segment [11]
##    segment            sum_lo sum_hi  lo_r  hi_r  lo_f  hi_f  lo_m  hi_m description
##    <chr>               <dbl>  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <chr>
##  1 Champions              12     15     4     5     4     5     4     5 Bought rec…
##  2 Loyal Customers         8     15     2     5     3     5     3     5 Spend good…
##  3 Potential Loyalist      5     11     3     5     1     3     1     3 Recent cus…
##  4 New Customers           4      7     4     5     0     1     0     1 Bought mor…
##  5 Promising               3      6     3     4     0     1     0     1 Recent sho…
##  6 Need Attention          6      9     2     3     2     3     2     3 Above aver…
##  7 About To Sleep          2      7     2     3     0     2     0     2 Below aver…
##  8 At Risk                 4     12     0     2     2     5     2     5 Spent big …
##  9 Can't Lose Them         8     11     0     1     4     5     4     5 Made big p…
## 10 Hibernating             3      6     1     2     1     2     1     2 Low spende…
## 11 Lost                    0      6     0     2     0     2     0     2 Lowest rec…
```

What about Stata?

Using haven

Haven is the Tidyverse's package for reading and managing files from Stata, SPSS, and SAS. You should prefer it to the older Base R package `foreign`, which has similar functionality.

We're going to import a General Social Survey dataset that's in Stata's `.dta` format.

```
library(haven)

# This will take a moment
gss_panel <- read_stata(here("data", "gss_panel_long.dta"))
```

We'll do some of the common recoding and reorganizing tasks that accompany this.

The GSS panel

The data:

```
gss_panel
```

```
## # A tibble: 14,610 × 2,757
##    firstyear   firstid  year    id   vpsu vstrat adults  ballot dateintv  famgen
##        <dbl> <dbl+lbl> <dbl> <dbl> <dbl+> <dbl+> <dbl+> <dbl+l> <dbl+lb> <dbl+l>
##  1      2006         9  2006     9      2   1957      1 3 [BAL…      709 1 [1 G…
##  2      2006         9  2008  3001     NA     NA      2 3 [BAL…      503 1 [1 G…
##  3      2006         9  2010  6001  NA(i)     NA      2 3 [BAL…      508 1 [1 G…
##  4      2006        10  2010  6002  NA(i)     NA      1 1 [BAL…      408 1 [1 G…
##  5      2006        10  2006    10      2   1957      2 1 [BAL…      630 2 [2 G…
##  6      2006        10  2008  3002     NA     NA      2 1 [BAL…      426 2 [2 G…
##  7      2006        11  2008  3003     NA     NA      2 3 [BAL…      718 4 [2 G…
##  8      2006        11  2010  6003  NA(i)     NA  NA(n) 3 [BAL…      518 2 [2 G…
##  9      2006        11  2006    11      2   1957      2 3 [BAL…      630 4 [2 G…
## 10      2006        12  2010  6004  NA(i)     NA      4 1 [BAL…      324 2 [2 G…
## # … with 14,600 more rows, and 2,747 more variables: form <dbl+lbl>,
## #   formwt <dbl>, gender1 <dbl+lbl>, hompop <dbl+lbl>, intage <dbl+lbl>,
## #   intid <dbl+lbl>, intyrs <dbl+lbl>, mode <dbl+lbl>, oversamp <dbl>,
## #   phase <dbl+lbl>, race <dbl+lbl>, reg16 <dbl+lbl>, region <dbl+lbl>,
## #   relate1 <dbl+lbl>, relhh1 <dbl+lbl>, relhhd1 <dbl+lbl>, respnum <dbl+lbl>,
## #   rvisitor <dbl+lbl>, sampcode <dbl+lbl>, sample <dbl+lbl>, sex <dbl+lbl>,
## #   size <dbl+lbl>, spaneng <dbl+lbl>, srcbelt <dbl+lbl>, version <dbl>, …
```

The GSS panel

Many variables.

Stata's missing value types are preserved

Data types are things like `dbl+lbl` indicating that Stata's numeric values and variable labels have been preserved.

The GSS panel

You can see the labeling system at work:

```
gss_panel %>%
  select(degree) %>%
  group_by(degree) %>%
  tally()
```

```
## # A tibble: 6 × 2
##                  degree     n
##               <dbl+lbl> <int>
## 1     0 [LT HIGH SCHOOL]  1850
## 2     1 [HIGH SCHOOL]     7274
## 3     2 [JUNIOR COLLEGE]  1161
## 4     3 [bachelor]        2767
## 5     4 [graduate]        1556
## 6 NA(d)                      2
```

The GSS panel

Values get pivoted, not labels, though.

```
gss_panel %>%
  select(sex, degree) %>%
  group_by(sex, degree) %>%
  tally() %>%
  pivot_wider(names_from = sex, values_from = n)
```

```
## # A tibble: 6 × 3
##                    degree   `1`   `2`
##                 <dbl+lbl> <int> <int>
## 1     0 [LT HIGH SCHOOL]   814  1036
## 2     1 [HIGH SCHOOL]     3131  4143
## 3     2 [JUNIOR COLLEGE]   440   721
## 4     3 [bachelor]        1293  1474
## 5     4 [graduate]         696   860
## 6 NA(d)                     NA     2
```

The GSS panel

Option 1: Just drop all the labels.

```
gss_panel %>%
  zap_missing() %>%
  zap_labels()
```

```
## # A tibble: 14,610 × 2,757
##    firstyear firstid  year    id  vpsu vstrat adults ballot dateintv famgen
##        <dbl>   <dbl> <dbl> <dbl> <dbl>  <dbl>  <dbl>  <dbl>    <dbl>  <dbl>
##  1      2006       9  2006     9     2   1957      1      3      709      1
##  2      2006       9  2008  3001    NA     NA      2      3      503      1
##  3      2006       9  2010  6001    NA     NA      2      3      508      1
##  4      2006      10  2010  6002    NA     NA      1      1      408      1
##  5      2006      10  2006    10     2   1957      2      1      630      2
##  6      2006      10  2008  3002    NA     NA      2      1      426      2
##  7      2006      11  2008  3003    NA     NA      2      3      718      4
##  8      2006      11  2010  6003    NA     NA     NA      3      518      2
##  9      2006      11  2006    11     2   1957      2      3      630      4
## 10      2006      12  2010  6004    NA     NA      4      1      324      2
## # … with 14,600 more rows, and 2,747 more variables: form <dbl>, formwt <dbl>,
## #   gender1 <dbl>, hompop <dbl>, intage <dbl>, intid <dbl>, intyrs <dbl>,
## #   mode <dbl>, oversamp <dbl>, phase <dbl>, race <dbl>, reg16 <dbl>,
## #   region <dbl>, relate1 <dbl>, relhh1 <dbl>, relhhd1 <dbl>, respnum <dbl>,
## #   rvisitor <dbl>, sampcode <dbl>, sample <dbl>, sex <dbl>, size <dbl>,
## #   spaneng <dbl>, srcbelt <dbl>, version <dbl>, visitors <dbl>, wtss <dbl>,
## #   wtssall <dbl>, wtssnr <dbl>, xnorcsiz <dbl>, hispanic <dbl>, …
```

The GSS panel

Option 2: Convert the labels

Let's focus on a few measures of interest, and do some recoding.

```r
## Categorical vars
cat_vars <- c("race", "sex", "degree", "relig", "income", "polviews", "fefam")

## Integer vars
int_vars <- c("year", "id", "ballot", "age", "tvhours")


## Survey design
wt_vars <- c("vpsu",
             "vstrat",
             "oversamp",
             "formwt",              # weight to deal with experimental randomization
             "wtssall",             # weight variable
             "sampcode",            # sampling error code
             "sample")              # sampling frame and method


my_gss_vars <- c(int_vars, cat_vars, wt_vars)
```

Now we're ready to go ...

Cut down the dataset

```
gss_sub <- gss_panel %>%
  select(all_of(my_gss_vars))

gss_sub
```

```
## # A tibble: 14,610 × 19
##     year    id       ballot       age     tvhours    race     sex  degree   relig
##    <dbl> <dbl>    <dbl+lbl> <dbl+lbl>   <dbl+lbl> <dbl+l> <dbl+l> <dbl+l> <dbl+l>
##  1  2006     9 3 [BALLOT C]        23 NA(a) [iap] 2 [bla… 2 [fem… 3 [bac… 4 [non…
##  2  2008  3001 3 [BALLOT C]        25 NA(i)       3 [oth… 2 [fem… 3 [bac… 4 [non…
##  3  2010  6001 3 [BALLOT C]        27 NA(i)       2 [bla… 2 [fem… 3 [bac… 4 [non…
##  4  2010  6002 1 [BALLOT A]        36     3       1 [whi… 2 [fem… 4 [gra… 4 [non…
##  5  2006    10 1 [BALLOT A]        32     3       3 [oth… 2 [fem… 4 [gra… 4 [non…
##  6  2008  3002 1 [BALLOT A]        34     3       3 [oth… 2 [fem… 4 [gra… 4 [non…
##  7  2008  3003 3 [BALLOT C]        83 NA(i)       2 [bla… 2 [fem… 0 [LT … 1 [pro…
##  8  2010  6003 3 [BALLOT C]        85 NA(i)       2 [bla… 2 [fem… 0 [LT … 1 [pro…
##  9  2006    11 3 [BALLOT C]        81 NA(a) [iap] 2 [bla… 2 [fem… 0 [LT … 1 [pro…
## 10  2010  6004 1 [BALLOT A]        51    10       3 [oth… 1 [mal… 1 [HIG… 2 [cat…
## # … with 14,600 more rows, and 10 more variables: income <dbl+lbl>,
## #   polviews <dbl+lbl>, fefam <dbl+lbl>, vpsu <dbl+lbl>, vstrat <dbl+lbl>,
## #   oversamp <dbl>, formwt <dbl>, wtssall <dbl+lbl>, sampcode <dbl+lbl>,
## #   sample <dbl+lbl>
```

The GSS Panel: Recoding

```
gss_sub %>%
  mutate(across(everything(), zap_missing)) %>%
  mutate(across(all_of(wt_vars), as.numeric)) %>%
  mutate(across(all_of(int_vars), as.integer)) %>%
  mutate(across(all_of(cat_vars), as_factor)) %>%
  mutate(across(all_of(cat_vars), fct_relabel, tolower)) %>%
  mutate(across(all_of(cat_vars), fct_relabel, tools::toTitleCase)) %>%
  mutate(income = stringr::str_replace(income, " - ", "-"))
```

```
## # A tibble: 14,610 × 19
##     year    id ballot   age tvhours race  sex    degree  relig  income  polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>   <fct>  <chr>   <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None   $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None   $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None   $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None   $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None   <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None   $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 8 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>
```

How we'd actually write this

```
gss_sub <- gss_sub %>%
  mutate(across(everything(), zap_missing),
         across(all_of(wt_vars), as.numeric),
         across(all_of(int_vars), as.integer),
         across(all_of(cat_vars), as_factor),
         across(all_of(cat_vars), fct_relabel, tolower),
         across(all_of(cat_vars), fct_relabel, tools::toTitleCase),
         income = stringr::str_replace(income, " - ", "-"))
```

The GSS panel: more recoding

Age quintiles: find the cutpoints

```r
# seq can make all kinds of sequences
seq(from = 0, to = 1, by = 0.2)
```

```
## [1] 0.0 0.2 0.4 0.6 0.8 1.0
```

```r
age_quintiles <- quantile(as.numeric(gss_panel$age),
                          probs = seq(0, 1, 0.2),
                          na.rm = TRUE)

## These are the quintile cutpoints
age_quintiles
```

```
##   0%  20%  40%  60%  80% 100%
##   18   33   43   53   65   89
```

The GSS panel: more recoding

Age quintiles: create the quintile variable

```
## Apply the cut
gss_sub %>%
  mutate(agequint = cut(x = age,
                        breaks = unique(age_quintiles),
                        include.lowest = TRUE)) %>%
  pull(agequint) %>% # grab a column and make it an ordinary vector
  table()
```

```
## .
## [18,33] (33,43] (43,53] (53,65] (65,89]
##    3157    2680    2851    3057    2720
```

We'll need to clean up those labels.

The GSS panel: more recoding

I told you that regexp stuff would pay off.

```
convert_agegrp <- function(x){
    x <- stringr::str_remove(x, "\\(")  # Remove open paren
    x <- stringr::str_remove(x, "\\[")  # Remove open bracket
    x <- stringr::str_remove(x, "\\]")  # Remove close bracket
    x <- stringr::str_replace(x, ",", "-") # Replace comma with dash
    x <- stringr::str_replace(x, "-89", "+") # Replace -89 with +
    regex <- "^(.*$)" # Matches everything in string to end of line
    x <- stringr::str_replace(x, regex, "Age \\1") # Preface string with "Age"
    x
}
```

The GSS panel: more recoding

```r
gss_sub
```

```
## # A tibble: 14,610 × 19
##     year    id ballot   age tvhours race  sex    degree  relig  income  polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>   <fct>  <chr>   <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None   $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None   $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None   $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None   $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None   <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None   $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 8 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>
```

The GSS panel: more recoding

```r
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE))
```

```
## # A tibble: 14,610 × 20
##     year    id ballot   age tvhours race  sex    degree  relig  income  polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>   <fct>  <chr>   <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None   $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None   $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None   $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None   $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None   <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None   $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 9 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>
```

The GSS panel: more recoding

```
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE)) %>%
    mutate(agequint = fct_relabel(agequint, convert_agegrp))
```

```
## # A tibble: 14,610 × 20
##     year    id ballot   age tvhours race  sex    degree relig income polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>  <fct> <chr>  <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None  $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None  $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None  $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None  $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None  <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None  $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 9 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>
```

The GSS panel: more recoding

```r
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE)) %>%
    mutate(agequint = fct_relabel(agequint, convert_agegrp)) %>%
    mutate(year_f = droplevels(factor(year)))
```

```
## # A tibble: 14,610 × 21
##     year    id ballot   age tvhours race  sex    degree relig income polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>  <fct> <chr>  <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None  $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None  $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None  $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None  $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None  <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None  $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 10 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>, year_f <fct>
```

The GSS panel: more recoding

```r
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE)) %>%
    mutate(agequint = fct_relabel(agequint, convert_agegrp)) %>%
    mutate(year_f = droplevels(factor(year))) %>%
    mutate(young = ifelse(age < 26, "Yes", "No"))
```

```
## # A tibble: 14,610 × 22
##     year    id ballot   age tvhours race  sex    degree relig income polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>  <fct> <chr>  <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None  $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None  $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None  $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None  $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None  <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None  $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 11 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>, year_f <fct>, young <chr>
```

The GSS panel: more recoding

```r
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE)) %>%
    mutate(agequint = fct_relabel(agequint, convert_agegrp)) %>%
    mutate(year_f = droplevels(factor(year))) %>%
    mutate(young = ifelse(age < 26, "Yes", "No")) %>%
    mutate(fefam_d = fct_recode(fefam,
                                Agree = "Strongly Agree",
                                Disagree = "Strongly Disagree"))
```

```
## # A tibble: 14,610 × 23
##     year    id ballot   age tvhours race  sex    degree relig income polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <fct>  <fct> <chr>  <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None  $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None  $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None  $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None  $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None  <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None  $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 12 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>, year_f <fct>, young <chr>, fefam_d <fct>
```

The GSS panel: more recoding

```r
gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE)) %>%
    mutate(agequint = fct_relabel(agequint, convert_agegrp)) %>%
    mutate(year_f = droplevels(factor(year))) %>%
    mutate(young = ifelse(age < 26, "Yes", "No")) %>%
    mutate(fefam_d = fct_recode(fefam,
                                Agree = "Strongly Agree",
                                Disagree = "Strongly Disagree")) %>%
    mutate(degree = factor(degree,
                           levels = levels(gss_sub$degree),  #<<
                           ordered = TRUE))  #<<
```

```
## # A tibble: 14,610 × 23
##     year    id ballot   age tvhours race  sex    degree  relig  income  polviews
##    <int> <int>  <int> <int>   <int> <fct> <fct>  <ord>   <fct>  <chr>   <fct>
##  1  2006     9      3    23      NA Black Female Bachel… None   $25000… Conserv…
##  2  2008  3001      3    25      NA Other Female Bachel… None   $25000… Extreme…
##  3  2010  6001      3    27      NA Black Female Bachel… None   $25000… Extreme…
##  4  2010  6002      1    36       3 White Female Gradua… None   $25000… Liberal
##  5  2006    10      1    32       3 Other Female Gradua… None   <NA>    Slightl…
##  6  2008  3002      1    34       3 Other Female Gradua… None   $25000… Moderate
##  7  2008  3003      3    83      NA Black Female Lt Hig… Prote… $20000… Liberal
##  8  2010  6003      3    85      NA Black Female Lt Hig… Prote… <NA>    Moderate
##  9  2006    11      3    81      NA Black Female Lt Hig… Prote… <NA>    Moderate
## 10  2010  6004      1    51      10 Other Male   High S… Catho… Lt $10… Liberal
## # … with 14,600 more rows, and 12 more variables: fefam <fct>, vpsu <dbl>,
## #   vstrat <dbl>, oversamp <dbl>, formwt <dbl>, wtssall <dbl>, sampcode <dbl>,
## #   sample <dbl>, agequint <fct>, year_f <fct>, young <chr>, fefam_d <fct>
```

How we'd actually write this

```
gss_sub <- gss_sub %>%
      mutate(agequint = cut(x = age,
                            breaks = unique(age_quintiles),
                            include.lowest = TRUE),
             agequint = fct_relabel(agequint, convert_agegrp),
             year_f = droplevels(factor(year)),
             young = ifelse(age < 26, "Yes", "No"),
             fefam_d = fct_recode(fefam,
                                  Agree = "Strongly Agree",
                                  Disagree = "Strongly Disagree"),
             degree = factor(degree,
                             levels = levels(gss_sub$degree),
                             ordered = TRUE))
```

How we'd actually write this

```
gss_sub <- gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE),
            agequint = fct_relabel(agequint, convert_agegrp),
            year_f = factor(year),
            young = ifelse(age < 26, "Yes", "No"),
            fefam_d = fct_recode(fefam,
                                Agree = "Strongly Agree",
                                Disagree = "Strongly Disagree"),
            degree = factor(degree,
                           levels = levels(gss_sub$degree),
                           ordered = TRUE))
```

How we'd actually write this

```
gss_sub <- gss_sub %>%
    mutate(agequint = cut(x = age,
                          breaks = unique(age_quintiles),
                          include.lowest = TRUE),
            agequint = fct_relabel(agequint, convert_agegrp),
            year_f = droplevels(factor(year)),
            young = ifelse(age < 26, "Yes", "No"),
            fefam_d = fct_recode(fefam,
                                Agree = "Strongly Agree",
                                Disagree = "Strongly Disagree"),
            degree = factor(degree,
                           levels = levels(gss_sub$degree),
                           ordered = TRUE))
```

GSS Panel

```
gss_sub %>%
  select(sex, year, year_f, age, young, fefam, fefam_d) %>%
  sample_n(15)
```

```
## # A tibble: 15 × 7
##    sex     year year_f   age young fefam                fefam_d
##    <fct>  <int> <fct>  <int> <chr> <fct>                <fct>
##  1 Male    2010 2010      64 No    Agree                Agree
##  2 Male    2010 2010      22 Yes   Strongly Disagree Disagree
##  3 Female  2010 2010      56 No    Agree                Agree
##  4 Male    2008 2008      73 No    <NA>                 <NA>
##  5 Female  2012 2012      32 No    Disagree             Disagree
##  6 Female  2008 2008      56 No    <NA>                 <NA>
##  7 Female  2008 2008      30 No    Disagree             Disagree
##  8 Female  2014 2014      33 No    Agree                Agree
##  9 Male    2010 2010      41 No    Disagree             Disagree
## 10 Female  2008 2008      86 No    <NA>                 <NA>
## 11 Female  2012 2012      34 No    Disagree             Disagree
## 12 Male    2012 2012      49 No    Strongly Agree       Agree
## 13 Male    2010 2010      30 No    Strongly Agree       Agree
## 14 Male    2008 2008      72 No    Agree                Agree
## 15 Female  2012 2012      35 No    Disagree             Disagree
```

GSS Panel

```
gss_sub %>%
  select(sex, degree) %>%
  group_by(sex, degree) %>%
  tally() %>%
  pivot_wider(names_from = sex, values_from = n)
```

```
## # A tibble: 6 × 3
##   degree          Male Female
##   <ord>          <int>  <int>
## 1 Lt High School   814   1036
## 2 High School     3131   4143
## 3 Junior College   440    721
## 4 Bachelor        1293   1474
## 5 Graduate         696    860
## 6 <NA>              NA      2
```

More on **factors**

We've already seen **`fct_relabel()`** and **`fct_recode()`** from forcats.

There are numerous other convenience functions for factors.

More on **factors**

We've already seen **`fct_relabel()`** and **`fct_recode()`** from forcats.

There are numerous other convenience functions for factors.

```
gss_sub %>%
  count(degree)
```

```
## # A tibble: 6 × 2
##   degree              n
##   <ord>           <int>
## 1 Lt High School   1850
## 2 High School      7274
## 3 Junior College   1161
## 4 Bachelor         2767
## 5 Graduate         1556
## 6 <NA>                2
```

```
levels(gss_sub$degree)
```

```
## [1] "Lt High School" "High School"    "Junior College" "Bachelor"
## [5] "Graduate"
```

More on factors

Make the NA values an explicit level

```
gss_sub %>%
  mutate(degree_na = fct_explicit_na(degree)) %>%
  count(degree_na)
```

```
## # A tibble: 6 × 2
##   degree_na          n
##   <ord>          <int>
## 1 Lt High School  1850
## 2 High School     7274
## 3 Junior College  1161
## 4 Bachelor        2767
## 5 Graduate        1556
## 6 (Missing)          2
```

More on **factors**

Relevel by frequency

```
gss_sub %>%
  mutate(degree_freq = fct_infreq(degree)) %>%
  count(degree_freq)
```

```
## # A tibble: 6 × 2
##   degree_freq        n
##   <ord>          <int>
## 1 High School     7274
## 2 Bachelor        2767
## 3 Lt High School  1850
## 4 Graduate        1556
## 5 Junior College  1161
## 6 <NA>               2
```

More on factors

Relevel manually

```r
is.ordered(gss_sub$sex)
```

```
## [1] FALSE
```

```r
levels(gss_sub$sex)
```

```
## [1] "Male"   "Female"
```

```r
summary(lm(age ~ sex, data = gss_sub))
```

```
##
## Call:
## lm(formula = age ~ sex, data = gss_sub)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -31.431 -13.972  -0.431  12.569  40.028
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  48.9720     0.2149 227.846   <2e-16 ***
## sexFemale     0.4594     0.2864   1.604    0.109
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 17.08 on 14463 degrees of freedom
##   (145 observations deleted due to missingness)
## Multiple R-squared:  0.0001779,    Adjusted R-squared:  0.0001088
## F-statistic: 2.573 on 1 and 14463 DF,  p-value: 0.1087
```

More on factors

Relevel manually

```
gss_sub <- gss_sub %>%
  mutate(sex = fct_relevel(sex, "Female"))

levels(gss_sub$sex)
```

```
## [1] "Female" "Male"
```

More on factors

Relevel manually

```
gss_sub <- gss_sub %>%
  mutate(sex = fct_relevel(sex, "Female"))

levels(gss_sub$sex)
```

```
## [1] "Female" "Male"
```

```
summary(lm(age ~ sex, data = gss_sub))
```

```
##
## Call:
## lm(formula = age ~ sex, data = gss_sub)
##
## Residuals:
##     Min      1Q  Median      3Q     Max
## -31.431 -13.972  -0.431  12.569  40.028
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  49.4313     0.1892 261.233   <2e-16 ***
## sexMale      -0.4594     0.2864  -1.604    0.109
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

More on **factors**

Interact or cross factors

```
gss_sub <- gss_sub %>%
  mutate(degree_by_race = fct_cross(race, degree))

gss_sub %>%
  count(degree_by_race)
```

```
## # A tibble: 16 × 2
##    degree_by_race           n
##    <fct>                <int>
##  1 White:Lt High School  1188
##  2 Black:Lt High School   379
##  3 Other:Lt High School   283
##  4 White:High School     5548
##  5 Black:High School     1180
##  6 Other:High School      546
##  7 White:Junior College   885
##  8 Black:Junior College   206
##  9 Other:Junior College    70
## 10 White:Bachelor        2334
## 11 Black:Bachelor         233
## 12 Other:Bachelor         200
## 13 White:Graduate        1293
## 14 Black:Graduate         116
## 15 Other:Graduate         147
## 16 <NA>                     2
```

More on factors

Relevel manually by lumping ... the least frequent n

```
gss_sub %>%
  mutate(degree_n = fct_lump_n(degree, n = 3)) %>%
  count(degree_n)
```

```
## # A tibble: 5 × 2
##   degree_n           n
##   <ord>          <int>
## 1 Lt High School  1850
## 2 High School     7274
## 3 Bachelor        2767
## 4 Other           2717
## 5 <NA>               2
```

More on factors

Relevel manually by lumping ...to other, manually

```
gss_sub %>%
  mutate(degree_o = fct_other(degree,
                              keep = c("Lt High School",
                                       "High School"))) %>%
  count(degree_o)
```

```
## # A tibble: 4 × 2
##   degree_o           n
##   <ord>          <int>
## 1 Lt High School  1850
## 2 High School     7274
## 3 Other           5484
## 4 <NA>               2
```