

Iterating on data with `purrr` and `map`

Data Wrangling: Session 7

Kieran Healy
Statistical Horizons, April 2022

Load the packages, as always

```
library(here)      # manage file paths

## here() starts at /Users/kjhealy/Documents/courses/data_wrangling

library(socviz)    # data and some useful functions

##
## Attaching package: 'socviz'

## The following object is masked from 'package:kjhutils':
## 
##     %nin%

library(tidyverse) # your friend and mine

## — Attaching packages ————— tidyverse 1.3.1 —
## ✓ ggplot2 3.3.5     ✓ purrr   0.3.4
## ✓ tibble  3.1.6     ✓ dplyr   1.0.8
## ✓ tidyr   1.2.0     ✓ stringr 1.4.0
## ✓ readr   2.1.2     ✓forcats 0.5.1

## — Conflicts ————— tidyverse_conflicts() —
## x readr::edition_get() masks testthat::edition_get()
## x dplyr::filter()     masks stats::filter()
## x purrr::is_null()    masks testthat::is_null()
## x dplyr::lag()       masks stats::lag()
## x readr::local_edition() masks testthat::local_edition()
## x dplyr::matches()   masks tidyR::matches(), testthat::matches()
```

Moar Data

More than one data file

Inside the data/ folder of the course packet is a folder named congress/

```
# A little trick from the fs package:  
fs::dir_tree(here("data", "congress"))  
  
## /Users/kjhealy/Documents/courses/data_wrangling/data/congress  
##   ├── 01_79_congress.csv  
##   ├── 02_80_congress.csv  
##   ├── 03_81_congress.csv  
##   ├── 04_82_congress.csv  
##   ├── 05_83_congress.csv  
##   ├── 06_84_congress.csv  
##   ├── 07_85_congress.csv  
##   ├── 08_86_congress.csv  
##   ├── 09_87_congress.csv  
##   ├── 10_88_congress.csv  
##   ├── 11_89_congress.csv  
##   ├── 12_90_congress.csv  
##   ├── 13_91_congress.csv  
##   ├── 14_92_congress.csv  
##   ├── 15_93_congress.csv  
##   ├── 16_94_congress.csv  
##   ├── 17_95_congress.csv  
##   ├── 18_96_congress.csv  
##   ├── 19_97_congress.csv  
##   ├── 20_98_congress.csv  
##   ├── 21_99_congress.csv  
##   ├── 22_100_congress.csv  
##   ├── 23_101_congress.csv  
##   ├── 24_102_congress.csv  
##   ├── 25_103_congress.csv  
##   ├── 26_104_congress.csv  
##   ├── 27_105_congress.csv  
##   ├── 28_106_congress.csv  
##   └── 29_107_congress.csv
```

More than one data file

Let's look at one.

```
read_csv(here("data", "congress", "17_95_congress.csv")) %>%
  janitor::clean_names() %>%
  head()

## # A tibble: 6 × 25
##   last     first    middle suffix nickname born   death sex   position party state
##   <chr>    <chr>    <chr>  <chr>  <chr>  <chr> <chr> <chr> <chr> <chr>
## 1 Abdnor   James    <NA>   <NA>   <NA>   02/1... 11/0... M   U.S. Re... Repu...
## 2 Abourezk  James    George  <NA>   <NA>   02/2... <NA>   M   U.S. Se... Demo...
## 3 Adams     Brockm... <NA>   <NA>   Brock   01/1... 09/1... M   U.S. Re... Demo...
## 4 Addabbo   Joseph   Patri... <NA>   <NA>   03/1... 04/1... M   U.S. Re... Demo...
## 5 Aiken     George   David   <NA>   <NA>   08/2... 11/1... M   U.S. Se... Repu...
## 6 Akaka     Daniel   Kahik... <NA>   <NA>   09/1... 04/0... M   U.S. Re... Demo...
## # ... with 14 more variables: district <chr>, start <chr>, end <chr>,
## #   religion <chr>, race <chr>, educational_attainment <chr>, job_type1 <chr>,
## #   job_type2 <chr>, job_type3 <chr>, job_type4 <chr>, job_type5 <lgl>,
## #   mil1 <chr>, mil2 <chr>, mil3 <chr>
```

We often find ourselves in this situation. We know each file has the same structure, and we would like to use them all at once.

Loops?

How to read them all in?

One traditional way, which we could do in R, is to write an explicit *loop* that iterated over a vector of filenames, read each file, and then joined the results together in a tall rectangle.

```
# Pseudocode

filenames <- c("01_79_congress.csv", "02_80_congress.csv", "03_81_congress.csv",
             "04_82_congress.csv" [etc etc])

collected_files <- NULL

for(i in 1:length(filenames)) {
  new_file <- read_file(filenames[i])
  collected_files <- append_to(collected_files, new_file)
}
```

Loops?

You may have noticed we have not written any loops, however.

While loops are still lurking there underneath the surface, what we will do instead is to take advantage of the combination of vectors and functions and *map* one to the other in order to generate results.

Speaking loosely, think of `map()` as a way of *iterating* without writing loops. You start with a vector of things. You feed it one thing at a time to some function. The function does whatever it does. You get back output that is the same length as your input, and of a specific type.

Mapping is just a kind of iteration

The `purrr` package provides a big family of mapping functions. One reason there are a lot of them is that `purrr`, like the rest of the tidyverse, is picky about data types.

Mapping is just a kind of iteration

The `purrr` package provides a big family of mapping functions. One reason there are a lot of them is that `purrr`, like the rest of the tidyverse, is picky about data types.

So in addition to the basic `map()`, which always returns a *list*, we also have `map_chr()`, `map_int()`, `map_dbl()`, `map_lgl()` and others. They always return the data type indicated by their suffix, or die trying.

Vectorized arithmetic again

The simplest cases are not that different from the vectorized arithmetic we're already familiar with.

```
a <- c(1:10)  
b <- 1  
  
# You know what R will do here  
a + b  
  
## [1] 2 3 4 5 6 7 8 9 10 11
```

Vectorized arithmetic again

The simplest cases are not that different from the vectorized arithmetic we're already familiar with.

```
a <- c(1:10)  
b <- 1  
  
# You know what R will do here  
a + b  
  
## [1] 2 3 4 5 6 7 8 9 10 11
```

R's vectorized rules add b to every element of a. In a sense, the **+** operation can be thought of as a function that takes each element of a and does something with it. In this case "add b".

Vectorized arithmetic again

We can make this explicit by writing a function:

```
add_b <- function(x) {  
  b <- 1  
  x + b # for any x  
}
```

Now:

```
add_b(x = a)  
## [1] 2 3 4 5 6 7 8 9 10 11
```

Vectorized arithmetic again

Again, R's vectorized approach means it automatically adds b to every element of the x we give it.

```
add_b(x = 10)
## [1] 11

add_b(x = c(1, 99, 1000))
## [1]    2   100 1001
```

Iterating in a pipeline

Some operations can't directly be vectorized in this way, which is why we need to manually iterate, or will want to write loops.

```
library(gapminder)
gapminder %>%
  summarize(country_n = n_distinct(country),
           continent_n = n_distinct(continent),
           year_n = n_distinct(year),
           lifeExp_n = n_distinct(lifeExp),
           population_n = n_distinct(population))

## # A tibble: 1 × 5
##   country_n continent_n year_n lifeExp_n population_n
##       <int>      <int>    <int>      <int>
## 1        142          5     12      1626        4060
```

That's tedious to write! Computers are supposed to allow us to avoid that sort of thing.

Iterating in a pipeline

So how would we iterate this? What we want is to apply the `n_distinct()` function to each column of `gapminder`, but in a way that still allows us to use pipelines and so on.

```
library(gapminder)
gapminder %>%
  summarize(n_distinct(country),
            n_distinct(continent),
            n_distinct(year),
            n_distinct(lifeExp),
            n_distinct(population))

## # A tibble: 1 × 5
##   `n_distinct(country)` `n_distinct(continen...` `n_distinct(ye...` `n_distinct(li...` 
##   <int>           <int>           <int>           <int>
## 1        142             5             12           1626
## # ... with 1 more variable: `n_distinct(population)` <int>
```

Using `n_distinct()` in this context is an idea I got from Rebecca Barter's discussion of `purrr`.

Iterating in a pipeline

You'd use **across()**, like this:

```
gapminder %>%
  summarize(across(everything(), n_distinct))

## # A tibble: 1 × 6
##   country continent year lifeExp   pop gdpPercap
##   <int>     <int> <int>   <int> <int>      <int>
## 1     142       5    12    1626  1704      1704
```

Iterating in a pipeline

But you could also do this ...

```
map(gapminder, n_distinct)

## $country
## [1] 142
##
## $continent
## [1] 5
##
## $year
## [1] 12
##
## $lifeExp
## [1] 1626
##
## $pop
## [1] 1704
##
## $gdpPercap
## [1] 1704
```

Read it as "Feed each column of gapminder to the **n_distinct()** function.

(This is pretty much what **across()** is doing more nicely.)

Iterating in a pipeline

Or, in pipeline form:

```
gapminder %>%  
  map(n_distinct)
```

```
## $country  
## [1] 142  
##  
## $continent  
## [1] 5  
##  
## $year  
## [1] 12  
##  
## $lifeExp  
## [1] 1626  
##  
## $pop  
## [1] 1704  
##  
## $gdpPercap  
## [1] 1704
```

You can see we are getting a *list* back.

Iterating in a pipeline

Or, in pipeline form:

```
result <- gapminder %>%
  map(n_distinct)

class(result)

## [1] "list"

result$continent

## [1] 5

result[[2]]

## [1] 5
```

Iterating in a pipeline

But we know `n_distinct()` should always return an integer. So we use `map_int()` instead of the generic `map()`.

```
gapminder %>%  
  map_int(n_distinct)  
  
## #> #> #> #> #>
```

	country	continent	year	lifeExp	pop	gdpPerCap
##	142	5	12	1626	1704	1704

The thing about the `map()` family is that they can deal with all kinds of input types and output types.

Get a vector of filenames

```
filenames <- dir(path = here("data", "congress"),
                  pattern = "*.csv",
                  full.names = TRUE)

filenames[1:15] # Just displaying the first 15, to save slide space

## [1] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/01_79_congress.csv"
## [2] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/02_80_congress.csv"
## [3] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/03_81_congress.csv"
## [4] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/04_82_congress.csv"
## [5] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/05_83_congress.csv"
## [6] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/06_84_congress.csv"
## [7] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/07_85_congress.csv"
## [8] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/08_86_congress.csv"
## [9] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/09_87_congress.csv"
## [10] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/10_88_congress.csv"
## [11] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/11_89_congress.csv"
## [12] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/12_90_congress.csv"
## [13] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/13_91_congress.csv"
## [14] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/14_92_congress.csv"
## [15] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/15_93_congress.csv"
```

And feed it to `read_csv()`

... using the variant of `map()` that returns data frames and tibbles.

```
df <- filenames %>%
  map_dfr(read_csv, .id = "congress") %>%
  janitor::clean_names()

df

## # A tibble: 20,580 × 26
##   congress last   first middle suffix nickname born   death sex   position party
##   <chr>     <chr>  <chr> <chr>  <chr>    <chr> <chr> <chr> <chr>    <chr>
## 1 1         Abern... Thom... Gerst... <NA>    <NA>   05/1... 01/2... M    U.S. Re... Demo...
## 2 1         Adams   Sher... <NA>    <NA>    <NA>   01/0... 10/2... M    U.S. Re... Repu...
## 3 1         Aiken   Geor... David   <NA>    <NA>   08/2... 11/1... M    U.S. Se... Repu...
## 4 1         Allen   Asa    Leona... <NA>    <NA>   01/0... 01/0... M    U.S. Re... Demo...
## 5 1         Allen   Leo    Elwood  <NA>    <NA>   10/0... 01/1... M    U.S. Re... Repu...
## 6 1         Almond  J.    Linds... Jr.   <NA>    <NA>   06/1... 04/1... M    U.S. Re... Demo...
## 7 1         Ander... Herm... Carl   <NA>    <NA>   01/2... 07/2... M    U.S. Re... Repu...
## 8 1         Ander... Clin... Presba <NA>    <NA>   10/2... 11/1... M    U.S. Re... Demo...
## 9 1         Ander... John   Zuing... <NA>    <NA>   03/2... 02/0... M    U.S. Re... Repu...
## 10 1        Andre... Augu... Herman <NA>    <NA>   10/1... 01/1... M    U.S. Re... Repu...
## # ... with 20,570 more rows, and 15 more variables: state <chr>, district <chr>,
## #   start <chr>, end <chr>, religion <chr>, race <chr>,
## #   educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
## #   job_type3 <chr>, job_type4 <chr>, job_type5 <chr>, mil1 <chr>, mil2 <chr>,
## #   mil3 <chr>
```

Now witness the firepower of this fully armed and operational



method of type-safe functional iteration

read_csv() can do this directly now!

No map_df() required

```
tmp <- read_csv(filenames, id = "path",
                 name_repair = janitor::make_clean_names)

tmp %>%
  mutate(congress = stringr::str_extract(path, "_\\d{2,3}_congress"),
         congress = stringr::str_extract(congress, "\\d{2,3}")) %>%
  relocate(congress)

## # A tibble: 20,580 × 27
##   congress path   last   first middle suffix nickname born   death sex   position
##   <chr>    <chr>  <chr> <chr> <chr>  <chr>   <chr> <chr> <chr> <chr>
## 1 79       /User... Aber... Thom... Gerst... <NA>   <NA>   05/1... 01/2... M   U.S. Re...
## 2 79       /User... Adams Sher... <NA>   <NA>   <NA>   01/0... 10/2... M   U.S. Re...
## 3 79       /User... Aiken Geor... David   <NA>   <NA>   08/2... 11/1... M   U.S. Se...
## 4 79       /User... Allen Asa   Leona... <NA>   <NA>   01/0... 01/0... M   U.S. Re...
## 5 79       /User... Allen Leo   Elwood <NA>   <NA>   10/0... 01/1... M   U.S. Re...
## 6 79       /User... Almo... J.   Linds... Jr.  <NA>   <NA>   06/1... 04/1... M   U.S. Re...
## 7 79       /User... Ande... Herm... Carl   <NA>   <NA>   01/2... 07/2... M   U.S. Re...
## 8 79       /User... Ande... Clin... Presba <NA>   <NA>   10/2... 11/1... M   U.S. Re...
## 9 79       /User... Ande... John   Zuing... <NA>   <NA>   03/2... 02/0... M   U.S. Re...
## 10 79      /User... Andr... Augu... Herman <NA>   <NA>   10/1... 01/1... M   U.S. Re...
## # ... with 20,570 more rows, and 16 more variables: party <chr>, state <chr>,
## #   district <chr>, start <chr>, end <chr>, religion <chr>, race <chr>,
## #   educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
## #   job_type3 <chr>, job_type4 <chr>, job_type5 <chr>, mil1 <chr>, mil2 <chr>,
## #   mil3 <chr>
```

Example: Iterating on the US Census

But mapped iteration is not just for local files!

```
## Register for a free Census API key
library(tidycensus)

out <- get_acs(geography = "county",
                variables = "B19013_001",
                state = "NY",
                county = "New York",
                survey = "acs1",
                year = 2005)

out

## # A tibble: 1 × 5
##   GEOID NAME           variable estimate    moe
##   <chr> <chr>          <chr>     <dbl> <dbl>
## 1 36061 New York County, New York B19013_001  55973  1462
```

Example: Iterating on the US Census

All counties in New York State for a specific year

```
out <- get_acs(geography = "county",
                variables = "B19013_001",
                state = "NY",
                survey = "acs1",
                year = 2005)

out

## # A tibble: 38 × 5
##   GEOID NAME           variable   estimate    moe
##   <chr> <chr>          <chr>      <dbl> <dbl>
## 1 36001 Albany County, New York B19013_001  50054  2030
## 2 36005 Bronx County, New York B19013_001 29228   853
## 3 36007 Broome County, New York B19013_001 36394  2340
## 4 36009 Cattaraugus County, New York B19013_001 37580  2282
## 5 36011 Cayuga County, New York B19013_001 42057  2406
## 6 36013 Chautauqua County, New York B19013_001 35495  2077
## 7 36015 Chemung County, New York B19013_001 37418  3143
## 8 36019 Clinton County, New York B19013_001 44757  3500
## 9 36027 Dutchess County, New York B19013_001 61889  2431
## 10 36029 Erie County, New York B19013_001 41967  1231
## # ... with 28 more rows
```

Example: Iterating on the US Census

What if we want the results for *every* available year?

First, a handy function: `set_names()`

```
x <- c(1:10)  
x  
## [1] 1 2 3 4 5 6 7 8 9 10  
  
x <- set_names(x, nm = letters[1:10])  
x  
## a b c d e f g h i j  
## 1 2 3 4 5 6 7 8 9 10
```

Example: Iterating on the US Census

By default, `set_names()` will label a vector with that vector's values:

```
c(1:10) %>%  
  set_names()  
  
##  1  2  3  4  5  6  7  8  9 10  
## 1  2  3  4  5  6  7  8  9 10
```

Example: Iterating on the US Census

Meanwhile, `map_dfr()` has an `.id` argument that lets you add a row-identifier to whatever you are binding. Like this:

```
df <- 2005:2019 %>%
  map_dfr(~ get_acs(geography = "county",
                     variables = "B19013_001",
                     state = "NY",
                     county = "New York",
                     survey = "acs1",
                     year = .x),
         .id = "id")

df

## # A tibble: 15 × 6
##   id    GEOID NAME          variable estimate    moe
##   <chr> <chr> <chr>        <chr>     <dbl> <dbl>
## 1 1    36061 New York County, New York B19013_001 55973 1462
## 2 2    36061 New York County, New York B19013_001 60017 1603
## 3 3    36061 New York County, New York B19013_001 64217 2002
## 4 4    36061 New York County, New York B19013_001 69017 1943
## 5 5    36061 New York County, New York B19013_001 68706 1471
## 6 6    36061 New York County, New York B19013_001 63832 2125
## 7 7    36061 New York County, New York B19013_001 66299 1783
## 8 8    36061 New York County, New York B19013_001 67099 1640
## 9 9    36061 New York County, New York B19013_001 72190 2200
## 10 10  36061 New York County, New York B19013_001 76089 2016
## # ... with 5 more rows, and 1 more variable: estimate.error <dbl>
```

Example: Iterating on the US Census

Our `id` column tracks the year. But we'd like it to *be* the year. So, we use `set_names()`, and also name it `year` when we create it:

```
df <- 2005:2019 %>%
  set_names() %>%
  map_dfr(~ get_acs(geography = "county",
                     variables = "B19013_001",
                     state = "NY",
                     county = "New York",
                     survey = "acs1",
                     year = .x),
         .id = "year") %>%
  mutate(year = as.integer(year))
```

Example: Iterating on the US Census

df

```
## # A tibble: 15 × 6
##   year GEOID NAME      variable estimate    moe
##   <int> <chr> <chr>     <chr>      <dbl> <dbl>
## 1 2005 36061 New York County, New York B19013_001  55973 1462
## 2 2006 36061 New York County, New York B19013_001  60017 1603
## 3 2007 36061 New York County, New York B19013_001  64217 2002
## 4 2008 36061 New York County, New York B19013_001  69017 1943
## 5 2009 36061 New York County, New York B19013_001  68706 1471
## 6 2010 36061 New York County, New York B19013_001  63832 2125
## 7 2011 36061 New York County, New York B19013_001  66299 1783
## 8 2012 36061 New York County, New York B19013_001  67099 1640
## 9 2013 36061 New York County, New York B19013_001  72190 2200
## 10 2014 36061 New York County, New York B19013_001  76089 2016
## 11 2015 36061 New York County, New York B19013_001  75575 2566
## 12 2016 36061 New York County, New York B19013_001  77559 2469
## 13 2017 36061 New York County, New York B19013_001  85071 3520
## 14 2018 36061 New York County, New York B19013_001  85066 3480
## 15 2019 36061 New York County, New York B19013_001  93651 3322
```

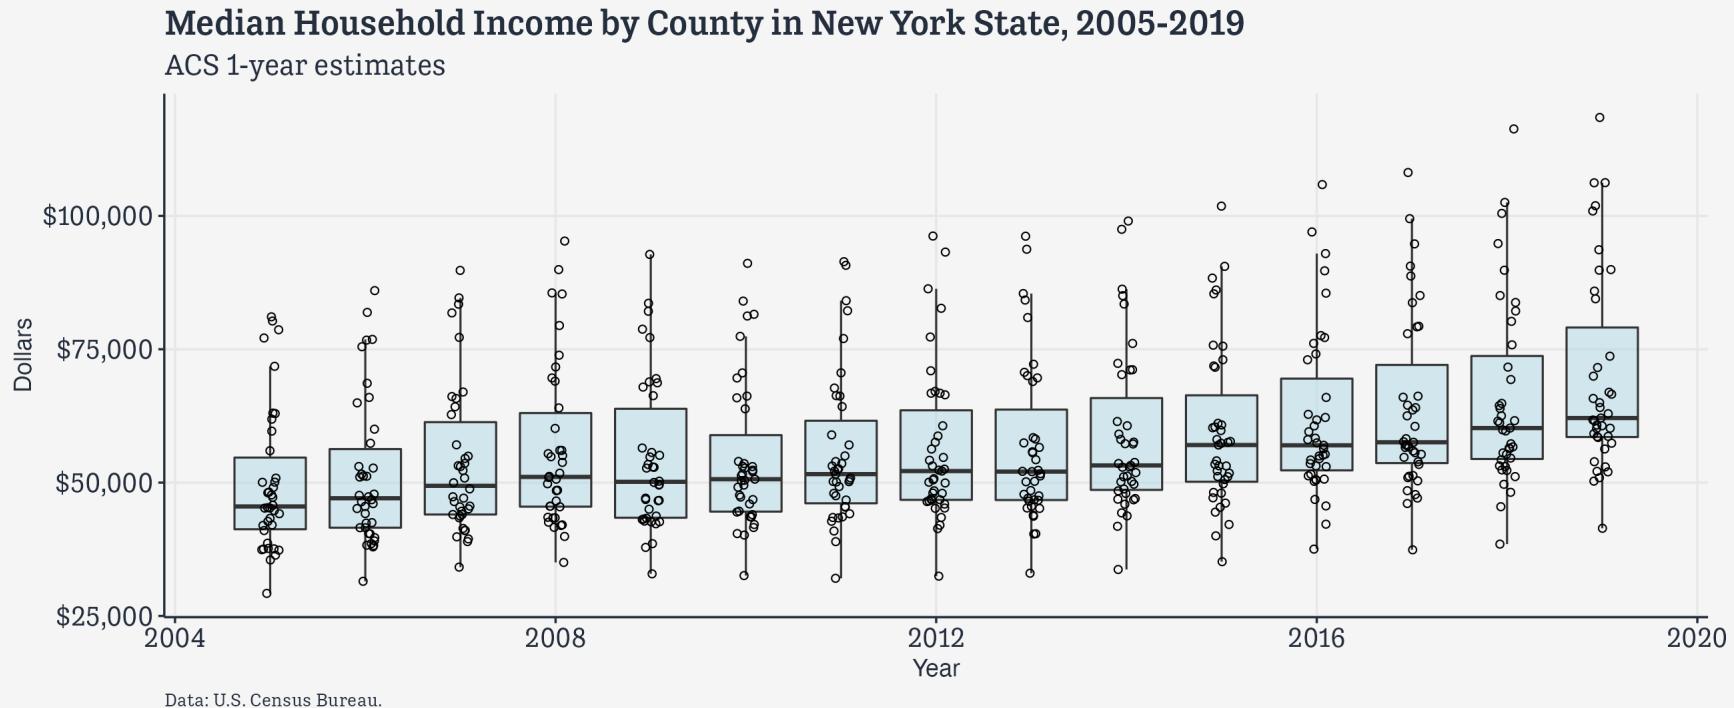
Now `year` is just the year. The `year` column will be created as a character vector, so we converted it back to an integer again at the end.

Example: Iterating on the US Census

```
p_out <- 2005:2019 %>%
  set_names() %>%
  map_dfr(~ get_acs(geography = "county", variables = "B19013_001",
                     state = "NY", survey = "acs1", year = .x), .id = "year") %>%
  mutate(year = as.integer(year)) %>%
  ggplot(mapping = aes(x = year, y = estimate, group = year)) +
  geom_boxplot(fill = "lightblue", alpha = 0.5, outlier.alpha = 0) +
  geom_jitter(position = position_jitter(width = 0.1), shape = 1) +
  scale_y_continuous(labels = scales::label_dollar()) + labs(x = "Year", y = "Dollars",
  title = "Median Household Income by County in New York State, 2005-2019",
  subtitle = "ACS 1-year estimates", caption = "Data: U.S. Census Bureau.")
```

Example: Iterating on the US Census

```
print(p_out)
```



Cleaning up congress

```
df <- filenames %>%
  map_dfr(read_csv, .id = "congress") %>%
  janitor::clean_names()

df %>%
  select(born, death, start, end)

## # A tibble: 20,580 × 4
##   born      death     start     end
##   <chr>     <chr>     <chr>     <chr>
## 1 05/16/1903 01/23/1953 01/03/1945 01/03/1953
## 2 01/08/1899 10/27/1986 01/03/1945 01/03/1947
## 3 08/20/1892 11/19/1984 01/03/1945 01/03/1979
## 4 01/05/1891 01/05/1969 01/03/1945 01/03/1953
## 5 10/05/1898 01/19/1973 01/03/1945 01/02/1949
## 6 06/15/1898 04/14/1986 02/04/1946 04/17/1948
## 7 01/27/1897 07/26/1978 01/03/1945 01/03/1963
## 8 10/23/1895 11/11/1975 01/03/1941 06/30/1945
## 9 03/22/1904 02/09/1981 01/03/1945 01/03/1953
## 10 10/11/1890 01/14/1958 01/03/1945 01/14/1958
## # ... with 20,570 more rows
```

We'll use the **lubridate** package to sort these out.

Lubridate has a wide range of functions to handle dates, times, and durations.

Cleaning up congress

```
library(lubridate)

date_recodes <- c("born", "death", "start", "end")
df <- df %>%
  mutate(across(any_of(date_recodes), mdy),
        congress = as.double(congress) + 78)

df

## # A tibble: 20,580 × 26
##   congress last     first   middle suffix nickname born       death      sex
##   <dbl> <chr>    <chr>   <chr>  <chr>  <chr> <date>    <date>    <chr>
## 1      79 Abernethy Thomas Gerst... <NA>   <NA>  1903-05-16 1953-01-23 M
## 2      79 Adams     Sherman <NA>   <NA>   <NA>  1899-01-08 1986-10-27 M
## 3      79 Aiken     George David   <NA>  <NA>  1892-08-20 1984-11-19 M
## 4      79 Allen     Asa    Leona... <NA>  <NA>  1891-01-05 1969-01-05 M
## 5      79 Allen     Leo    Elwood <NA>  <NA>  1898-10-05 1973-01-19 M
## 6      79 Almond    J.    Linds... Jr. <NA>  <NA>  1898-06-15 1986-04-14 M
## 7      79 Andersen Herman Carl   <NA>  <NA>  1897-01-27 1978-07-26 M
## 8      79 Anderson Clinton Presba <NA>  <NA>  1895-10-23 1975-11-11 M
## 9      79 Anderson John   Zuинг... <NA>  <NA>  1904-03-22 1981-02-09 M
## 10     79 Andresen August Herman <NA>  <NA>  1890-10-11 1958-01-14 M
## # ... with 20,570 more rows, and 17 more variables: position <chr>, party <chr>,
## # state <chr>, district <chr>, start <date>, end <date>, religion <chr>,
## # race <chr>, educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
## # job_type3 <chr>, job_type4 <chr>, job_type5 <chr>, mil1 <chr>, mil2 <chr>,
## # mil3 <chr>
```

Cleaning up congress

```
sessions <- tibble(congress = 79:116,
                    start_year = seq(1945, 2019, by = 2),
                    end_year = seq(1947, 2021, by = 2)) %>%
  mutate(start_year = ymd(paste(start_year, "01", "03", sep = "-")),
        end_year = ymd(paste(end_year, "01", "03", sep = "-")))

sessions

## # A tibble: 38 × 3
##   congress start_year end_year
##       <int>     <date>    <date>
## 1       79 1945-01-03 1947-01-03
## 2       80 1947-01-03 1949-01-03
## 3       81 1949-01-03 1951-01-03
## 4       82 1951-01-03 1953-01-03
## 5       83 1953-01-03 1955-01-03
## 6       84 1955-01-03 1957-01-03
## 7       85 1957-01-03 1959-01-03
## 8       86 1959-01-03 1961-01-03
## 9       87 1961-01-03 1963-01-03
## 10      88 1963-01-03 1965-01-03
## # ... with 28 more rows
```

We're going to join these tables

The big table

```
df %>%  
  select(congress, last, born)  
  
## # A tibble: 20,580 × 3  
##   congress last    born  
##   <dbl> <chr>   <date>  
## 1      79 Abernethy 1903-05-16  
## 2      79 Adams    1899-01-08  
## 3      79 Aiken    1892-08-20  
## 4      79 Allen    1891-01-05  
## 5      79 Allen    1898-10-05  
## 6      79 Almond   1898-06-15  
## 7      79 Andersen 1897-01-27  
## 8      79 Anderson 1895-10-23  
## 9      79 Anderson 1904-03-22  
## 10     79 Andresen 1890-10-11  
## # ... with 20,570 more rows
```

The smaller table

```
sessions  
## # A tibble: 38 × 3  
##   congress start_year end_year  
##   <int> <date>    <date>  
## 1      79 1945-01-03 1947-01-03  
## 2      80 1947-01-03 1949-01-03  
## 3      81 1949-01-03 1951-01-03  
## 4      82 1951-01-03 1953-01-03  
## 5      83 1953-01-03 1955-01-03  
## 6      84 1955-01-03 1957-01-03  
## 7      85 1957-01-03 1959-01-03  
## 8      86 1959-01-03 1961-01-03  
## 9      87 1961-01-03 1963-01-03  
## 10     88 1963-01-03 1965-01-03  
## # ... with 28 more rows
```

We're going to **join** these tables

We will use **left_join()** which is what you want most of the time when you are looking to merge a smaller table with additional information into a larger main one.

```
df <- left_join(df, sessions) %>%
  relocate(start_year:end_year, .after = congress)

## Joining, by = "congress"

df

## # A tibble: 20,580 × 28
##   congress start_year end_year   last   first middle suffix nickname born
##       <dbl>    <date>    <date>   <chr>  <chr>  <chr>  <chr>  <chr>    <date>
## 1       79 1945-01-03 1947-01-03 Abern... Thom... Gerst... <NA>  <NA>  1903-05-16
## 2       79 1945-01-03 1947-01-03 Adams  Sher... <NA>  <NA>  <NA>  1899-01-08
## 3       79 1945-01-03 1947-01-03 Aiken   Geor... David  <NA>  <NA>  1892-08-20
## 4       79 1945-01-03 1947-01-03 Allen   Asa   Leona... <NA>  <NA>  1891-01-05
## 5       79 1945-01-03 1947-01-03 Allen   Leo   Elwood <NA>  <NA>  1898-10-05
## 6       79 1945-01-03 1947-01-03 Almond J. Linds... Jr.  <NA>  <NA>  1898-06-15
## 7       79 1945-01-03 1947-01-03 Ander... Herm... Carl  <NA>  <NA>  1897-01-27
## 8       79 1945-01-03 1947-01-03 Ander... Clin... Presba <NA>  <NA>  1895-10-23
## 9       79 1945-01-03 1947-01-03 Ander... John  Zuing... <NA>  <NA>  1904-03-22
## 10      79 1945-01-03 1947-01-03 Andre... Augu... Herman <NA>  <NA>  1890-10-11
## ... with 20,570 more rows, and 19 more variables: death <date>, sex <chr>,
## #   position <chr>, party <chr>, state <chr>, district <chr>, start <date>,
## #   end <date>, religion <chr>, race <chr>, educational_attainment <chr>,
## #   job_type1 <chr>, job_type2 <chr>, job_type3 <chr>, job_type4 <chr>,
## #   job_type5 <chr>, mil1 <chr>, mil2 <chr>, mil3 <chr>
```

Table joins

X	y
1	x1
2	x2
3	x3

1	y1
2	y2
4	y4

*Spiffy Join Animations courtesy [Garrick Aden-Buie](#)

Left join, `left_join()`

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x, and all columns from x and y. Rows in x with no match in y will have NA values in the new columns.

Left join (contd), `left_join()`

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4
		2	y5

If there are multiple matches between x and y, all combinations of the matches are returned.

Inner join, `inner_join()`

inner_join(x, y)

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x where there are matching values in y, and all columns from x and y.

Full join, **full_join()**

`full_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows and all columns from both x and y. Where there are not matching values, returns NA for the one missing.

Semi join, **semi_join()**

`semi_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x where there are matching values in y, keeping just columns from x.

Anti join, **anti_join()**

anti_join(x, y)	
1	x1
2	x2
3	x3
1	y1
2	y2
4	y4

All rows from x where there are not matching values in y, keeping just columns from x.

Left join, `left_join()`

Most of the time you will be looking to make a `left_join()`

Missing Data

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
df <- tribble(  
  ~subject, ~age,  
  "A", 20,  
  "B", 25,  
  "C", NA,  
  "D", 34  
)  
  
df  
  
## # A tibble: 4 × 2  
##   subject    age  
##   <chr>     <dbl>  
## 1 A          20  
## 2 B          25  
## 3 C          NA  
## 4 D          34
```

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# OK
df %>%
  filter(age == 25)

## # A tibble: 1 × 2
##   subject    age
##   <chr>     <dbl>
## 1 B          25
```

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# OK
df %>%
  filter(age == 25)

## # A tibble: 1 × 2
##   subject    age
##   <chr>     <dbl>
## 1 B          25

# Nope
df %>%
  filter(age == NA)

## # A tibble: 0 × 2
## # ... with 2 variables: subject <chr>, age <dbl>

# E.g.
23 == NA

## [1] NA
```

Never test for missingness with ==

Always use **is.na()** instead

```
# Yes
df %>%
  filter(is.na(age))

## # A tibble: 1 × 2
##   subject    age
##   <chr>     <dbl>
## 1 C          NA
```

A quick plug for **naniar** and **visdat**

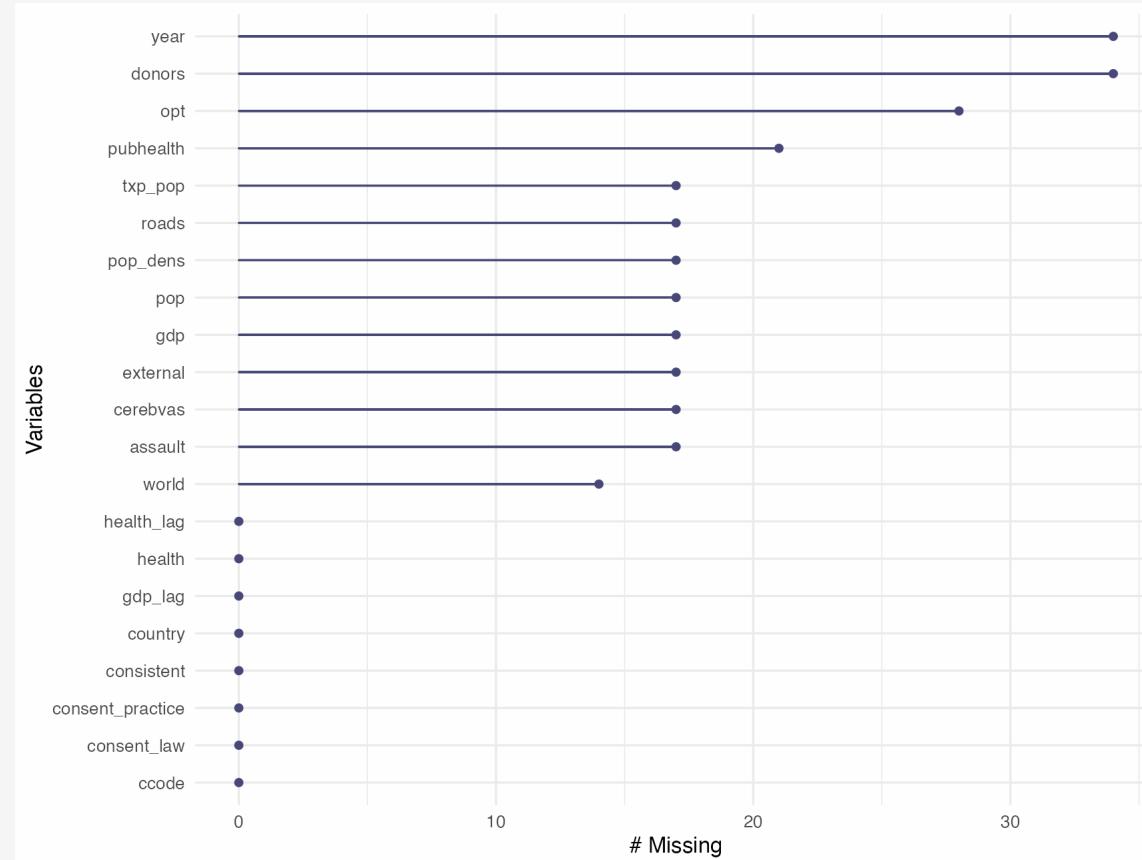
```
library(naniar)
library(visdat)

organdata

## # A tibble: 238 × 21
##   country     year    donors    pop  pop_dens    gdp  gdp_lag  health  health_lag
##   <chr>      <date>   <dbl>   <int>    <dbl>   <int>    <dbl>    <dbl>
## 1 Australia NA        NA     17065    0.220  16774   16591    1300     1224
## 2 Australia 1991-01-01 12.1    17284    0.223  17171   16774    1379     1300
## 3 Australia 1992-01-01 12.4    17495    0.226  17914   17171    1455     1379
## 4 Australia 1993-01-01 12.5    17667    0.228  18883   17914    1540     1455
## 5 Australia 1994-01-01 10.2    17855    0.231  19849   18883    1626     1540
## 6 Australia 1995-01-01 10.2    18072    0.233  21079   19849    1737     1626
## 7 Australia 1996-01-01 10.6    18311    0.237  21923   21079    1846     1737
## 8 Australia 1997-01-01 10.3    18518    0.239  22961   21923    1948     1846
## 9 Australia 1998-01-01 10.5    18711    0.242  24148   22961    2077     1948
## 10 Australia 1999-01-01 8.67   18926    0.244  25445   24148    2231     2077
## # ... with 228 more rows, and 12 more variables: pubhealth <dbl>, roads <dbl>,
## # cerebvas <int>, assault <int>, external <int>, txp_pop <dbl>, world <chr>,
## # opt <chr>, consent_law <chr>, consent_practice <chr>, consistent <chr>,
## # ccode <chr>
```

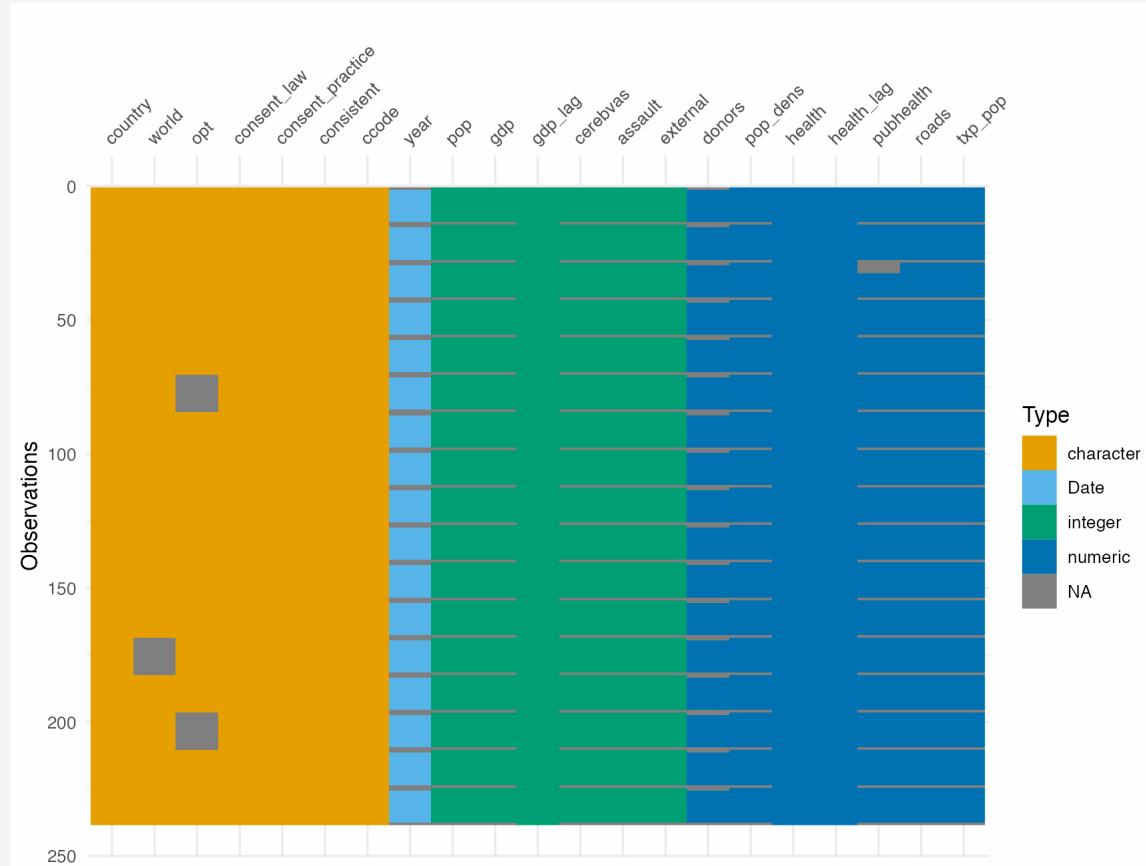
A quick plug for **naniar** and **visdat**

```
gg_miss_var(organdata)
```



A quick plug for **naniar** and **visdat**

```
vis_dat(organdata)
```



A quick plug for **naniar** and **visdat**

```
miss_var_summary(organdata)

## # A tibble: 21 × 3
##   variable n_miss pct_miss
##   <chr>     <int>    <dbl>
## 1 year       34     14.3
## 2 donors     34     14.3
## 3 opt        28     11.8
## 4 pubhealth   21     8.82
## 5 pop         17     7.14
## 6 pop_dens    17     7.14
## 7 gdp         17     7.14
## 8 roads        17     7.14
## 9 cerebvas    17     7.14
## 10 assault     17     7.14
## # ... with 11 more rows
```

A quick plug for **naniar** and **visdat**

```
miss_case_summary(organdata)

## # A tibble: 238 × 3
##   case n_miss pct_miss
##   <int>    <int>     <dbl>
## 1     84      12     57.1
## 2    182      12     57.1
## 3    210      12     57.1
## 4     14      11     52.4
## 5     28      11     52.4
## 6     42      11     52.4
## 7     56      11     52.4
## 8     70      11     52.4
## 9     98      11     52.4
## 10   112      11     52.4
## # ... with 228 more rows
```

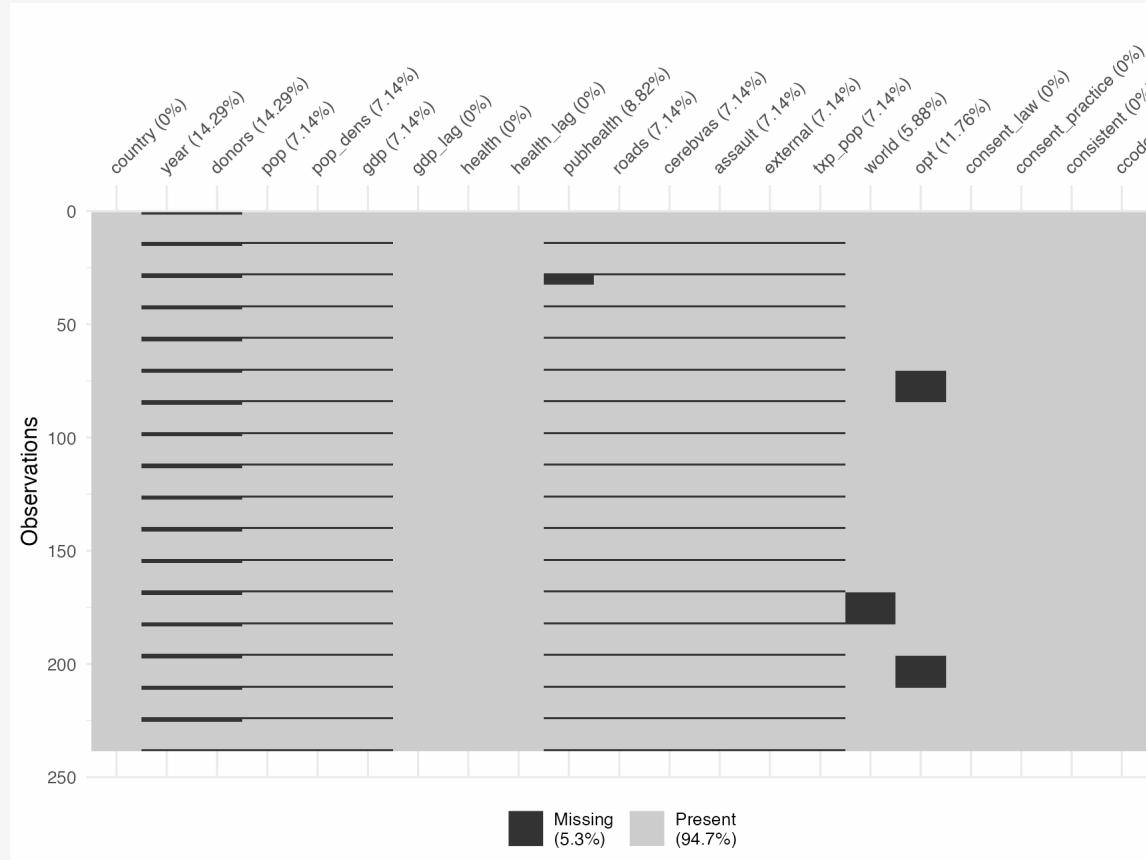
A quick plug for **naniar** and **visdat**

```
organdata %>%
  select(consent_law, year, pubhealth, roads) %>%
  group_by(consent_law) %>%
  miss_var_summary()

## # A tibble: 6 × 4
## # Groups: consent_law [2]
##   consent_law variable n_miss pct_miss
##   <chr>       <chr>     <int>    <dbl>
## 1 Informed    year        16     14.3
## 2 Informed    pubhealth    8      7.14
## 3 Informed    roads        8      7.14
## 4 Presumed    year        18     14.3
## 5 Presumed    pubhealth   13     10.3
## 6 Presumed    roads        9      7.14
```

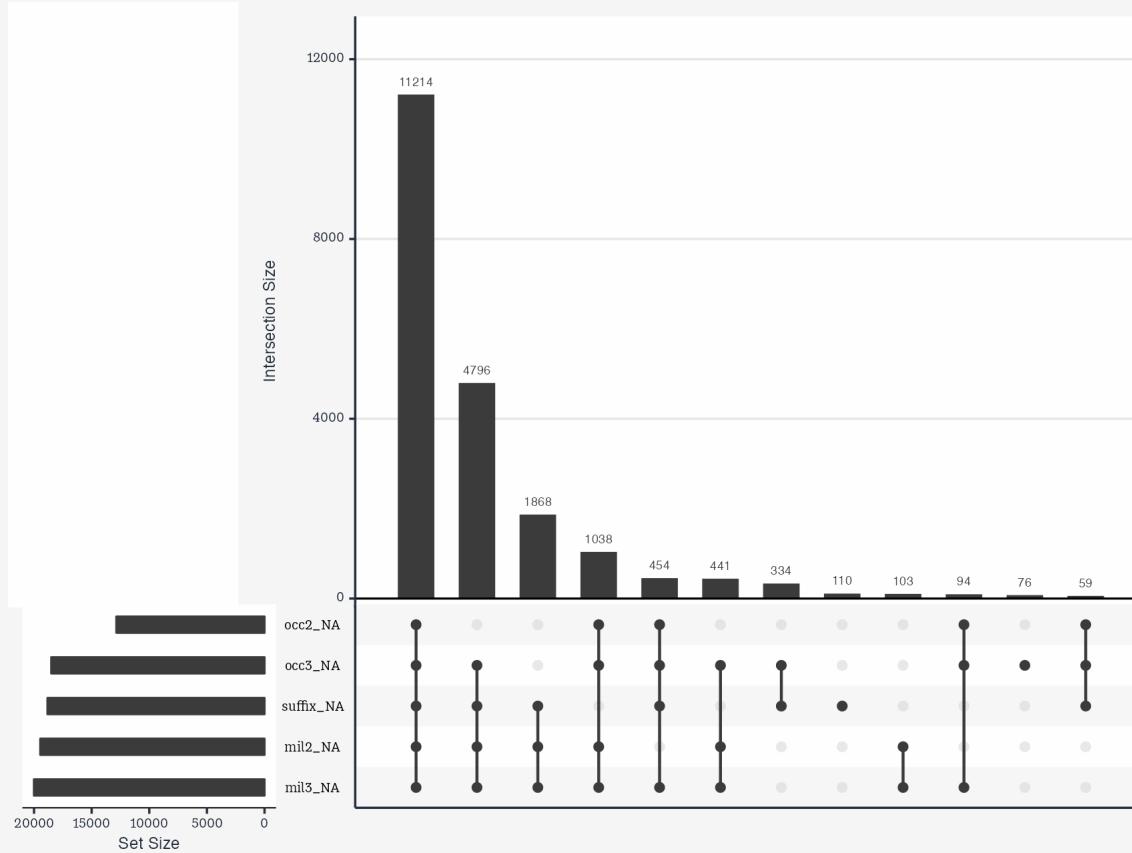
A quick plug for **naniar** and **visdat**

```
vis_miss(organdata)
```



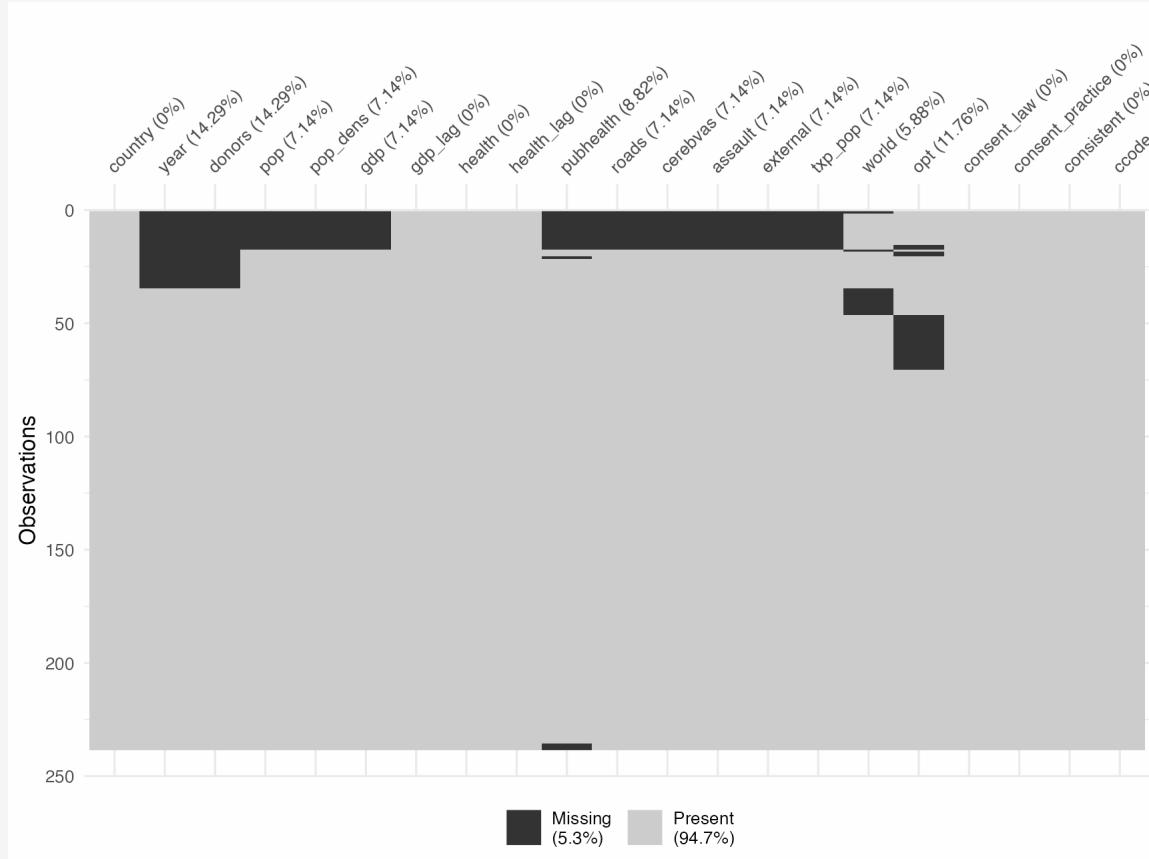
A quick plug for **naniar** and **visdat**

```
library(congress)
gg_miss_upset(congress)
```



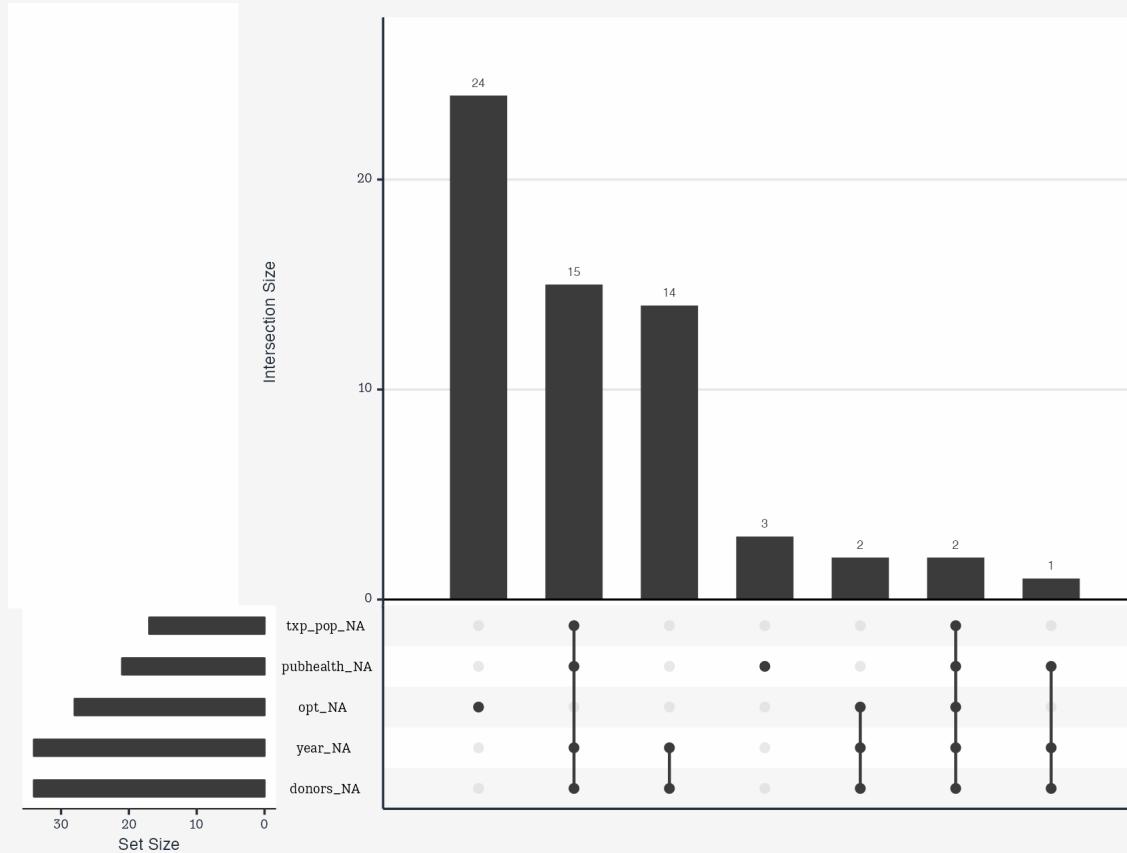
A quick plug for **naniar** and **visdat**

```
vis_miss(organdata, cluster = TRUE)
```



A quick plug for **naniar** and **visdat**

```
gg_miss_upset(organdata)
```

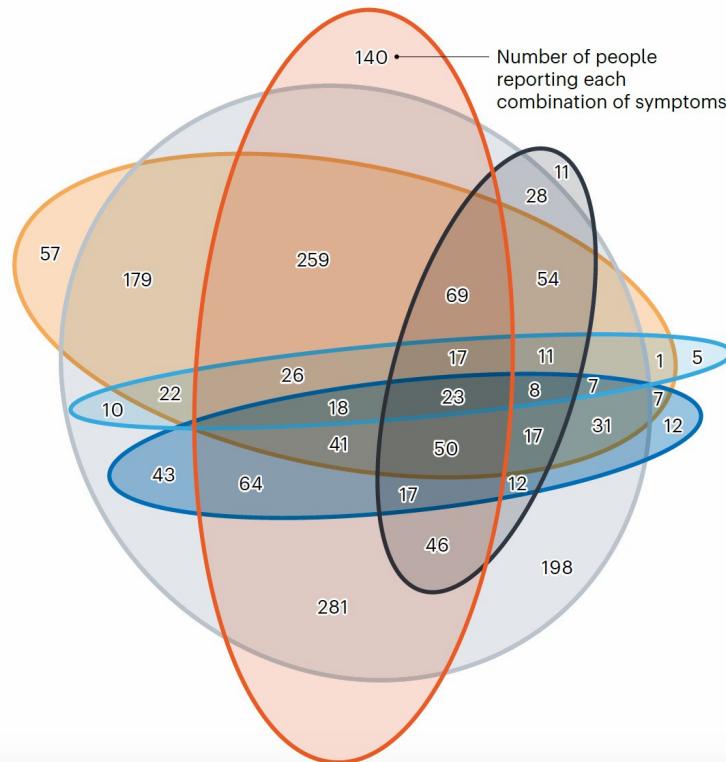


Upset plots and a bit of wrangling

TRACKING SYMPTOMS

On 7 April, around 60% of app users who tested positive for COVID-19 and reported symptoms had lost their sense of smell.

Anosmia (loss of smell) Cough Fatigue
Diarrhoea Shortness of breath Fever



PHOTOGRAPH: PIER MARCO; DATA SOURCE: COVID SYMPTOM TRACKER TEAM

Upset plots and a bit of wrangling

```
symptoms <- c("Anosmia", "Cough", "Fatigue",
             "Diarrhea", "Breath", "Fever")
names(symptoms) <- symptoms
symptoms

##      Anosmia      Cough      Fatigue      Diarrhea      Breath      Fever
## "Anosmia"    "Cough"    "Fatigue"   "Diarrhea"    "Breath"    "Fever"
```

Upset plots and a bit of wrangling

```
# An Excel file!
dat <- readxl::read_xlsx(here("data", "symptoms.xlsx"))
dat %>% print(n = nrow(dat))
```

```
## # A tibble: 32 × 2
##   combination           count
##   <chr>                <dbl>
## 1 Anosmia                  140
## 2 Cough                      57
## 3 Fatigue                   198
## 4 Diarrhea                   12
## 5 Breath                      5
## 6 Fever                      11
## 7 Cough&Fatigue              179
## 8 Fatigue&Fever                 28
## 9 Breath&Fatigue                 10
## 10 Diarrhea&Fatigue               43
## 11 Anosmia&Fatigue              281
## 12 Breath&Cough                   1
## 13 Anosmia&Diarrhea&Fatigue            64
## 14 Breath&Cough&Fatigue                22
## 15 Anosmia&Cough&Fatigue              259
## 16 Anosmia&Fever&Fatigue                 46
## 17 Cough&Fever&Fatigue                  54
## 18 Cough&Diarrhea                     7
## 19 Cough&Diarrhea&Fatigue                31
## 20 Anosmia&Breath&Cough&Fatigue            26
## 21 Anosmia&Cough&Fatigue&Fever             69
## 22 Anosmia&Breath&Cough&Diarrhea&Fatigue          18
## 23 Anosmia&Breath&Cough&Fatigue&Fever            17
```

Upset plots and a bit of wrangling

```
subsets <- dat %>%
  pull(combination)

## Check if each subset mentions each symptom or not
symptom_mat <- map_dfc(subsets, str_detect, symptoms) %>%
  data.frame() %>%
  t() %>% # transpose the result, this is a little gross, sorry
  as_tibble(.name_repair = "unique")

colnames(symptom_mat) <- symptoms
symptom_mat$count <- dat$count
```

Upset plots and a bit of wrangling

Now we have a table we can do something with.

```
symptom_mat %>% print(n = nrow(symptom_mat))

## # A tibble: 32 × 7
##   Anosmia Cough Fatigue Diarrhea Breath Fever count
##   <lgl>    <lgl> <lgl>    <lgl>   <lgl> <dbl>
## 1 TRUE     FALSE  FALSE    FALSE   FALSE  140
## 2 FALSE    TRUE   FALSE    FALSE   FALSE  57
## 3 FALSE    FALSE  TRUE     FALSE   FALSE  198
## 4 FALSE    FALSE  FALSE    TRUE    FALSE  12
## 5 FALSE    FALSE  FALSE    FALSE   TRUE   5
## 6 FALSE    FALSE  FALSE    FALSE   FALSE  TRUE  11
## 7 FALSE    TRUE   TRUE     FALSE   FALSE  FALSE 179
## 8 FALSE    FALSE  TRUE     FALSE   FALSE  TRUE  28
## 9 FALSE    FALSE  TRUE     FALSE   TRUE   FALSE 10
## 10 FALSE   FALSE  TRUE     TRUE    FALSE  FALSE 43
## 11 TRUE    FALSE  TRUE     FALSE   FALSE  FALSE 281
## 12 FALSE   TRUE   FALSE    FALSE   TRUE   FALSE 1
## 13 TRUE    FALSE  TRUE     TRUE    FALSE  FALSE 64
## 14 FALSE   TRUE   TRUE     FALSE   TRUE   FALSE 22
## 15 TRUE    TRUE   TRUE     FALSE   FALSE  FALSE 259
## 16 TRUE    FALSE  TRUE     FALSE   FALSE  TRUE  46
## 17 FALSE   TRUE   TRUE     FALSE   FALSE  TRUE  54
## 18 FALSE   TRUE   FALSE    TRUE    FALSE  FALSE 7
## 19 FALSE   TRUE   TRUE     TRUE    FALSE  FALSE 31
## 20 TRUE    TRUE   TRUE     FALSE   TRUE   FALSE 26
## 21 TRUE    TRUE   TRUE     FALSE   FALSE  TRUE  69
## 22 TRUE    TRUE   TRUE     TRUE    TRUE   FALSE 18
## 23 TRUE    TRUE   TRUE     FALSE   TRUE   TRUE  17
```

Upset plots and a bit of wrangling

Uncounting tables

```
indvs <- symptom_mat %>%
  uncount(count)

indvs

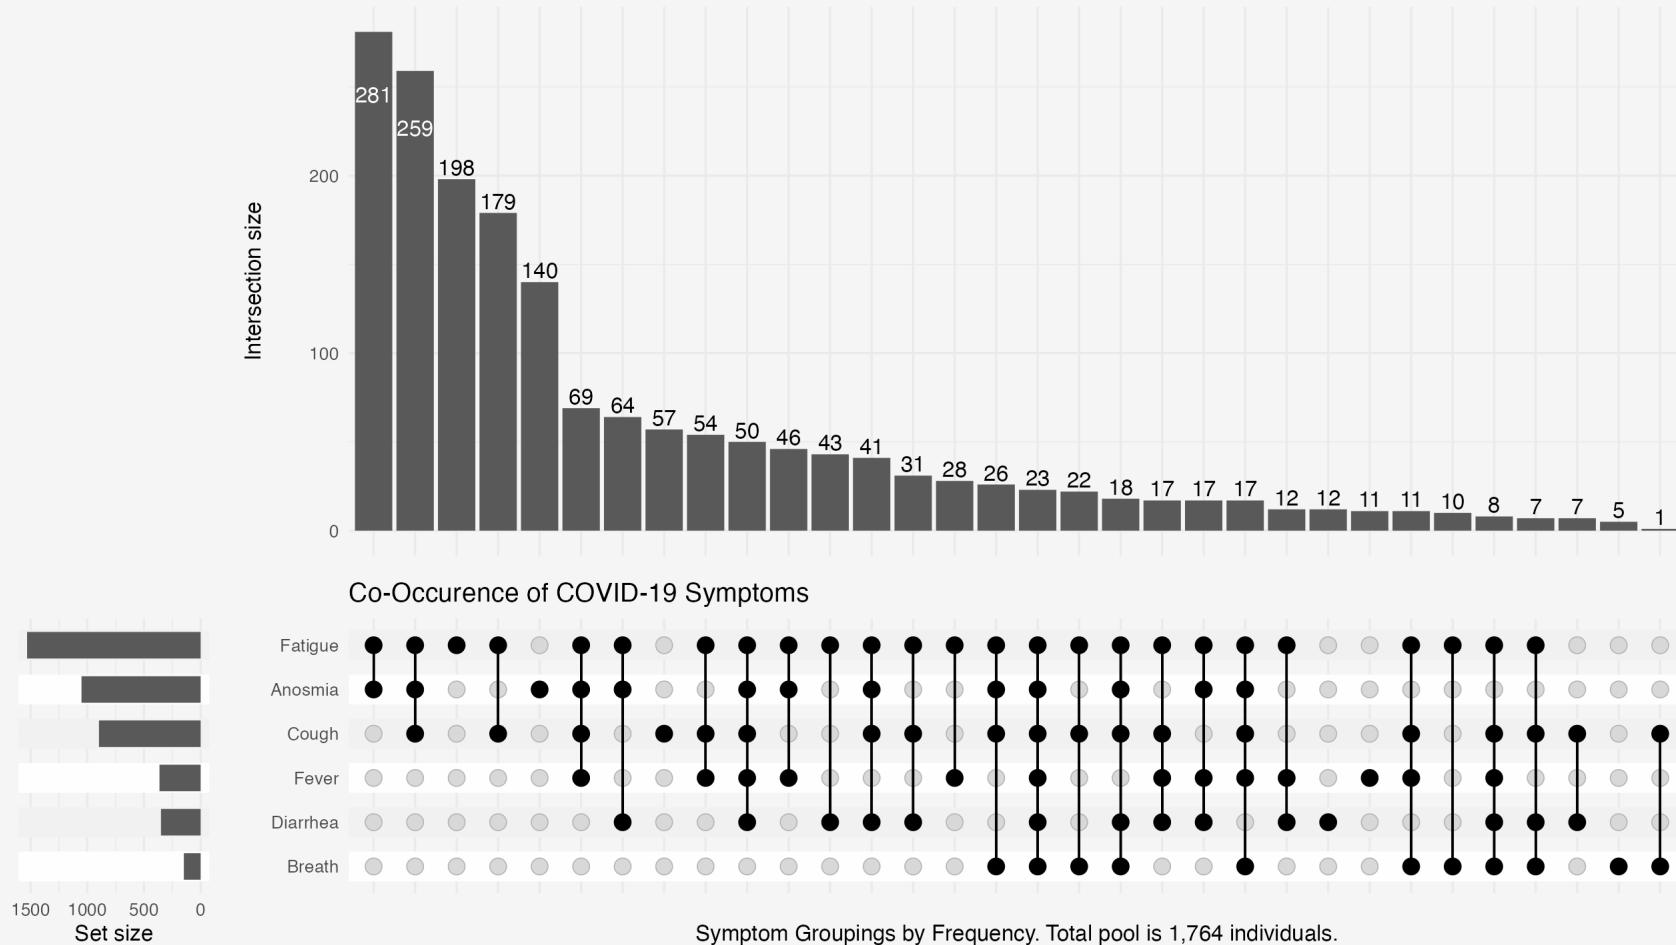
## # A tibble: 1,764 × 6
##   Anosmia Cough Fatigue Diarrhea Breath Fever
##   <lgl>    <lgl> <lgl>    <lgl>    <lgl>
## 1 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 2 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 3 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 4 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 5 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 6 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 7 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 8 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 9 TRUE     FALSE  FALSE    FALSE    FALSE  FALSE
## 10 TRUE    FALSE  FALSE    FALSE    FALSE  FALSE
## # ... with 1,754 more rows
```

Now we've reconstructed the individual-level observations.

Upset plots and a bit of wrangling

```
# devtools::install_github("krassowski/complex-upset")  
  
library(ComplexUpset)  
  
upset(data = indvs, intersect = symptoms,  
      name="Symptom Groupings by Frequency. Total pool is 1,764 individuals.",  
      min_size = 0,  
      width_ratio = 0.125) +  
  labs(title = "Co-Occurrence of COVID-19 Symptoms",  
       caption = "Data: covid.joinzoe.com/us | Graph: @kjhealy")
```

Upset plots and a bit of wrangling



Models

This is not a **statistics** seminar!

I'll just give you an example of the sort of thing that many other modeling packages implement for all kinds of modeling techniques.

Again, the principle is tidy incorporation of models and their output.

Tidy regression output with **broom**

```
library(broom)
library(gapminder)

out <- lm(formula = lifeExp ~ gdpPercap + pop + continent,
          data = gapminder)
```

Tidy regression output with broom

We can't *do* anything with this, programatically.

```
summary(out)

##
## Call:
## lm(formula = lifeExp ~ gdpPercap + pop + continent, data = gapminder)
##
## Residuals:
##     Min      1Q  Median      3Q     Max 
## -49.161  -4.486   0.297   5.110  25.175 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept) 4.781e+01 3.395e-01 140.819 < 2e-16 ***
## gdpPercap   4.495e-04 2.346e-05 19.158 < 2e-16 ***
## pop         6.570e-09 1.975e-09  3.326 0.000901 *** 
## continentAmericas 1.348e+01 6.000e-01 22.458 < 2e-16 ***
## continentAsia    8.193e+00 5.712e-01 14.342 < 2e-16 ***
## continentEurope   1.747e+01 6.246e-01 27.973 < 2e-16 ***
## continentOceania  1.808e+01 1.782e+00 10.146 < 2e-16 *** 
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.365 on 1697 degrees of freedom
## Multiple R-squared:  0.5821,    Adjusted R-squared:  0.5806 
## F-statistic: 393.9 on 6 and 1697 DF,  p-value: < 2.2e-16
```

Tidy regression output with **broom**

```
library(broom)

tidy(out)

## # A tibble: 7 × 5
##   term      estimate std.error statistic p.value
##   <chr>     <dbl>     <dbl>     <dbl>     <dbl>
## 1 (Intercept) 4.78e+1  0.340     141.     0
## 2 gdpPercap    4.50e-4  0.0000235    19.2    3.24e- 74
## 3 pop         6.57e-9  0.00000000198    3.33   9.01e- 4
## 4 continentAmericas 1.35e+1  0.600     22.5    5.19e- 98
## 5 continentAsia  8.19e+0  0.571     14.3    4.06e- 44
## 6 continentEurope 1.75e+1  0.625     28.0    6.34e-142
## 7 continentOceania 1.81e+1  1.78      10.1   1.59e- 23
```

That's a *lot* nicer. Now it's just a tibble. We know those.

Tidy regression output with broom

```
out_conf <- tidy(out, conf.int = TRUE)
out_conf

## # A tibble: 7 × 7
##   term      estimate    std.error statistic   p.value conf.low conf.high
##   <chr>     <dbl>        <dbl>     <dbl>      <dbl>    <dbl>       <dbl>
## 1 (Intercept) 4.78e+1    3.40e-1    141.     0        4.71e+1    4.85e+1
## 2 gdpPercap   4.50e-4    2.35e-5    19.2    3.24e-74   4.03e-4    4.96e-4
## 3 pop        6.57e-9    1.98e-9    3.33    9.01e- 4   2.70e-9    1.04e-8
## 4 continentAmericas 1.35e+1    6.00e-1    22.5    5.19e- 98  1.23e+1    1.47e+1
## 5 continentAsia  8.19e+0    5.71e-1    14.3    4.06e- 44  7.07e+0    9.31e+0
## 6 continentEurope 1.75e+1    6.25e-1    28.0    6.34e-142  1.62e+1    1.87e+1
## 7 continentOceania 1.81e+1    1.78e+0    10.1    1.59e- 23  1.46e+1    2.16e+1
```

Tidy regression output with broom

```
out_conf %>%  
  filter(term %nin% "(Intercept)") %>%  
  mutate(nicelabs = prefix_strip(term, "continent")) %>%  
  select(nicelabs, everything())  
  
## # A tibble: 6 × 8  
##   nicelabs term      estimate std.error statistic  p.value conf.low conf.high  
##   <chr>     <chr>      <dbl>     <dbl>      <dbl>     <dbl>    <dbl>    <dbl>  
## 1 gdpPercap gdpPercap  4.50e-4  2.35e-5     19.2  3.24e- 74  4.03e-4  4.96e-4  
## 2 Pop       pop        6.57e-9  1.98e-9     3.33  9.01e- 4   2.70e-9  1.04e-8  
## 3 Americas   continent... 1.35e+1  6.00e-1     22.5  5.19e- 98  1.23e+1  1.47e+1  
## 4 Asia      continent...  8.19e+0  5.71e-1     14.3  4.06e- 44  7.07e+0  9.31e+0  
## 5 Europe    continent...  1.75e+1  6.25e-1     28.0  6.34e-142 1.62e+1  1.87e+1  
## 6 Oceania   continent...  1.81e+1  1.78e+0     10.1  1.59e- 23  1.46e+1  2.16e+1
```

Grouped analysis and **list columns**

```
eu77 <- gapminder %>% filter(continent == "Europe", year == 1977)
fit <- lm(lifeExp ~ log(gdpPercap), data = eu77)
```

```
summary(fit)
```

```
##
## Call:
## lm(formula = lifeExp ~ log(gdpPercap), data = eu77)
##
## Residuals:
##     Min      1Q      Median      3Q      Max
## -7.4956 -1.0306  0.0935  1.1755  3.7125
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 29.489     7.161   4.118 0.000306 ***
## log(gdpPercap) 4.488     0.756   5.936 2.17e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.114 on 28 degrees of freedom
## Multiple R-squared:  0.5572,    Adjusted R-squared:  0.5414
## F-statistic: 35.24 on 1 and 28 DF,  p-value: 2.173e-06
```

Grouped analysis and **list columns**

```
out_le <- gapminder %>%
  group_by(continent, year) %>%
  nest()

out_le

## # A tibble: 60 × 3
## # Groups:   continent, year [60]
##   continent   year     data
##   <fct>      <int> <list>
## 1 Asia        1952 <tibble [33 × 4]>
## 2 Asia        1957 <tibble [33 × 4]>
## 3 Asia        1962 <tibble [33 × 4]>
## 4 Asia        1967 <tibble [33 × 4]>
## 5 Asia        1972 <tibble [33 × 4]>
## 6 Asia        1977 <tibble [33 × 4]>
## 7 Asia        1982 <tibble [33 × 4]>
## 8 Asia        1987 <tibble [33 × 4]>
## 9 Asia        1992 <tibble [33 × 4]>
## 10 Asia       1997 <tibble [33 × 4]>
## # ... with 50 more rows
```

Think of nesting as a kind of "super-grouping". Look in the object inspector.

Grouped analysis and **list columns**

It's still in there.

```
out_le %>% filter(continent == "Europe" & year == 1977) %>%
  unnest(cols = c(data))

## # A tibble: 30 × 6
## # Groups:   continent, year [1]
##   continent  year country      lifeExp     pop gdpPercap
##   <fct>     <int> <fct>      <dbl>    <int>     <dbl>
## 1 Europe     1977 Albania     68.9  2509048     3533.
## 2 Europe     1977 Austria     72.2  7568430     19749.
## 3 Europe     1977 Belgium     72.8  9821800     19118.
## 4 Europe     1977 Bosnia and Herzegovina 69.9  4086000     3528.
## 5 Europe     1977 Bulgaria    70.8  8797022     7612.
## 6 Europe     1977 Croatia     70.6  4318673     11305.
## 7 Europe     1977 Czech Republic 70.7  10161915    14800.
## 8 Europe     1977 Denmark     74.7  5088419     20423.
## 9 Europe     1977 Finland     72.5  4738902     15605.
## 10 Europe    1977 France      73.8  53165019    18293.
## # ... with 20 more rows
```

Grouped analysis and **list columns**

Here we **map()** a custom function to every row in the data column.

```
fit_ols <- function(df) {  
  lm(lifeExp ~ log(gdpPerCap), data = df)  
}  
  
out_le <- gapminder %>%  
  group_by(continent, year) %>%  
  nest() %>%  
  mutate(model = map(data, fit_ols))
```

Grouped analysis and **list** columns

out_le

```
## # A tibble: 60 × 4
## # Groups: continent, year [60]
##   continent year data          model
##   <fct>     <int> <list>        <list>
## 1 Asia       1952 <tibble [33 × 4]> <lm>
## 2 Asia       1957 <tibble [33 × 4]> <lm>
## 3 Asia       1962 <tibble [33 × 4]> <lm>
## 4 Asia       1967 <tibble [33 × 4]> <lm>
## 5 Asia       1972 <tibble [33 × 4]> <lm>
## 6 Asia       1977 <tibble [33 × 4]> <lm>
## 7 Asia       1982 <tibble [33 × 4]> <lm>
## 8 Asia       1987 <tibble [33 × 4]> <lm>
## 9 Asia       1992 <tibble [33 × 4]> <lm>
## 10 Asia      1997 <tibble [33 × 4]> <lm>
## # ... with 50 more rows
```

Grouped analysis and **list columns**

We can tidy the nested models, too.

```
fit_ols <- function(df) {  
  lm(lifeExp ~ log(gdpPerCap), data = df)  
}  
  
out_tidy <- gapminder %>%  
  group_by(continent, year) %>%  
  nest() %>%  
  mutate(model = map(data, fit_ols),  
        tidied = map(model, tidy)) %>%  
  unnest(cols = c(tidied)) %>%  
  filter(term %nin% "(Intercept)" &  
        continent %nin% "Oceania")
```

Grouped analysis and **list columns**

out_tidy

```
## # A tibble: 48 × 9
## # Groups: continent, year [48]
##   continent year data    model term      estimate std.error statistic p.value
##   <fct>     <int> <list> <list> <chr>      <dbl>     <dbl>     <dbl>     <dbl>
## 1 Asia       1952 <tibble> <lm> log(gdp...  4.16      1.25      3.33  2.28e-3
## 2 Asia       1957 <tibble> <lm> log(gdp...  4.17      1.28      3.26  2.71e-3
## 3 Asia       1962 <tibble> <lm> log(gdp...  4.59      1.24      3.72  7.94e-4
## 4 Asia       1967 <tibble> <lm> log(gdp...  4.50      1.15      3.90  4.77e-4
## 5 Asia       1972 <tibble> <lm> log(gdp...  4.44      1.01      4.41  1.16e-4
## 6 Asia       1977 <tibble> <lm> log(gdp...  4.87      1.03      4.75  4.42e-5
## 7 Asia       1982 <tibble> <lm> log(gdp...  4.78      0.852     5.61  3.77e-6
## 8 Asia       1987 <tibble> <lm> log(gdp...  5.17      0.727     7.12  5.31e-8
## 9 Asia       1992 <tibble> <lm> log(gdp...  5.09      0.649     7.84  7.60e-9
## 10 Asia      1997 <tibble> <lm> log(gdp...  5.11      0.628     8.15  3.35e-9
## # ... with 38 more rows
```

Grouped analysis and **list** columns

```
out_tidy %>%  
  ungroup() %>%  
  sample_n(5)  
  
## # A tibble: 5 × 9  
##   continent year data    model term      estimate std.error statistic p.value  
##   <fct>     <int> <list> <list> <chr>       <dbl>     <dbl>     <dbl>     <dbl>  
## 1 Europe      2002 <tibble> <lm> log(gdpP... 3.74      0.445      8.40 3.91e-9  
## 2 Europe      1992 <tibble> <lm> log(gdpP... 3.48      0.545      6.39 6.44e-7  
## 3 Africa      1977 <tibble> <lm> log(gdpP... 4.51      0.920      4.90 1.04e-5  
## 4 Africa      2007 <tibble> <lm> log(gdpP... 4.26      1.19       3.58 7.79e-4  
## 5 Americas    1952 <tibble> <lm> log(gdpP... 10.4      2.72       3.84 8.27e-4
```