

Iterating on Data

Data Wrangling, Session 7

Kieran Healy

Code Horizons

April 2025

Iterating on data with purrr and map

Load the packages, as always

```
library(here)      # manage file paths  
library(socviz)    # data and some useful functions  
library(tidyverse) # your friend and mine
```

Moar Data

More than one data file

Inside the `data/` folder of the course packet is a folder named `congress/`

```
# A little trick from the fs package:  
fs::dir_tree(here("data", "congress"))
```

```
/Users/kjhealy/Documents/courses/data_wrangling/data/congress
```

```
├── 01_79_congress.csv  
├── 02_80_congress.csv  
├── 03_81_congress.csv  
├── 04_82_congress.csv  
├── 05_83_congress.csv  
├── 06_84_congress.csv  
├── 07_85_congress.csv  
├── 08_86_congress.csv  
├── 09_87_congress.csv  
├── 10_88_congress.csv  
├── 11_89_congress.csv  
├── 12_90_congress.csv  
├── 13_91_congress.csv  
├── 14_92_congress.csv  
├── 15_93_congress.csv  
├── 16_94_congress.csv  
├── 17_95_congress.csv  
├── 18_96_congress.csv  
├── 19_97_congress.csv  
├── 20_98_congress.csv  
├── 21_99_congress.csv  
├── 22_100_congress.csv  
├── 23_101_congress.csv  
└── 24_102_congress.csv
```

More than one data file

Let's look at one.

```
read_csv(here("data", "congress", "17_95_congress.csv")) ▷  
janitor::clean_names() ▷  
head()
```

```
# A tibble: 6 × 25  
  last      first  middle suffix nickname born  death sex  position party state  
  <chr>   <chr>   <chr>  <chr>   <chr>   <chr> <chr> <chr>  <chr>   <chr> <chr>  
1 Abdnor   James  <NA>   <NA>   <NA>    02/1... 11/0... M    U.S. Re... Repu... SD  
2 Abourezk James  George <NA>   <NA>    02/2... <NA>   M    U.S. Se... Demo... SD  
3 Adams    Brockm... <NA>   <NA>   Brock   01/1... 09/1... M    U.S. Re... Demo... WA  
4 Addabbo  Joseph  Patri... <NA>   <NA>    03/1... 04/1... M    U.S. Re... Demo... NY  
5 Aiken    George  David  <NA>   <NA>    08/2... 11/1... M    U.S. Se... Repu... VT  
6 Akaka    Daniel  Kahik... <NA>   <NA>    09/1... 04/0... M    U.S. Re... Demo... HI  
# i 14 more variables: district <chr>, start <chr>, end <chr>, religion <chr>,  
#   race <chr>, educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,  
#   job_type3 <chr>, job_type4 <chr>, job_type5 <lgl>, mil1 <chr>, mil2 <chr>,  
#   mil3 <chr>
```

We often find ourselves in this situation. We know each file has the same structure, and we would like to use them all at once.

Loops?

How to read them all in?

One traditional way, which we could do in R, is to write an explicit *loop* that iterated over a vector of filenames, read each file, and then joined the results together in a tall rectangle.

```
# Pseudocode

filenames ← c("01_79_congress.csv", "02_80_congress.csv", "03_81_congress.csv",
              "04_82_congress.csv" [etc etc])

collected_files ← NULL

for(i in 1:length(filenames)) {
  new_file ← read_file(filenames[i])
  collected_files ← append_to(collected_files, new_files)
}
```

Loops?

Mapping is just a kind of iteration

Vectorized arithmetic again

The simplest cases are not that different from the vectorized arithmetic we're already familiar with.

```
a ← c(1:10)
```

```
b ← 1
```

```
# You know what R will do here
```

```
a + b
```

```
[1] 2 3 4 5 6 7 8 9 10 11
```

Vectorized arithmetic again

We can make this explicit by writing a function:

```
add_b ← function(x) {  
  b ← 1  
  x + b # for any x  
}
```

Vectorized arithmetic again

We can make this explicit by writing a function:

```
add_b ← function(x) {  
  b ← 1  
  x + b # for any x  
}
```

Now:

```
add_b(x = a)
```

```
[1] 2 3 4 5 6 7 8 9 10 11
```

Vectorized arithmetic again

Again, R's vectorized approach means it automatically adds **b** to every element of the **x** we give it.

```
add_b(x = 10)
```

```
[1] 11
```

```
add_b(x = c(1, 99, 1000))
```

```
[1]    2  100 1001
```

Iterating in a pipeline

Some operations can't directly be vectorized in this way, which is why we need to manually iterate, or will want to write loops.

```
library(gapminder)
gapminder >
  summarize(country_n = n_distinct(country),
            continent_n = n_distinct(continent),
            year_n = n_distinct(year),
            lifeExp_n = n_distinct(lifeExp),
            population_n = n_distinct(population))
```

```
# A tibble: 1 × 5
  country_n continent_n year_n lifeExp_n population_n
  <int>      <int>    <int>    <int>      <int>
1      142         5      12     1626      4060
```

That's tedious to write! Computers are supposed to allow us to avoid that sort of thing.

Iterating in a pipeline

So how would we iterate this? What we want is to apply the `n_distinct()` function to each column of `gapminder`, but in a way that still allows us to use pipelines and so on.

```
library(gapminder)
gapminder >
  summarize(n_distinct(country),
            n_distinct(continent),
            n_distinct(year),
            n_distinct(lifeExp),
            n_distinct(population))

# A tibble: 1 × 5
#   `n_distinct(country)` `n_distinct(continent)` `n_distinct(year)`
#   <int>                <int>                <int>
# 1         142           5                    12
# i 2 more variables: `n_distinct(lifeExp)` <int>,
#   `n_distinct(population)` <int>
```

Using `n_distinct()` in this context is an idea I got from Rebecca Barter's discussion of `purrr`.

Iterating in a pipeline

You'd use `across()`, like this:

```
gapminder ►  
  summarize(across(everything(), n_distinct))
```

```
# A tibble: 1 × 6  
  country continent  year lifeExp   pop gdpPercap  
  <int>      <int> <int>   <int> <int>    <int>  
1    142         5    12    1626  1704    1704
```


Iterating in a pipeline

But you could also do this ...

```
map(gapminder, n_distinct)
```

```
$country
```

```
[1] 142
```

```
$continent
```

```
[1] 5
```

```
$year
```

```
[1] 12
```

```
$lifeExp
```

```
[1] 1626
```

```
$pop
```

```
[1] 1704
```

```
$gdpPercap
```

```
[1] 1704
```

Read it as “Feed each column of `gapminder` to the `n_distinct()` function.

(This is pretty much what `across()` is doing more nicely.)

Iterating in a pipeline

Or, in pipeline form:

You can see we are getting a *list* back.

```
gapminder ►  
  map(n_distinct)
```

```
$country  
[1] 142
```

```
$continent  
[1] 5
```

```
$year  
[1] 12
```

```
$lifeExp  
[1] 1626
```

```
$pop  
[1] 1704
```

```
$gdpPercap  
[1] 1704
```

Iterating in a pipeline

Or, in pipeline form:

```
result <- gapminder >  
  map(n_distinct)  
  
class(result)
```

```
[1] "list"
```

```
result$continent
```

```
[1] 5
```

```
result[[2]]
```

```
[1] 5
```

Iterating in a pipeline

But we know `n_distinct()` should always return an integer. So we use `map_int()` instead of the generic `map()`.

```
gapminder ►  
  map_int(n_distinct)
```

country	continent	year	lifeExp	pop	gdpPercap
142	5	12	1626	1704	1704

Get a vector of **filenames**

```
filenames ← dir(path = here("data", "congress"),  
                pattern = "*.csv",  
                full.names = TRUE)
```

```
filenames[1:15] # Just displaying the first 15, to save slide space
```

```
[1] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/01_79_congress.csv"  
[2] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/02_80_congress.csv"  
[3] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/03_81_congress.csv"  
[4] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/04_82_congress.csv"  
[5] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/05_83_congress.csv"  
[6] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/06_84_congress.csv"  
[7] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/07_85_congress.csv"  
[8] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/08_86_congress.csv"  
[9] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/09_87_congress.csv"  
[10] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/10_88_congress.csv"  
[11] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/11_89_congress.csv"  
[12] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/12_90_congress.csv"  
[13] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/13_91_congress.csv"  
[14] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/14_92_congress.csv"  
[15] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/15_93_congress.csv"
```

And feed it to `read_csv()`

... using `map()` and binding the resulting list into a tibble.

```
df <- filenames >
map(read_csv) >
list_rbind(names_to = "congress") >
janitor::clean_names()
```

df

A tibble: 20,580 × 26

	congress	last	first	middle	suffix	nickname	born	death	sex	position	party
	<int>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1	1	Abern...	Thom...	Gerst...	<NA>	<NA>	05/1...	01/2...	M	U.S. Re...	Demo...
2	1	Adams	Sher...	<NA>	<NA>	<NA>	01/0...	10/2...	M	U.S. Re...	Repu...
3	1	Aiken	Geor...	David	<NA>	<NA>	08/2...	11/1...	M	U.S. Se...	Repu...
4	1	Allen	Asa	Leona...	<NA>	<NA>	01/0...	01/0...	M	U.S. Re...	Demo...
5	1	Allen	Leo	Elwood	<NA>	<NA>	10/0...	01/1...	M	U.S. Re...	Repu...
6	1	Almond	J.	Linds...	Jr.	<NA>	06/1...	04/1...	M	U.S. Re...	Demo...
7	1	Ander...	Herm...	Carl	<NA>	<NA>	01/2...	07/2...	M	U.S. Re...	Repu...
8	1	Ander...	Clin...	Presba	<NA>	<NA>	10/2...	11/1...	M	U.S. Re...	Demo...
9	1	Ander...	John	Zuing...	<NA>	<NA>	03/2...	02/0...	M	U.S. Re...	Repu...
10	1	Andre...	Augu...	Herman	<NA>	<NA>	10/1...	01/1...	M	U.S. Re...	Repu...

i 20,570 more rows

i 15 more variables: state <chr>, district <chr>, start <chr>, end <chr>,
religion <chr>, race <chr>, educational_attainment <chr>, job_type1 <chr>,
job_type2 <chr>, job_type3 <chr>, job_type4 <chr>, job_type5 <chr>,
mil1 <chr>, mil2 <chr>, mil3 <chr>

Now witness the firepower of this fully armed and operational



method of type-safe functional iteration

read_csv() can do this directly

In fact `map()` is not required for this particular use:

```
tmp ← read_csv(filenames, id = "path",
               name_repair = janitor::make_clean_names)

tmp ▷
  mutate(congress = str_extract(path, "\\d{2,3}_congress"),
         congress = str_extract(congress, "\\d{2,3}")) ▷
  relocate(congress)
```

```
# A tibble: 20,580 × 27
  congress path    last first middle suffix nickname born death sex position
  <chr>    <chr>  <chr> <chr> <chr>  <chr>  <chr>  <chr> <chr> <chr> <chr>
1 79      /User... Aber... Thom... Gerst... <NA>  <NA>  05/1... 01/2... M    U.S. Re...
2 79      /User... Adams Sher... <NA>  <NA>  <NA>  01/0... 10/2... M    U.S. Re...
3 79      /User... Aiken Geor... David  <NA>  <NA>  08/2... 11/1... M    U.S. Se...
4 79      /User... Allen Asa   Leona... <NA>  <NA>  01/0... 01/0... M    U.S. Re...
5 79      /User... Allen Leo   Elwood  <NA>  <NA>  10/0... 01/1... M    U.S. Re...
6 79      /User... Almo... J.     Linds... Jr.    <NA>  06/1... 04/1... M    U.S. Re...
7 79      /User... Ande... Herm... Carl   <NA>  <NA>  01/2... 07/2... M    U.S. Re...
8 79      /User... Ande... Clin... Presba <NA>  <NA>  10/2... 11/1... M    U.S. Re...
9 79      /User... Ande... John  Zuing... <NA>  <NA>  03/2... 02/0... M    U.S. Re...
10 79     /User... Andr... Augu... Herman <NA>  <NA>  10/1... 01/1... M    U.S. Re...

# i 20,570 more rows
# i 16 more variables: party <chr>, state <chr>, district <chr>, start <chr>,
#   end <chr>, religion <chr>, race <chr>, educational_attainment <chr>,
#   job_type1 <chr>, job_type2 <chr>, job_type3 <chr>, job_type4 <chr>,
#   job_type5 <chr>, mil1 <chr>, mil2 <chr>, mil3 <chr>
```


Example: Iterating on the US Census

Iterating on the US Census

Mapped iteration is very general, and not just for local files

```
## Register for a free Census API key
library(tidycensus)

out <- get_acs(geography = "county",
               variables = "B19013_001",
               state = "NY",
               county = "New York",
               survey = "acs1",
               year = 2005)
```

out

```
# A tibble: 1 × 5
  GEOID NAME          variable estimate moe
  <chr> <chr>          <chr>      <dbl> <dbl>
1 36061 New York County, New York B19013_001 55973 1462
```

Iterating on the US Census

All counties in New York State for a specific year

```
out ← get_acs(geography = "county",
               variables = "B19013_001",
               state = "NY",
               survey = "acs1",
               year = 2005)
```

out

A tibble: 38 × 5

	GEOID	NAME	variable	estimate	moe
	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	36001	Albany County, New York	B19013_001	50054	2030
2	36005	Bronx County, New York	B19013_001	29228	853
3	36007	Broome County, New York	B19013_001	36394	2340
4	36009	Cattaraugus County, New York	B19013_001	37580	2282
5	36011	Cayuga County, New York	B19013_001	42057	2406
6	36013	Chautauqua County, New York	B19013_001	35495	2077
7	36015	Chemung County, New York	B19013_001	37418	3143
8	36019	Clinton County, New York	B19013_001	44757	3500
9	36027	Dutchess County, New York	B19013_001	61889	2431
10	36029	Erie County, New York	B19013_001	41967	1231

i 28 more rows

Iterating on the US Census

What if we want the results for *every* available year? First, a handy function: `set_names()`

```
x ← c(1:10)
```

```
x
```

```
[1] 1 2 3 4 5 6 7 8 9 10
```

```
x ← set_names(x, nm = letters[1:10])
```

```
x
```

```
a b c d e f g h i j  
1 2 3 4 5 6 7 8 9 10
```

Iterating on the US Census

By default, `set_names()` will label a vector with that vector's values:

```
c(1:10) ►  
  set_names()
```

```
1 2 3 4 5 6 7 8 9 10  
1 2 3 4 5 6 7 8 9 10
```

Iterating on the US Census

This works with `map()` just fine:

```
df <- 2005:2019 >
  map(\(x) get_acs(geography = "county",
                  variables = "B19013_001",
                  state = "NY",
                  survey = "acs1",
                  year = x)) >
list_rbind(names_to = "year")
```

df

```
# A tibble: 580 × 6
  year GEOID NAME          variable estimate moe
  <int> <chr> <chr>          <chr>      <dbl> <dbl>
1     1  1 36001 Albany County, New York B19013_001  50054  2030
2     1  1 36005 Bronx County, New York   B19013_001  29228   853
3     1  1 36007 Broome County, New York  B19013_001  36394  2340
4     1  1 36009 Cattaraugus County, New York B19013_001  37580  2282
5     1  1 36011 Cayuga County, New York   B19013_001  42057  2406
6     1  1 36013 Chautauqua County, New York B19013_001  35495  2077
7     1  1 36015 Chemung County, New York   B19013_001  37418  3143
8     1  1 36019 Clinton County, New York  B19013_001  44757  3500
9     1  1 36027 Dutchess County, New York B19013_001  61889  2431
10    1  1 36029 Erie County, New York   B19013_001  41967  1231
# i 570 more rows
```

Iterating on the US Census

Our `id` column *tracks* the year. But we'd like it to *be* the year. So, we use `set_names()`:

```
df <- 2005:2019 >
  set_names() >
  map(\(x) get_acs(geography = "county",
                  variables = "B19013_001",
                  state = "NY",
                  survey = "acs1",
                  year = x)) >
  list_rbind(names_to = "year") >
  mutate(year = as.integer(year))
```

Iterating on the US Census

```
df
```

```
# A tibble: 580 × 6
  year GEOID NAME          variable estimate moe
  <int> <chr> <chr>          <chr>      <dbl> <dbl>
1  2005 36001 Albany County, New York B19013_001  50054  2030
2  2005 36005 Bronx County, New York  B19013_001  29228   853
3  2005 36007 Broome County, New York B19013_001  36394  2340
4  2005 36009 Cattaraugus County, New York B19013_001  37580  2282
5  2005 36011 Cayuga County, New York  B19013_001  42057  2406
6  2005 36013 Chautauqua County, New York B19013_001  35495  2077
7  2005 36015 Chemung County, New York  B19013_001  37418  3143
8  2005 36019 Clinton County, New York  B19013_001  44757  3500
9  2005 36027 Dutchess County, New York B19013_001  61889  2431
10 2005 36029 Erie County, New York     B19013_001  41967  1231
# i 570 more rows
```

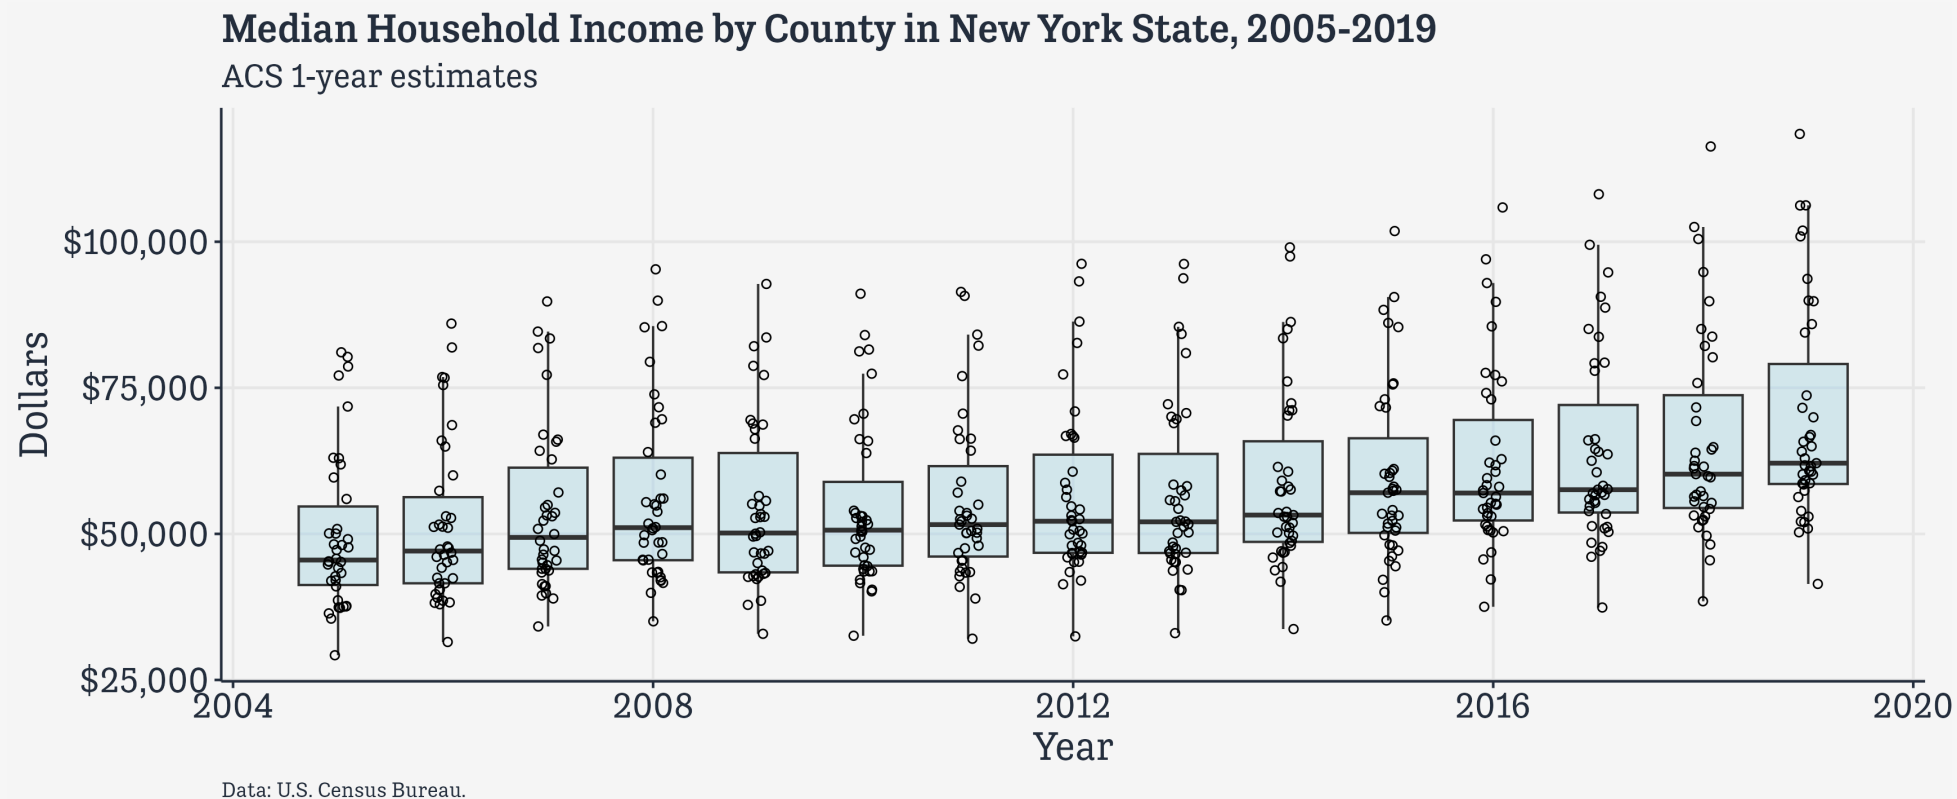
Now **year** is just the year. The **year** column will be created as a character vector, so we converted it back to an integer again at the end.

Iterating on the US Census

```
p_out <- 2005:2019 >
  set_names() >
  map(\(x) get_acs(geography = "county",
                  variables = "B19013_001",
                  state = "NY",
                  survey = "acs1",
                  year = x)) >
  list_rbind(names_to = "year") >
  mutate(year = as.integer(year)) >
  ggplot(mapping = aes(x = year, y = estimate, group = year)) +
  geom_boxplot(fill = "lightblue", alpha = 0.5, outlier.alpha = 0) +
  geom_jitter(position = position_jitter(width = 0.1), shape = 1) +
  scale_y_continuous(labels = scales::label_dollar()) +
  labs(x = "Year", y = "Dollars",
       title = "Median Household Income by County in New York State, 2005-2019",
       subtitle = "ACS 1-year estimates", caption = "Data: U.S. Census Bureau.")
```

Iterating on the US Census

```
print(p_out)
```



Example: cleaning up congress

Cleaning up congress

```
df <- filenames >
  map(read_csv) >
  list_rbind(names_to = "congress") >
  janitor::clean_names()

df >
  select(born, death, start, end)
```

```
# A tibble: 20,580 × 4
  born      death      start      end
  <chr>      <chr>      <chr>      <chr>
1 05/16/1903 01/23/1953 01/03/1945 01/03/1953
2 01/08/1899 10/27/1986 01/03/1945 01/03/1947
3 08/20/1892 11/19/1984 01/03/1945 01/03/1979
4 01/05/1891 01/05/1969 01/03/1945 01/03/1953
5 10/05/1898 01/19/1973 01/03/1945 01/02/1949
6 06/15/1898 04/14/1986 02/04/1946 04/17/1948
7 01/27/1897 07/26/1978 01/03/1945 01/03/1963
8 10/23/1895 11/11/1975 01/03/1941 06/30/1945
9 03/22/1904 02/09/1981 01/03/1945 01/03/1953
10 10/11/1890 01/14/1958 01/03/1945 01/14/1958
# i 20,570 more rows
```

We'll use the **lubridate** package to sort these out.

Lubridate has a wide range of functions to handle dates, times, and durations.

Cleaning up congress

```
library(lubridate)

date_recodes ← c("born", "death", "start", "end")
df ← df ►
  mutate(across(any_of(date_recodes), mdy),
         congress = as.integer(congress) + 78)

df
```

A tibble: 20,580 × 26

	congress	last	first	middle	suffix	nickname	born	death	sex
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<date>	<date>	<chr>
1	79	Abernethy	Thomas	Gerst...	<NA>	<NA>	1903-05-16	1953-01-23	M
2	79	Adams	Sherman	<NA>	<NA>	<NA>	1899-01-08	1986-10-27	M
3	79	Aiken	George	David	<NA>	<NA>	1892-08-20	1984-11-19	M
4	79	Allen	Asa	Leona...	<NA>	<NA>	1891-01-05	1969-01-05	M
5	79	Allen	Leo	Elwood	<NA>	<NA>	1898-10-05	1973-01-19	M
6	79	Almond	J.	Linds...	Jr.	<NA>	1898-06-15	1986-04-14	M
7	79	Andersen	Herman	Carl	<NA>	<NA>	1897-01-27	1978-07-26	M
8	79	Anderson	Clinton	Presba	<NA>	<NA>	1895-10-23	1975-11-11	M
9	79	Anderson	John	Zuing...	<NA>	<NA>	1904-03-22	1981-02-09	M
10	79	Andresen	August	Herman	<NA>	<NA>	1890-10-11	1958-01-14	M

i 20,570 more rows

i 17 more variables: position <chr>, party <chr>, state <chr>,
district <chr>, start <date>, end <date>, religion <chr>, race <chr>,
educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
job_type3 <chr>, job_type4 <chr>, job_type5 <chr>, mil1 <chr>, mil2 <chr>,
mil3 <chr>

Cleaning up congress

```
sessions ← tibble(congress = 79:116,  
                  start_year = seq(1945, 2019, by = 2),  
                  end_year = seq(1947, 2021, by = 2)) ▷  
  mutate(start_year = ymd(paste(start_year, "01", "03", sep = "-")),  
         end_year = ymd(paste(end_year, "01", "03", sep = "-")))
```

sessions

```
# A tibble: 38 × 3  
  congress start_year end_year  
    <int>   <date>     <date>  
1      79 1945-01-03 1947-01-03  
2      80 1947-01-03 1949-01-03  
3      81 1949-01-03 1951-01-03  
4      82 1951-01-03 1953-01-03  
5      83 1953-01-03 1955-01-03  
6      84 1955-01-03 1957-01-03  
7      85 1957-01-03 1959-01-03  
8      86 1959-01-03 1961-01-03  
9      87 1961-01-03 1963-01-03  
10     88 1963-01-03 1965-01-03  
# i 28 more rows
```

We're going to join these tables

The big table:

```
df ►  
select(congress, last, born)
```

```
# A tibble: 20,580 × 3  
  congress last      born  
    <dbl> <chr>    <date>  
1      79 Abernethy 1903-05-16  
2      79 Adams    1899-01-08  
3      79 Aiken    1892-08-20  
4      79 Allen    1891-01-05  
5      79 Allen    1898-10-05  
6      79 Almond   1898-06-15  
7      79 Andersen  1897-01-27  
8      79 Anderson  1895-10-23  
9      79 Anderson  1904-03-22  
10     79 Andresen  1890-10-11  
# i 20,570 more rows
```

The smaller table

```
sessions
```

```
# A tibble: 38 × 3  
  congress start_year end_year  
    <int> <date>    <date>  
1      79 1945-01-03 1947-01-03  
2      80 1947-01-03 1949-01-03  
3      81 1949-01-03 1951-01-03  
4      82 1951-01-03 1953-01-03  
5      83 1953-01-03 1955-01-03  
6      84 1955-01-03 1957-01-03  
7      85 1957-01-03 1959-01-03  
8      86 1959-01-03 1961-01-03  
9      87 1961-01-03 1963-01-03  
10     88 1963-01-03 1965-01-03  
# i 28 more rows
```

We're going to **join** these tables

We will use `left_join()` which is what you want most of the time when you are looking to merge a smaller table with additional information into a larger main one.

```
df ← left_join(df, sessions) ▷  
  relocate(start_year:end_year, .after = congress)
```

Joining with `by = join_by(congress)`

df

```
# A tibble: 20,580 × 28  
  congress start_year end_year last first middle suffix nickname born  
    <dbl> <date> <date> <chr> <chr> <chr> <chr> <chr> <date>  
1      79 1945-01-03 1947-01-03 Abern... Thom... Gerst... <NA> <NA> 1903-05-16  
2      79 1945-01-03 1947-01-03 Adams  Sher... <NA> <NA> <NA> 1899-01-08  
3      79 1945-01-03 1947-01-03 Aiken  Geor... David <NA> <NA> 1892-08-20  
4      79 1945-01-03 1947-01-03 Allen  Asa    Leona... <NA> <NA> 1891-01-05  
5      79 1945-01-03 1947-01-03 Allen  Leo    Elwood <NA> <NA> 1898-10-05  
6      79 1945-01-03 1947-01-03 Almond J.      Linds... Jr.    <NA> <NA> 1898-06-15  
7      79 1945-01-03 1947-01-03 Ander... Herm... Carl  <NA> <NA> 1897-01-27  
8      79 1945-01-03 1947-01-03 Ander... Clin... Presba <NA> <NA> 1895-10-23  
9      79 1945-01-03 1947-01-03 Ander... John  Zuing... <NA> <NA> 1904-03-22  
10     79 1945-01-03 1947-01-03 Andre... Augu... Herman <NA> <NA> 1890-10-11  
# i 20,570 more rows  
# i 19 more variables: death <date>, sex <chr>, position <chr>, party <chr>,  
# state <chr>, district <chr>, start <date>, end <date>, religion <chr>,  
# race <chr>, educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
```


Table joins

x		y	
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Spiffy Join Animatations courtesy [Garrick Aden-Buie](#)

Left join, `left_join()`

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x, and all columns from x and y. Rows in x with no match in y will have NA values in the new columns.

Left join (contd), `left_join()`

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4
		2	y5

If there are multiple matches between x and y, all combinations of the matches are returned.

Inner join, `inner_join()`

`inner_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x where there are matching values in y, and all columns from x and y.

Full join, **full_join()**

`full_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows and all columns from both x and y. Where there are not matching values, returns NA for the one missing.

Semi join, `semi_join()`

`semi_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x where there are matching values in y, keeping just columns from x.

Anti join, **anti_join()**

`anti_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

All rows from x where there are not matching values in y, keeping just columns from x.

Left join, `left_join()`

Most of the time you will be looking to make a `left_join()`

More on Missing Data

Never test for missingness with `=`

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
df <- tribble(  
  ~subject, ~age,  
  "A", 20,  
  "B", 25,  
  "C", NA,  
  "D", 34  
)
```

```
df
```

```
# A tibble: 4 × 2  
  subject    age  
  <chr>    <dbl>  
1 A         20  
2 B         25  
3 C          NA  
4 D         34
```

Never test for missingness with `=`

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# OK  
df ▶  
  filter(age = 25)
```

```
# A tibble: 1 × 2  
  subject    age  
  <chr>    <dbl>  
1 B         25
```

Never test for missingness with `=`

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# Nope  
df ▶  
  filter(age = NA)
```

```
# A tibble: 0 × 2  
#   i 2 variables: subject <chr>, age <dbl>
```

Never test for missingness with `=`

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# E.g.  
23 = NA
```

```
[1] NA
```

Never test for missingness with `=`

Always use `is.na()` instead

```
# Yes  
df ▶  
  filter(is.na(age))
```

```
# A tibble: 1 × 2  
  subject    age  
  <chr>    <dbl>  
1 C         NA
```

A quick plug for **naniar** and **visdat**

```
library(naniar)
library(visdat)
```

```
organdata
```

```
# A tibble: 238 × 21
```

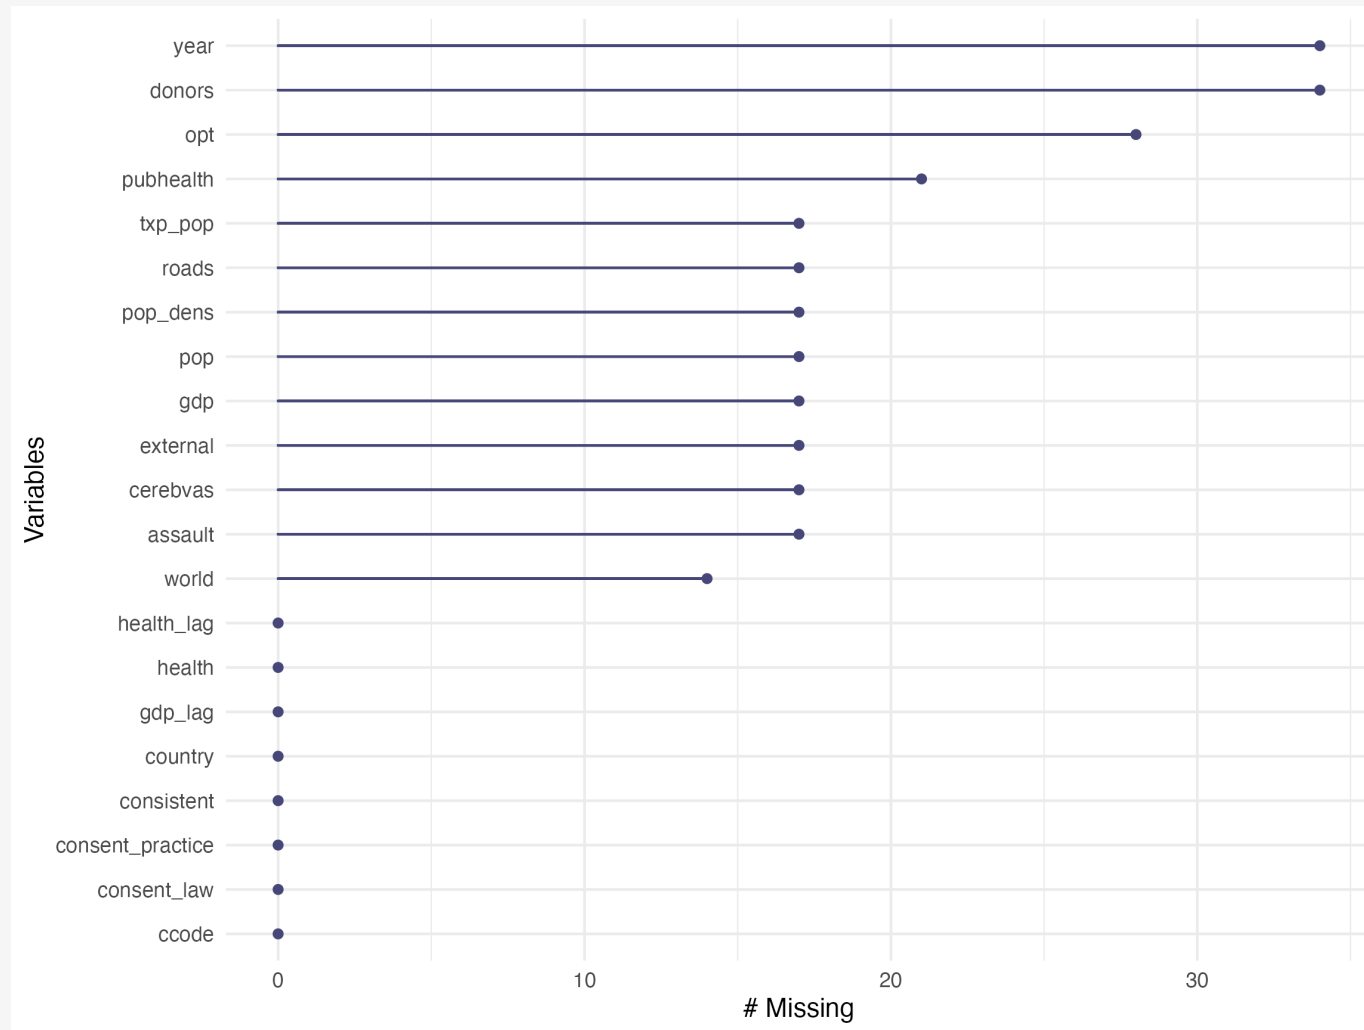
	country	year	donors	pop	pop_dens	gdp	gdp_lag	health	health_lag
	<chr>	<date>	<dbl>	<int>	<dbl>	<int>	<int>	<dbl>	<dbl>
1	Australia	NA	NA	17065	0.220	16774	16591	1300	1224
2	Australia	1991-01-01	12.1	17284	0.223	17171	16774	1379	1300
3	Australia	1992-01-01	12.4	17495	0.226	17914	17171	1455	1379
4	Australia	1993-01-01	12.5	17667	0.228	18883	17914	1540	1455
5	Australia	1994-01-01	10.2	17855	0.231	19849	18883	1626	1540
6	Australia	1995-01-01	10.2	18072	0.233	21079	19849	1737	1626
7	Australia	1996-01-01	10.6	18311	0.237	21923	21079	1846	1737
8	Australia	1997-01-01	10.3	18518	0.239	22961	21923	1948	1846
9	Australia	1998-01-01	10.5	18711	0.242	24148	22961	2077	1948
10	Australia	1999-01-01	8.67	18926	0.244	25445	24148	2231	2077

```
# i 228 more rows
```

```
# i 12 more variables: pubhealth <dbl>, roads <dbl>, cerebvas <int>,  
# assault <int>, external <int>, txp_pop <dbl>, world <chr>, opt <chr>,  
# consent_law <chr>, consent_practice <chr>, consistent <chr>, ccode <chr>
```

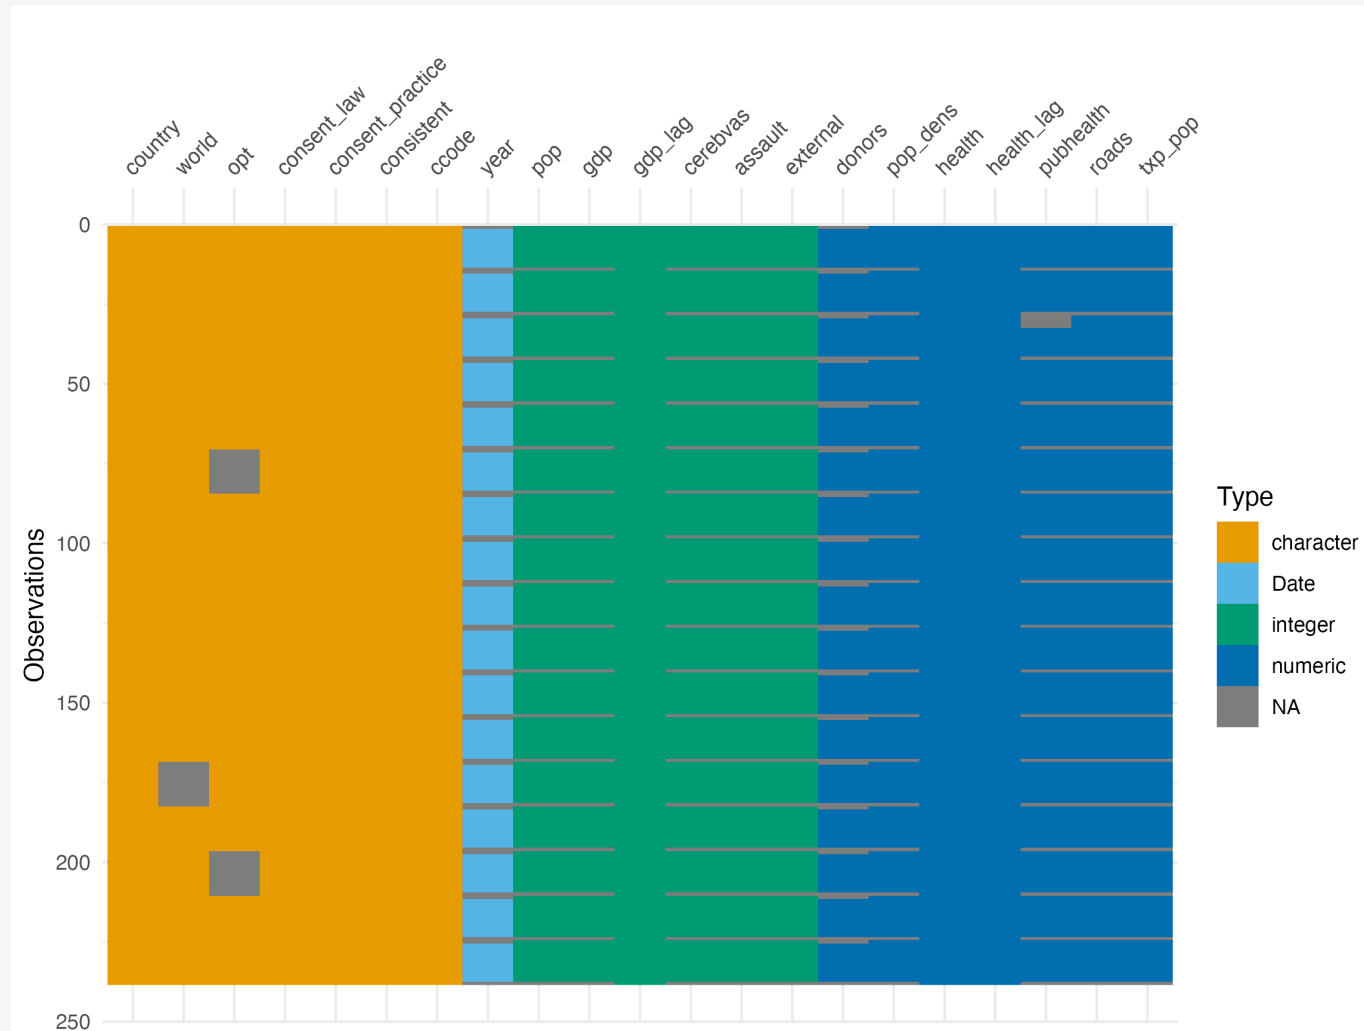
A quick plug for **naniar** and **visdat**

```
gg_miss_var(organdata)
```



A quick plug for **naniar** and **visdat**

```
vis_dat(organdata)
```



A quick plug for **naniar** and **visdat**

```
miss_var_summary(organdata)
```

```
# A tibble: 21 × 3
  variable n_miss pct_miss
  <chr>      <int>    <num>
1 year         34     14.3
2 donors        34     14.3
3 opt          28     11.8
4 pubhealth     21      8.82
5 pop          17      7.14
6 pop_dens      17      7.14
7 gdp           17      7.14
8 roads         17      7.14
9 cerebvas      17      7.14
10 assault      17      7.14
# i 11 more rows
```

A quick plug for **naniar** and **visdat**

```
miss_case_summary(organdata)
```

```
# A tibble: 238 × 3
  case n_miss pct_miss
  <int> <int>   <dbl>
1    84     12    57.1
2   182     12    57.1
3   210     12    57.1
4    14     11    52.4
5    28     11    52.4
6    42     11    52.4
7    56     11    52.4
8    70     11    52.4
9    98     11    52.4
10  112     11    52.4
# i 228 more rows
```

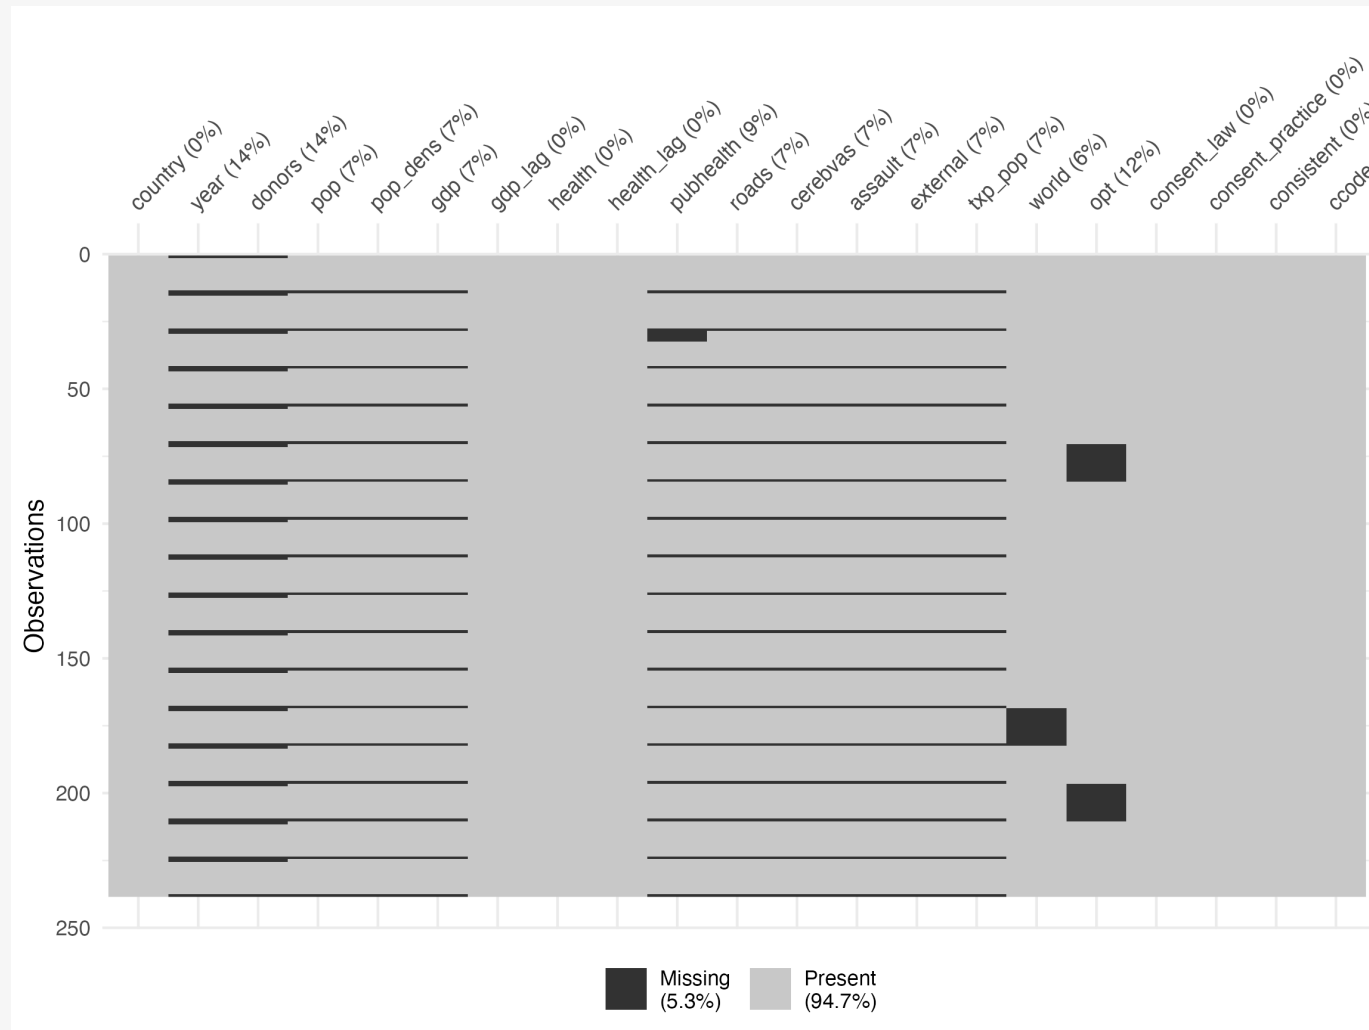
A quick plug for **naniar** and **visdat**

```
organdata >
  select(consent_law, year, pubhealth, roads) >
  group_by(consent_law) >
  miss_var_summary()
```


```
# A tibble: 6 × 4
# Groups:   consent_law [2]
  consent_law variable  n_miss pct_miss
  <chr>      <chr>      <int>   <num>
1 Informed   year         16    14.3
2 Informed   pubhealth     8     7.14
3 Informed   roads         8     7.14
4 Presumed   year         18    14.3
5 Presumed   pubhealth    13    10.3
6 Presumed   roads         9     7.14
```

A quick plug for **naniar** and **visdat**

```
vis_miss(organdata)
```




A quick plug for **naniar** and **visdat**



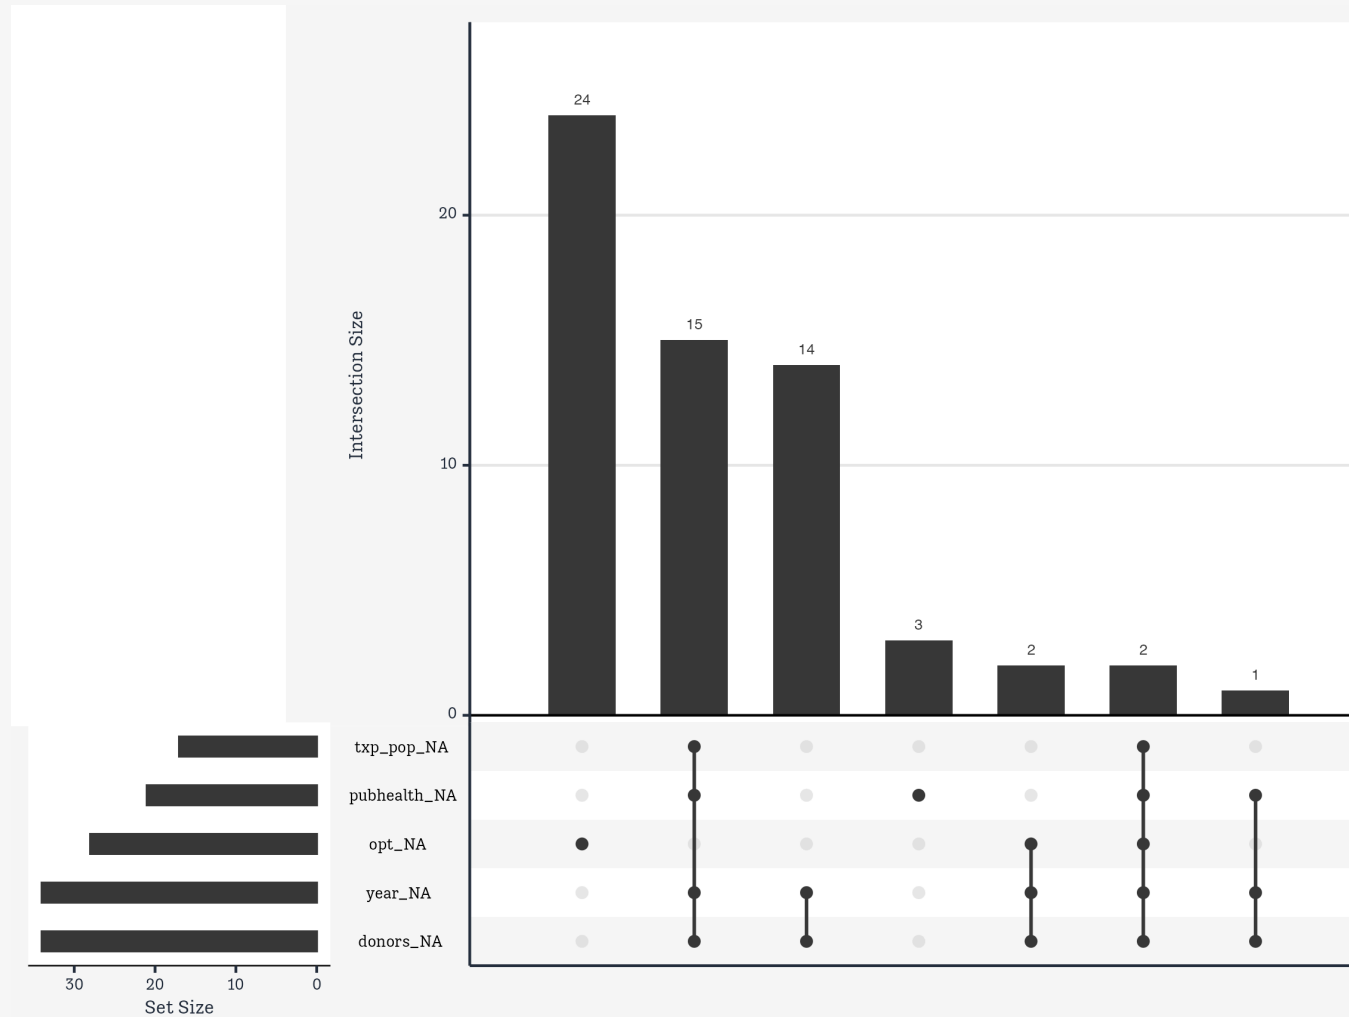
A quick plug for **naniar** and **visdat**

```
vis_miss(organdata, cluster = TRUE)
```



A quick plug for **naniar** and **visdat**

```
gg_miss_upset(organdata)
```



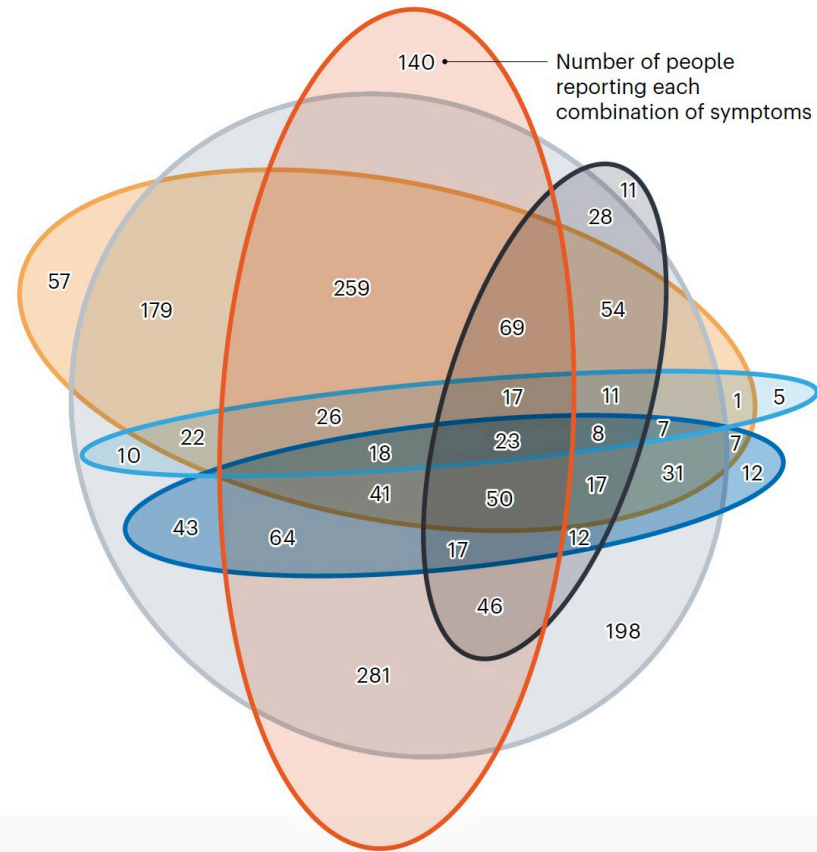
Example: Upset Plots

Upset plots and a bit of wrangling

TRACKING SYMPTOMS

On 7 April, around 60% of app users who tested positive for COVID-19 and reported symptoms had lost their sense of smell.

— Anosmia (loss of smell) — Cough — Fatigue
— Diarrhoea — Shortness of breath — Fever



:scale 35%

Upset plots and a bit of wrangling

```
symptoms ← c("Anosmia", "Cough", "Fatigue",  
             "Diarrhea", "Breath", "Fever")  
names(symptoms) ← symptoms  
symptoms
```

Anosmia	Cough	Fatigue	Diarrhea	Breath	Fever
"Anosmia"	"Cough"	"Fatigue"	"Diarrhea"	"Breath"	"Fever"

Upset plots and a bit of wrangling

```
# An Excel file!  
dat ← readxl::read_xlsx(here("data", "symptoms.xlsx"))  
dat ► print(n = nrow(dat))
```

```
# A tibble: 32 × 2
```

	combination	count
	<chr>	<dbl>
1	Anosmia	140
2	Cough	57
3	Fatigue	198
4	Diarrhea	12
5	Breath	5
6	Fever	11
7	Cough&Fatigue	179
8	Fatigue&Fever	28
9	Breath&Fatigue	10
10	Diarrhea&Fatigue	43
11	Anosmia&Fatigue	281
12	Breath&Cough	1
13	Anosmia&Diarrhea&Fatigue	64
14	Breath&Cough&Fatigue	22
15	Anosmia&Cough&Fatigue	259
16	Anosmia&Fever&Fatigue	46

Upset plots and a bit of wrangling

```
subsets ← dat ▷  
  pull(combination)  
  
## Check if each subset mentions each symptom or not  
symptom_mat ← map(subsets, \(x) str_detect(x, symptoms)) ▷  
  set_names(nm = subsets) ▷  
  map(\(x) set_names(x, nm = symptoms)) ▷  
  bind_rows(.id = "subset") ▷  
  left_join(dat, join_by(subset = combination))
```

Upset plots and a bit of wrangling

Now we have a table we can do something with.

```
symptom_mat > print(n = nrow(symptom_mat))
```

```
# A tibble: 32 × 8
```

subset <chr>	Anosmia <lgl>	Cough <lgl>	Fatigue <lgl>	Diarrhea <lgl>	Breath <lgl>	Fever <lgl>	count <dbl>
1 Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE	140
2 Cough	FALSE	TRUE	FALSE	FALSE	FALSE	FALSE	57
3 Fatigue	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE	198
4 Diarrhea	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	12
5 Breath	FALSE	FALSE	FALSE	FALSE	TRUE	FALSE	5
6 Fever	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	11
7 Cough&Fatigue	FALSE	TRUE	TRUE	FALSE	FALSE	FALSE	179
8 Fatigue&Fever	FALSE	FALSE	TRUE	FALSE	FALSE	TRUE	28
9 Breath&Fatigue	FALSE	FALSE	TRUE	FALSE	TRUE	FALSE	10
10 Diarrhea&Fatigue	FALSE	FALSE	TRUE	TRUE	FALSE	FALSE	43
11 Anosmia&Fatigue	TRUE	FALSE	TRUE	FALSE	FALSE	FALSE	281
12 Breath&Cough	FALSE	TRUE	FALSE	FALSE	TRUE	FALSE	1
13 Anosmia&Diarrhea&Fatigue	TRUE	FALSE	TRUE	TRUE	FALSE	FALSE	64
14 Breath&Cough&Fatigue	FALSE	TRUE	TRUE	FALSE	TRUE	FALSE	22
15 Anosmia&Cough&Fatigue	TRUE	TRUE	TRUE	FALSE	FALSE	FALSE	259
16 Anosmia&Fever&Fatigue	TRUE	FALSE	TRUE	FALSE	FALSE	TRUE	46

Upset plots and a bit of wrangling

Uncounting tables:

```
indvs ← symptom_mat ►  
  uncount(count)
```

```
indvs
```

```
# A tibble: 1,764 × 7
```

	subset	Anosmia	Cough	Fatigue	Diarrhea	Breath	Fever
	<chr>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>	<lgl>
1	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
2	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
3	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
4	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
5	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
6	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
7	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
8	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
9	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE
10	Anosmia	TRUE	FALSE	FALSE	FALSE	FALSE	FALSE

```
# i 1,754 more rows
```

Now we've reconstructed the individual-level observations.


Upset plots and a bit of wrangling

```
# devtools::install_github("krassowski/complex-upset")

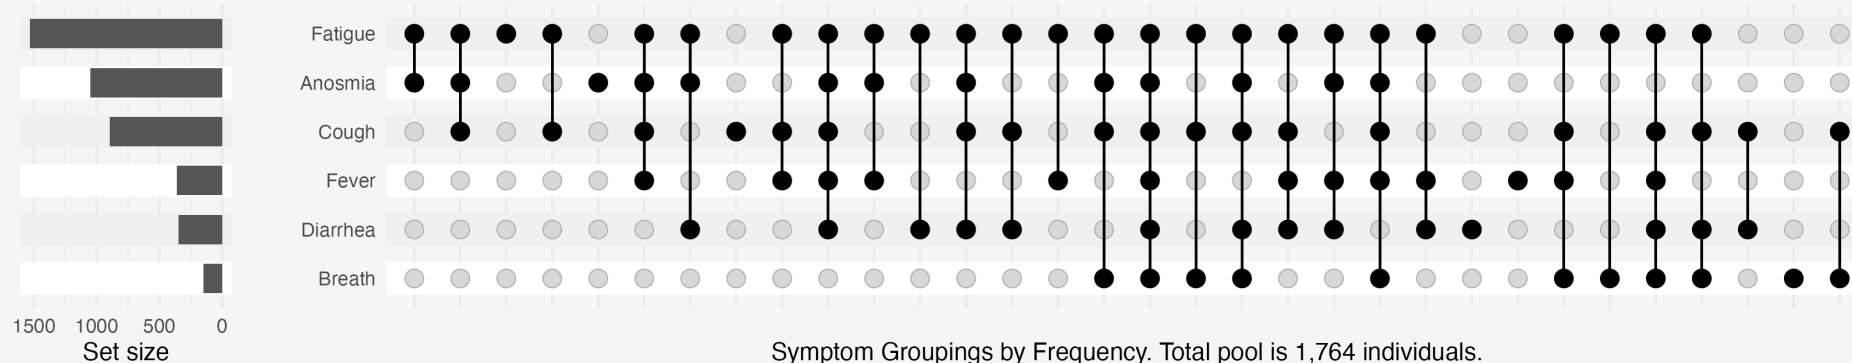
library(ComplexUpset)

upset(data = indivs, intersect = symptoms,
      name="Symptom Groupings by Frequency. Total pool is 1,764 individuals.",
      min_size = 0,
      width_ratio = 0.125) +
labs(title = "Co-Occurrence of COVID-19 Symptoms",
     caption = "Data: covid.joinzoe.com/us | Graph: @kjhealy")
```


Upset plots and a bit of wrangling



Co-Occurrence of COVID-19 Symptoms



Wrangling Models

This is not a **statistics** seminar!

I'll just give you an example of the sort of thing that many other modeling packages implement for all kinds of modeling techniques.

Again, the principle is tidy incorporation of models and their output.

Tidy regression output with **broom**

```
library(broom)  
library(gapminder)
```

```
out ← lm(formula = lifeExp ~ gdpPercap + pop + continent,  
          data = gapminder)
```

Tidy regression output with **broom**

We can't *do* anything with this, programmatically.

```
summary(out)
```

Call:

```
lm(formula = lifeExp ~ gdpPercap + pop + continent, data = gapminder)
```

Residuals:

Min	1Q	Median	3Q	Max
-49.161	-4.486	0.297	5.110	25.175

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.781e+01	3.395e-01	140.819	< 2e-16	**
gdpPercap	4.495e-04	2.346e-05	19.158	< 2e-16	**
pop	6.570e-09	1.975e-09	3.326	0.000901	**
continentAmericas	1.348e+01	6.000e-01	22.458	< 2e-16	**
continentAsia	8.193e+00	5.712e-01	14.342	< 2e-16	**
continentEurope	1.747e+01	6.246e-01	27.973	< 2e-16	**
continentOceania	1.808e+01	1.782e+00	10.146	< 2e-16	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Tidy regression output with **broom**

```
library(broom)
```

```
tidy(out)
```

```
# A tibble: 7 × 5
```

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>
1	(Intercept)	4.78e+1	0.340	141.	0
2	gdpPercap	4.50e-4	0.0000235	19.2	3.24e- 74
3	pop	6.57e-9	0.00000000198	3.33	9.01e- 4
4	continentAmericas	1.35e+1	0.600	22.5	5.19e- 98
5	continentAsia	8.19e+0	0.571	14.3	4.06e- 44
6	continentEurope	1.75e+1	0.625	28.0	6.34e-142
7	continentOceania	1.81e+1	1.78	10.1	1.59e- 23

That's a *lot* nicer. Now it's just a tibble. We know those.

Tidy regression output with **broom**

```
out_conf ← tidy(out, conf.int = TRUE)
out_conf
```

```
# A tibble: 7 × 7
```

	term <chr>	estimate <dbl>	std.error <dbl>	statistic <dbl>	p.value <dbl>	conf.low <dbl>	conf.high <dbl>
1	(Intercept)	4.78e+1	3.40e-1	141.	0	4.71e+1	4.85e+1
2	gdpPercap	4.50e-4	2.35e-5	19.2	3.24e- 74	4.03e-4	4.96e-4
3	pop	6.57e-9	1.98e-9	3.33	9.01e- 4	2.70e-9	1.04e-8
4	continentAmericas	1.35e+1	6.00e-1	22.5	5.19e- 98	1.23e+1	1.47e+1
5	continentAsia	8.19e+0	5.71e-1	14.3	4.06e- 44	7.07e+0	9.31e+0
6	continentEurope	1.75e+1	6.25e-1	28.0	6.34e-142	1.62e+1	1.87e+1
7	continentOceania	1.81e+1	1.78e+0	10.1	1.59e- 23	1.46e+1	2.16e+1

Tidy regression output with **broom**

```
out_conf >
  filter(term %nin% "(Intercept)") >
  mutate(nicelabs = prefix_strip(term, "continent")) >
  select(nicelabs, everything())
```

```
# A tibble: 6 × 8
  nicelabs term      estimate std.error statistic  p.value conf.low conf.high
  <chr>    <chr>      <dbl>    <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 gdpPercap gdpPercap  4.50e-4  2.35e-5    19.2  3.24e- 74  4.03e-4  4.96e-4
2 Pop      pop        6.57e-9  1.98e-9     3.33  9.01e- 4  2.70e-9  1.04e-8
3 Americas continent... 1.35e+1  6.00e-1    22.5  5.19e- 98  1.23e+1  1.47e+1
4 Asia     continent... 8.19e+0  5.71e-1    14.3  4.06e- 44  7.07e+0  9.31e+0
5 Europe   continent... 1.75e+1  6.25e-1    28.0  6.34e-142  1.62e+1  1.87e+1
6 Oceania  continent... 1.81e+1  1.78e+0    10.1  1.59e- 23  1.46e+1  2.16e+1
```


Grouped analysis and list columns

```
eu77 <- gapminder > filter(continent == "Europe", year == 1977)
fit <- lm(lifeExp ~ log(gdpPercap), data = eu77)
```

```
summary(fit)
```

Call:

```
lm(formula = lifeExp ~ log(gdpPercap), data = eu77)
```

Residuals:

Min	1Q	Median	3Q	Max
-7.4956	-1.0306	0.0935	1.1755	3.7125

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	29.489	7.161	4.118	0.000306	**
log(gdpPercap)	4.488	0.756	5.936	2.17e-06	**

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.114 on 28 degrees of freedom

Multiple R-squared: 0.5572, Adjusted R-squared: 0.5414

F-statistic: 35.24 on 1 and 28 DF, p-value: 2.173e-06

Grouped analysis and **list columns**

```
out_le ← gapminder ▷  
  group_by(continent, year) ▷  
  nest()
```

```
out_le
```

```
# A tibble: 60 × 3  
# Groups:   continent, year [60]  
  continent year data  
  <fct>      <int> <list>  
1 Asia      1952 <tibble [33 × 4]>  
2 Asia      1957 <tibble [33 × 4]>  
3 Asia      1962 <tibble [33 × 4]>  
4 Asia      1967 <tibble [33 × 4]>  
5 Asia      1972 <tibble [33 × 4]>  
6 Asia      1977 <tibble [33 × 4]>  
7 Asia      1982 <tibble [33 × 4]>  
8 Asia      1987 <tibble [33 × 4]>  
9 Asia      1992 <tibble [33 × 4]>  
10 Asia     1997 <tibble [33 × 4]>  
# i 50 more rows
```

Think of nesting as a kind of “super-grouping”. Look in the object inspector.

Grouped analysis and **list columns**

It's still in there.

```
out_le > filter(continent = "Europe" & year = 1977) >  
  unnest(cols = c(data))
```

```
# A tibble: 30 × 6
```

```
# Groups:   continent, year [1]
```

	continent	year	country	lifeExp	pop	gdpPercap
	<fct>	<int>	<fct>	<dbl>	<int>	<dbl>
1	Europe	1977	Albania	68.9	2509048	3533.
2	Europe	1977	Austria	72.2	7568430	19749.
3	Europe	1977	Belgium	72.8	9821800	19118.
4	Europe	1977	Bosnia and Herzegovina	69.9	4086000	3528.
5	Europe	1977	Bulgaria	70.8	8797022	7612.
6	Europe	1977	Croatia	70.6	4318673	11305.
7	Europe	1977	Czech Republic	70.7	10161915	14800.
8	Europe	1977	Denmark	74.7	5088419	20423.
9	Europe	1977	Finland	72.5	4738902	15605.
10	Europe	1977	France	73.8	53165019	18293.

```
# i 20 more rows
```

Grouped analysis and **list columns**

Here we **map()** a custom function to every row in the **data** column.

```
fit_ols ← function(df) {  
  lm(lifeExp ~ log(gdpPercap), data = df)  
}  
  
out_le ← gapminder ▷  
  group_by(continent, year) ▷  
  nest() ▷  
  mutate(model = map(data, fit_ols))
```

Grouped analysis and **list columns**

```
out_le
```

```
# A tibble: 60 × 4
# Groups:   continent, year [60]
  continent year data          model
  <fct>      <int> <list>          <list>
1 Asia      1952 <tibble [33 × 4]> <lm>
2 Asia      1957 <tibble [33 × 4]> <lm>
3 Asia      1962 <tibble [33 × 4]> <lm>
4 Asia      1967 <tibble [33 × 4]> <lm>
5 Asia      1972 <tibble [33 × 4]> <lm>
6 Asia      1977 <tibble [33 × 4]> <lm>
7 Asia      1982 <tibble [33 × 4]> <lm>
8 Asia      1987 <tibble [33 × 4]> <lm>
9 Asia      1992 <tibble [33 × 4]> <lm>
10 Asia     1997 <tibble [33 × 4]> <lm>
# i 50 more rows
```

Grouped analysis and **list columns**

We can tidy the nested models, too.

```
fit_ols ← function(df) {  
  lm(lifeExp ~ log(gdpPercap), data = df)  
}  
  
out_tidy ← gapminder ▷  
  group_by(continent, year) ▷  
  nest() ▷  
  mutate(model = map(data, fit_ols),  
         tidied = map(model, tidy)) ▷  
  unnest(cols = c(tidied)) ▷  
  filter(term %nin% "(Intercept)" &  
         continent %nin% "Oceania")
```

Grouped analysis and **list columns**

```
out_tidy
```

```
# A tibble: 48 × 9
```

```
# Groups:   continent, year [48]
```

	continent	year	data	model	term	estimate	std.error	statistic	p.value
	<fct>	<int>	<list>	<list>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Asia	1952	<tibble>	<lm>	log(gdp...	4.16	1.25	3.33	2.28e-3
2	Asia	1957	<tibble>	<lm>	log(gdp...	4.17	1.28	3.26	2.71e-3
3	Asia	1962	<tibble>	<lm>	log(gdp...	4.59	1.24	3.72	7.94e-4
4	Asia	1967	<tibble>	<lm>	log(gdp...	4.50	1.15	3.90	4.77e-4
5	Asia	1972	<tibble>	<lm>	log(gdp...	4.44	1.01	4.41	1.16e-4
6	Asia	1977	<tibble>	<lm>	log(gdp...	4.87	1.03	4.75	4.42e-5
7	Asia	1982	<tibble>	<lm>	log(gdp...	4.78	0.852	5.61	3.77e-6
8	Asia	1987	<tibble>	<lm>	log(gdp...	5.17	0.727	7.12	5.31e-8
9	Asia	1992	<tibble>	<lm>	log(gdp...	5.09	0.649	7.84	7.60e-9
10	Asia	1997	<tibble>	<lm>	log(gdp...	5.11	0.628	8.15	3.35e-9

```
# i 38 more rows
```

Grouped analysis and **list columns**

```
out_tidy ▷  
  ungroup() ▷  
  sample_n(5)
```

```
# A tibble: 5 × 9
```

	continent	year	data	model	term	estimate	std.error	statistic	p.value
	<fct>	<int>	<list>	<list>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	Africa	1957	<tibble>	<lm>	log(gdpP...	2.69	1.06	2.55	1.40e-2
2	Europe	2002	<tibble>	<lm>	log(gdpP...	3.74	0.445	8.40	3.91e-9
3	Europe	1982	<tibble>	<lm>	log(gdpP...	4.37	0.791	5.52	6.63e-6
4	Americas	2007	<tibble>	<lm>	log(gdpP...	4.49	0.752	5.98	4.26e-6
5	Europe	1992	<tibble>	<lm>	log(gdpP...	3.48	0.545	6.39	6.44e-7