Iterating on Data

Data Wrangling, Session 7

Kieran Healy

Code Horizons

October 2024

Load the packages, as always

```r
library(here)      # manage file paths
library(socviz)    # data and some useful functions
library(tidyverse) # your friend and mine
```

Moar Data

More than one data file

Inside the `data/` folder of the course packet is a folder named `congress/`

```
# A little trick from the fs package:
fs::dir_tree(here("data", "congress"))
```

```
/Users/kjhealy/Documents/courses/data_wrangling/data/congress
├── 01_79_congress.csv
├── 02_80_congress.csv
├── 03_81_congress.csv
├── 04_82_congress.csv
├── 05_83_congress.csv
├── 06_84_congress.csv
├── 07_85_congress.csv
├── 08_86_congress.csv
├── 09_87_congress.csv
├── 10_88_congress.csv
├── 11_89_congress.csv
├── 12_90_congress.csv
├── 13_91_congress.csv
├── 14_92_congress.csv
├── 15_93_congress.csv
├── 16_94_congress.csv
├── 17_95_congress.csv
├── 18_96_congress.csv
├── 19_97_congress.csv
├── 20_98_congress.csv
├── 21_99_congress.csv
├── 22_100_congress.csv
├── 23_101_congress.csv
├── 24_102_congress.csv
```

More than one data file

Let's look at one.

```r
read_csv(here("data", "congress", "17_95_congress.csv")) ▷
  janitor::clean_names() ▷
  head()
```

```
# A tibble: 6 × 25
  last     first   middle suffix nickname born  death sex   position party state
  <chr>    <chr>   <chr>  <chr>  <chr>    <chr> <chr> <chr> <chr>    <chr> <chr>
1 Abdnor   James   <NA>   <NA>   <NA>     02/1… 11/0… M     U.S. Re… Repu… SD
2 Abourezk James   George <NA>   <NA>     02/2… <NA>  M     U.S. Se… Demo… SD
3 Adams    Brockm… <NA>   <NA>   Brock    01/1… 09/1… M     U.S. Re… Demo… WA
4 Addabbo  Joseph  Patri… <NA>   <NA>     03/1… 04/1… M     U.S. Re… Demo… NY
5 Aiken    George  David  <NA>   <NA>     08/2… 11/1… M     U.S. Se… Repu… VT
6 Akaka    Daniel  Kahik… <NA>   <NA>     09/1… 04/0… M     U.S. Re… Demo… HI
# ℹ 14 more variables: district <chr>, start <chr>, end <chr>, religion <chr>,
#   race <chr>, educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
#   job_type3 <chr>, job_type4 <chr>, job_type5 <lgl>, mil1 <chr>, mil2 <chr>,
#   mil3 <chr>
```

We often find ourselves in this situation. We know each file has the same structure, and we would like to use them all at once.

Loops?

How to read them all in?

One traditional way, which we could do in R, is to write an explicit *loop* that iterated over a vector of filenames, read each file, and then joined the results together in a tall rectangle.

```
# Pseudocode

filenames ← c("01_79_congress.csv", "02_80_congress.csv", "03_81_congress.csv",
              "04_82_congress.csv" [etc etc])

collected_files ← NULL

for(i in 1:length(filenames)) {
    new_file ← read_file(filenames[i])
    collected_files ← append_to(collected_files, new_files)
}
```

Loops?

Mapping is just a kind of iteration

Vectorized arithmetic again

The simplest cases are not that different from the vectorized arithmetic we're already familiar with.

```r
a ← c(1:10)

b ← 1

# You know what R will do here
a + b
```

```
[1]  2  3  4  5  6  7  8  9 10 11
```

Vectorized arithmetic again

We can make this explicit by writing a function:

```
add_b ← function(x) {
  b ← 1
  x + b # for any x
}
```

Vectorized arithmetic again

We can make this explicit by writing a function:

```
add_b ← function(x) {
  b ← 1
  x + b # for any x
}
```

Now:

```
add_b(x = a)
```

```
 [1]  2  3  4  5  6  7  8  9 10 11
```

Vectorized arithmetic again

Again, R's vectorized approach means it automatically adds b to every element of the x we give it.

```r
add_b(x = 10)
```
```
[1] 11
```

```r
add_b(x = c(1, 99, 1000))
```
```
[1]    2  100 1001
```

Iterating in a pipeline

Some operations can't directly be vectorized in this way, which is why we need to manually iterate, or will want to write loops.

```
library(gapminder)
gapminder ▷
  summarize(country_n = n_distinct(country),
            continent_n = n_distinct(continent),
            year_n = n_distinct(year),
            lifeExp_n = n_distinct(lifeExp),
            population_n = n_distinct(population))
```

```
# A tibble: 1 × 5
  country_n continent_n year_n lifeExp_n population_n
      <int>       <int>  <int>     <int>        <int>
1       142           5     12      1626         4060
```

That's tedious to write! Computers are supposed to allow us to avoid that sort of thing.

Iterating in a pipeline

So how would we iterate this? What we want is to apply the `n_distinct()` function to each column of `gapminder`, but in a way that still allows us to use pipelines and so on.

```r
library(gapminder)
gapminder |>
  summarize(n_distinct(country),
            n_distinct(continent),
            n_distinct(year),
            n_distinct(lifeExp),
            n_distinct(population))
```

```
# A tibble: 1 × 5
  `n_distinct(country)` `n_distinct(continent)` `n_distinct(year)`
                  <int>                   <int>              <int>
1                   142                       5                 12
# ℹ 2 more variables: `n_distinct(lifeExp)` <int>,
#   `n_distinct(population)` <int>
```

Using `n_distinct()` in this context is an idea I got from Rebecca Barter's discussion of `purrr`.

Iterating in a pipeline

You'd use **across()**, like this:

```r
gapminder |>
  summarize(across(everything(), n_distinct))
```

```
# A tibble: 1 × 6
  country continent  year lifeExp   pop gdpPercap
    <int>     <int> <int>   <int> <int>     <int>
1     142         5    12    1626  1704      1704
```

Iterating in a pipeline

But you could also do this ...

```
map(gapminder, n_distinct)
```

```
$country
[1] 142

$continent
[1] 5

$year
[1] 12

$lifeExp
[1] 1626

$pop
[1] 1704

$gdpPercap
[1] 1704
```

Read it as "Feed each column of
gapminder to the n_distinct()
function.

(This is pretty much what across()
is doing more nicely.)

Iterating in a pipeline

Or, in pipeline form:

```
gapminder ▷
  map(n_distinct)
```

```
$country
[1] 142

$continent
[1] 5

$year
[1] 12

$lifeExp
[1] 1626

$pop
[1] 1704

$gdpPercap
[1] 1704
```

You can see we are getting a *list* back.

Iterating in a pipeline

Or, in pipeline form:

```
result ← gapminder ▷
  map(n_distinct)

class(result)
```

```
[1] "list"
```

```
result$continent
```

```
[1] 5
```

```
result[[2]]
```

```
[1] 5
```

Iterating in a pipeline

But we know `n_distinct()` should always return an integer. So we use `map_int()` instead of the generic `map()`.

```
gapminder ▷
  map_int(n_distinct)
```

```
  country continent     year   lifeExp      pop gdpPercap
      142         5       12      1626     1704      1704
```

Get a vector of **filenames**

```r
filenames ← dir(path = here("data", "congress"),
                pattern = "*.csv",
                full.names = TRUE)

filenames[1:15] # Just displaying the first 15, to save slide space
```

```
 [1] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/01_79_congress.csv"
 [2] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/02_80_congress.csv"
 [3] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/03_81_congress.csv"
 [4] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/04_82_congress.csv"
 [5] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/05_83_congress.csv"
 [6] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/06_84_congress.csv"
 [7] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/07_85_congress.csv"
 [8] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/08_86_congress.csv"
 [9] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/09_87_congress.csv"
[10] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/10_88_congress.csv"
[11] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/11_89_congress.csv"
[12] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/12_90_congress.csv"
[13] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/13_91_congress.csv"
[14] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/14_92_congress.csv"
[15] "/Users/kjhealy/Documents/courses/data_wrangling/data/congress/15_93_congress.csv"
```

And feed it to `read_csv()`

… using `map()` and binding the resulting list into a tibble.

```
df ← filenames ▷
  map(read_csv) ▷
  list_rbind(names_to = "congress") ▷
  janitor::clean_names()

df
```

```
# A tibble: 20,580 × 26
   congress last   first middle suffix nickname born   death sex   position party
      <int> <chr>  <chr> <chr>  <chr>  <chr>    <chr>  <chr> <chr> <chr>    <chr>
 1        1 Abern… Thom… Gerst… <NA>   <NA>     05/1… 01/2… M     U.S. Re… Demo…
 2        1 Adams  Sher… <NA>   <NA>   <NA>     01/0… 10/2… M     U.S. Re… Repu…
 3        1 Aiken  Geor… David  <NA>   <NA>     08/2… 11/1… M     U.S. Se… Repu…
 4        1 Allen  Asa   Leona… <NA>   <NA>     01/0… 01/0… M     U.S. Re… Demo…
 5        1 Allen  Leo   Elwood <NA>   <NA>     10/0… 01/1… M     U.S. Re… Repu…
 6        1 Almond J.    Linds… Jr.    <NA>     06/1… 04/1… M     U.S. Re… Demo…
 7        1 Ander… Herm… Carl   <NA>   <NA>     01/2… 07/2… M     U.S. Re… Repu…
 8        1 Ander… Clin… Presba <NA>   <NA>     10/2… 11/1… M     U.S. Re… Demo…
 9        1 Ander… John  Zuing… <NA>   <NA>     03/2… 02/0… M     U.S. Re… Repu…
10        1 Andre… Augu… Herman <NA>   <NA>     10/1… 01/1… M     U.S. Re… Repu…
# ℹ 20,570 more rows
# ℹ 15 more variables: state <chr>, district <chr>, start <chr>, end <chr>,
#   religion <chr>, race <chr>, educational_attainment <chr>, job_type1 <chr>,
#   job_type2 <chr>, job_type3 <chr>, job_type4 <chr>, job_type5 <chr>,
#   mil1 <chr>, mil2 <chr>, mil3 <chr>
```

Now witness the firepower of this fully armed and operational

method of type-safe functional iteration

`read_csv()` can do this directly

In fact `map()` is not required for this particular use:

```r
tmp ← read_csv(filenames, id = "path",
               name_repair = janitor::make_clean_names)

tmp ▷
  mutate(congress = str_extract(path, "_\\d{2,3}_congress"),
         congress = str_extract(congress, "\\d{2,3}")) ▷
  relocate(congress)
```

```
# A tibble: 20,580 × 27
   congress path    last  first middle suffix nickname born   death sex   position
   <chr>    <chr>   <chr> <chr> <chr>  <chr>  <chr>    <chr>  <chr> <chr> <chr>
 1 79       /User… Aber… Thom… Gerst… <NA>   <NA>     05/1… 01/2… M     U.S. Re…
 2 79       /User… Adams Sher… <NA>   <NA>   <NA>     01/0… 10/2… M     U.S. Re…
 3 79       /User… Aiken Geor… David  <NA>   <NA>     08/2… 11/1… M     U.S. Se…
 4 79       /User… Allen Asa   Leona… <NA>   <NA>     01/0… 01/0… M     U.S. Re…
 5 79       /User… Allen Leo   Elwood <NA>   <NA>     10/0… 01/1… M     U.S. Re…
 6 79       /User… Almo… J.    Linds… Jr.    <NA>     06/1… 04/1… M     U.S. Re…
 7 79       /User… Ande… Herm… Carl   <NA>   <NA>     01/2… 07/2… M     U.S. Re…
 8 79       /User… Ande… Clin… Presba <NA>   <NA>     10/2… 11/1… M     U.S. Re…
 9 79       /User… Ande… John  Zuing… <NA>   <NA>     03/2… 02/0… M     U.S. Re…
10 79       /User… Andr… Augu… Herman <NA>   <NA>     10/1… 01/1… M     U.S. Re…
# ℹ 20,570 more rows
# ℹ 16 more variables: party <chr>, state <chr>, district <chr>, start <chr>,
#   end <chr>, religion <chr>, race <chr>, educational_attainment <chr>,
#   job_type1 <chr>, job_type2 <chr>, job_type3 <chr>, job_type4 <chr>,
#   job_type5 <chr>, mil1 <chr>, mil2 <chr>, mil3 <chr>
```

Example: Iterating on the US Census

Iterating on the US Census

Mapped iteration is very general, and not just for local files

```r
## Register for a free Census API key
library(tidycensus)
```

```r
out ← get_acs(geography = "county",
              variables = "B19013_001",
              state = "NY",
              county = "New York",
              survey = "acs1",
              year = 2005)
```

```r
out
```

```
# A tibble: 1 × 5
  GEOID NAME                          variable    estimate   moe
  <chr> <chr>                         <chr>          <dbl> <dbl>
1 36061 New York County, New York B19013_001       55973  1462
```

Iterating on the US Census

All counties in New York State for a specific year

```
out ← get_acs(geography = "county",
              variables = "B19013_001",
              state = "NY",
              survey = "acs1",
              year = 2005)
```

`out`

```
# A tibble: 38 × 5
   GEOID NAME                       variable    estimate   moe
   <chr> <chr>                      <chr>          <dbl> <dbl>
 1 36001 Albany County, New York    B19013_001     50054  2030
 2 36005 Bronx County, New York     B19013_001     29228   853
 3 36007 Broome County, New York    B19013_001     36394  2340
 4 36009 Cattaraugus County, New York B19013_001   37580  2282
 5 36011 Cayuga County, New York    B19013_001     42057  2406
 6 36013 Chautauqua County, New York B19013_001    35495  2077
 7 36015 Chemung County, New York   B19013_001     37418  3143
 8 36019 Clinton County, New York   B19013_001     44757  3500
 9 36027 Dutchess County, New York  B19013_001     61889  2431
10 36029 Erie County, New York      B19013_001     41967  1231
# ℹ 28 more rows
```

Iterating on the US Census

What if we want the results for *every* available year? First, a handy function: **set_names()**

```
x ← c(1:10)

x
```

```
 [1]  1  2  3  4  5  6  7  8  9 10
```

```
x ← set_names(x, nm = letters[1:10])

x
```

```
 a  b  c  d  e  f  g  h  i  j
 1  2  3  4  5  6  7  8  9 10
```

Iterating on the US Census

By default, `set_names()` will label a vector with that vector's values:

```
c(1:10) ▷
  set_names()
```

```
 1  2  3  4  5  6  7  8  9 10
 1  2  3  4  5  6  7  8  9 10
```

Iterating on the US Census

This works with `map()` just fine:

```r
df <- 2005:2019 |>
  map(\(x) get_acs(geography = "county",
                   variables = "B19013_001",
                   state = "NY",
                   survey = "acs1",
                   year = x)) |>
  list_rbind(names_to = "year")
```

`df`

```
# A tibble: 580 × 6
    year GEOID NAME                      variable   estimate   moe
   <int> <chr> <chr>                     <chr>         <dbl> <dbl>
 1     1 36001 Albany County, New York   B19013_001    50054  2030
 2     1 36005 Bronx County, New York    B19013_001    29228   853
 3     1 36007 Broome County, New York   B19013_001    36394  2340
 4     1 36009 Cattaraugus County, New York B19013_001 37580  2282
 5     1 36011 Cayuga County, New York   B19013_001    42057  2406
 6     1 36013 Chautauqua County, New York B19013_001  35495  2077
 7     1 36015 Chemung County, New York  B19013_001    37418  3143
 8     1 36019 Clinton County, New York  B19013_001    44757  3500
 9     1 36027 Dutchess County, New York B19013_001    61889  2431
10     1 36029 Erie County, New York     B19013_001    41967  1231
# i 570 more rows
```

Iterating on the US Census

Our `id` column *tracks* the year. But we'd like it to *be* the year. So, we use
`set_names()`:

```r
df ← 2005:2019 ▷
  set_names() ▷
  map(\(x) get_acs(geography = "county",
                   variables = "B19013_001",
                   state = "NY",
                   survey = "acs1",
                   year = x)) ▷
  list_rbind(names_to = "year") ▷
  mutate(year = as.integer(year))
```

Iterating on the US Census

```
df
```

```
# A tibble: 580 × 6
    year GEOID NAME                        variable    estimate   moe
   <int> <chr> <chr>                       <chr>          <dbl> <dbl>
 1  2005 36001 Albany County, New York     B19013_001     50054  2030
 2  2005 36005 Bronx County, New York      B19013_001     29228   853
 3  2005 36007 Broome County, New York     B19013_001     36394  2340
 4  2005 36009 Cattaraugus County, New York B19013_001    37580  2282
 5  2005 36011 Cayuga County, New York     B19013_001     42057  2406
 6  2005 36013 Chautauqua County, New York B19013_001     35495  2077
 7  2005 36015 Chemung County, New York    B19013_001     37418  3143
 8  2005 36019 Clinton County, New York    B19013_001     44757  3500
 9  2005 36027 Dutchess County, New York   B19013_001     61889  2431
10  2005 36029 Erie County, New York       B19013_001     41967  1231
# ℹ 570 more rows
```

Now `year` is just the year. The `year` column will be created as a character vector, so we converted it back to an integer again at the end.

Iterating on the US Census

```r
p_out ← 2005:2019 ▷
  set_names() ▷
  map(\(x) get_acs(geography = "county",
                   variables = "B19013_001",
                   state = "NY",
                   survey = "acs1",
                   year = x)) ▷
  list_rbind(names_to = "year") ▷
  mutate(year = as.integer(year)) ▷
  ggplot(mapping = aes(x = year, y = estimate, group = year)) +
  geom_boxplot(fill = "lightblue", alpha = 0.5, outlier.alpha = 0) +
  geom_jitter(position = position_jitter(width = 0.1), shape = 1) +
  scale_y_continuous(labels = scales::label_dollar()) +
  labs(x = "Year", y = "Dollars",
       title = "Median Household Income by County in New York State, 2005-2019",
       subtitle = "ACS 1-year estimates", caption = "Data: U.S. Census Bureau.")
```

Iterating on the US Census

```
print(p_out)
```

Median Household Income by County in New York State, 2005-2019
ACS 1-year estimates

Data: U.S. Census Bureau.
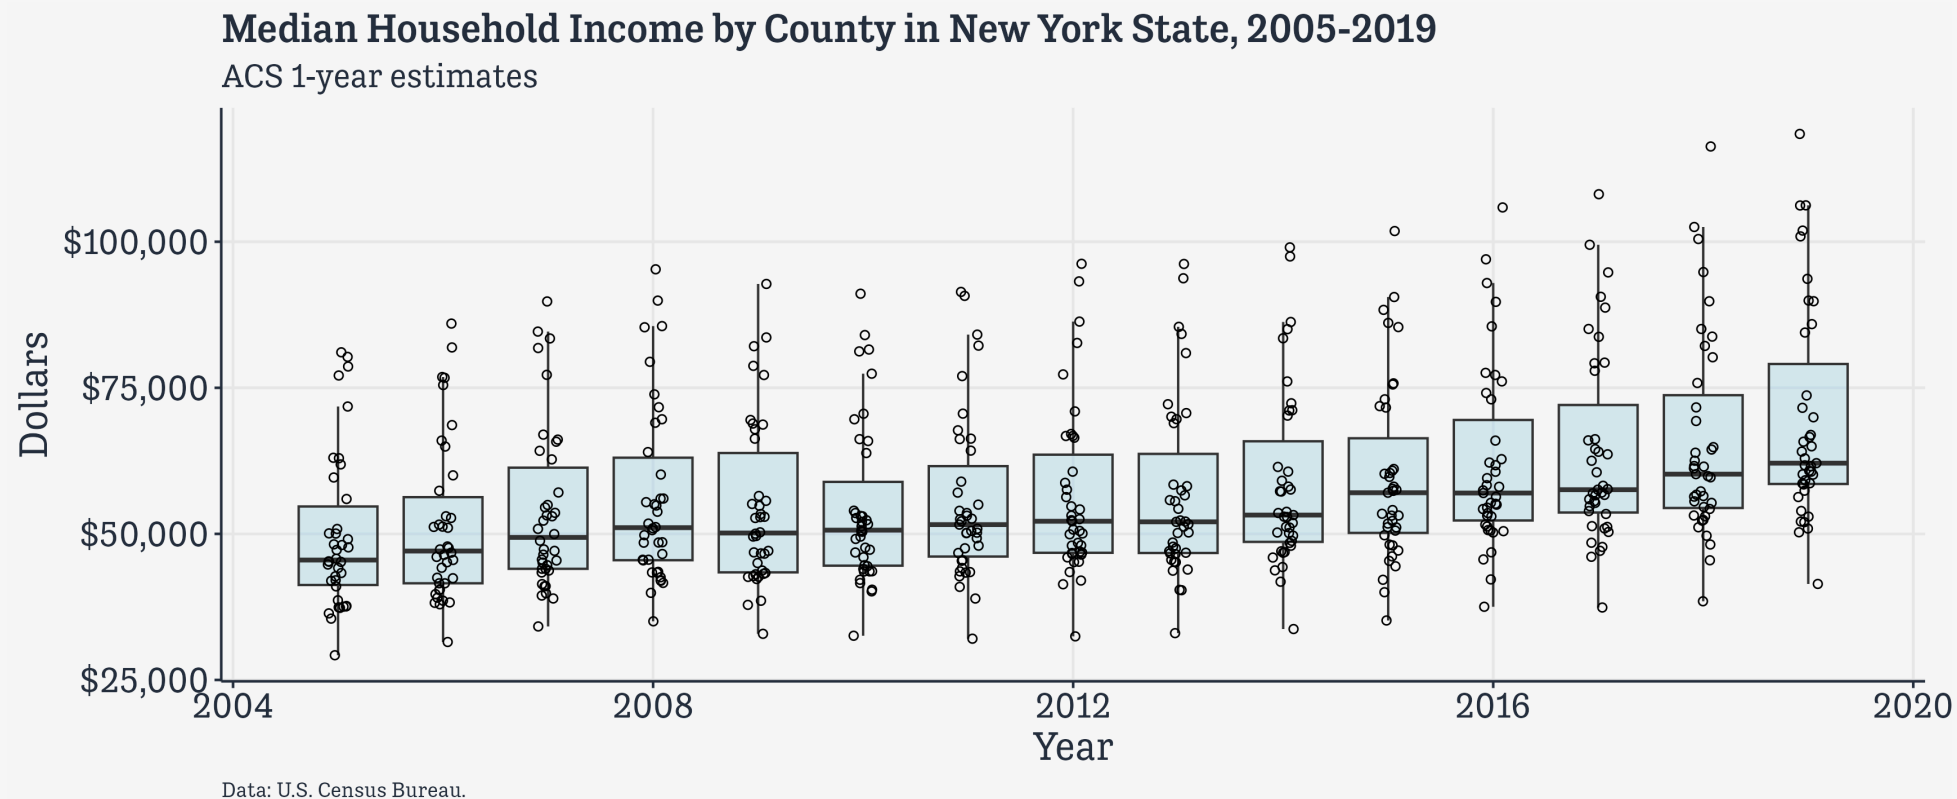
Example: cleaning up congress

Cleaning up congress

```r
df ← filenames ▷
  map(read_csv) ▷
  list_rbind(names_to = "congress") ▷
  janitor::clean_names()

df ▷
  select(born, death, start, end)
```

```
# A tibble: 20,580 × 4
   born       death      start      end
   <chr>      <chr>      <chr>      <chr>
 1 05/16/1903 01/23/1953 01/03/1945 01/03/1953
 2 01/08/1899 10/27/1986 01/03/1945 01/03/1947
 3 08/20/1892 11/19/1984 01/03/1945 01/03/1979
 4 01/05/1891 01/05/1969 01/03/1945 01/03/1953
 5 10/05/1898 01/19/1973 01/03/1945 01/02/1949
 6 06/15/1898 04/14/1986 02/04/1946 04/17/1948
 7 01/27/1897 07/26/1978 01/03/1945 01/03/1963
 8 10/23/1895 11/11/1975 01/03/1941 06/30/1945
 9 03/22/1904 02/09/1981 01/03/1945 01/03/1953
10 10/11/1890 01/14/1958 01/03/1945 01/14/1958
# ℹ 20,570 more rows
```

We'll use the **lubridate** package to sort these out.

Lubridate has a wide range of functions to handle dates, times, and durations.

Cleaning up congress

```r
library(lubridate)

date_recodes ← c("born", "death", "start", "end")
df ← df ▷
    mutate(across(any_of(date_recodes), mdy),
           congress = as.integer(congress) + 78)

df
```

```
# A tibble: 20,580 × 26
   congress last      first   middle  suffix nickname born       death       sex
      <dbl> <chr>     <chr>   <chr>   <chr>  <chr>    <date>     <date>      <chr>
 1       79 Abernethy Thomas  Gerst…  <NA>   <NA>     1903-05-16 1953-01-23  M
 2       79 Adams     Sherman <NA>    <NA>   <NA>     1899-01-08 1986-10-27  M
 3       79 Aiken     George  David   <NA>   <NA>     1892-08-20 1984-11-19  M
 4       79 Allen     Asa     Leona…  <NA>   <NA>     1891-01-05 1969-01-05  M
 5       79 Allen     Leo     Elwood  <NA>   <NA>     1898-10-05 1973-01-19  M
 6       79 Almond    J.      Linds…  Jr.    <NA>     1898-06-15 1986-04-14  M
 7       79 Andersen  Herman  Carl    <NA>   <NA>     1897-01-27 1978-07-26  M
 8       79 Anderson  Clinton Presba  <NA>   <NA>     1895-10-23 1975-11-11  M
 9       79 Anderson  John    Zuing…  <NA>   <NA>     1904-03-22 1981-02-09  M
10       79 Andresen  August  Herman  <NA>   <NA>     1890-10-11 1958-01-14  M
# ℹ 20,570 more rows
# ℹ 17 more variables: position <chr>, party <chr>, state <chr>,
#   district <chr>, start <date>, end <date>, religion <chr>, race <chr>,
#   educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
#   job_type3 <chr>, job_type4 <chr>, job_type5 <chr>, mil1 <chr>, mil2 <chr>,
#   mil3 <chr>
```

Cleaning up **congress**

```
sessions ← tibble(congress = 79:116,
                  start_year = seq(1945, 2019, by = 2),
                  end_year = seq(1947, 2021, by = 2)) ▷
  mutate(start_year = ymd(paste(start_year, "01", "03", sep = "-")),
         end_year = ymd(paste(end_year, "01", "03", sep = "-")))


sessions
```

```
# A tibble: 38 × 3
   congress start_year end_year
      <int> <date>     <date>
 1       79 1945-01-03 1947-01-03
 2       80 1947-01-03 1949-01-03
 3       81 1949-01-03 1951-01-03
 4       82 1951-01-03 1953-01-03
 5       83 1953-01-03 1955-01-03
 6       84 1955-01-03 1957-01-03
 7       85 1957-01-03 1959-01-03
 8       86 1959-01-03 1961-01-03
 9       87 1961-01-03 1963-01-03
10       88 1963-01-03 1965-01-03
# i 28 more rows
```

We're going to join these tables

The big table:

The smaller table

```
df ▷
  select(congress, last, born)
```

```
# A tibble: 20,580 × 3
   congress last      born
      <dbl> <chr>     <date>
 1       79 Abernethy 1903-05-16
 2       79 Adams     1899-01-08
 3       79 Aiken     1892-08-20
 4       79 Allen     1891-01-05
 5       79 Allen     1898-10-05
 6       79 Almond    1898-06-15
 7       79 Andersen  1897-01-27
 8       79 Anderson  1895-10-23
 9       79 Anderson  1904-03-22
10       79 Andresen  1890-10-11
# ℹ 20,570 more rows
```

```
sessions
```

```
# A tibble: 38 × 3
   congress start_year end_year
      <int> <date>     <date>
 1       79 1945-01-03 1947-01-03
 2       80 1947-01-03 1949-01-03
 3       81 1949-01-03 1951-01-03
 4       82 1951-01-03 1953-01-03
 5       83 1953-01-03 1955-01-03
 6       84 1955-01-03 1957-01-03
 7       85 1957-01-03 1959-01-03
 8       86 1959-01-03 1961-01-03
 9       87 1961-01-03 1963-01-03
10       88 1963-01-03 1965-01-03
# ℹ 28 more rows
```

We're going to join these tables

We will use `left_join()` which is what you want most of the time when you are looking to merge a smaller table with additional information into a larger main one.

```
df ← left_join(df, sessions) ▷
  relocate(start_year:end_year, .after = congress)
```

Joining with `by = join_by(congress)`

df

```
# A tibble: 20,580 × 28
   congress start_year end_year   last    first middle suffix nickname born
      <dbl> <date>     <date>     <chr>   <chr> <chr>  <chr>  <chr>    <date>
 1       79 1945-01-03 1947-01-03 Abern… Thom… Gerst… <NA>    <NA>     1903-05-16
 2       79 1945-01-03 1947-01-03 Adams  Sher… <NA>   <NA>    <NA>     1899-01-08
 3       79 1945-01-03 1947-01-03 Aiken  Geor… David  <NA>    <NA>     1892-08-20
 4       79 1945-01-03 1947-01-03 Allen  Asa   Leona… <NA>    <NA>     1891-01-05
 5       79 1945-01-03 1947-01-03 Allen  Leo   Elwood <NA>    <NA>     1898-10-05
 6       79 1945-01-03 1947-01-03 Almond J.    Linds… Jr.    <NA>     1898-06-15
 7       79 1945-01-03 1947-01-03 Ander… Herm… Carl   <NA>    <NA>     1897-01-27
 8       79 1945-01-03 1947-01-03 Ander… Clin… Presba <NA>    <NA>     1895-10-23
 9       79 1945-01-03 1947-01-03 Ander… John  Zuing… <NA>    <NA>     1904-03-22
10       79 1945-01-03 1947-01-03 Andre… Augu… Herman <NA>    <NA>     1890-10-11
# ℹ 20,570 more rows
# ℹ 19 more variables: death <date>, sex <chr>, position <chr>, party <chr>,
#   state <chr>, district <chr>, start <date>, end <date>, religion <chr>,
#   race <chr>, educational_attainment <chr>, job_type1 <chr>, job_type2 <chr>,
```

Table joins

Spiffy Join Animatations courtesy Garrick Aden-Buie

Left join, left_join()

left_join(x, y)

All rows from x, and all columns from x and y. Rows in x with no match in y will have NA values in the new columns.

Left join (contd), left_join()

If there are multiple matches between x and y, all combinations of the matches are returned.

Inner join, inner_join()

All rows from x where there are matching values in y, and all columns from x and y.

Full join, full_join()

full_join(x, y)

All rows and all columns from both x and y. Where there are not matching values, returns NA for the one missing.

Semi join, semi_join()

All rows from x where there are matching values in y, keeping just columns from x.

Anti join, anti_join()

All rows from x where there are not matching values in y, keeping just columns from x.

Left join, left_join()

Most of the time you will be looking to make a `left_join()`

More on Missing Data

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```r
df ← tribble(
  ~subject, ~age,
  "A", 20,
  "B", 25,
  "C", NA,
  "D", 34
)

df
```

```
# A tibble: 4 × 2
  subject   age
  <chr>   <dbl>
1 A          20
2 B          25
3 C          NA
4 D          34
```

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# OK
df ▷
  filter(age == 25)
```

```
# A tibble: 1 × 2
  subject    age
  <chr>    <dbl>
1 B           25
```

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# Nope
df |>
  filter(age == NA)
```

```
# A tibble: 0 × 2
# i 2 variables: subject <chr>, age <dbl>
```

Never test for missingness with ==

The result of almost any operation involving a missing/unknown value will be missing/unknown.

```
# E.g.
23 == NA
```

```
[1] NA
```

Never test for missingness with ==

Always use `is.na()` instead

```r
# Yes
df |>
  filter(is.na(age))
```

```
# A tibble: 1 × 2
  subject   age
  <chr>   <dbl>
1 C          NA
```

A quick plug for naniar and visdat

```
library(naniar)
library(visdat)

organdata
```

```
# A tibble: 238 × 21
   country   year       donors   pop pop_dens   gdp gdp_lag health health_lag
   <chr>     <date>      <dbl> <int>    <dbl> <int>   <int>  <dbl>      <dbl>
 1 Australia NA            NA   17065    0.220 16774   16591   1300       1224
 2 Australia 1991-01-01  12.1   17284    0.223 17171   16774   1379       1300
 3 Australia 1992-01-01  12.4   17495    0.226 17914   17171   1455       1379
 4 Australia 1993-01-01  12.5   17667    0.228 18883   17914   1540       1455
 5 Australia 1994-01-01  10.2   17855    0.231 19849   18883   1626       1540
 6 Australia 1995-01-01  10.2   18072    0.233 21079   19849   1737       1626
 7 Australia 1996-01-01  10.6   18311    0.237 21923   21079   1846       1737
 8 Australia 1997-01-01  10.3   18518    0.239 22961   21923   1948       1846
 9 Australia 1998-01-01  10.5   18711    0.242 24148   22961   2077       1948
10 Australia 1999-01-01   8.67  18926    0.244 25445   24148   2231       2077
# ℹ 228 more rows
# ℹ 12 more variables: pubhealth <dbl>, roads <dbl>, cerebvas <int>,
#   assault <int>, external <int>, txp_pop <dbl>, world <chr>, opt <chr>,
#   consent_law <chr>, consent_practice <chr>, consistent <chr>, ccode <chr>
```
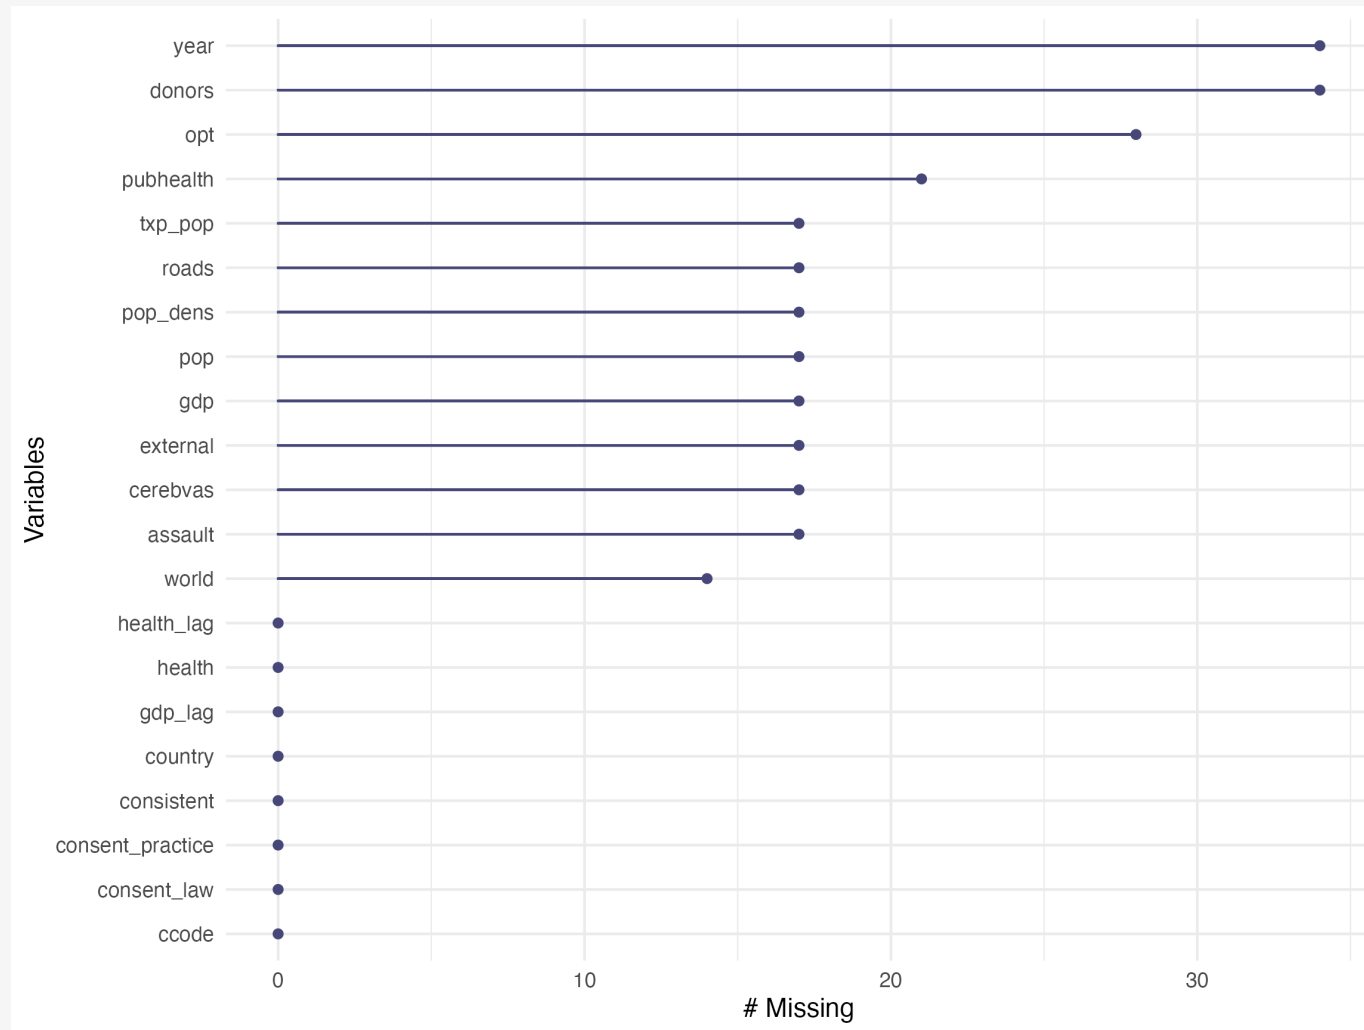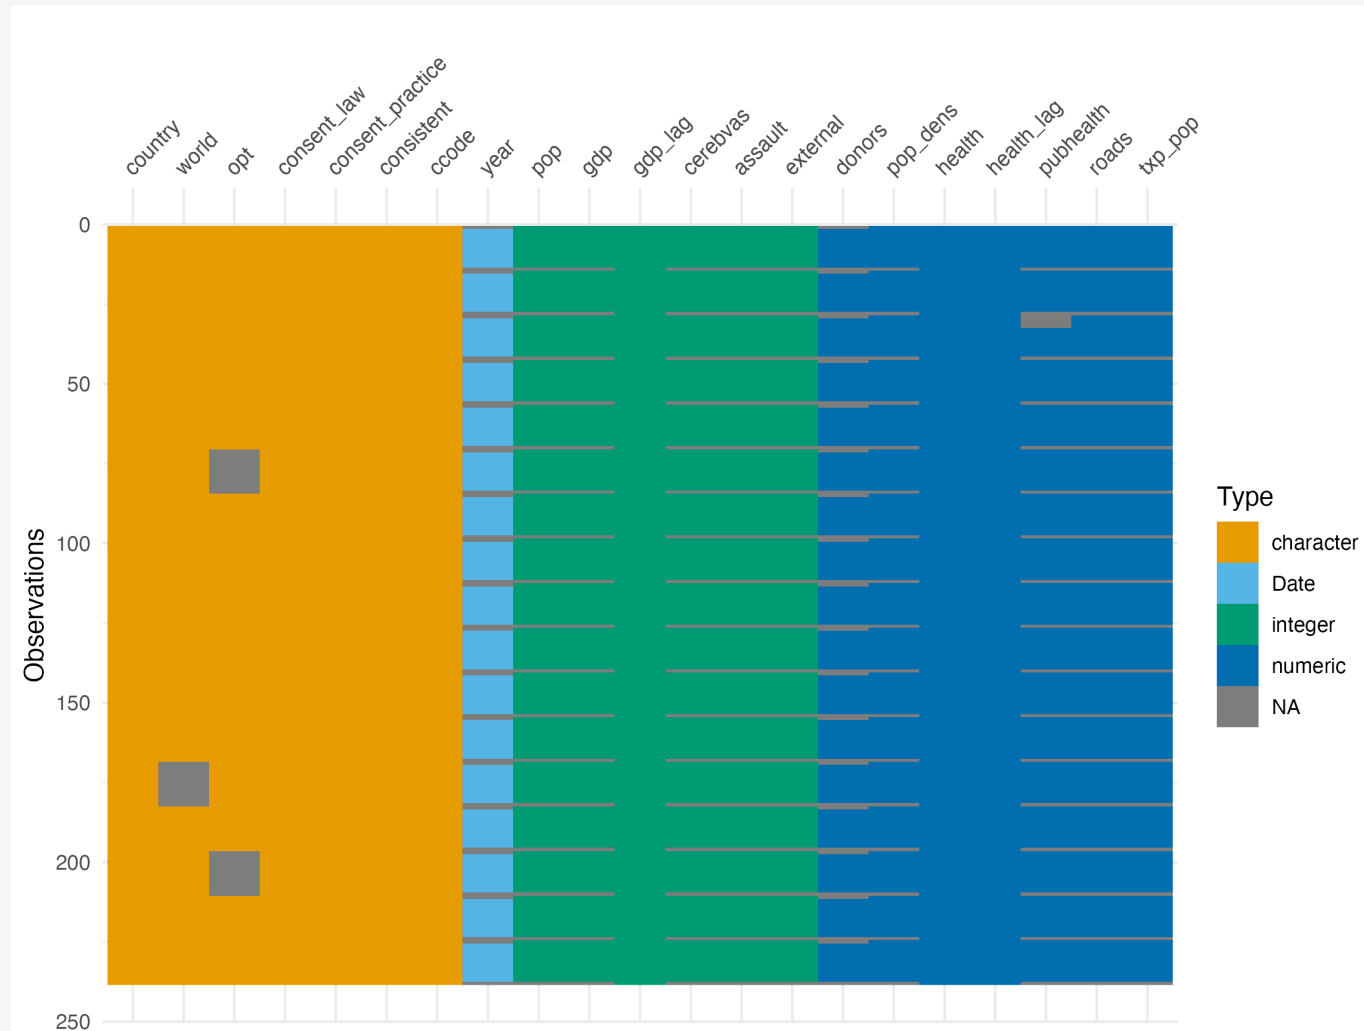
A quick plug for naniar and visdat

`gg_miss_var(organdata)`

A quick plug for **naniar** and **visdat**

```
vis_dat(organdata)
```

A quick plug for naniar and visdat

```
miss_var_summary(organdata)
```

```
# A tibble: 21 × 3
   variable   n_miss pct_miss
   <chr>       <int>    <num>
 1 year          34    14.3
 2 donors        34    14.3
 3 opt           28    11.8
 4 pubhealth     21     8.82
 5 pop           17     7.14
 6 pop_dens      17     7.14
 7 gdp           17     7.14
 8 roads         17     7.14
 9 cerebvas      17     7.14
10 assault       17     7.14
# ℹ 11 more rows
```

A quick plug for naniar and visdat

```
miss_case_summary(organdata)
```

```
# A tibble: 238 × 3
    case n_miss pct_miss
   <int>  <int>    <dbl>
 1    84     12     57.1
 2   182     12     57.1
 3   210     12     57.1
 4    14     11     52.4
 5    28     11     52.4
 6    42     11     52.4
 7    56     11     52.4
 8    70     11     52.4
 9    98     11     52.4
10   112     11     52.4
# ℹ 228 more rows
```
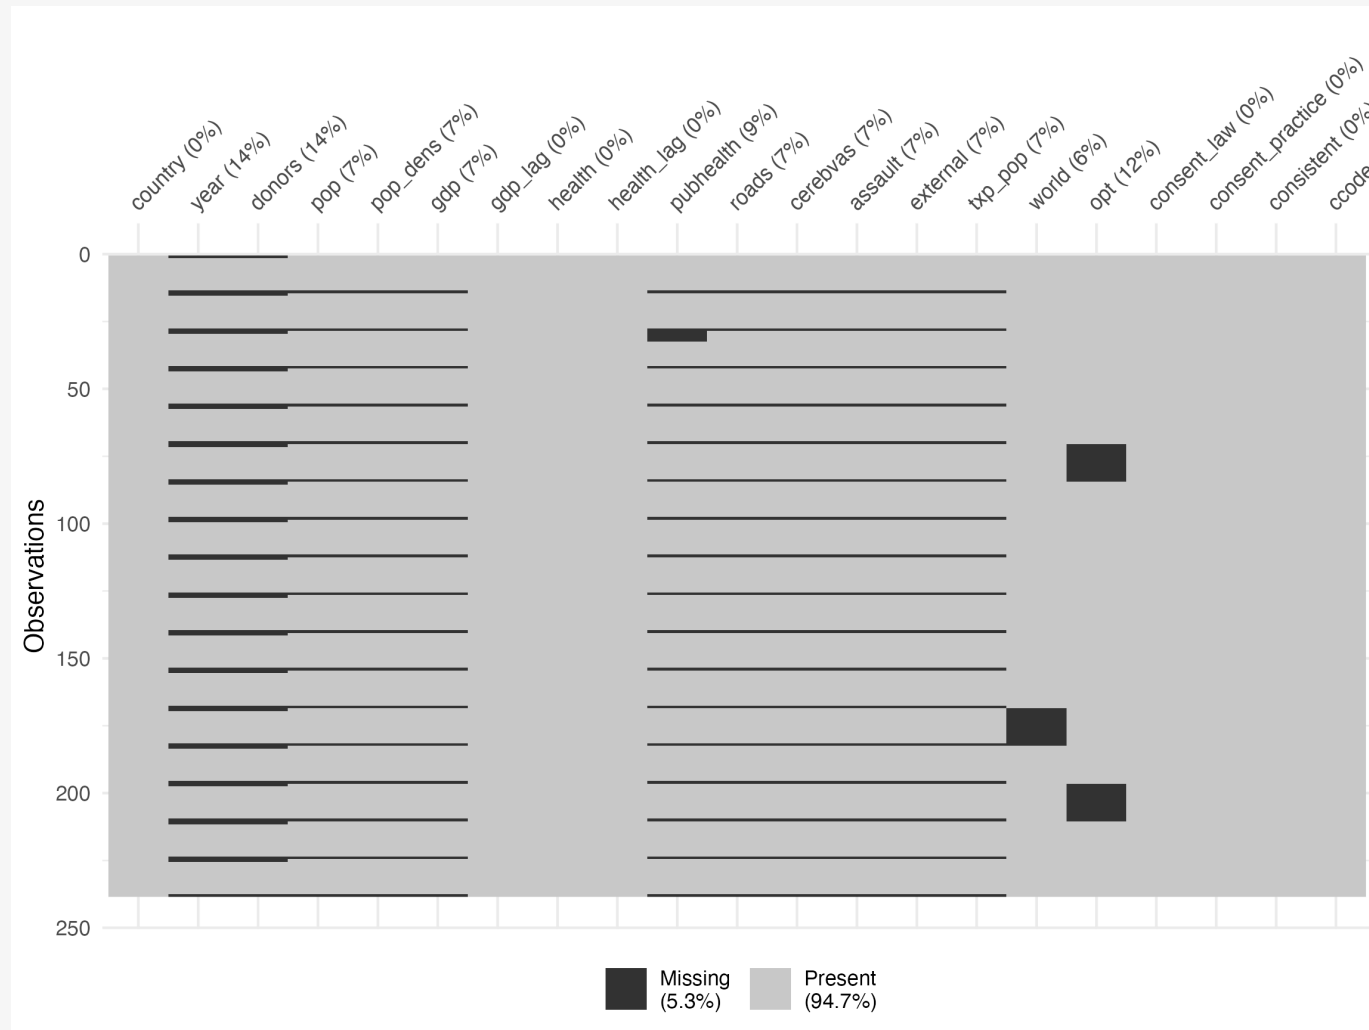
A quick plug for naniar and visdat

```
organdata ▷
  select(consent_law, year, pubhealth, roads) ▷
  group_by(consent_law) ▷
  miss_var_summary()
```

```
# A tibble: 6 × 4
# Groups:   consent_law [2]
  consent_law variable   n_miss pct_miss
  <chr>       <chr>       <int>    <num>
1 Informed    year           16    14.3
2 Informed    pubhealth       8     7.14
3 Informed    roads           8     7.14
4 Presumed    year           18    14.3
5 Presumed    pubhealth      13    10.3
6 Presumed    roads           9     7.14
```

A quick plug for naniar and visdat

```
vis_miss(organdata)
```

A quick plug for naniar and visdat

```
library(dwcongress)
gg_miss_upset(congress)
```

A quick plug for naniar and visdat

```
vis_miss(organdata, cluster = TRUE)
```

A quick plug for naniar and visdat

`gg_miss_upset(organdata)`

Example: Upset Plots

Upset plots and a bit of wrangling

TRACKING SYMPTOMS

On 7 April, around 60% of app users who tested positive for COVID-19 and reported symptoms had lost their sense of smell.

— Anosmia (loss of smell) — Cough — Fatigue
— Diarrhoea — Shortness of breath — Fever

140 •———— Number of people reporting each combination of symptoms

PHOTOGRAPH: PIER MARCO; DATA SOURCE: COVID SYMPTOM TRACKER TEAM

:scale 35%

Upset plots and a bit of wrangling

```r
symptoms <- c("Anosmia", "Cough", "Fatigue",
              "Diarrhea", "Breath", "Fever")
names(symptoms) <- symptoms
symptoms
```

```
   Anosmia       Cough     Fatigue    Diarrhea      Breath       Fever
 "Anosmia"     "Cough"   "Fatigue"  "Diarrhea"    "Breath"     "Fever"
```

Upset plots and a bit of wrangling

```r
# An Excel file!
dat ← readxl::read_xlsx(here("data", "symptoms.xlsx"))
dat ▷ print(n = nrow(dat))
```

```
# A tibble: 32 × 2
   combination                     count
   <chr>                           <dbl>
 1 Anosmia                           140
 2 Cough                              57
 3 Fatigue                           198
 4 Diarrhea                           12
 5 Breath                              5
 6 Fever                              11
 7 Cough&Fatigue                     179
 8 Fatigue&Fever                      28
 9 Breath&Fatigue                     10
10 Diarrhea&Fatigue                   43
11 Anosmia&Fatigue                   281
12 Breath&Cough                        1
13 Anosmia&Diarrhea&Fatigue           64
14 Breath&Cough&Fatigue               22
15 Anosmia&Cough&Fatigue             259
16 Anosmia&Fever&Fatigue              46
```

Upset plots and a bit of wrangling

```r
subsets ← dat ▷
  pull(combination)

## Check if each subset mentions each symptom or not
symptom_mat ← map(subsets, \(x) str_detect(x, symptoms)) ▷
  set_names(nm = subsets) ▷
  map(\(x) set_names(x, nm = symptoms)) ▷
  bind_rows(.id = "subset") ▷
  left_join(dat, join_by(subset == combination))
```

Upset plots and a bit of wrangling

Now we have a table we can do something with.

```
symptom_mat ▷ print(n = nrow(symptom_mat))
```

```
# A tibble: 32 × 8
   subset                Anosmia Cough Fatigue Diarrhea Breath Fever count
   <chr>                 <lgl>   <lgl> <lgl>   <lgl>    <lgl>  <lgl> <dbl>
 1 Anosmia               TRUE    FALSE FALSE   FALSE    FALSE  FALSE   140
 2 Cough                 FALSE   TRUE  FALSE   FALSE    FALSE  FALSE    57
 3 Fatigue               FALSE   FALSE TRUE    FALSE    FALSE  FALSE   198
 4 Diarrhea              FALSE   FALSE FALSE   TRUE     FALSE  FALSE    12
 5 Breath                FALSE   FALSE FALSE   FALSE    TRUE   FALSE     5
 6 Fever                 FALSE   FALSE FALSE   FALSE    FALSE  TRUE     11
 7 Cough&Fatigue         FALSE   TRUE  TRUE    FALSE    FALSE  FALSE   179
 8 Fatigue&Fever         FALSE   FALSE TRUE    FALSE    FALSE  TRUE     28
 9 Breath&Fatigue        FALSE   FALSE TRUE    FALSE    TRUE   FALSE    10
10 Diarrhea&Fatigue      FALSE   FALSE TRUE    TRUE     FALSE  FALSE    43
11 Anosmia&Fatigue       TRUE    FALSE TRUE    FALSE    FALSE  FALSE   281
12 Breath&Cough          FALSE   TRUE  FALSE   FALSE    TRUE   FALSE     1
13 Anosmia&Diarrhea&Fatigue  TRUE  FALSE TRUE  TRUE     FALSE  FALSE    64
14 Breath&Cough&Fatigue  FALSE   TRUE  TRUE    FALSE    TRUE   FALSE    22
15 Anosmia&Cough&Fatigue TRUE    TRUE  TRUE    FALSE    FALSE  FALSE   259
16 Anosmia&Fever&Fatigue TRUE    FALSE TRUE    FALSE    FALSE  TRUE     46
```

Upset plots and a bit of wrangling

Uncounting tables:

```r
indvs ← symptom_mat ▷
    uncount(count)

indvs
```

```
# A tibble: 1,764 × 7
   subset  Anosmia Cough Fatigue Diarrhea Breath Fever
   <chr>   <lgl>   <lgl> <lgl>   <lgl>    <lgl>  <lgl>
 1 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 2 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 3 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 4 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 5 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 6 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 7 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 8 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
 9 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
10 Anosmia TRUE    FALSE FALSE   FALSE    FALSE  FALSE
# ℹ 1,754 more rows
```

Now we've reconstructed the individual-level observations.

Upset plots and a bit of wrangling

```r
# devtools::install_github("krassowski/complex-upset")

library(ComplexUpset)

upset(data = indvs, intersect = symptoms,
      name="Symptom Groupings by Frequency. Total pool is 1,764 individuals.",
      min_size = 0,
      width_ratio = 0.125) +
   labs(title = "Co-Occurence of COVID-19 Symptoms",
        caption = "Data: covid.joinzoe.com/us | Graph: @kjhealy")
```

Wrangling Models

This is not a **statistics** seminar!

I'll just give you an example of the sort of thing that many other modeling packages implement for all kinds of modeling techniques.

Again, the principle is tidy incorporation of models and their output.

Tidy regression output with broom

```r
library(broom)
library(gapminder)
```

```r
out ← lm(formula = lifeExp ~ gdpPercap + pop + continent,
         data = gapminder)
```

Tidy regression output with broom

We can't *do* anything with this, programatically.

```
summary(out)
```

```
Call:
lm(formula = lifeExp ~ gdpPercap + pop + continent, data = gapminder)

Residuals:
    Min      1Q  Median      3Q     Max
-49.161  -4.486   0.297   5.110  25.175

Coefficients:
                   Estimate Std. Error t value Pr(>|t|)
(Intercept)       4.781e+01  3.395e-01 140.819  < 2e-16 ***
gdpPercap         4.495e-04  2.346e-05  19.158  < 2e-16 ***
pop               6.570e-09  1.975e-09   3.326 0.000901 ***
continentAmericas 1.348e+01  6.000e-01  22.458  < 2e-16 ***
continentAsia     8.193e+00  5.712e-01  14.342  < 2e-16 ***
continentEurope   1.747e+01  6.246e-01  27.973  < 2e-16 ***
continentOceania  1.808e+01  1.782e+00  10.146  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Tidy regression output with broom

```
library(broom)
```

```
tidy(out)
```

```
# A tibble: 7 × 5
  term              estimate   std.error statistic  p.value
  <chr>                <dbl>       <dbl>     <dbl>    <dbl>
1 (Intercept)        4.78e+1 0.340           141.   0
2 gdpPercap          4.50e-4 0.0000235        19.2 3.24e- 74
3 pop                6.57e-9 0.00000000198     3.33 9.01e-  4
4 continentAmericas  1.35e+1 0.600            22.5 5.19e- 98
5 continentAsia      8.19e+0 0.571            14.3 4.06e- 44
6 continentEurope    1.75e+1 0.625            28.0 6.34e-142
7 continentOceania   1.81e+1 1.78             10.1 1.59e- 23
```

That's a *lot* nicer. Now it's just a tibble. We know those.

Tidy regression output with broom

```
out_conf <- tidy(out, conf.int = TRUE)
out_conf
```

```
# A tibble: 7 × 7
  term               estimate   std.error statistic   p.value conf.low conf.high
  <chr>                 <dbl>       <dbl>     <dbl>     <dbl>    <dbl>     <dbl>
1 (Intercept)         4.78e+1      3.40e-1    141.     0         4.71e+1   4.85e+1
2 gdpPercap           4.50e-4      2.35e-5     19.2   3.24e- 74  4.03e-4   4.96e-4
3 pop                 6.57e-9      1.98e-9      3.33  9.01e-  4  2.70e-9   1.04e-8
4 continentAmericas   1.35e+1      6.00e-1     22.5   5.19e- 98  1.23e+1   1.47e+1
5 continentAsia       8.19e+0      5.71e-1     14.3   4.06e- 44  7.07e+0   9.31e+0
6 continentEurope     1.75e+1      6.25e-1     28.0   6.34e-142  1.62e+1   1.87e+1
7 continentOceania    1.81e+1      1.78e+0     10.1   1.59e- 23  1.46e+1   2.16e+1
```

Tidy regression output with broom

```
out_conf ▷
    filter(term %nin% "(Intercept)") ▷
    mutate(nicelabs = prefix_strip(term, "continent")) ▷
    select(nicelabs, everything())
```

```
# A tibble: 6 × 8
  nicelabs   term      estimate std.error statistic  p.value conf.low conf.high
  <chr>      <chr>        <dbl>     <dbl>     <dbl>    <dbl>    <dbl>     <dbl>
1 gdpPercap  gdpPercap  4.50e-4   2.35e-5     19.2  3.24e- 74  4.03e-4   4.96e-4
2 Pop        pop        6.57e-9   1.98e-9      3.33 9.01e-  4  2.70e-9   1.04e-8
3 Americas   continent… 1.35e+1   6.00e-1     22.5  5.19e- 98  1.23e+1   1.47e+1
4 Asia       continent… 8.19e+0   5.71e-1     14.3  4.06e- 44  7.07e+0   9.31e+0
5 Europe     continent… 1.75e+1   6.25e-1     28.0  6.34e-142  1.62e+1   1.87e+1
6 Oceania    continent… 1.81e+1   1.78e+0     10.1  1.59e- 23  1.46e+1   2.16e+1
```

Grouped analysis and list columns

```r
eu77 <- gapminder |> filter(continent == "Europe", year == 1977)
fit <- lm(lifeExp ~ log(gdpPercap), data = eu77)
```

```r
summary(fit)
```

```
Call:
lm(formula = lifeExp ~ log(gdpPercap), data = eu77)

Residuals:
    Min      1Q  Median      3Q     Max
-7.4956 -1.0306  0.0935  1.1755  3.7125

Coefficients:
               Estimate Std. Error t value Pr(>|t|)
(Intercept)      29.489      7.161   4.118 0.000306 ***
log(gdpPercap)    4.488      0.756   5.936 2.17e-06 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.114 on 28 degrees of freedom
Multiple R-squared:  0.5572,    Adjusted R-squared:  0.5414
F-statistic: 35.24 on 1 and 28 DF,  p-value: 2.173e-06
```

Grouped analysis and list columns

```
out_le ← gapminder ▷
    group_by(continent, year) ▷
    nest()

out_le
```

```
# A tibble: 60 × 3
# Groups:   continent, year [60]
   continent  year data
   <fct>     <int> <list>
 1 Asia       1952 <tibble [33 × 4]>
 2 Asia       1957 <tibble [33 × 4]>
 3 Asia       1962 <tibble [33 × 4]>
 4 Asia       1967 <tibble [33 × 4]>
 5 Asia       1972 <tibble [33 × 4]>
 6 Asia       1977 <tibble [33 × 4]>
 7 Asia       1982 <tibble [33 × 4]>
 8 Asia       1987 <tibble [33 × 4]>
 9 Asia       1992 <tibble [33 × 4]>
10 Asia       1997 <tibble [33 × 4]>
# ℹ 50 more rows
```

Think of nesting as a kind of "super-grouping". Look in the object inspector.

Grouped analysis and list columns

It's still in there.

```
out_le ▷ filter(continent == "Europe" & year == 1977) ▷
    unnest(cols = c(data))
```

```
# A tibble: 30 × 6
# Groups:   continent, year [1]
   continent  year country                lifeExp      pop gdpPercap
   <fct>     <int> <fct>                    <dbl>    <int>     <dbl>
 1 Europe     1977 Albania                   68.9  2509048     3533.
 2 Europe     1977 Austria                   72.2  7568430    19749.
 3 Europe     1977 Belgium                   72.8  9821800    19118.
 4 Europe     1977 Bosnia and Herzegovina    69.9  4086000     3528.
 5 Europe     1977 Bulgaria                  70.8  8797022     7612.
 6 Europe     1977 Croatia                   70.6  4318673    11305.
 7 Europe     1977 Czech Republic            70.7 10161915    14800.
 8 Europe     1977 Denmark                   74.7  5088419    20423.
 9 Europe     1977 Finland                   72.5  4738902    15605.
10 Europe     1977 France                    73.8 53165019    18293.
# ℹ 20 more rows
```

Grouped analysis and list columns

Here we `map()` a custom function to every row in the `data` column.

```r
fit_ols ← function(df) {
    lm(lifeExp ~ log(gdpPercap), data = df)
}

out_le ← gapminder ▷
    group_by(continent, year) ▷
    nest() ▷
    mutate(model = map(data, fit_ols))
```

Grouped analysis and list columns

```
out_le
```

```
# A tibble: 60 × 4
# Groups:   continent, year [60]
   continent  year data             model
   <fct>     <int> <list>           <list>
 1 Asia       1952 <tibble [33 × 4]> <lm>
 2 Asia       1957 <tibble [33 × 4]> <lm>
 3 Asia       1962 <tibble [33 × 4]> <lm>
 4 Asia       1967 <tibble [33 × 4]> <lm>
 5 Asia       1972 <tibble [33 × 4]> <lm>
 6 Asia       1977 <tibble [33 × 4]> <lm>
 7 Asia       1982 <tibble [33 × 4]> <lm>
 8 Asia       1987 <tibble [33 × 4]> <lm>
 9 Asia       1992 <tibble [33 × 4]> <lm>
10 Asia       1997 <tibble [33 × 4]> <lm>
# ℹ 50 more rows
```

Grouped analysis and list columns

We can tidy the nested models, too.

```r
fit_ols ← function(df) {
    lm(lifeExp ~ log(gdpPercap), data = df)
}

out_tidy ← gapminder ▷
    group_by(continent, year) ▷
    nest() ▷
    mutate(model = map(data, fit_ols),
           tidied = map(model, tidy)) ▷
    unnest(cols = c(tidied)) ▷
    filter(term %nin% "(Intercept)" &
           continent %nin% "Oceania")
```

Grouped analysis and list columns

```
# A tibble: 48 × 9
# Groups:   continent, year [48]
   continent  year data     model  term      estimate std.error statistic p.value
   <fct>     <int> <list>   <list> <chr>         <dbl>     <dbl>     <dbl>   <dbl>
 1 Asia       1952 <tibble> <lm>   log(gdp…       4.16     1.25       3.33 2.28e-3
 2 Asia       1957 <tibble> <lm>   log(gdp…       4.17     1.28       3.26 2.71e-3
 3 Asia       1962 <tibble> <lm>   log(gdp…       4.59     1.24       3.72 7.94e-4
 4 Asia       1967 <tibble> <lm>   log(gdp…       4.50     1.15       3.90 4.77e-4
 5 Asia       1972 <tibble> <lm>   log(gdp…       4.44     1.01       4.41 1.16e-4
 6 Asia       1977 <tibble> <lm>   log(gdp…       4.87     1.03       4.75 4.42e-5
 7 Asia       1982 <tibble> <lm>   log(gdp…       4.78     0.852      5.61 3.77e-6
 8 Asia       1987 <tibble> <lm>   log(gdp…       5.17     0.727      7.12 5.31e-8
 9 Asia       1992 <tibble> <lm>   log(gdp…       5.09     0.649      7.84 7.60e-9
10 Asia       1997 <tibble> <lm>   log(gdp…       5.11     0.628      8.15 3.35e-9
# ℹ 38 more rows
```

Grouped analysis and list columns

```
out_tidy ▷
    ungroup() ▷
    sample_n(5)
```

```
# A tibble: 5 × 9
  continent  year data     model  term       estimate std.error statistic p.value
  <fct>     <int> <list>   <list> <chr>          <dbl>     <dbl>     <dbl>   <dbl>
1 Africa     1957 <tibble> <lm>   log(gdpP…       2.69     1.06       2.55 1.40e-2
2 Europe     2002 <tibble> <lm>   log(gdpP…       3.74     0.445      8.40 3.91e-9
3 Europe     1982 <tibble> <lm>   log(gdpP…       4.37     0.791      5.52 6.63e-6
4 Americas   2007 <tibble> <lm>   log(gdpP…       4.49     0.752      5.98 4.26e-6
5 Europe     1992 <tibble> <lm>   log(gdpP…       3.48     0.545      6.39 6.44e-7
```