Iterating Safely

Data Wrangling, Session 7b

Kieran Healy

Code Horizons

July 29, 2024

Safely iterating with purrr and map

Load the packages, as always

```r
library(here)      # manage file paths
library(socviz)    # data and some useful functions
library(tidyverse) # your friend and mine
```

Additional libraries

```r
library(survey)
library(srvyr)
library(broom)
library(gssr) # https://kjhealy.github.io/gssr
```

The complete GSS

```
data(gss_all)

gss_all
```

```
# A tibble: 72,390 × 6,694
   year         id wrkstat      hrs1         hrs2         evwork         occ prestige
   <dbl+lbl> <dbl> <dbl+lbl>    <dbl+lbl>    <dbl+lbl>    <dbl+lbl>    <dbl> <dbl+lb>
 1 1972          1 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   205   50
 2 1972          2 5 [retire… NA(i) [iap] NA(i) [iap]     1 [yes]   441   45
 3 1972          3 2 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   270   44
 4 1972          4 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]     1   57
 5 1972          5 7 [keepin… NA(i) [iap] NA(i) [iap]     1 [yes]   385   40
 6 1972          6 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   281   49
 7 1972          7 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   522   41
 8 1972          8 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   314   36
 9 1972          9 2 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   912   26
10 1972         10 1 [workin… NA(i) [iap] NA(i) [iap] NA(i) [iap]   984   18
# ℹ 72,380 more rows
# ℹ 6,686 more variables: wrkslf <dbl+lbl>, wrkgovt <dbl+lbl>,
#   commute <dbl+lbl>, industry <dbl+lbl>, occ80 <dbl+lbl>, prestg80 <dbl+lbl>,
#   indus80 <dbl+lbl>, indus07 <dbl+lbl>, occonet <dbl+lbl>, found <dbl+lbl>,
#   occ10 <dbl+lbl>, occindv <dbl+lbl>, occstatus <dbl+lbl>, occtag <dbl+lbl>,
#   prestg10 <dbl+lbl>, prestg105plus <dbl+lbl>, indus10 <dbl+lbl>,
```

Set up our analysis

```r
cont_vars ← c("year", "id", "ballot", "age")
cat_vars ← c("race", "sex", "fefam")
wt_vars ← c("vpsu",
            "vstrat",
            "oversamp",
            "formwt",          # weight to deal with experimental randomization
            "wtssall",         # main weight variable
            "sampcode",        # sampling error code
            "sample")          # sampling frame and method
my_vars ← c(cont_vars, cat_vars, wt_vars)
```

Clean the labeled variables

```r
gss_df ← gss_all ▷
  filter(year > 1974 & year < 2021) ▷
  select(all_of(my_vars)) ▷
  mutate(across(everything(), haven::zap_missing), # Convert labeled missing to regular NA
         across(all_of(wt_vars), as.numeric),
         across(all_of(cat_vars), as_factor),
         across(all_of(cat_vars), fct_relabel, tolower),
         across(all_of(cat_vars), fct_relabel, tools::toTitleCase),
         compwt = oversamp * formwt * wtssall)
```

Working dataset

```
gss_df
```

```
# A tibble: 60,213 × 15
   year          id ballot   age   race  sex   fefam vpsu vstrat oversamp formwt
   <dbl+lbl> <dbl> <dbl+lb> <dbl> <fct> <fct> <fct> <dbl>  <dbl>    <dbl>  <dbl>
 1 1975       1 NA        38    White Male  <NA>     1   7001        1     NA
 2 1975       2 NA        20    White Fema… <NA>     1   7001        1     NA
 3 1975       3 NA        61    White Fema… <NA>     1   7001        1     NA
 4 1975       4 NA        19    White Male  <NA>     1   7001        1     NA
 5 1975       5 NA        28    White Male  <NA>     1   7001        1     NA
 6 1975       6 NA        28    White Fema… <NA>     1   7002        1     NA
 7 1975       7 NA        35    White Fema… <NA>     1   7002        1     NA
 8 1975       8 NA        64    White Fema… <NA>     1   7002        1     NA
 9 1975       9 NA        53    White Male  <NA>     1   7002        1     NA
10 1975      10 NA        34    White Fema… <NA>     1   7002        1     NA
# ℹ 60,203 more rows
# ℹ 4 more variables: wtssall <dbl>, sampcode <dbl>, sample <dbl>, compwt <dbl>
```

The `fefam` question

```
gss_df ▷
  count(fefam)
```

```
# A tibble: 5 × 2
  fefam                 n
  <fct>             <int>
1 Strongly Agree     2543
2 Agree              8992
3 Disagree          13061
4 Strongly Disagree  5479
5 <NA>              30138
```

Recoding

```r
gss_df <- gss_df |>
  mutate(fefam_d = forcats::fct_recode(fefam,
                                       Agree = "Strongly Agree",
                                       Disagree = "Strongly Disagree"),
    fefam_n = recode(fefam_d, "Agree" = 1, "Disagree" = 0))

# factor version
gss_df |>
  count(fefam_d)
```

```
# A tibble: 3 × 2
  fefam_d        n
  <fct>      <int>
1 Agree      11535
2 Disagree   18540
3 <NA>       30138
```

```r
# numeric version, 1 is "Agree"
gss_df |>
  count(fefam_n)
```

```
# A tibble: 3 × 2
  fefam_n        n
    <dbl> <int>
1       0 18540
2       1 11535
3      NA 30138
```

Unweighted model

```r
out_all ← glm(fefam_n ~ age + sex + race,
              data = gss_df,
              family="binomial",
              na.action = na.omit)


summary(out_all)
```

```
Call:
glm(formula = fefam_n ~ age + sex + race, family = "binomial",
    data = gss_df, na.action = na.omit)

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.9185878  0.0399581 -48.015  < 2e-16 ***
age          0.0323648  0.0007275  44.486  < 2e-16 ***
sexFemale   -0.2247518  0.0248741  -9.036  < 2e-16 ***
raceBlack    0.0668275  0.0363201   1.840   0.0658 .
raceOther    0.3659411  0.0493673   7.413 1.24e-13 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 39921  on 29980  degrees of freedom
Residual deviance: 37746  on 29976  degrees of freedom
```

Tidied output

```
tidy(out_all)
```

```
# A tibble: 5 × 5
  term          estimate std.error statistic  p.value
  <chr>            <dbl>     <dbl>     <dbl>    <dbl>
1 (Intercept)  -1.92     0.0400      -48.0  0
2 age           0.0324   0.000728     44.5  0
3 sexFemale    -0.225    0.0249       -9.04 1.63e-19
4 raceBlack     0.0668   0.0363        1.84 6.58e- 2
5 raceOther     0.366    0.0494        7.41 1.24e-13
```

group_map() and possibly()

Model each year

```r
out_yr ← gss_df ▷
  group_by(year) ▷
  group_map_dfr(possibly(~ tidy(glm(fefam_n ~ age + sex + race,
                    data = .x,
                    family = "binomial",
                    na.action = na.omit),
               conf.int = TRUE),
           otherwise = NULL))

out_yr
```

```
# A tibble: 105 × 8
   year      term       estimate std.error statistic  p.value conf.low conf.high
   <dbl+lbl> <chr>         <dbl>     <dbl>     <dbl>    <dbl>    <dbl>     <dbl>
 1 1977      (Intercep…  -1.20     0.178      -6.75   1.47e-11  -1.55    -0.854
 2 1977      age          0.0483   0.00388    12.4    1.56e-35   0.0408   0.0561
 3 1977      sexFemale   -0.341    0.118      -2.90   3.77e- 3  -0.572   -0.111
 4 1977      raceBlack   -0.0613   0.180      -0.340  7.34e- 1  -0.412    0.295
 5 1977      raceOther    0.188    0.576       0.326  7.44e- 1  -0.912    1.40
 6 1985      (Intercep…  -1.89     0.168     -11.2    2.89e-29  -2.23    -1.56
 7 1985      age          0.0432   0.00332    13.0    1.03e-38   0.0368   0.0498
 8 1985      sexFemale   -0.261    0.112      -2.34   1.94e- 2  -0.481   -0.0426
 9 1985      raceBlack    0.148    0.189       0.782  4.34e- 1  -0.223    0.519
10 1985      raceOther   -0.319    0.338      -0.944  3.45e- 1  -1.00     0.329
```

group_map() and possibly()

```r
possibly(~ tidy(glm(...)), otherwise = NULL)
```

group_map() and possibly()

```
out_yr ▷
  filter(term == "sexFemale") ▷
  ggplot(mapping = aes(x = year, y = estimate,
                       ymin = conf.low, ymax = conf.high)) +
  geom_hline(yintercept = 0, linetype = "dotted") +
  geom_line() +
  geom_pointrange()
```
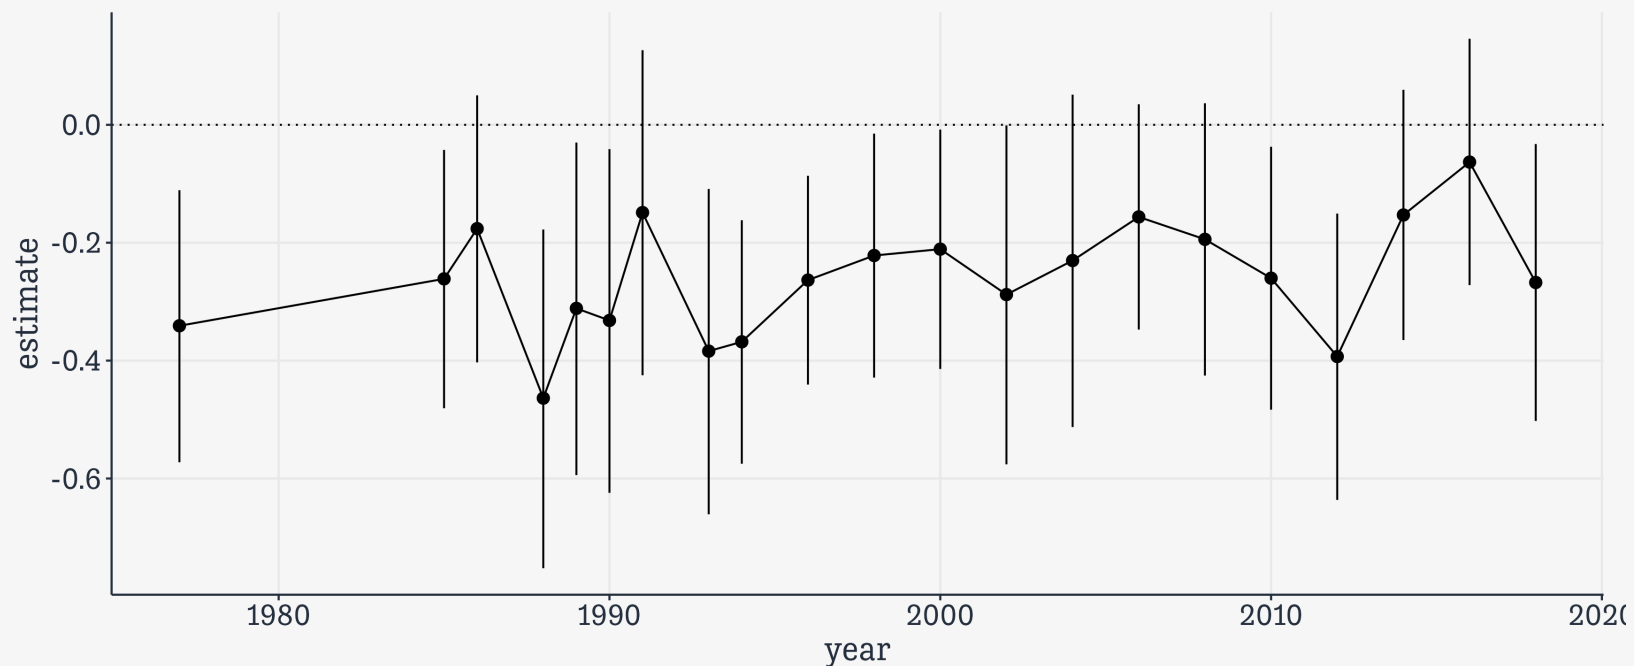
Survey-weighted estimates

```r
options(survey.lonely.psu = "adjust")
options(na.action="na.pass")


gss_svy ← gss_df ▷
  filter(year > 1974) ▷
  mutate(stratvar = interaction(year, vstrat)) ▷
  as_survey_design(id = vpsu,
                   strata = stratvar,
                   weights = wtssall,
                   nest = TRUE)
gss_svy
```

```
Stratified 1 - level Cluster Sampling design (with replacement)
With (4555) clusters.
Called via srvyr
Sampling variables:
 - ids: vpsu
 - strata: stratvar
 - weights: wtssall
Data variables: year (dbl+lbl), id (dbl), ballot (dbl+lbl), age (dbl+lbl), race
  (fct), sex (fct), fefam (fct), vpsu (dbl), vstrat (dbl), oversamp (dbl),
  formwt (dbl), wtssall (dbl), sampcode (dbl), sample (dbl), compwt (dbl),
  fefam_d (fct), fefam_n (dbl), stratvar (fct)
```

Survey-weighted estimates

```
gss_svy ▷
  drop_na(fefam_d) ▷
  group_by(year, sex, race, fefam_d) ▷
  summarize(prop = survey_mean(na.rm = TRUE,
                               vartype = "ci"))
```

```
# A tibble: 252 × 7
# Groups:   year, sex, race [126]
   year      sex    race   fefam_d    prop prop_low prop_upp
   <dbl+lbl> <fct>  <fct>  <fct>     <dbl>    <dbl>    <dbl>
 1 1977      Male   White  Agree     0.694   0.655    0.732
 2 1977      Male   White  Disagree  0.306   0.268    0.345
 3 1977      Male   Black  Agree     0.686   0.564    0.807
 4 1977      Male   Black  Disagree  0.314   0.193    0.436
 5 1977      Male   Other  Agree     0.632   0.357    0.906
 6 1977      Male   Other  Disagree  0.368   0.0936   0.643
 7 1977      Female White  Agree     0.640   0.601    0.680
 8 1977      Female White  Disagree  0.360   0.320    0.399
 9 1977      Female Black  Agree     0.553   0.472    0.634
10 1977      Female Black  Disagree  0.447   0.366    0.528
# i 242 more rows
```

Survey-weighted estimates

```
out_svy_all ← svyglm(fefam_n ~ age + sex + race,
                     design = gss_svy,
                     family = quasibinomial(),
                     na.action = na.omit)


tidy(out_svy_all)
```

```
# A tibble: 5 × 5
  term         estimate std.error statistic    p.value
  <chr>           <dbl>     <dbl>     <dbl>      <dbl>
1 (Intercept)   -1.83     0.0478     -38.3   6.34e-234
2 age            0.0310   0.000852    36.4   9.99e-217
3 sexFemale     -0.235    0.0277      -8.48  4.55e- 17
4 raceBlack      0.0282   0.0432       0.653 5.14e-  1
5 raceOther      0.382    0.0588       6.50  1.06e- 10
```

Survey-weighted estimates

```r
out_svy_yrs ← gss_svy ▷
  group_by(year) ▷
  group_map_dfr(possibly(~ tidy(svyglm(fefam_n ~ age + sex + race,
                          design = .x,
                          family = quasibinomial(),
                          na.action = na.omit),
                    conf.int = TRUE),
                otherwise = NULL))

out_svy_yrs
```

```
# A tibble: 105 × 8
   year       term       estimate std.error statistic  p.value conf.low conf.high
   <dbl+lbl>  <chr>         <dbl>     <dbl>     <dbl>    <dbl>    <dbl>     <dbl>
 1 1977       (Intercep…   -1.09     0.184     -5.93  3.74e- 7   -1.46    -0.720
 2 1977       age           0.0469   0.00403   11.6   2.63e-15    0.0388   0.0550
 3 1977       sexFemale    -0.344    0.126     -2.73  9.05e- 3   -0.599   -0.0901
 4 1977       raceBlack    -0.144    0.215     -0.669 5.07e- 1   -0.576    0.288
 5 1977       raceOther     0.276    0.552      0.500 6.19e- 1   -0.835    1.39
 6 1985       (Intercep…   -1.89     0.199     -9.49  9.05e-13   -2.29    -1.49
 7 1985       age           0.0431   0.00369   11.7   6.47e-16    0.0357   0.0505
 8 1985       sexFemale    -0.174    0.123     -1.42  1.61e- 1   -0.421    0.0720
 9 1985       raceBlack     0.157    0.228      0.688 4.95e- 1   -0.301    0.614
10 1985       raceOther    -0.533    0.268     -1.99  5.24e- 2   -1.07     0.00573
# ℹ 95 more rows
```

Survey-weighted estimates

```
out_svy_yrs ▷
  filter(term == "sexFemale") ▷
  ggplot(mapping = aes(x = year,
                       y = estimate,
                       ymin = conf.low,
                       ymax = conf.high)) +
  geom_hline(yintercept = 0, linetype = "dotted") +
  geom_line() +
  geom_pointrange()
```